The Logtalk Handbook

Release v3.21.0-b01

Paulo Moura

Oct 13, 2018

CONTENTS:

USER MANUAL

1.1 Main features

Several years ago, I decided that the best way to learn object-oriented programming was to build my own object-oriented language. Prolog being always my favorite language, I chose to extend it with object-oriented capabilities. Strong motivation also come from my frustration with Prolog shortcomings for writing large applications. Eventually this work has led to the Logtalk programming language as its know today. The first system to use the name Logtalk appeared in February 1995. At that time, Logtalk was mainly an experiment in computational reflection with a rudimentary runtime and no compiler. Based on feedback by users and on the author subsequent work, the name was retained and Logtalk as created as a full programming language focusing on using object-oriented concepts for code encapsulation and reuse. Its first public alpha version was released in July 1998 with the first stable release (2.0) following in February 1999. Development of the third generation of Logtalk started in 2012 with the first public alpha version in August 2012 and the first stable release (3.0.0) in January 2015.

Logtalk provides the following features:

1.1.1 Integration of logic and object-oriented programming

Logtalk tries to bring together the main advantages of these two programming paradigms. On one hand, the object orientation allows us to work with the same set of entities in the successive phases of application development, giving us a way of organizing and encapsulating the knowledge of each entity within a given domain. On the other hand, logic programming allows us to represent, in a declarative way, the knowledge we have of each entity. Together, these two advantages allow us to minimize the distance between an application and its problem domain, turning the writing and maintenance of programming easier and more productive.

In a more pragmatically view, Logtalk objects provide Prolog with the possibility of defining several namespaces, instead of the traditional Prolog single database, addressing some of the needs of large software projects.

1.1.2 Integration of event-driven and object-oriented programming

Event-driven programming enables the building of reactive systems, where computing which takes place at each moment is a result of the observation of occurring events. This integration complements object-oriented programming, in which each computing is initiated by the explicit sending of a message to an object. The user dynamically defines what events are to be observed and establishes monitors for these events. This is specially useful when representing relationships between objects that imply constraints in the state of participating objects [Rumbaugh87], [Rumbaugh88], [Fornarino_et_al_89], [Razek92]. Other common uses are reflective applications like code debugging or profiling [Maes87]. Predicates can be implicitly called when a spied event occurs, allowing programming solutions which minimize object

coupling. In addition, events provide support for behavioral reflection and can be used to implement the concepts of *pointcut* and *advice* found on Aspect-Oriented Programming.

1.1.3 Support for component-based programming

Predicates can be encapsulated inside *categories* which can be imported by any object, without any code duplication and irrespective of object hierarchies. A category is a first-class encapsulation entity, at the same level as objects and protocols, which can be used as a component when building new objects. Thus, objects may be defined through composition of categories, which act as fine-grained units of code reuse. Categories may also extend existing objects. Categories can be used to implement *mixins* and *aspects*. Categories allows for code reuse between non-related objects, independent of hierarchy relations, in the same vein as protocols allow for interface reuse.

1.1.4 Support for both prototype and class-based systems

Almost any (if not all) object-oriented languages available today are either class-based or prototype-based [Lieberman86], with a strong predominance of class-based languages. Logtalk provides support for both hierarchy types. That is, we can have both prototype and class hierarchies in the same application. Prototypes solve a problem of class-based systems where we sometimes have to define a class that will have only one instance in order to reuse a piece of code. Classes solves a dual problem in prototype based systems where it is not possible to encapsulate some code to be reused by other objects but not by the encapsulating object. Stand-alone objects, that is, objects that do not belong to any hierarchy, are a convenient solution to encapsulate code that will be reused by several unrelated objects.

1.1.5 Support for multiple object hierarchies

Languages like Smalltalk-80 [Goldberg83], Objective-C [Cox86] and Java [Joy_et_al_00] define a single hierarchy rooted in a class usually named `Object`. This makes it easy to ensure that all objects share a common behavior but also tends to result in lengthy hierarchies where it is difficult to express objects which represent exceptions to default behavior. In Logtalk we can have multiple, independent, object hierarchies. Some of them can be prototype-based while others can be class-based. Furthermore, stand-alone objects provide a simple way to encapsulate utility predicates that do not need or fit in an object hierarchy.

1.1.6 Separation between interface and implementation

This is an expected (should we say standard ?) feature of almost any modern programming language. Logtalk provides support for separating interface from implementation in a flexible way: predicate directives can be contained in an object, a category or a protocol (first-order entities in Logtalk) or can be spread in both objects, categories and protocols.

1.1.7 Private, protected and public inheritance

Logtalk supports private, protected and public inheritance in a similar way to C++ [Stroustrup86], enabling us to restrict the scope of inherited, imported or implemented predicates (by default inheritance is public).

1.1.8 Private, protected and public object predicates

Logtalk supports data hiding by implementing private, protected and public object predicates in a way similar to C++ [Stroustrup86]. Private predicates can only be called from the container object. Protected predicates can be called by the container object or by the container descendants. Public predicates can be called from any object.

1.1.9 Parametric objects

Object names can be compound terms (instead of atoms), providing a way to parameterize object predicates. Parametric objects are implemented in a similar way to `L&O` [McCabe92], `OL(P)` [Fromherz93] or `SICStus Objects` [SICStus95] (however, access to parameter values is done via a built-in method instead of making the parameters scope global over the whole object). Parametric objects allows us to treat any predicate clause as defining an *instantiation* of a parametric object. Thus, a parametric object allows us to encapsulate and associate any number of predicates with a compound term.

1.1.10 High level multi-threading programming support

High level multi-threading programming is available when running Logtalk with selected back-end Prolog compilers, allowing objects to support both synchronous and asynchronous messages. Logtalk allows programmers to take advantage of modern multi-processor and multi-core computers without bothering with the details of creating and destroying threads, implement thread communication, or synchronizing threads.

1.1.11 Smooth learning curve

Logtalk has a smooth learning curve, by adopting standard Prolog syntax and by enabling an incremental learning and use of most of its features.

1.1.12 Compatibility with most Prolog systems and the ISO standard

The Logtalk system has been designed to be compatible with most Prolog compilers and, in particular, with the ISO Prolog standard [ISO95]. It runs in almost any computer system with a modern Prolog compiler.

1.1.13 Performance

The current Logtalk implementation works as a trans-compiler: Logtalk source files are first compiled to Prolog source files, which are then compiled by the chosen Prolog compiler. Therefore, Logtalk performance necessarily depends on the back-end Prolog compiler. The Logtalk compiler preserves the programmers choices when writing efficient code that takes advantage of tail recursion and first-argument indexing.

As an object-oriented language, Logtalk can use both *static binding* and *dynamic binding* for matching messages and methods. Furthermore, Logtalk entities (objects, protocols, and categories) are independently compiled, allowing for a very flexible programming development. Entities can be edited, compiled, and loaded at runtime, without necessarily implying recompilation of all related entities.

When dynamic binding is used, the Logtalk runtime engine implements caching of method lookups (including messages to *self* and *super* calls), ensuring a performance level close to what could be achieved when using static binding.

1.1.14 Logtalk scope

Logtalk, being a superset of Prolog, shares with it the same preferred areas of application but also extends them with those areas where object-oriented features provide an advantage compared to plain Prolog. Among these areas we have:

Logic and object-oriented programming teaching and researching Logtalk smooth learning curve, combined with support for both prototype and class-based programming, protocols, components or aspects via category-based composition, and other advanced object-oriented features allow a smooth introduction to object-oriented programming to people with a background in Prolog programming. The distribution of Logtalk source code using an open-source license provides a framework for people to learn and then modify to try out new ideas on object-oriented programming research. In addition, the Logtalk distribution includes plenty of programming examples that can be used in the classroom for teaching logic and object-oriented programming concepts.

Structured knowledge representations and knowledge-based systems Logtalk objects, coupled with event-driven programming features, enable easy implementation of frame-like systems and similar structured knowledge representations.

Blackboard systems, agent-based systems, and systems with complex object relationships Logtalk support for event-driven programming can provide a basis for the dynamic and reactive nature of blackboard type applications.

Highly portable applications Logtalk is compatible with most modern Prolog systems that support official and de facto standards. Used as a way to provide Prolog with namespaces, it avoids the porting problems of most Prolog module systems. Platform, operating system, or compiler specific code can be isolated from the rest of the code by encapsulating it in objects with well-defined interfaces.

Alternative to a Prolog module system Logtalk can be used as an alternative to a Prolog compiler module system. Most Prolog applications that use modules can be converted into Logtalk applications, improving portability across Prolog systems and taking advantage of the stronger encapsulation and reuse framework provided by Logtalk object-oriented features.

Integration with other programming languages Logtalk support for most key object-oriented features helps users integrating Prolog with object-oriented languages like C++, Java, or Smalltalk by facilitating a high-level mapping between the two languages.

1.2 Nomenclature

Depending on your Object-oriented Programming background (or lack of it), you may find Logtalk nomenclature either familiar or at odds with the terms used in other languages. In addition, being a superset of Prolog, terms such as *predicate* and *method* are often used interchangeably. Logtalk inherits most of its nomenclature from Smalltalk, arguably (and somehow sadly) not the most popular OOP language nowadays. In this section, we map nomenclatures from popular OOP languages such as C++ and Java to the Logtalk nomenclature.

1.2.1 C++ nomenclature

There are several C++ glossaries available on the Internet. The list that follows relates the most commonly used C++ terms with their Logtalk equivalents.

abstract class Logtalk uses an *operational* definition of abstract class: any class that does not inherit a method for creating new instances can be considered an abstract class. Moreover, Logtalk supports *interfaces/protocols*, which are often a better way to provide the functionality of C++ abstract classes.

base class Logtalk uses the term *superclass* with the same meaning.

data member Logtalk uses *predicates* for representing both behavior and data.

constructor function There are no special methods for creating new objects in Logtalk. Instead, Logtalk provides a built-in predicate, *create_object/4*, which can be used as a building block to define more sophisticated object creation predicates.

derived class Logtalk uses the term *subclass* with the same meaning.

destructor function There are no special methods for deleting new objects in Logtalk. Instead, Logtalk provides a built-in predicate, *abolish_object/1*, which is often used to define more sophisticated object deletion predicates.

friend function Not supported in Logtalk. Nevertheless, see the manual section on *meta-predicates*.

instance In Logtalk, an instance can be either created dynamically at runtime or defined statically in a source file in the same way as classes.

member Logtalk uses the term *predicate*.

member function Logtalk uses *predicates* for representing both behavior and data.

namespace Logtalk does not support multiple identifier namespaces. All Logtalk entity identifiers share the same namespace (Logtalk entities are objects, categories, and protocols).

nested class Logtalk does not support nested classes.

template Logtalk supports *parametric objects*, which allows you to get the similar functionality of templates at runtime.

this Logtalk uses the built-in context method *self/1* for retrieving the current instance. Logtalk also provides a *this/1* method but for returning the class containing the method being executed. Why the name clashes? Well, the notion of *self* was inherited from Smalltalk, which predates C++.

virtual member function There is no `virtual` keyword in Logtalk. Any inherited or imported predicate can be redefined (either overridden or specialized). Logtalk can use *static binding* or *dynamic binding* for locating both method declarations and method definitions. Moreover, methods that are declared but not defined simply fail when called (as per closed-world assumption).

1.2.2 Java nomenclature

There are several Java glossaries available on the Internet. The list that follows relates the most commonly used Java terms with their Logtalk equivalents.

abstract class Logtalk uses an *operational* definition of abstract class: any class that does not inherit a method for creating new instances is an abstract class. I.e. there is no `abstract` keyword in Logtalk.

abstract method In Logtalk, you may simply declare a method (*predicate*) in a class without defining it, leaving its definition to some descendant sub-class.

assertion There is no `assertion` keyword in Logtalk. Assertions are supported using Logtalk compilation hooks and developer tools.

extends There is no `extends keyword` in Logtalk. Class inheritance is indicated using *specialization relations*. Moreover, the *extends relation* is used in Logtalk to indicate protocol, category, or prototype extension.

interface Logtalk uses the term *protocol* with the same meaning.

callback method Logtalk supports *event-driven programming*, the most common use context of callback methods.

class method Class methods may be implemented in Logtalk by using a *metaclass* for the class and defining the class methods in the metaclass. I.e. class methods are simply instance methods of the class metaclass.

class variable True class variables may be implemented in Logtalk by using a *metaclass* for the class and defining the class variables in the class. I.e. class variables are simply instance variables of the class metaclass. Shared instance variables may be implemented by using the built-in database methods (which can be used to implement

variable assignment) to access and updated a single occurrence of the variable stored in the class (there is no `static` keyword in Logtalk).

constructor There are no special methods for creating new objects in Logtalk. Instead, Logtalk provides a built-in predicate, *create_object/4*, which is often used to define more sophisticated object creation predicates.

final There is no `final` keyword in Logtalk; methods may always be redefined in subclasses (and instances!).

inner class Inner classes are not supported in Logtalk.

instance In Logtalk, an instance can be either created dynamically at runtime or defined statically in a source file in the same way as classes.

method Logtalk uses the term *predicate* interchangeably with the term *method*.

method call Logtalk usually uses the expression *message sending* for method calls, true to its Smalltalk heritage.

method signature Logtalk selects the method/predicate to execute in order to answer a method call based only on the method name and number of arguments. Logtalk (and Prolog) are not typed languages in the same sense as Java.

package There is no concept of packages in Logtalk. All Logtalk entities (objects, protocols, categories) share a single namespace. But Logtalk does support a concept of *library* that allows grouping of entities whose source files share a common prefix.

reflection Logtalk supports both *structural reflection* (using a set of built-in predicates and built-in methods) and *behavioral reflection* (using *event-driven programming*).

static There is no `static` keyword in Logtalk. See the entries on *class methods* and *class variables*.

super Instead of a `super` keyword, Logtalk provides a super operator, *^^/1*, for calling overridden methods.

synchronized Logtalk supports *multi-threading programming* in selected Prolog compilers, including a *synchronized/1* predicate directive. Logtalk allows you to synchronize a predicate or a set of predicates using per-predicate or per-predicate-set *mutexes*.

this Logtalk uses the built-in context method *self/1* for retrieving the current instance. Logtalk also provides a *this/1* method but for returning the class containing the predicate clause being executed. Why the name clashes? Well, the notion of *self* was inherited from Smalltalk, which predates Java.

1.3 Messages

Messages allows us to call object predicates. Logtalk uses the same nomenclature found in other object-oriented programming languages such as Smalltalk. Therefore, the terms *predicate* and *method* are often used interchangeably when referring to predicates defined inside objects and categories. A message must always match a predicate within the scope of the sender object.

Note that message sending is only the same as calling an object's predicate if the object does not inherit (or import) predicate definitions from other objects (or categories). Otherwise, the predicate definition that will be executed may depend on the relationships of the object with other Logtalk entities.

1.3.1 Operators used in message sending

Logtalk declares the following operators for the message sending control constructs:

```
:- op(600, xfy, ::).
:- op(600,  fy, ::).
:- op(600,  fy, ^^).
```

It is assumed that these operators remain active (once the Logtalk compiler and runtime files are loaded) until the end of the Prolog session (this is the usual behavior of most Prolog compilers). Note that these operator definitions are compatible with the predefined operators in the Prolog ISO standard.

1.3.2 Sending a message to an object

Sending a message to an object is accomplished by using the *::/2* control construct:

```
| ?- Object::Message.
```

The message must match a public predicate declared for the receiving object. The message may also correspond to a protected or private predicate if the *sender* matches the predicate scope container. If the predicate is declared but not defined, the message simply fails (as per the closed-world assumption).

1.3.3 Delegating a message to an object

It is also possible to send a message to an object while preserving the original *sender* by using the *[]/1* delegation control construct:

```
..., [Object::Message], ....
```

This control construct can only be used within objects and categories (at the interpreter top-level, the *sender* is always the pseudo-object *user* so using this control construct would be equivalent to use the `::/2` message sending control construct).

1.3.4 Sending a message to *self*

While defining a predicate, we sometimes need to send a message to *self*, i.e., to the same object that has received the original message. This is done in Logtalk through the *::/1* control construct:

```
::Message
```

The message must match either a public or protected predicate declared for the receiving object or a private predicate within the scope of the *sender* otherwise an error will be thrown (see the Reference Manual for details). If the message is sent from inside a category or if we are using private inheritance, then the message may also match a private predicate. Again, if the predicate is declared but not defined, the message simply fails (as per the closed-world assumption).

1.3.5 Broadcasting

In the Logtalk context, broadcasting is interpreted as the sending of several messages to the same object. This can be achieved by using the message sending method described above. However, for convenience, Logtalk implements an extended syntax for message sending that may improve program readability in some cases. This extended syntax uses the `(,)/2`, `(;)/2`, and `(->)/2` control constructs. For example, if we wish to send several messages to the same object, we can write:

```
| ?- Object::(Message1, Message2, ...).
```

This is semantically equivalent to:

```
| ?- Object::Message1, Object::Message2, ... .
```

This extended syntax may also be used with the `::/1` message sending control construct.

1.3.6 Calling imported and inherited predicate definitions

When redefining a predicate, sometimes we need to call the inherited definition in the new code. This functionality, introduced by the Smalltalk language through the `super` primitive, is available in Logtalk using the *^^/1* control construct:

```
^^Predicate
```

Most of the time we will use this control construct by instantiating the pattern:

```
Predicate :-
    ...,            % do something
    ^^Predicate,    % call inherited definition
    ... .           % do something more
```

This control construct is generalized in Logtalk where it may be used to call any imported or inherited predicate definition. This control construct may be used within objects and categories. When combined with *static binding*, this control construct allows imported and inherited predicates to be called with the same performance of local predicates. As with the message sending control constructs, the `^^/1` call simply fails when the predicate is declared but not defined (as per the closed-world assumption).

1.3.7 Message sending and event generation

Every message sent using the *::/2* control construct generates two events, one before and one after the message execution. Messages that are sent using the *::/1* (message to *self*) control construct or the *^^/1* super mechanism described above do not generate any events. The rationale behind this distinction is that messages to *self* and *super* calls are only used internally in the definition of methods or to execute additional messages with the same target object (represented by *self*). In other words, events are only generated when using an object's public interface; they cannot be used to break object encapsulation.

If we need to generate events for a public message sent to *self*, then we just need to write something like:

```
Predicate :-
    ...,
    % get self reference
    self(Self),
    % send a message to self using ::/2
    Self::Message,
    ... .
```

If we also need the sender of the message to be other than the object containing the predicate definition, we can write:

```
Predicate :-
    ...,
    % send a message to self using ::/2
    % sender will be the pseudo-object user
    self(Self),
    {Self::Message},
    ... .
```

When events are not used, is possible to turn off event generation on a per object basis by using the *events/1* compiler flag (see the *Event-driven programming* section for more details).

1.3.8 Message sending performance

Logtalk supports both *static binding* and *dynamic binding*. Static binding is used whenever messages are sent (using the `::/2` control construct) to static objects already loaded and with the `optimize` compiler flag turned on. When that is not the case (or when using the `::/1` control construct), Logtalk uses dynamic binding coupled with a caching mechanism that avoids repeated lookups of predicate declarations and predicate definitions. This is a solution common to other programming languages supporting dynamic binding. Message lookups are automatically cached the first time a message is sent. Cache entries are automatically removed when loading entities or using Logtalk dynamic features that invalidate the cached lookups.

Whenever static binding is used, message sending performance is roughly the same as a predicate call in plain Prolog. When discussing Logtalk dynamic binding performance, two distinct cases should be considered: messages sent by the user from the top-level interpreter and messages sent from compiled objects. In addition, the message declaration and definition lookups may, or may not be already cached by the runtime engine. In what follows, we will assume that the message lookups are already cached.

Translating message processing to predicate calls

In order to better understand the performance trade-offs of using Logtalk dynamic binding when compared to plain Prolog or to Prolog module systems, is useful to translate message processing in terms of predicate calls. However, in doing this, we should keep in mind that the number of predicate calls is not necessarily proportional to the time taken to execute them.

With event-support turned on, a message sent from a compiled object (or category) to another object translates to a minimum of five predicate calls:

checking for *before* events one call to the built-in predicate `\+/1` and a call to its argument, assuming that no events are defined

method call using the cached lookup one call to a dynamic predicate (the cache entry)

checking for *after* events one call to the built-in predicate `\+/1` and a call to its argument, assuming that no events are defined

Given that events can be dynamically defined at runtime, there is no room for reducing the number of predicate calls without turning off support for event-driven programming. When events are defined, the number of predicate calls grows proportional to the number of events and event handlers (monitors). Event-driven programming support can be switched off for specific object using the *events/1* compiler flag. Doing so, reduces the number of predicate calls from three to just one.

Messages to *self* are transparent regarding events and, as such, imply only one predicate call (to the cache entry, a dynamic predicate).

When a message is sent by the user from the top-level interpreter, Logtalk needs to perform a runtime translation of the message term in order to prove the corresponding goal. Thus, while sending a message from a compiled object corresponds to either three predicate calls (event-support on) or one predicate call (event-support off), the same message sent by the user from the top-level interpreter necessarily implies an overhead. Considering the time taken for the user to type the goal and read the reply, this overhead is of no practical consequence.

When a message is not cached, the number of predicate calls depends on the number of steps needed for the Logtalk runtime engine to lookup the corresponding predicate scope declaration (to check if the message is valid) and then to lookup a predicate definition for answering the message.

Processing time

Not all predicate calls take the same time. Moreover, the time taken to process a specific predicate call depends on the Prolog compiler implementation details. As such, the only valid performance measure is the time taken for processing

a message.

The usual way of measuring the time taken by a predicate call is to repeat the call a number of times and than to calculate the average time. A sufficient large number of repetitions would hopefully lead to an accurate measure. Care should be taken to subtract the time taken by the repetition code itself. In addition, we should be aware of any limitations of the predicates used to measure execution times. One way to make sense of numbers we get is to repeat the test with the same predicate using plain Prolog and with the predicate encapsulated in a module.

A simple predicate for helping benchmarking predicate calls could be:

```
benchmark(N, Goal) :-
    repeat(N),
        call(Goal),
    fail.

benchmark(_, _).
```

The rational of using a failure-driven loop is to try to avoid any interference on our timing measurements from garbage-collection or memory expansion mechanisms. Based on the predicate `benchmark/2`, we may define a more convenient predicate for performing our benchmarks. For example:

```
benchmark(Goal) :-
    % some sufficiently large number of repetitions
    N = 10000000,
    write('Number of repetitions: '), write(N), nl,
    % replace by your Prolog-specific predicate
    get_cpu_time(Seconds1),
    benchmark(N, Goal),
    get_cpu_time(Seconds2),
    Average is (Seconds2 - Seconds1)/N,
    write('Average time per call: '), write(Average), write(' seconds'), nl,
    Speed is 1.0/Average,
    write('Number of calls per second: '), write(Speed), nl.
```

We can get a baseline for our timings by doing:

```
| ?- benchmark(true).
```

For comparing message sending performance across several Prolog compilers, we would call the `benchmark/1` predicate with a suitable argument. For example:

```
| ?- benchmark(list::length([1,2,3,4,5,6,7,8,9,0], _)).
```

For comparing message sending performance with predicate calls in plain Prolog and with calls to predicates encapsulated in modules, we should use exactly the same predicate definition in the three cases.

It should be stressed that message sending is only one of the factors affecting the performance of a Logtalk application (and often not the most important one). The strengths and limitations of the chosen Prolog compiler play a crucial role on all aspects of the development, reliability, usability, and performance of a Logtalk application. It is advisable to take advantage of the Logtalk wide compatibility with most Prolog compilers to test for the best match for developing your Logtalk applications.

1.4 Objects

The main goal of Logtalk objects is the encapsulation and reuse of predicates. Instead of a single database containing all your code, Logtalk objects provide separated namespaces or databases allowing the partitioning of code in

more manageable parts. Logtalk does not aim to bring some sort of new dynamic state change concept to Logic Programming or Prolog.

In Logtalk, the only pre-defined objects are the built-in objects `user` and `logtalk`, which are described at the end of this section.

1.4.1 Objects, prototypes, classes, and instances

There are only three kinds of encapsulation entities in Logtalk: objects, protocols, and categories. Logtalk uses the term *object* in a broad sense. The terms *prototype*, *parent*, *class*, *subclass*, *superclass*, *metaclass*, and *instance* always designate an object. Different names are used to emphasize the *role* played by an object in a particular context. I.e. we use a term other than object when we want to make the relationship with other objects explicit. For example, an object with an *instantiation* relation with other object plays the role of an *instance*, while the instantiated object plays the role of a *class*; an object with a *specialization* relation with other object plays the role of a *subclass*, while the specialized object plays the role of a *superclass*; an object with an *extension* relation with other object plays the role of a *prototype*, the same for the extended object. A *stand-alone* object, i.e. an object with no relations with other objects, is always interpreted as a prototype. In Logtalk, entity relations essentially define *patterns* of code reuse. An entity is compiled accordingly to the roles it plays.

Logtalk allows you to work from standalone objects to any kind of hierarchy, either class-based or prototype-based. You may use single or multiple inheritance, use or forgo metaclasses, implement reflective designs, use parametric objects, and take advantage of protocols and categories (think components).

Prototypes

Prototypes are either self-defined objects or objects defined as extensions to other prototypes with whom they share common properties. Prototypes are ideal for representing one-of-a-kind objects. Prototypes usually represent concrete objects in the application domain. When linking prototypes using *extension* relations, Logtalk uses the term *prototype hierarchies* although most authors prefer to use the term *hierarchy* only with class generalization/specialization relations. In the context of logic programming, prototypes are often the ideal replacement for modules.

Classes

Classes are used to represent abstractions of common properties of sets of objects. Classes provide an ideal structuring solution when you want to express hierarchies of abstractions or work with many similar objects. Classes are used indirectly through *instantiation*. Contrary to most object-oriented programming languages, instances can be created both dynamically at runtime or defined in a source file like other objects.

1.4.2 Defining a new object

We can define a new object in the same way we write Prolog code: by using a text editor. Logtalk source files may contain one or more objects, categories, or protocols. If you prefer to define each entity in its own source file, it is recommended that the file be named after the object. By default, all Logtalk source files use the extension `.lgt` but this is optional and can be set in the adapter files. Intermediate Prolog source files (generated by the Logtalk compiler) have, by default, a `_lgt` suffix and a `.pl` extension. Again, this can be set to match the needs of a particular Prolog compiler in the corresponding adapter file. For instance, we may define an object named `vehicle` and save it in a `vehicle.lgt` source file which will be compiled to a `vehicle_lgt.pl` Prolog file (depending on the backend compiler, the names of the intermediate Prolog files may include a directory hash).

Object names can be atoms or compound terms (when defining parametric objects, see below). Objects, categories, and protocols share the same name space: we cannot have an object with the same name as a protocol or a category.

Object code (directives and predicates) is textually encapsulated by using two Logtalk directives: *object/1-5* and *end_object/0*. The most simple object will be one that is self-contained, not depending on any other Logtalk entity:

```
:- object(Object).
    ...
:- end_object.
```

If an object implements one or more protocols then the opening directive will be:

```
:- object(Object,
    implements([Protocol1, Protocol2, ...])).
    ...
:- end_object.
```

An object can import one or more categories:

```
:- object(Object,
    imports([Category1, Category2, ...])).
    ...
:- end_object.
```

If an object both implements protocols and imports categories then we will write:

```
:- object(Object,
    implements([Protocol1, Protocol2, ...]),
    imports([Category1, Category2, ...])).
    ...
:- end_object.
```

In object-oriented programming objects are usually organized in hierarchies that enable interface and code sharing by inheritance. In Logtalk, we can construct prototype-based hierarchies by writing:

```
:- object(Prototype,
    extends(Parent)).
    ...
:- end_object.
```

We can also have class-based hierarchies by defining instantiation and specialization relations between objects. To define an object as a class instance we will write:

```
:- object(Object,
    instantiates(Class)).
    ...
:- end_object.
```

A class may specialize another class, its superclass:

```
:- object(Class,
    specializes(Superclass)).
    ...
:- end_object.
```

If we are defining a reflexive system where every class is also an instance, we will probably be using the following pattern:

```
:- object(Class,
    instantiates(Metaclass),
```

```
        specializes(Superclass)).
        ...
:- end_object.
```

In short, an object can be a *stand-alone* object or be part of an object hierarchy. The hierarchy can be prototype-based (defined by extending other objects) or class-based (with instantiation and specialization relations). An object may also implement one or more protocols or import one or more categories.

A *stand-alone* object (i.e. an object with no extension, instantiation, or specialization relations with other objects) is always compiled as a prototype, that is, a self-describing object. If we want to use classes and instances, then we will need to specify at least one instantiation or specialization relation. The best way to do this is to define a set of objects that provide the basis of a reflective system [Cointe87], [Moura94]. For example:

```
% default root of the inheritance graph
% predicates common to all objects

:- object(object,
    instantiates(class)).
    ...
:- end_object.


% default metaclass for all classes
% predicates common to all instantiable classes

:- object(class,
    instantiates(class),
    specializes(abstract_class)).
    ...
:- end_object.


% default metaclass for all abstract classes
% predicates common to all classes

:- object(abstract_class,
    instantiates(class),
    specializes(object)).
    ...
:- end_object.
```

Note that with these instantiation and specialization relations, `object`, `class`, and `abstract_class` are, at the same time, classes and instances of some class. In addition, each object inherits its own predicates and the predicates of the other two objects without any inheritance loop problems.

When a full-blown reflective system solution is not needed, the above scheme can be simplified by making an object an instance of itself, i.e. by making a class its own metaclass. For example:

```
:- object(class,
    instantiates(class)).
    ...
:- end_object.
```

We can use, in the same application, both prototype and class-based hierarchies (and freely exchange messages between all objects). We cannot however mix the two types of hierarchies by, e.g., specializing an object that extends another object in this current Logtalk version.

1.4.3 Parametric objects

Parametric objects have a compound term for name instead of an atom. This compound term usually contains free variables that can be instantiated when sending or as a consequence of sending a message to the object, thus acting as object parameters. The object predicates can then be coded to depend on those parameters, which are logical variables shared by all object predicates. When an object state is set at object creation and never changed, parameters provide a better solution than using the object's database via asserts. Parametric objects can also be used to associate a set of predicates to terms that share a common functor and arity.

In order to give access to an object parameter, Logtalk provides a *parameter/2* built-in local method:

```
:- object(Functor(Arg1, Arg2, ...)).

    ...

    Predicate :-
        ...,
        parameter(Number, Value),
        ... .
```

An alternative solution is to use the built-in local method *this/1*. For example:

```
:- object(foo(Arg)).

    ...

    bar :-
        ...,
        this(foo(Arg)),
        ... .
```

Both solutions are equally efficient as calls to the methods `this/1` and `parameter/2` are usually compiled inline into a clause head unification. The drawback of this second solution is that we must check all calls of `this/1` if we change the object name. Note that we can't use these method with the message sending operators (*::/2*, *::/1*, or *^^/1*).

A third alternative to access object parameters is to use *parameter variables*. Although parameter variables introduce a concept of entity global variables, they allow object parameters to be added, rearranged, or removed without requiring any changes to the clauses that refer to them. Note that using parameter variables doesn't change the fact that entity parameters are logical variables.

When storing a parametric object in its own source file, the convention is to name the file after the object, with the object arity appended. For instance, when defining an object named `sort(Type)`, we may save it in a `sort_1.lgt` text file. This way it is easy to avoid file name clashes when saving Logtalk entities that have the same functor but different arity.

Compound terms with the same functor and with the same number of arguments as a parametric object identifier may act as *proxies* to a parametric object. Proxies may be stored on the database as Prolog facts and be used to represent different instantiations of a parametric object identifier. Logtalk provides a convenient notation for accessing proxies represented as Prolog facts when sending a message:

```
{Proxy}::Message
```

In this context, the proxy argument is proved as a plain Prolog goal. If successful, the message is sent to the corresponding parametric object. Typically, the proof allows retrieving of parameter instantiations. This construct can either be used with a proxy argument that is sufficiently instantiated in order to unify with a single Prolog fact or with a proxy argument that unifies with several facts on backtracking.

1.4.4 Finding defined objects

We can find, by backtracking, all defined objects by calling the *current_object/1* built-in predicate with a non-instantiated variable:

```
| ?- current_object(Object).
```

This predicate can also be used to test if an object is defined by calling it with a valid object identifier (an atom or a compound term).

1.4.5 Creating a new object in runtime

An object can be dynamically created at runtime by using the *create_object/4* built-in predicate:

```
| ?- create_object(Object, Relations, Directives, Clauses).
```

The first argument should be either a variable or the name of the new object (a Prolog atom or compound term, which must not match any existing entity name). The remaining three arguments correspond to the relations described in the opening object directive and to the object code contents (directives and clauses).

For instance, the call:

```
| ?- create_object(foo, [extends(bar)], [public(foo/1)], [foo(1), foo(2)]).
```

is equivalent to compiling and loading the object:

```
:- object(foo,
    extends(bar)).

    :- dynamic.

    :- public(foo/1).

    foo(1).
    foo(2).

:- end_object.
```

If we need to create a lot of (dynamic) objects at runtime, then is best to define a metaclass or a prototype with a predicate that will call this built-in predicate to make new objects. This predicate may provide automatic object name generation, name checking, and accept object initialization options.

1.4.6 Abolishing an existing object

Dynamic objects can be abolished using the *abolish_object/1* built-in predicate:

```
| ?- abolish_object(Object).
```

The argument must be an identifier of a defined dynamic object, otherwise an error will be thrown.

1.4.7 Object directives

Object directives are used to set initialization goals, define object properties, to document an object dependencies on other Logtalk entities, and to load the contents of files into an object.

Object initialization

We can define a goal to be executed as soon as an object is (compiled and) loaded to memory with the *initialization/1* directive:

```
:- initialization(Goal).
```

The argument can be any valid Prolog or Logtalk goal, including a message to other object. For example:

```
:- object(foo).

    :- initialization(init).
    :- private(init/0).

    init :-
        ... .

    ...

:- end_object.
```

Or:

```
:- object(assembler).

    :- initialization(control::start).
    ...

:- end_object.
```

The initialization goal can also be a message to *self* in order to call an inherited or imported predicate. For example, assuming that we have a `monitor` category defining a `reset/0` predicate:

```
:- object(profiler,
    imports(monitor)).

    :- initialization(::reset).
    ...

:- end_object.
```

Note, however, that descendant objects do not inherit initialization directives. In this context, *self* denotes the object that contains the directive. Also note that by initialization we do not necessarily mean setting an object dynamic state.

Dynamic objects

Similar to Prolog predicates, an object can be either static or dynamic. An object created during the execution of a program is always dynamic. An object defined in a file can be either dynamic or static. Dynamic objects are declared by using the *dynamic/0* directive in the object source code:

```
:- dynamic.
```

The directive must precede any predicate directives or clauses. Please be aware that using dynamic code results in a performance hit when compared to static code. We should only use dynamic objects when these need to be abolished during program execution. In addition, note that we can declare and define dynamic predicates within a static object.

Object documentation

An object can be documented with arbitrary user-defined information by using the *info/1* directive:

```
:- info(List).
```

See the *Documenting applications* section for details.

Loading files into an object

The *include/1* directive can be used to load the contents of a file into an object. A typical usage scenario is to load a plain Prolog database into an object thus providing a simple way to encapsulate it. For example, assume a cities. pl file defining facts for a city/4 predicate. We could define a wrapper for this database by writing:

```
:- object(cities).

    :- public(city/4).

    :- include(dbs('cities.pl')).

:- end_object.
```

The include/1 directive can also be used when creating an object dynamically. For example:

```
| ?- create_object(cities, [], [public(city/4), include(dbs('cities.pl'))], []).
```

1.4.8 Object relationships

Logtalk provides six sets of built-in predicates that enable us to query the system about the possible relationships that an object may have with other entities.

The built-in predicates *instantiates_class/2-3* and *instantiates_class/2-3* can be used to query all instantiation relations:

```
| ?- instantiates_class(Instance, Class).
```

or, if we want to know the instantiation scope:

```
| ?- instantiates_class(Instance, Class, Scope).
```

Specialization relations can be found by using either the *specializes_class/2-3* or the *specializes_class/2-3* built-in predicates:

```
| ?- specializes_class(Class, Superclass).
```

or, if we want to know the specialization scope:

```
| ?- specializes_class(Class, Superclass, Scope).
```

For prototypes, we can query extension relations with the *extends_object/2-3* or the *extends_object/2-3* built-in predicates:

```
| ?- extends_object(Object, Parent).
```

or, if we want to know the extension scope:

```
| ?- extends_object(Object, Parent, Scope).
```

In order to find which objects import which categories we can use the built-in predicates *imports_category/2-3* or *imports_category/2-3*:

```
| ?- imports_category(Object, Category).
```

or, if we want to know the importation scope:

```
| ?- imports_category(Object, Category, Scope).
```

To find which objects implements which protocols we can use the *implements_protocol/2-3* and *conforms_to_protocol/2-3* built-in predicates:

```
| ?- implements_protocol(Object, Protocol, Scope).
```

or, if we also want inherited protocols:

```
| ?- conforms_to_protocol(Object, Protocol, Scope).
```

Note that, if we use a non-instantiated variable for the first argument, we will need to use the *current_object/1* built-in predicate to ensure that the entity returned is an object and not a category.

To find which objects are explicitly complemented by categories we can use the *complements_object/2* built-in predicate:

```
| ?- complements_object(Category, Object).
```

Note that more than one category may explicitly complement a single object.

1.4.9 Object properties

We can find the properties of defined objects by calling the built-in predicate *object_property/2*:

```
| ?- object_property(Object, Property).
```

The following object properties are supported:

static The object is static

dynamic The object is dynamic (and thus can be abolished in runtime by calling the *abolish_object/1* built-in predicate)

built_in The object is a built-in object (and thus always available)

threaded The object supports/makes multi-threading calls

file(Path) Absolute path of the source file defining the object (if applicable)

file(Basename, Directory) Basename and directory of the source file defining the object (if applicable)

lines(BeginLine, EndLine) Source file begin and end lines of the object definition (if applicable)

context_switching_calls The object supports context switching calls (i.e. can be used with the <</2 debugging control construct)

dynamic_declarations The object supports dynamic declarations of predicates

events Messages sent from the object generate events

source_data Source data available for the object

complements(Permission) The object supports complementing categories with the specified permission (`allow` or `restrict`)

complements The object supports complementing categories

public(Predicates) List of public predicates declared by the object

protected(Predicates) List of protected predicates declared by the object

private(Predicates) List of private predicates declared by the object

declares(Predicate, Properties) List of properties for a predicate declared by the object

defines(Predicate, Properties) List of properties for a predicate defined by the object

includes(Predicate, Entity, Properties) List of properties for an object multifile predicate that are defined in the specified entity (the properties include `number_of_clauses(Number)`, `number_of_rules(Number)`, and `line_count(Line)` with `Line` being the begin line of the multifile predicate clause)

provides(Predicate, Entity, Properties) List of properties for other entity multifile predicate that are defined in the object (the properties include `number_of_clauses(Number)`, `number_of_rules(Number)`, and `line_count(Line)` with `Line` being the begin line of the multifile predicate clause)

alias(Predicate, Properties) List of properties for a predicate alias declared by the object (the properties include `for(Original)`, `from(Entity)`, `non_terminal(NonTerminal)`, and `line_count(Line)` with `Line` being the begin line of the alias directive)

calls(Call, Properties) List of properties for predicate calls made by the object (`Call` is either a predicate indicator or a control construct such as `::/1-2` or `^^/1` with a predicate indicator as argument; note that `Call` may not be ground in case of a call to a control construct where its argument is only know at runtime; the properties include `caller(Caller)`, `alias(Alias)`, and `line_count(Line)` with both `Caller` and `Alias` being predicate indicators and `Line` being the begin line of the predicate clause or directive making the call)

updates(Predicate, Properties) List of properties for dynamic predicate updates (and also access using the `clause/2` predicate) made by the object (`Predicate` is either a predicate indicator or a control construct such as `::/1-2` or `:/2` with a predicate indicator as argument; note that `Predicate` may not be ground in case of a control construct argument only know at runtime; the properties include `updater(Updater)`, `alias(Alias)`, and `line_count(Line)` with `Updater` being a (possibly multifile) predicate indicator, `Alias` being a predicate indicator, and `Line` being the begin line of the predicate clause or directive updating the predicate)

number_of_clauses(Number) Total number of predicate clauses defined in the object at compilation time (includes both user-defined clauses and auxiliary clauses generated by the compiler or by the expansion hooks)

number_of_rules(Number) Total number of predicate rules defined in the object at compilation time (includes both user-defined rules and auxiliary rules generated by the compiler or by the expansion hooks)

number_of_user_clauses(Number) Total number of user-defined predicate clauses defined in the object at compilation time

number_of_user_rules(Number) Total number of user-defined predicate rules defined in the object at compilation time

debugging The object is compiled in debug mode

module The object resulted from the compilation of a Prolog module

When a predicate is called from an `initialization/1` directive, the argument of the `caller/1` property is `:-/1`.

Some of the properties such as line numbers are only available when the object is defined in a source file compiled with the `source_data` flag turned on.

The properties that return the number of clauses (rules) report the clauses (rules) *textually defined in the object* for both multifile and non-multifile predicates. Thus, these numbers exclude clauses (rules) for multifile predicates *contributed* by other entities.

1.4.10 Built-in objects

Logtalk defines some built-in objects that are always available for any application.

The built-in pseudo-object *user*

Logtalk defines a built-in, pseudo-object named `user` that virtually contains all user predicate definitions not encapsulated in a Logtalk entity. These predicates are assumed to be implicitly declared public. Messages sent from this pseudo-object, which includes messages sent from the top-level interpreter, generate events when the default value of the `events` flag is set to `allow`. Defining complementing categories for this pseudo-object is not supported.

With some of the backend Prolog compilers that support a module system, it is possible to load (the) Logtalk (compiler/runtime) into a module other than the pseudo-module *user*. In this case, the Logtalk pseudo-object *user* virtually contains all user predicate definitions defined in the module where Logtalk was loaded.

The built-in object *logtalk*

Logtalk defines a built-in object named `logtalk` that provides structured message printing mechanism predicates, structured question asking predicates, debugging event predicates, predicates for accessing the internal database of loaded files and their properties, and also a set of low-level utility predicates normally used when defining hook objects.

The following predicates are defined:

`expand_library_path(Library, Path)` Expands a file specification in library notation to a full operating-system path.

`loaded_file(Path)` Returns the full path of a currently loaded source file.

`loaded_file_property(Path, Property)` Returns a property for a currently loaded source file. Valid properties are `basename/1`, `directory/1`, `flags/1` (explicit flags used when the file was loaded), `text_properties/1` (list, possibly empty, whose possible elements are `encoding/1` and `bom/1`), `target/1` (full path for the Prolog file generated by the compilation of the loaded source file), `modified/1` (time stamp that should be treated as an opaque term but that may be used for comparisons), `parent/1` (parent file, if it exists, that loaded the file; a file may have multiple parents), and `library/1` (library name when there is a library whose location is the same as the loaded file directory).

`compile_aux_clauses(Clauses)` Compiles a list of clauses in the context of the entity under compilation. This method is usually called from `goal_expansion/2` hooks in order to compile auxiliary clauses generated for supporting an expanded goal. The compilation of the clauses avoids the risk of making the predicate whose clause is being goal-expanded discontiguous by accident.

`entity_prefix(Entity, Prefix)` Converts an entity identifier into its internal prefix or an internal prefix into an entity identifier.

compile_predicate_heads(Heads, Entity, TranslatedHeads, ContextArgument)
Compiles a predicate head or a list of predicate heads in the context of the specified entity or in the context of the entity being compiled when `Entity` is not instantiated.

compile_predicate_indicators(PredicateIndicators, Entity, TranslatedPredicateIndicators)
Compiles a predicate indicator or a list of predicate indicators in the context of the specified entity or in the context of the entity being compiled when `Entity` is not instantiated.

decompile_predicate_heads(TranslatedHeads, Entity, EntityType, Heads) Decompiles a compiled predicate head or a list of compiled predicate heads returning the entity, entity type, and source level heads. Requires the entity to be currently loaded.

decompile_predicate_indicators(TranslatedPredicateIndicators, Entity, EntityType, Predicate
Decompiles a compiled predicate indicator or a list of compiled predicate indicators returning the entity, entity type, and source level predicate indicators. Requires the entity to be currently loaded.

execution_context(ExecutionContext, Entity, Sender, This, Self, MetaCallContext, Stack)
Allows constructing and accessing execution context components.

print_message(Kind, Component, Term) Prints a message term after converting it into a list of tokens using the `message_tokens//2` hook non-terminal. When the conversion fails, the message term itself is printed.

print_message_tokens(Stream, Prefix, Tokens) Prints a list of message tokens to the specified stream and prefixing each line with the specified prefix.

print_message_token(Stream, Prefix, Token, Tokens) Hook predicate, declared multifile and dynamic, allowing the default printing of a token to be overridden.

message_tokens(Term, Component) Hook non-terminal, declared multifile and dynamic, allowing the translation of a message into a list of tokens for printing.

message_prefix_stream(Kind, Component, Prefix, Stream) Hook predicate, declared multifile and dynamic, allowing the definition of line prefix and output stream for messages.

message_hook(Term, Kind, Component, Tokens) Hook predicate, declared multifile and dynamic, allowing the overriding the default printing of a message.

trace_event(Event, EventExecutionContext) Hook predicate, declared multifile and dynamic, for handling trace events generated by the execution of source code compiled in debug mode. The Logtalk runtime calls all defined handlers using a failure-driven loop. Thus, care must be taken that the handlers are deterministic to avoid potential termination issues.

debug_handler_provider(Provider) Multifile predicate for declaring an object that provides a debug handler. There can only be one debug handler provider loaded at the same time. The Logtalk runtime uses this hook predicate for detecting multiple instances of the handler and for better error reporting.

debug_handler(Event, EventExecutionContext) Multifile predicate for handling debug events generated by the execution of source code compiled in debug mode.

To use these predicates, simply send the corresponding message to the `logtalk` object.

1.5 Protocols

Protocols enable the separation between interface and implementation: several objects can implement the same protocol and an object can implement several protocols. Protocols may contain only predicate declarations. In some languages the term *interface* is used with similar meaning. Logtalk allows predicate declarations of any scope within protocols, contrary to some languages that only allow public declarations.

Logtalk defines three built-in protocols, `monitoring`, `expanding`, and `forwarding`, which are described at the end of this section.

1.5.1 Defining a new protocol

We can define a new protocol in the same way we write Prolog code: by using a text editor. Logtalk source files may contain one or more objects, categories, or protocols. If you prefer to define each entity in its own source file, it is recommended that the file be named after the protocol. By default, all Logtalk source files use the extension `.lgt` but this is optional and can be set in the adapter files. Intermediate Prolog source files (generated by the Logtalk compiler) have, by default, a `_lgt` suffix and a `.pl` extension. Again, this can be set to match the needs of a particular Prolog compiler in the corresponding adapter file. For example, we may define a protocol named `listp` and save it in a `listp.lgt` source file that will be compiled to a `listp_lgt.pl` Prolog file (depending on the backend compiler, the names of the intermediate Prolog files may include a directory hash).

Protocol names must be atoms. Objects, categories and protocols share the same name space: we cannot have a protocol with the same name as an object or a category.

Protocol directives are textually encapsulated by using two Logtalk directives: *protocol/1-2* and *end_protocol/0*. The most simple protocol will be one that is self-contained, not depending on any other Logtalk entity:

```
:- protocol(Protocol).
    ...
:- end_protocol.
```

If a protocol extends one or more protocols, then the opening directive will be:

```
:- protocol(Protocol,
    extends([Protocol1, Protocol2, ...])).
    ...
:- end_protocol.
```

In order to maximize protocol reuse, all predicates specified in a protocol should relate to the same functionality. Therefore, the only recommended use of protocol extension is when you need both a minimal protocol and an extended version of the same protocol with additional, useful predicates.

1.5.2 Finding defined protocols

We can find, by backtracking, all defined protocols by using the *current_protocol/1* built-in predicate with a non-instantiated variable:

```
| ?- current_protocol(Protocol).
```

This predicate can also be used to test if a protocol is defined by calling it with a valid protocol identifier (an atom).

1.5.3 Creating a new protocol in runtime

We can create a new (dynamic) protocol in runtime by calling the Logtalk built-in predicate *create_protocol/3*:

```
| ?- create_protocol(Protocol, Relations, Directives).
```

The first argument should be either a variable or the name of the new protocol (a Prolog atom, which must not match an existing entity name). The remaining two arguments correspond to the relations described in the opening protocol directive and to the protocol directives.

For instance, the call:

```
| ?- create_protocol(ppp, [extends(qqq)], [public([foo/1, bar/1])]).
```

is equivalent to compiling and loading the protocol:

```
:- protocol(ppp,
    extends(qqq)).

    :- dynamic.

    :- public([foo/1, bar/1]).

:- end_protocol.
```

If we need to create a lot of (dynamic) protocols at runtime, then is best to define a metaclass or a prototype with a predicate that will call this built-in predicate in order to provide more sophisticated behavior.

1.5.4 Abolishing an existing protocol

Dynamic protocols can be abolished using the *abolish_protocol/1* built-in predicate:

```
| ?- abolish_protocol(Protocol).
```

The argument must be an identifier of a defined dynamic protocol, otherwise an error will be thrown.

1.5.5 Protocol directives

Protocol directives are used to define protocol properties and documentation.

Dynamic protocols

As usually happens with Prolog code, a protocol can be either static or dynamic. A protocol created during the execution of a program is always dynamic. A protocol defined in a file can be either dynamic or static. Dynamic protocols are declared by using the *dynamic/0* directive in the protocol source code:

```
:- dynamic.
```

The directive must precede any predicate directives. Please be aware that using dynamic code results in a performance hit when compared to static code. We should only use dynamic protocols when these need to be abolished during program execution.

Protocol documentation

A protocol can be documented with arbitrary user-defined information by using the *info/1* directive:

```
:- info(List).
```

See the *Documenting applications* section for details.

Loading files into a protocol

The *include/1* directive can be used to load the contents of a file into a protocol. See the *Objects* section for an example of using this directive.

1.5.6 Protocol relationships

Logtalk provides two sets of built-in predicates that enable us to query the system about the possible relationships that a protocol have with other entities.

The built-in predicates *extends_protocol/2-3* and *extends_protocol/2-3* return all pairs of protocols so that the first one extends the second:

```
| ?- extends_protocol(Protocol1, Protocol2).
```

or, if we want to know the extension scope:

```
| ?- extends_protocol(Protocol1, Protocol2, Scope).
```

To find which objects or categories implement which protocols we can call the *implements_protocol/2-3* or *implements_protocol/2-3* built-in predicates:

```
| ?- implements_protocol(ObjectOrCategory, Protocol).
```

or, if we want to know the implementation scope:

```
| ?- implements_protocol(ObjectOrCategory, Protocol, Scope).
```

Note that, if we use a non-instantiated variable for the first argument, we will need to use the *current_object/1* or *current_category/1* built-in predicates to identify the kind of entity returned.

1.5.7 Protocol properties

We can find the properties of defined protocols by calling the *protocol_property/2* built-in predicate:

```
| ?- protocol_property(Protocol, Property).
```

A protocol may have the property `static`, `dynamic`, or `built_in`. Dynamic protocols can be abolished in runtime by calling the *abolish_protocol/1* built-in predicate. Depending on the back-end Prolog compiler, a protocol may have additional properties related to the source file where it is defined.

The following protocol properties are supported:

static The protocol is static

dynamic The protocol is dynamic (and thus can be abolished in runtime by calling the *abolish_category/1* built-in predicate)

built_in The protocol is a built-in protocol (and thus always available)

source_data Source data available for the protocol

file(Path) Absolute path of the source file defining the protocol (if applicable)

file(Basename, Directory) Basename and directory of the source file defining the protocol (if applicable)

lines(BeginLine, EndLine) Source file begin and end lines of the protocol definition (if applicable)

public(Predicates) List of public predicates declared by the protocol

protected(Predicates) List of protected predicates declared by the protocol

private(Predicates) List of private predicates declared by the protocol

declares(Predicate, Properties) List of properties for a predicate declared by the protocol

`alias(Predicate, Properties)` List of properties for a predicate alias declared by the protocol (the properties include `for(Original)`, `from(Entity)`, `non_terminal(NonTerminal)`, and `line_count(Line)` with `Line` being the begin line of the alias directive)

Some of the properties such as line numbers are only available when the protocol is defined in a source file compiled with the `source_data` flag turned on.

1.5.8 Implementing protocols

Any number of objects or categories can implement a protocol. The syntax is very simple:

```
:- object(Object,
    implements(Protocol)).
    ...
:- end_object.
```

or, in the case of a category:

```
:- category(Object,
    implements(Protocol)).
    ...
:- end_category.
```

To make all public predicates declared via an implemented protocol protected or to make all public and protected predicates private we prefix the protocol's name with the corresponding keyword. For instance:

```
:- object(Object,
    implements(private::Protocol)).
    ...
:- end_object.
```

or:

```
:- object(Object,
    implements(protected::Protocol)).
    ...
:- end_object.
```

Omitting the scope keyword is equivalent to writing:

```
:- object(Object,
    implements(public::Protocol)).
    ...
:- end_object.
```

The same rules applies to protocols implemented by categories.

1.5.9 Built-in protocols

Logtalk defines a set of built-in protocols that are always available for any application.

The built-in protocol *expanding*

Logtalk defines a built-in protocol named `expanding` that declares the `term_expansion/2` and `goal_expansion/2` predicates. See the description of the *hook/1* compiler flag for more details.

The built-in protocol *monitoring*

Logtalk defines a built-in protocol named `monitoring` declares the *before/3* and *after/3* public event handler predicates. See the *Event-driven programming* section for more details.

The built-in protocol *forwarding*

Logtalk defines a built-in protocol named `forwarding` that declares the *forward/1* user-defined message forwarding handler, which is automatically called (if defined) by the runtime for any message that the receiving object does not understand. See also the *[]/1* control construct.

1.6 Categories

Categories are fine-grained units of code reuse and can be regarded as a dual concept of protocols. Categories provide a way to encapsulate a set of related predicate declarations and definitions that do not represent a complete object and that only make sense when composed with other predicates. Categories may also be used to break a complex object in functional units. A category can be imported by several objects (without code duplication), including objects participating in prototype or class-based hierarchies. This concept of categories shares some ideas with Smalltalk-80 functional categories [Goldberg83], Flavors mix-ins [Moon86] (without necessarily implying multi-inheritance), and Objective-C categories [Cox86]. Categories may also *complement* existing objects, thus providing a hot patching mechanism inspired by the Objective-C categories functionality.

1.6.1 Defining a new category

We can define a new category in the same way we write Prolog code: by using a text editor. Logtalk source files may contain one or more objects, categories, or protocols. If you prefer to define each entity in its own source file, it is recommended that the file be named after the category. By default, all Logtalk source files use the extension `.lgt` but this is optional and can be set in the adapter files. Intermediate Prolog source files (generated by the Logtalk compiler) have, by default, a `_lgt` suffix and a `.pl` extension. Again, this can be set to match the needs of a particular Prolog compiler in the corresponding adapter file. For instance, we may define a category named `documenting` and save it in a `documenting.lgt` source file that will be compiled to a `documenting_lgt.pl` Prolog file (depending on the backend compiler, the names of the intermediate Prolog files may include a directory hash).

Category names can be atoms or compound terms (when defining parametric categories). Objects, categories, and protocols share the same name space: we cannot have a category with the same name as an object or a protocol.

Category code (directives and predicates) is textually encapsulated by using two Logtalk directives: *category/1-3* and *end_category/0*. The most simple category will be one that is self-contained, not depending on any other Logtalk entity:

```
:- category(Category).
    ...
:- end_category.
```

If a category implements one or more protocols then the opening directive will be:

```
:- category(Category,
    implements([Protocol1, Protocol2, ...])).
    ...
:- end_category.
```

A category may be defined as a composition of other categories by writing:

```logtalk
:- category(Category,
    extends([Category1, Category2, ...])).
    ...
:- end_category.
```

This feature should only be used when extending a category without breaking its functional cohesion (for example, when a modified version of a category is needed for importing on several unrelated objects). The preferred way of composing several categories is by importing them into an object. When a category overrides a predicate defined in an extended category, the overridden definition can still be used by using the *alias/2* predicate directive.

Categories cannot inherit from objects. In addition, categories cannot define clauses for dynamic predicates. This restriction applies because a category can be imported by several objects and because we cannot use the database handling built-in methods with categories (messages can only be sent to objects). However, categories may contain declarations for dynamic predicates and they can contain predicates which handle dynamic predicates. For example:

```logtalk
:- category(attributes).

    :- public(attribute/2).
    :- public(set_attribute/2).
    :- public(del_attribute/2).

    :- private(attribute_/2).
    :- dynamic(attribute_/2).

    attribute(Attribute, Value) :-
        % called in the context of "self"
        ::attribute_(Attribute, Value).

    set_attribute(Attribute, Value) :-
        % retract old clauses in "self"
        ::retractall(attribute_(Attribute, _)),
        % assert new clause in "self"
        ::assertz(attribute_(Attribute, Value)).

    del_attribute(Attribute, Value) :-
        % retract clause in "self"
        ::retract(attribute_(Attribute, Value)).

:- end_category.
```

Each object importing this category will have its own `attribute_/2` private, dynamic predicate. The predicates `attribute/2`, `set_attribute/2`, and `del_attribute/2` always access and modify the dynamic predicate contained in the object receiving the corresponding messages (i.e. *self*). But it's also possible to define predicates that handle dynamic predicates in the context of *this* instead of *self*. For example:

```logtalk
:- category(attributes).

    :- public(attribute/2).
    :- public(set_attribute/2).
    :- public(del_attribute/2).

    :- private(attribute_/2).
    :- dynamic(attribute_/2).

    attribute(Attribute, Value) :-
        % call in the context of "this"
        attribute_(Attribute, Value).
```

(continues on next page)

```
    set_attribute(Attribute, Value) :-
        % retract old clauses in "this"
        retractall(attribute_(Attribute, _)),
        % asserts clause in "this"
        assertz(attribute_(Attribute, Value)).

    del_attribute(Attribute, Value) :-
        % retract clause in "this"
        retract(attribute_(Attribute, Value)).

:- end_category.
```

When defining a category that declares and handles dynamic predicates, working in the context of *this* ties those dynamic predicates to the object importing the category while working in the context of *self* allows each object inheriting from the object that imports the category to have its own set of clauses for those dynamic predicates.

Hot patching

A category may explicitly complement one or more existing objects, thus providing hot patching functionality inspired by Objective-C categories:

```
:- category(Category,
    complements([Object1, Object2, ....])).
    ...
:- end_category.
```

This allows us to add missing directives (e.g. to define aliases for complemented object predicates), replace broken predicate definitions, add new predicates, and add protocols and categories to existing objects without requiring access or modifications to their source code. Common scenarios are adding logging or debugging predicates to a set of objects. Complemented objects need to be compiled with the `complements` compiler flag set `allow` (to allow both patching and adding functionality) or `restrict` (to allow only adding new functionality). A complementing category takes preference over a previously loaded complementing category for the same object thus allowing patching a previous patch if necessary.

Note that *super calls* from predicates defined in complementing categories lookup inherited definitions as if the calls were made from the complemented object instead of the category ancestors. This allows more comprehensive object patching. But it also means that, if you want to patch an object so that it imports a category that extends another category and uses super calls to access the extended category predicates, you will need to define a (possibly empty) complementing category that extends the category that you want to add.

An unfortunate consequence of allowing an object to be patched at runtime using a complementing category is that it disables the use of *static binding* optimizations for messages sent to the complemented object as it can always be later patched, thus rendering the static binding optimizations invalid.

Another important caveat is that, while a complementing category can replace a predicate definition, local callers of the replaced predicate will still call the unpatched version of the predicate. This is a consequence of the lack of a portable solution at the backend Prolog compiler level for destructively replacing static predicates.

1.6.2 Finding defined categories

We can find, by backtracking, all defined categories by using the *current_category/1* Logtalk built-in predicate with a non-instantiated variable:

```
| ?- current_category(Category).
```

This predicate can also be used to test if a category is defined by calling it with a valid category identifier (an atom or a compound term).

1.6.3 Creating a new category in runtime

A category can be dynamically created at runtime by using the *create_category/4* built-in predicate:

```
| ?- create_category(Category, Relations, Directives, Clauses).
```

The first argument should be either a variable or the name of the new category (a Prolog atom, which must not match with an existing entity name). The remaining three arguments correspond to the relations described in the opening category directive and to the category code contents (directives and clauses).

For instance, the call:

```
| ?- create_category(
        ccc,
        [implements(ppp)],
        [private(bar/1)],
        [(foo(X):-bar(X)), bar(1), bar(2)]
     ).
```

is equivalent to compiling and loading the category:

```
:- category(ccc,
    implements(ppp)).

    :- dynamic.

    :- private(bar/1).

    foo(X) :-
        bar(X).

    bar(1).
    bar(2).

:- end_category.
```

If we need to create a lot of (dynamic) categories at runtime, then is best to define a metaclass or a prototype with a predicate that will call this built-in predicate in order to provide more sophisticated behavior.

1.6.4 Abolishing an existing category

Dynamic categories can be abolished using the *abolish_category/1* built-in predicate:

```
| ?- abolish_category(Category).
```

The argument must be an identifier of a defined dynamic category, otherwise an error will be thrown.

1.6.5 Category directives

Category directives are used to define category properties, to document a category dependencies on other Logtalk entities, and to load the contents of files into a category.

Dynamic categories

As usually happens with Prolog code, a category can be either static or dynamic. A category created during the execution of a program is always dynamic. A category defined in a file can be either dynamic or static. Dynamic categories are declared by using the *dynamic/0* directive in the category source code:

```
:- dynamic.
```

The directive must precede any predicate directives or clauses. Please be aware that using dynamic code results in a performance hit when compared to static code. We should only use dynamic categories when these need to be abolished during program execution.

Category documentation

A category can be documented with arbitrary user-defined information by using the *info/1* directive:

```
:- info(List).
```

See the *Documenting applications* section for details.

Loading files into a category

The *include/1* directive can be used to load the contents of a file into a category. See the *Objects* section for an example of using this directive.

1.6.6 Category relationships

Logtalk provides two sets of built-in predicates that enable us to query the system about the possible relationships that a category can have with other entities.

The built-in predicates *implements_protocol/2-3* and *conforms_to_protocol/2-3* allows us to find which categories implements which protocols:

```
| ?- implements_protocol(Category, Protocol, Scope).
```

or, if we also want inherited protocols:

```
| ?- conforms_to_protocol(Category, Protocol, Scope).
```

Note that, if we use a non-instantiated variable for the first argument, we will need to use the *current_category/1* built-in predicate to ensure that the returned entity is a category and not an object.

To find which objects import which categories we can use the *imports_category/2-3* or *imports_category/2-3* built-in predicates:

```
| ?- imports_category(Object, Category).
```

or, if we want to know the importation scope:

```
| ?- imports_category(Object, Category, Scope).
```

Note that a category may be imported by several objects.

To find which categories extend other categories we can use the *extends_category/2-3* or *extends_category/2-3* built-in predicates:

```
| ?- extends_category(Category1, Category2).
```

or, if we want to know the extension scope:

```
| ?- extends_category(Category1, Category2, Scope).
```

Note that a category may be extended by several categories.

To find which categories explicitly complement existing objects we can use the *complements_object/2* built-in predicate:

```
| ?- complements_object(Category, Object).
```

Note that a category may explicitly complement several objects.

1.6.7 Category properties

We can find the properties of defined categories by calling the built-in predicate *category_property/2*:

```
| ?- category_property(Category, Property).
```

The following category properties are supported:

static The category is static

dynamic The category is dynamic (and thus can be abolished in runtime by calling the *abolish_category/1* built-in predicate)

built_in The category is a built-in category (and thus always available)

file(Path) Absolute path of the source file defining the category (if applicable)

file(Basename, Directory) Basename and directory of the source file defining the category (if applicable)

lines(BeginLine, EndLine) Source file begin and end lines of the category definition (if applicable)

events Messages sent from the category generate events

source_data Source data available for the category

public(Predicates) List of public predicates declared by the category

protected(Predicates) List of protected predicates declared by the category

private(Predicates) List of private predicates declared by the category

declares(Predicate, Properties) List of properties for a predicate declared by the category

defines(Predicate, Properties) List of properties for a predicate defined by the category

includes(Predicate, Entity, Properties) List of properties for an object multifile predicate that are defined in the specified entity (the properties include `number_of_clauses(Number)`, `number_of_rules(Number)`, and `line_count(Line)` with `Line` being the begin line of the multifile predicate clause)

provides(Predicate, Entity, Properties) List of properties for other entity multifile predicate that are defined in the category (the properties include `number_of_clauses(Number)`, `number_of_rules(Number)`, and `line_count(Line)` with `Line` being the begin line of the multifile predicate clause)

alias(Predicate, Properties) List of properties for a predicate alias declared by the category (the properties include `for(Original)`, `from(Entity)`, `non_terminal(NonTerminal)`, and `line_count(Line)` with `Line` being the begin line of the alias directive)

calls(Call, Properties) List of properties for predicate calls made by the category (`Call` is either a predicate indicator or a control construct such as `::/1-2` or `^^/1` with a predicate indicator as argument; note that `Call` may not be ground in case of a call to a control construct where its argument is only know at runtime; the properties include `caller(Caller)`, `alias(Alias)`, and `line_count(Line)` with both `Caller` and `Alias` being predicate indicators and `Line` being the begin line of the predicate clause or directive making the call)

updates(Predicate, Properties) List of properties for dynamic predicate updates (and also access using the `clause/2` predicate) made by the object (`Predicate` is either a predicate indicator or a control construct such as `::/1-2` or `:/2` with a predicate indicator as argument; note that `Predicate` may not be ground in case of a control construct argument only know at runtime; the properties include `updater(Updater)`, `alias(Alias)`, and `line_count(Line)` with `Updater` being a (possibly multifile) predicate indicator, `Alias` being a predicate indicator, and `Line` being the begin line of the predicate clause or directive updating the predicate)

number_of_clauses(Number) Total number of predicate clauses defined in the category (includes both user-defined clauses and auxiliary clauses generated by the compiler or by the expansion hooks)

number_of_rules(Number) Total number of predicate rules defined in the category (includes both user-defined rules and auxiliary rules generated by the compiler or by the expansion hooks)

number_of_user_clauses(Number) Total number of user-defined predicate clauses defined in the category

number_of_user_rules(Number) Total number of user-defined predicate rules defined in the category

Some properties such as line numbers are only available when the category is defined in a source file compiled with the `source_data` flag turned on.

The properties that return the number of clauses (rules) report the clauses (rules) *textually defined in the object* for both multifile and non-multifile predicates. Thus, these numbers exclude clauses (rules) for multifile predicates *contributed* by other entities.

1.6.8 Importing categories

Any number of objects can import a category. In addition, an object may import any number of categories. The syntax is very simple:

```
:- object(Object,
    imports([Category1, Category2, ...])).
    ...
:- end_object.
```

To make all public predicates imported via a category protected or to make all public and protected predicates private we prefix the category's name with the corresponding keyword:

```
:- object(Object,
    imports(private::Category)).
    ...
:- end_object.
```

or:

```
:- object(Object,
    imports(protected::Category)).
    ...
:- end_object.
```

Omitting the scope keyword is equivalent to writing:

```
:- object(Object,
    imports(public::Category)).
    ...
:- end_object.
```

1.6.9 Calling category predicates

Category predicates can be called from within an object by sending a message to *self* or using a *super* call. Consider the following category:

```
:- category(output).

    :- public(out/1).

    out(X) :-
        write(X), nl.

:- end_category.
```

The predicate out/1 can be called from within an object importing the category by simply sending a message to *self*. For example:

```
:- object(worker,
    imports(output)).

    ...
    do(Task) :-
        execute(Task, Result),
        ::out(Result).
    ...

:- end_object.
```

This is the recommended way of calling a category predicate that can be specialized/overridden in a descendant object as the predicate definition lookup will start from *self*.

A direct call the predicate definition found in an imported category can be made using the *^^/1* control construct. For example:

```
:- object(worker,
    imports(output)).

    ...
    do(Task) :-
        execute(Task, Result),
        ^^out(Result).
    ...
```

```
:- end_object.
```

This alternative should only be used when the user knows a priori that the category predicates will not be specialized or redefined by descendant objects of the object importing the category. Its advantage is that, when the `optimize` compiler flag is turned on, the Logtalk compiler will try to optimize the calls by using *static binding*. When *dynamic binding* is used due to e.g. the lack of sufficient information at compilation time, the performance is similar to calling the category predicate using a message to *self* (in both cases a predicate lookup caching mechanism is used).

1.6.10 Parametric categories

Category predicates can be parameterized in the same way as object predicates by using a compound term as the category identifier. The category parameters can be accessed by calling the *parameter/2* or *this/1* built-in local methods in the category predicate clauses or by using *parameter variables*. Category parameter values can be defined by the importing objects. For example:

```
:- object(speech(Season, Event),
    imports([dress(Season), speech(Event)])).
    ...
:- end_object.
```

Note that access to category parameters is only possible from within the category. In particular, calls to the *this/1* built-in local method from category predicates always access the importing object identifier (and thus object parameters, not category parameters).

1.7 Predicates

Predicate directives and clauses can be encapsulated inside objects and categories. Protocols can only contain predicate directives. From the point-of-view of an object-oriented language, predicates allows both object state and object behavior to be represented. Mutable object state can be represented using dynamic object predicates.

1.7.1 Reserved predicate names

For performance reasons, a few predicates have a fixed interpretation. These predicates are declared in the built-in protocols. They are: *goal_expansion/2* and *term_expansion/2*, declared in the `expanding` protocol; *before/3* and *after/3*, declared in the `monitoring` protocol; and *forward/1*, declared in the `forwarding` protocol. By default, the compiler prints a warning when a definition for one of these predicates is found but the reference to the corresponding built-in protocol is missing.

1.7.2 Declaring predicates

All object (or category) predicates that we want to access from other objects (or categories) must be explicitly declared. A predicate declaration must contain, at least, a *scope* directive. Other directives may be used to document the predicate or to ensure proper compilation of the predicate definitions.

Scope directives

A predicate scope directive specifies *from where* the predicate can be called, i.e. its *visibility*. Predicates can be *public*, *protected*, *private*, or *local*. Public predicates can be called from any object. Protected predicates can only

be called from the container object or from a container descendant. Private predicates can only be called from the container object. Local predicates, like private predicates, can only be called from the container object (or category) but they are *invisible* to the reflection built-in methods (*current_op/3*, *current_predicate/1*, and *predicate_property/2*) and to the message error handling mechanisms (i.e. sending a message corresponding to a local predicate results in a `predicate_declaration` existence error instead of a scope error).

The scope declarations are made using the directives *public/1*, *protected/1*, and *private/1*. For example:

```
:- public(init/1).

:- protected(valid_init_option/1).

:- private(process_init_options/1).
```

If a predicate does not have a (local or inherited) scope declaration, it is assumed that the predicate is local. Note that we do not need to write scope declarations for all defined predicates. One exception is local dynamic predicates: declaring them as private predicates may allow the Logtalk compiler to generate optimized code for asserting and retracting clauses.

Note that a predicate scope directive doesn't specify *where* a predicate is, or can be, defined. For example, a private predicate can only be called from an object holding its scope directive. But it can be defined in descendant objects. A typical example is an object playing the role of a class defining a private (possibly dynamic) predicate for its descendant instances. Only the class can call (and possibly assert/retract clauses for) the predicate but its clauses can be found/defined in the instances themselves.

Mode directive

Often predicates can only be called using specific argument patterns. The valid arguments and instantiation modes of those arguments can be documented by using the *mode/2* directive. For example:

```
:- mode(member(?term, ?list), zero_or_more).
```

The first directive argument describes a valid calling mode. The minimum information will be the instantiation mode of each argument. The first four possible values are described in [ISO95]). The remaining two can be found in use in some Prolog systems.

+ Argument must be instantiated (but not necessarily ground).

– Argument should be a free (non-instantiated) variable (when bound, the call will unify the returned term with the given term).

? Argument can either be instantiated or free.

@ Argument will not be further instantiated (modified).

++ Argument must be ground.

–– Argument must be unbound.

These four mode atoms are also declared as prefix operators by the Logtalk compiler. This makes it possible to include type information for each argument like in the example above. Some possible type values are: `event`, `object`, `category`, `protocol`, `callable`, `term`, `nonvar`, `var`, `atomic`, `atom`, `number`, `integer`, `float`, `compound`, and `list`. The first four are Logtalk specific. The remaining are common Prolog types. We can also use our own types that can be either atoms or ground compound terms.

The second directive argument documents the number of proofs (not necessarily distinct solutions) for the specified mode. Note that different modes for the same predicate often have different determinism. The possible values are:

zero Predicate always fails.

one Predicate always succeeds once.

zero_or_one Predicate either fails or succeeds.

zero_or_more Predicate has zero or more proofs.

one_or_more Predicate has one or more proofs.

one_or_error Predicate either succeeds once or throws an error (see below).

error Predicate will throw an error (see below).

Mode declarations can also be used to document that some call modes will throw an error. For instance, regarding the `arg/3` ISO Prolog built-in predicate, we may write:

```
:- mode(arg(-, -, +), error).
```

Note that most predicates have more than one valid mode implying several mode directives. For example, to document the possible use modes of the `atom_concat/3` ISO built-in predicate we would write:

```
:- mode(atom_concat(?atom, ?atom, +atom), one_or_more).
:- mode(atom_concat(+atom, +atom, -atom), zero_or_one).
```

Some old Prolog compilers supported some sort of mode directives to improve performance. To the best of my knowledge, there is no modern Prolog compiler supporting this kind of directive. The current version of the Logtalk compiler just parses and then discards this directive (however, see the description on *synchronized predicates* in the *multi-threading programming* section). Nevertheless, the use of mode directives is a good starting point for documenting your predicates.

Meta-predicate directive

Some predicates may have arguments that will be called as goals or closures that will be used for constructing goals. To ensure that these goals will be executed in the correct context (i.e. in the *calling context*, not in the meta-predicate *definition context*) we need to use the *meta_predicate/1* directive. For example:

```
:- meta_predicate(findall(*, 0, *)).
```

The meta-predicate mode arguments in this directive have the following meaning:

0 Meta-argument that will be called as a goal.

N Meta-argument that will be a closure used to construct a call by appending `N` arguments. The value of `N` must be a non-negative integer.

:: Argument that is context-aware but that will not be called as a goal.

^ Goal that may be existentially quantified (`Vars^Goal`).

* Normal argument.

The following meta-predicate mode arguments are for use only when writing backend Prolog adapter files to deal with proprietary built-in meta-predicates and meta-directives:

/ Predicate indicator (`Name/Arity`), list of predicate indicators, or conjunction of predicate indicators.

// Non-terminal indicator (`Name//Arity`), list of predicate indicators, or conjunction of predicate indicators.

[0] List of goals.

[N] List of closures.

[/] List of predicate indicators.

[//] List of non-terminal indicators.

To the best of my knowledge, the use of non-negative integers to specify closures has first introduced on Quintus Prolog for providing information for predicate cross-reference tools.

As each Logtalk entity is independently compiled, this directive must be included in every object or category that contains a definition for the described meta-predicate, even if the meta-predicate declaration is inherited from another entity, to ensure proper compilation of meta-arguments.

Discontiguous directive

The clause of an object (or category) predicate may not be contiguous. In that case, we must declare the predicate discontiguous by using the *discontiguous/1* directive:

```
:- discontiguous(foo/1).
```

This is a directive that we should avoid using: it makes your code harder to read and it is not supported by some Prolog compilers.

As each Logtalk entity is compiled independently of other entities, this directive must be included in every object or category that contains a definition for the described predicate (even if the predicate declaration is inherited from other entity).

Dynamic directive

An object predicate can be static or dynamic. By default, all object predicates are static. To declare a dynamic predicate we use the *dynamic/1* directive:

```
:- dynamic(foo/1).
```

This directive may also be used to declare dynamic grammar rule non-terminals. As each Logtalk entity is compiled independently from other entities, this directive must be included in every object that contains a definition for the described predicate (even if the predicate declaration is inherited from other object or imported from a category). If we omit the dynamic declaration then the predicate definition will be compiled static. In the case of dynamic objects, static predicates cannot be redefined using the database built-in methods (despite being internally compiled to dynamic code).

Dynamic predicates can be used to represent persistent mutable object state. Note that static objects may declare and define dynamic predicates.

Operator directive

An object (or category) predicate can be declared as an operator using the familiar *op/3* directive:

```
:- op(Priority, Specifier, Operator).
```

Operators are local to the object (or category) where they are declared. This means that, if you declare a public predicate as an operator, you cannot use operator notation when sending to an object (where the predicate is visible) the respective message (as this would imply visibility of the operator declaration in the context of the *sender* of the message). If you want to declare global operators and, at the same time, use them inside an entity, just write the corresponding directives at the top of your source file, before the entity opening directive.

When the same operators are used on several entities within the same source file, the corresponding directives must appear before any entity that uses them. However, this results in a global scope for the operators. If you prefer the operators to be local to the source file, just *undefine* them at the end of the file. For example:

```
% before any entity that uses the operator
:- op(400, xfx, results).


...


% after all entities that used the operator
:- op(0, xfx, results).
```

Uses directive

When a predicate makes heavy use of predicates defined on other objects, its predicate clauses can be verbose due to all the necessary message sending goals. Consider the following example:

```
foo :-
    ...,
    findall(X, list::member(X, L), A),
    list::append(A, B, C),
    list::select(Y, C, R),
    ...
```

Logtalk provides a directive, *uses/2*, which allows us to simplify the code above. The usage template for this directive is:

```
:- uses(Object, [Functor1/Arity1, Functor2/Arity2, ...]).
```

Rewriting the code above using this directive results in a simplified and more readable predicate definition:

```
:- uses(list, [
    append/3, member/2, select/3
]).

foo :-
    ...,
    findall(X, member(X, L), A),
    append(A, B, C),
    select(Y, C, R),
    ...
```

Logtalk also supports an extended version of this directive that allows the declaration of predicate aliases using the notation `Predicate as Alias` (or the alternative notation `Predicate::Alias`). For example:

```
:- uses(btrees, [new/1 as new_btree/1]).
:- uses(queues, [new/1 as new_queue/1]).
```

You may use this extended version for solving conflicts between predicates declared on several `uses/2` directives or just for giving new names to the predicates that will be more meaningful on their using context.

The `uses/2` directive allows simpler predicate definitions as long as there are no conflicts between the predicates declared in the directive and the predicates defined in the object (or category) containing the directive. A predicate (or its alias if defined) cannot be listed in more than one `uses/2` directive. In addition, a `uses/2` directive cannot list a predicate (or its alias if defined) which is defined in the object (or category) containing the directive. Any conflicts are reported by Logtalk as compilation errors.

Alias directive

Logtalk allows the definition of an alternative name for an inherited or imported predicate (or for an inherited or imported grammar rule non-terminal) through the use of the *alias/2* directive:

```
:- alias(Entity, [
    Predicate1 as Alias1,
    Predicate2 as Alias2,
    ...
]).
```

This directive can be used in objects, protocols, or categories. The first argument, `Entity`, must be an entity referenced in the opening directive of the entity containing the `alias/2` directive. It can be an extended or implemented protocol, an imported category, an extended prototype, an instantiated class, or a specialized class. The second argument is a list of pairs of predicate indicators (or grammar rule non-terminal indicators) using the `as` infix operator as connector.

A common use for the `alias/2` directive is to give an alternative name to an inherited predicate in order to improve readability. For example:

```
:- object(square,
    extends(rectangle)).

    :- alias(rectangle, [width/1 as side/1]).


    ...

:- end_object.
```

The directive allows both `width/1` and `side/1` to be used as messages to the object `square`. Thus, using this directive, there is no need to explicitly declare and define a "new" `side/1` predicate. Note that the `alias/2` directive does not rename a predicate, only provides an alternative, additional name; the original name continues to be available (although it may be masked due to the default inheritance conflict mechanism).

Another common use for this directive is to solve conflicts when two inherited predicates have the same functor and arity. We may want to call the predicate which is masked out by the Logtalk lookup algorithm (see the *Inheritance* section) or we may need to call both predicates. This is simply accomplished by using the `alias/2` directive to give alternative names to masked out or conflicting predicates. Consider the following example:

```
:- object(my_data_structure,
    extends(list, set)).

    :- alias(list, [member/2 as list_member/2]).
    :- alias(set,  [member/2 as set_member/2]).


    ...

:- end_object.
```

Assuming that both `list` and `set` objects define a `member/2` predicate, without the `alias/2` directives, only the definition of `member/2` predicate in the object `list` would be visible on the object `my_data_structure`, as a result of the application of the Logtalk predicate lookup algorithm. By using the `alias/2` directives, all the following messages would be valid (assuming a public scope for the predicates):

```
% uses list member/2
| ?- my_data_structure::list_member(X, L).
```

```
 % uses set member/2
| ?- my_data_structure::set_member(X, L).

% uses list member/2
| ?- my_data_structure::member(X, L).
```

When used this way, the `alias/2` directive provides functionality similar to programming constructs of other object-oriented languages that support multi-inheritance (the most notable example probably being the renaming of inherited features in Eiffel).

Note that the `alias/2` directive never hides a predicate which is visible on the entity containing the directive as a result of the Logtalk lookup algorithm. However, it may be used to make visible a predicate which otherwise would be masked by another predicate, as illustrated in the above example.

The `alias/2` directive may also be used to give access to an inherited predicate, which otherwise would be masked by another inherited predicate, while keeping the original name as follows:

```
:- object(my_data_structure,
    extends(list, set)).

    :- alias(list, [member/2 as list_member/2]).
    :- alias(set,  [member/2 as set_member/2]).

    member(X, L) :-
        ::set_member(X, L).

    ...

:- end_object.
```

Thus, when sending the message `member/2` to `my_data_structure`, the predicate definition in `set` will be used instead of the one contained in `list`.

Documenting directive

A predicate can be documented with arbitrary user-defined information by using the *info/2* directive:

```
:- info(Name/Arity, List).
```

The second argument is a list of `Key is Value` terms. See the *Documenting applications* section for details.

Multifile directive

A predicate can be declared *multifile* by using the *multifile/1* directive:

```
:- multifile(Name/Arity).
```

This allows clauses for a predicate to be defined in several objects and/or categories. This is a directive that should be used with care. Support for this directive have been added to Logtalk primarily to support migration of Prolog module code. Spreading clauses for a predicate among several Logtalk entities can be handy in some cases but can also make your code difficult to understand. Logtalk precludes using a multifile predicate for breaking object encapsulation by checking that the object (or category) declaring the predicate (using a scope directive) defines it also as multifile. This entity is said to contain the *primary declaration* for the multifile predicate. In addition, note that the `multifile/1` directive is mandatory when defining multifile predicates.

Consider the following simple example:

```
:- object(main).

    :- public(a/1).
    :- multifile(a/1).
    a(1).

:- end_object.
```

After compiling and loading the `main` object, we can define other objects (or categories) that contribute with clauses for the multifile predicate. For example:

```
:- object(other).

    :- multifile(main::a/1).
    main::a(2).
    main::a(X) :-
        b(X).

    b(3).
    b(4).

:- end_object.
```

After compiling and loading the above objects, you can use queries such as:

```
| ?- main::a(X).

X = 1 ;
X = 2 ;
X = 3 ;
X = 4
yes
```

Entities containing *primary multifile predicate declarations* must always be compiled before entities defining clauses for those multifile predicates. The Logtalk compiler will print a warning if the scope directive is missing.

Multifile predicates may also be declared dynamic using the same `Entity::Name/Arity` notation (multifile predicates are static by default).

When a clause of a multifile predicate is a rule, its body is compiled within the context of the object or category defining the clause. This allows clauses for multifile predicates to call local object or category predicates. But the values of the *sender*, *this*, and *self* in the implicit execution context are passed from the clause head to the clause body. This is necessary to ensure that these values are always valid and to allow multifile predicate clauses to be defined in categories. A call to the `parameter/2` execution context methods, however, retrieves parameters of the entity defining the clause, not from the entity for which the clause is defined. The parameters of the entity for which the clause is defined can be accessed by simple unification at the clause head.

Local calls to the database methods from multifile predicate clauses defined in an object take place in the object own database instead of the database of the entity holding the multifile predicate primary declaration. Similarly, local calls to the `expand_term/2` and `expand_goal/2` methods from a multifile predicate clause look for clauses of the `term_expansion/2` and `goal_expansion/2` hook predicates starting from the entity defining the clause instead of the entity holding the multifile predicate primary declaration. Local calls to the `current_predicate/1`, `predicate_property/2`, and `current_op/3` methods from multifile predicate clauses defined in an object also lookup predicates and their properties in the object own database instead of the database of the entity holding the multifile predicate primary declaration.

Coinductive directive

A predicate can be declared *coinductive* by using the *coinductive/1* directive. For example:

```
:- coinductive(comember/2).
```

Logtalk support for coinductive predicates is experimental and requires a back-end Prolog compiler with minimal support for cyclic terms.

1.7.3 Defining predicates

Object predicates

We define object predicates as we have always defined Prolog predicates, the only difference be that we have four more control structures (the three message sending operators plus the external call operator) to play with. For example, if we wish to define an object containing common utility list predicates like `append/2` or `member/2` we could write something like:

```
:- object(list).

    :- public(append/3).
    :- public(member/2).

    append([], L, L).
    append([H| T], L, [H| T2]) :-
        append(T, L, T2).

    member(H, [H| _]).
    member(H, [_| T]) :-
        member(H, T).

:- end_object.
```

Note that, abstracting from the opening and closing object directives and the scope directives, what we have written is plain Prolog. Calls in a predicate definition body default to the local predicates, unless we use the message sending operators or the external call operator. This enables easy conversion from Prolog code to Logtalk objects: we just need to add the necessary encapsulation and scope directives to the old code.

Category predicates

Because a category can be imported by multiple objects, dynamic private predicates must be called either in the context of *self*, using the *message to self* control structure, *::/1*, or in the context of *this* (i.e. in the context of the object importing the category). For example, if we want to define a category implementing variables using destructive assignment where the variable values are stored in *self* we could write:

```
:- category(variable).

    :- public(get/2).
    :- public(set/2).

    :- private(value_/2).
    :- dynamic(value_/2).

    get(Var, Value) :-
```

```
        ::value_(Var, Value).

    set(Var, Value) :-
        ::retractall(value_(Var, _)),
        ::asserta(value_(Var, Value).

:- end_category.
```

In this case, the `get/2` and `set/2` predicates will always access/update the correct definition, contained in the object receiving the messages. The alternative, storing the variable values in *this*, such that each object importing the category will have its own definition for the `value_/2` private predicate is simple: just omit the use of the `::/1` control construct in the code above.

A category can only contain clauses for static predicates. Nevertheless, as the example above illustrates, there are no restrictions in declaring and calling dynamic predicates from inside a category.

Meta-predicates

Meta-predicates may be defined inside objects (and categories) as any other predicate. A meta-predicate is declared using the `meta_predicate/1` directive as described earlier on this section. When defining a meta-predicate, the arguments in the clause heads corresponding to the meta-arguments must be variables. All meta-arguments are called in the context of the entity calling the meta-predicate.

Some meta-predicates have meta-arguments which are not goals but closures. Logtalk supports the definition of meta-predicates that are called with closures instead of goals as long as the definition uses the Logtalk built-in predicate *call/1-N* to call the closure with the additional arguments. For example:

```
:- public(all_true/2).
:- meta_predicate(all_true(1, *)).

all_true(_, []).
all_true(Closure, [Arg| Args]) :-
    call(Closure, Arg),
    all_true(Closure, Args).
```

Note that in this case the meta-predicate directive specifies that the closure will be extended with exactly one extra argument.

When calling a meta-predicate, a closure can correspond to a user-defined predicate, a built-in predicate, a *lambda expression*, or a control construct.

Lambda expressions

The use of lambda expressions as meta-predicate goal and closure arguments often saves writing auxiliary predicates for the sole purpose of calling the meta-predicates. A simple example of a lambda expression is:

```
| ?- meta::map([X,Y]>>(Y is 2*X), [1,2,3], Ys).
Ys = [2,4,6]
yes
```

In this example, a lambda expression, `[X,Y]>>(Y is 2*X)`, is used as an argument to the `map/3` list mapping predicate, defined in the library object `meta`, in order to double the elements of a list of integers. Using a lambda expression avoids writing an auxiliary predicate for the sole purpose of doubling the list elements. The lambda parameters are represented by the list `[X,Y]`, which is connected to the lambda goal, `(Y is 2*X)`, by the `(>>)/2` operator.

Currying is supported. I.e. it is possible to write a lambda expression whose goal is another lambda expression. The above example can be rewritten as:

```
| ?- meta::map([X]>>([Y]>>(Y is 2*X)), [1,2,3], Ys).
Ys = [2,4,6]
yes
```

Lambda expressions may also contain lambda free variables. I.e. variables that are global to the lambda expression. For example, using GNU Prolog as the back-end compiler, we can write:

```
| ?- meta::map({Z}/[X,Y]>>(Z#=X+Y), [1,2,3], Zs).
Z = _#22(3..268435455)
Zs = [_#3(2..268435454),_#66(1..268435453),_#110(0..268435452)]
yes
```

The ISO Prolog construct {}/1 for representing the lambda free variables as this representation is often associated with set representation. Note that the order of the free variables is of no consequence (on the other hand, a list is used for the lambda parameters as their order does matter).

Both lambda free variables and lambda parameters can be any Prolog term. Consider the following example by Markus Triska:

```
| ?- meta::map([A-B,B-A]>>true, [1-a,2-b,3-c], Zs).
Zs = [a-1,b-2,c-3]
yes
```

Lambda expressions can be used, as expected, in non-deterministic queries as in the following example using SWI-Prolog as the back-end compiler and Markus Triska's CLP(FD) library:

```
| ?- meta::map({Z}/[X,Y]>>(clpfd:(Z#=X+Y)), Xs, Ys).
Xs = [],
Ys = [] ;
Xs = [_G1369],
Ys = [_G1378],
_G1369+_G1378#=Z ;
Xs = [_G1579, _G1582],
Ys = [_G1591, _G1594],
_G1582+_G1594#=Z,
_G1579+_G1591#=Z ;
Xs = [_G1789, _G1792, _G1795],
Ys = [_G1804, _G1807, _G1810],
_G1795+_G1810#=Z,
_G1792+_G1807#=Z,
_G1789+_G1804#=Z ;
...
```

As illustrated by the above examples, lambda expression syntax reuses the ISO Prolog construct {}/1 and the standard operators (/)/2 and (>>)/2, thus avoiding defining new operators, which is always tricky for a portable system such as Logtalk. The operator (>>)/2 was chosen as it suggests an arrow, similar to the syntax used in other languages such as OCaml and Haskell to connect lambda parameters with lambda functions. This syntax was also chosen in order to simplify parsing, error checking, and compilation of lambda expressions. The full specification of the lambda expression syntax can be found in the *Lambda expressions* section of the language grammar.

The compiler checks whenever possible that all variables in a lambda expression are either classified as free variables or as lambda parameters. Non-classified variables in a lambda expression should be regarded as a programming error. Unfortunately, the dynamic features of the language and lack of sufficient information at compile time may prevent the compiler of checking all uses of lambda expressions. The compiler also checks if a variable is classified as both

a free variable and a lambda parameter. An optimizing meta-predicate and lambda expression compiler, based on the term-expansion mechanism, is provided for practical performance.

Definite clause grammar rules

Definite clause grammar rules provide a convenient notation to represent the rewrite rules common of most grammars in Prolog. In Logtalk, definite clause grammar rules can be encapsulated in objects and categories. Currently, the ISO/IEC WG17 group is working on a draft specification for a definite clause grammars Prolog standard. Therefore, in the mean time, Logtalk follows the common practice of Prolog compilers supporting definite clause grammars, extending it to support calling grammar rules contained in categories and objects. A common example of a definite clause grammar is the definition of a set of rules for parsing simple arithmetic expressions:

```
:- object(calculator).

    :- public(parse/2).

    parse(Expression, Value) :-
        phrase(expr(Value), Expression).

    expr(Z) --> term(X), "+", expr(Y), {Z is X + Y}.
    expr(Z) --> term(X), "-", expr(Y), {Z is X - Y}.
    expr(X) --> term(X).

    term(Z) --> number(X), "*", term(Y), {Z is X * Y}.
    term(Z) --> number(X), "/", term(Y), {Z is X / Y}.
    term(Z) --> number(Z).

    number(C) --> "+", number(C).
    number(C) --> "-", number(X), {C is -X}.
    number(X) --> [C], {0'0 =< C, C =< 0'9, X is C - 0'0}.

:- end_object.
```

The predicate *phrase/2* called in the definition of predicate `parse/2` above is a Logtalk built-in method, similar to the predicate with the same name found on most Prolog compilers that support definite clause grammars. After compiling and loading this object, we can test the grammar rules with calls such as the following one:

```
| ?- calculator::parse("1+2-3*4", Result).

Result = -9
yes
```

In most cases, the predicates resulting from the translation of the grammar rules to regular clauses are not declared. Instead, these predicates are usually called by using the built-in methods *phrase/2* and *phrase/3* as shown in the example above. When we want to use the built-in methods `phrase/2` and `phrase/3`, the non-terminal used as first argument must be within the scope of the *sender*. For the above example, assuming that we want the predicate corresponding to the `expr//1` non-terminal to be public, the corresponding scope directive would be:

```
:- public(expr//1).
```

The `//` infix operator used above tells the Logtalk compiler that the scope directive refers to a grammar rule non-terminal, not to a predicate. The idea is that the predicate corresponding to the translation of the `expr//1` non-terminal will have a number of arguments equal to one plus the number of additional arguments necessary for processing the implicit difference list of tokens.

In the body of a grammar rule, we can call rules that are inherited from ancestor objects, imported from categories, or contained in other objects. This is accomplished by using non-terminals as messages. Using a non-terminal as a

message to *self* allows us to call grammar rules in categories and ancestor objects. To call grammar rules encapsulated in other objects, we use a non-terminal as a message to those objects. Consider the following example, containing grammar rules for parsing natural language sentences:

```
:- object(sentence,
    imports(determiners, nouns, verbs)).

    :- public(parse/2).

    parse(List, true) :-
        phrase(sentence, List).
    parse(_, false).

    sentence --> noun_phrase, verb_phrase.

    noun_phrase --> ::determiner, ::noun.
    noun_phrase --> ::noun.

    verb_phrase --> ::verb.
    verb_phrase --> ::verb, noun_phrase.

:- end_object.
```

The categories imported by the object would contain the necessary grammar rules for parsing determiners, nouns, and verbs. For example:

```
:- category(determiners).

    :- private(determiner//0).

    determiner --> [the].
    determiner --> [a].

:- end_category.
```

Along with the message sending operators (::/1, ::/2, and ^^/1), we may also use other control constructs such as \+/1, !/0, ;/2, ->/2, and {}/1 in the body of a grammar. In addition, grammar rules may contain meta-calls (a variable taking the place of a non-terminal), which are translated to calls of the built-in method phrase/3.

You may have noticed that Logtalk defines *{}/1* as a control construct for bypassing the compiler when compiling a clause body goal. As exemplified above, this is the same control construct that is used in grammar rules for bypassing the expansion of rule body goals when a rule is converted into a clause. Both control constructs can be combined in order to call a goal from a grammar rule body, while bypassing at the same time the Logtalk compiler. Consider the following example:

```
bar :-
    write('bar predicate called'), nl.


:- object(bypass).

    :- public(foo//0).

    foo --> {{bar}}.

:- end_object.
```

After compiling and loading this code, we may try the following query:

```
| ?- logtalk << phrase(bypass::foo, _, _).

bar predicate called
yes
```

This is the expected result as the expansion of the grammar rule into a clause leaves the `{bar}` goal untouched, which, in turn, is converted into the goal `bar` when the clause is compiled.

A grammar rule non-terminal may be declared as dynamic or discontiguous, as any object predicate, using the same `Name//Arity` notation illustrated above for the scope directives. In addition, grammar rule non-terminals can be documented using the *info/2* directive, as in the following example:

```
:- public(sentence//0).

:- info(sentence//0, [
    comment is 'Rewrites sentence into noun and verb phrases.']).
```

1.7.4 Built-in object predicates (methods)

Logtalk defines a set of built-in object predicates or methods to access message execution context, to find sets of solutions, to inspect objects, for database handling, for term and goal expansion, and for printing messages. Similar to Prolog built-in predicates, these built-in methods should not be redefined.

Execution context methods

Logtalk defines five built-in private methods to access an object execution context. These methods are in the common usage scenarios translated to a single unification performed at compile time with a clause head context argument. Therefore, they can be freely used without worrying about performance penalties. When called from inside a category, these methods refer to the execution context of the object importing the category. These methods are private and cannot be used as messages to objects.

To find the object that received the message under execution we may use the *self/1* method. We may also retrieve the object that has sent the message under execution using the *sender/1* method.

The method *this/1* enables us to retrieve the name of the object for which the predicate clause whose body is being executed is defined instead of using the name directly. This helps to avoid breaking the code if we decide to change the object name and forget to change the name references. This method may also be used from within a category. In this case, the method returns the object importing the category on whose behalf the predicate clause is being executed.

Here is a short example including calls to these three object execution context methods:

```
:- object(test).

    :- public(test/0).

    test :-
        this(This),
        write('Calling predicate definition in '),
        writeq(This), nl,
        self(Self),
        write('to answer a message received by '),
        writeq(Self), nl,
        sender(Sender),
        write('that was sent by '),
        writeq(Sender), nl, nl.
```

(continues on next page)

```
:- end_object.


:- object(descendant,
    extends(test)).

:- end_object.
```

After compiling and loading these two objects, we can try the following goal:

```
| ?- descendant::test.

Calling predicate definition in test
to answer a message received by descendant
that was sent by user
yes
```

Note that the goals `self(Self)`, `sender(Sender)`, and `this(This)`, being translated to unifications with the clause head context arguments at compile time, are effectively removed from the clause body. Therefore, a clause such as:

```
predicate(Arg) :-
    self(Self),
    atom(Arg),
    ... .
```

is compiled with the goal `atom(Arg)` as the first condition on the clause body. As such, the use of these context execution methods do not interfere with the optimizations that some Prolog compilers perform when the first clause body condition is a call to a built-in type-test predicate or a comparison operator.

For parametric objects and categories, the method *parameter/2* enables us to retrieve current parameter values (see the section on parametric objects for a detailed description). For example:

```
:- object(block(_Color)).

    :- public(test/0).

    test :-
        parameter(1, Color),
        write('Color parameter value is '),
        writeq(Color), nl.

:- end_object.
```

An alternative to the `parameter/2` predicate is to use *parameter variables*:

```
:- object(block(_Color_)).

    :- public(test/0).

    test :-
        write('Color parameter value is '),
        writeq(_Color_), nl.

:- end_object.
```

After compiling and loading either version of the object, we can try the following goal:

```
| ?- block(blue)::test.

Color parameter value is blue
yes
```

Calls to the `parameter/2` method are translated to a compile time unification when the second argument is a variable. When the second argument is bound, the calls are translated to a call to the built-in predicate `arg/3`.

When type-checking predicate arguments, it is often useful to include the predicate execution context when reporting an argument error. The *context/1* method provides access to that context. For example, assume a predicate `foo/2` that takes an atom and an integer as arguments. We could type-check the arguments by writing (using the library `type` object):

```
foo(A, N) :-
    % type-check arguments
    context(Context),
    type::check(atom, A, Context),
    type::check(integer, N, Context),
    % arguments are fine; go ahead
    ... .
```

Error handling and throwing methods

Besides the *catch/3* and *throw/1* methods inherited from Prolog, Logtalk also provides a set of convenience methods to throw standard `error/2` exception terms: *instantiation_error/0*, *type_error/2*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *representation_error/1*, *evaluation_error/1*, *resource_error/1*, *syntax_error/1*, and *system_error/0*.

Database methods

Logtalk provides a set of built-in methods for *object database* handling similar to the usual database Prolog predicates: *abolish/1*, *asserta/1*, *assertz/1*, *clause/2*, *retract/1*, and *retractall/1*. These methods always operate on the database of the object receiving the corresponding message.

When working with dynamic grammar rule non-terminals, you may use the built-in method *expand_term/2* convert a grammar rule into a clause that can then be used with the database methods.

Meta-call methods

Logtalk supports the generalized *call/1-N* meta-predicate. This built-in private meta-predicate must be used in the implementation of meta-predicates which work with closures instead of goals. In addition, Logtalk supports the built-in private meta-predicates *ignore/1*, *once/1*, and *\+/1*. These methods cannot be used as messages to objects.

All solutions methods

The usual all solutions meta-predicates are built-in private methods in Logtalk: *bagof/3*, *findall/3*, *findall/4*, and *setof/3*. There is also a *forall/2* method that implements generate and test loops. These methods cannot be used as messages to objects.

Reflection methods

Logtalk provides a comprehensive set of built-in predicates and built-in methods for querying about entities and predicates. Some of the information, however, requires that the source files are compiled with the `source_data` flag turned on.

The reflection API supports two different views on entities and their contents, which we may call the *transparent box view* and the *black box view*. In the transparent box view, we look into an entity disregarding how it will be used and returning all information available on it, including predicate declarations and predicate definitions. This view is supported by the entity property built-in predicates. In the black box view, we look into an entity from a usage point-of-view using built-in methods for inspecting object operators and predicates that are within scope from where we are making the call: *current_op/3*, which returns operator specifications, *predicate_property/2*, which returns predicate properties, and *current_predicate/1*, which enables us to query about predicate definitions. See below for a more detailed description of these methods.

Definite clause grammar parsing methods and non-terminals

Logtalk supports two definite clause grammar parsing built-in private methods, *phrase/2* and *phrase/3*, with definitions similar to the predicates with the same name found on most Prolog compilers that support definite clause grammars. These methods cannot be used as messages to objects.

Logtalk also supports *phrase//1*, *call//1-N*, and *eos//0* built-in non-terminals. The `call//1-N` non-terminals takes a closure (which can be a lambda expression) plus zero or more additional arguments and are processed by appending the input list of tokens and the list of remaining tokens to the arguments.

Term and goal expansion methods

Logtalk supports an *expand_term/2* built-in method for expanding a term into another term or a list of terms. This method is mainly used to translate grammar rules into Prolog clauses. It can be customized, e.g. for bypassing the default Logtalk grammar rule translator, by defining clauses for the *term_expansion/2* hook predicate. Logtalk also supports an *expand_goal/2* built-in method for expanding a goal. This method can also be customized by defining clauses for the *goal_expansion/2* hook predicate.

Term and goal expansion may be performed either by calling the `expand_term/2` and `expand_goal/2` built-in methods explicitly or by using *hook objects*. A hook object is simply an object defining clauses for the term- and goal-expansion hook predicates. To compile a source file using a hook object for expanding its terms and goals, you can use the *hook/1* compiler flag in the second argument of the *logtalk_compile/2* or *logtalk_load/2* built-in predicates. In alternative, you can use a *set_logtalk_flag/2* directive in the source file itself. When compiling a source file, the compiler will first try the source file specific hook object, if defined. If that fails, it tries the default hook object, if defined. If that also fails, the compiler tries the Prolog dialect specific expansion predicate definitions if defined in the adapter file.

Sometimes we have multiple hook objects that we need to use in the compilation of a source file. The Logtalk library includes support for two basic expansion workflows: a pipeline of hook objects, where the expansion results from a hook object are feed to the next hook object in the pipeline, and a set of hook objects, where expansions are tried until one of is found that succeeds. These workflows are implemented as parametric objects allowing combining them to implement more sophisticated expansion workflows.

Clauses for the `term_expansion/2` and `goal_expansion/2` predicates defined within an object or a category are never used in the compilation of the object or the category itself, however. In order to use clauses for the `term_expansion/2` and `goal_expansion/2` predicates defined in plain Prolog, simply specify the pseudo-object `user` as the hook object when compiling source files. When using backend Prolog compilers that support a module system, it can also be specified a module containing clauses for the expanding predicates as long as the module name doesn't coincide with an object name. But note that Prolog module libraries may provide definitions of the expansion predicates that are not compatible with the Logtalk compiler. Specially when setting the hook object to

user, be aware of any Prolog library that is loaded, possibly by default or implicitly by the Prolog system, that may be contributing definitions of the expansion predicates. It is usually much safer to define a specific hook object for combining multiple expansions in a fully controlled way.

Logtalk provides a *logtalk_load_context/2* built-in predicate that can be used to access the compilation/loading context when performing term-expansion or goal-expansion.

Printing messages

Logtalk features a *structured message printing* mechanism. This mechanism gives the programmer full control of message printing, allowing it to filter, rewrite, or redirect any message. The origins of the message printing mechanism that inspired the Logtalk implementation goes back to Quintus Prolog, where it was apparently implemented by Dave Bowen (thanks to Richard O'Keefe for the historical bits). Variations of this mechanism can also be found currently on e.g. SICStus Prolog, SWI-Prolog, and YAP.

Why a mechanism for printing messages? Consider the different components in a Logtalk application development and execution. At the bottom level, you have the Logtalk compiler and runtime. The Logtalk compiler writes messages related to e.g. compiling and loading files, compiling entities, compilation warnings and errors. The Logtalk runtime may write banner messages and handles execution errors that may result in printing human-level messages. The development environment can be console-based or you may be using a GUI tool such as PDT. In the latter case, PDT needs to intercept the Logtalk compiler and runtime messages to present the relevant information using its GUI. Then you have all the other components in a typical application. For example, your own libraries and third-party libraries. The libraries may want to print messages on its own, e.g. banners, debugging information, or logging information. As you assemble all your application components, you want to have the final word on which messages are printed, where, an in what conditions. Uncontrolled message printing by libraries could potentially e.g. disturb application flow, expose implementation details, spam the user with irrelevant details, or break user interfaces.

The Logtalk message printing mechanism provides you with a set of predicates, defined in the `logtalk` built-in object, and some hook predicates. Two of the most important predicates are *print_message(Kind, Component, Term)*, which is used for printing a message, and *message_hook(Term, Kind, Component, Tokens)*, a user-defined hook predicate used for intercepting messages. The `Kind` argument is used to represent the nature of the message being printed. It can be e.g. a logging message, a warning message or an error message, In Logtalk this argument can be either an atom, e.g. `error`, or a compound term, e.g. `comment(loading)`. Using a compound term allows easy partitioning of messages of the same kind in different groups. The following kinds of message are recognized by default:

banner banner messages (used e.g. when loading tools or main application components; can be suppressed by setting the `report` flag to `warnings` or `off`)

help messages printed in reply for the user asking for help (mostly for helping port existing Prolog code)

information and information(Group) messages usually printed in reply to a user request for information

silent and silent(Group) not printed by default (but can be intercepted using the `message_hook/4` predicate)

comment and comment(Group) useful but usually not essential messages (can be suppressed by setting the `report` flag to `warnings` or `off`)

warning and warning(Group) warning messages (generated e.g. by the compiler; can be suppressed by turning off the `report` flag)

error and error(Group) error messages (generated e.g. by the compiler)

Note that you can define your own alternative message kind identifiers, for your own components, together with suitable definitions for their associated prefixes and output streams.

Messages are represented by atoms or compound terms, handy for machine-processing, and converted into a list of tokens, for human consumption. This conversion is performed using the multifile non-terminal *message_tokens(Term, Component)*. A simple example is:

```
:- multifile(logtalk::message_tokens//2).
:- dynamic(logtalk::message_tokens//2).

logtalk::message_tokens(loaded_settings_file(Path), core) -->
    ['Loaded settings file found on directory ~w'-[Path], nl, nl].
```

The `Component` argument is new in the Logtalk implementation and is useful to filter messages belonging to a specific component (e.g. the Logtalk compiler and runtime is identified by the atom `core`) and also to avoid conflicts when two components define the same message term (e.g. `banner`).

The following tokens can be used when translating a message:

at_same_line Signals a following part to a multi-part message with no line break in between; this token is ignored when it's not the first in the list of tokens

flush Flush the output stream (by calling the `flush_output/1` standard predicate)

nl Change line in the output stream

Format-Arguments `Format` must be an atom and `Arguments` must be a list of format arguments (the token arguments are passed to a call to the `format/3` de facto standard predicate)

term(Term, Options) `Term` can be any term and `Options` must be a list of valid `write_term/3` output options (the token arguments are passed to a call to the `write_term/3` standard predicate)

ansi(Attributes, Format, Arguments) Taken from SWI-Prolog; by default, do nothing; can be used for styled output

begin(Kind, Var) Taken from SWI-Prolog; by default, do nothing; can be used together with `end(Var)` to wrap a sequence of message tokens

end(Var) Taken from SWI-Prolog; by default, do nothing

There are also predicates for printing a list of tokens, *print_message_tokens(Stream, Prefix, Tokens)*, for hooking into printing an individual token, *print_message_token(Stream, Prefix, Token, Tokens)*, and for setting default output stream and message prefixes, *message_prefix_stream(Kind, Component, Prefix, Stream)*. For example, the SWI-Prolog adapter file uses the `print_message_token/4` hook predicate to enable coloring of messages printed on a console.

Using the message printing mechanism in your applications and libraries is easy. Simply chose a unique component name for your application (e.g. the name of the application itself), call `logtalk::print_message/3` for every message that you may want to print, and define default translations for your message terms using the multifile non-terminal `logtalk::message_tokens/2`. For a full programming example, see e.g. the `lgtunit` tool source code.

Asking questions

Logtalk features a *structured question asking* mechanism that complements the message printing mechanism. This feature provides an abstraction for the common task of asking a user a question and reading back its reply. By default, this mechanism writes the question, writes a prompt, and reads the answer from the current user input and output streams but allows both steps to be intercepted, filtered, rewritten, and redirected. Two typical examples are using a GUI dialog for asking questions and automatically providing answers to specific questions.

The question asking mechanism works in tandem with the message printing mechanism, using it to print the question text and a prompt. It provides an asking predicate and a hook predicate, both declared and defined in the `logtalk` built-in object. The asking predicate, *ask_question(Kind, Component, Question, Check, Answer)*, is used for ask a question and read the answer. The hook predicate, *question_hook(Question, Kind, Component, Tokens, Check, Answer)*, is used for intercepting questions. The `Kind` argument is used to represent the nature of the question being asked. Its default value is `question` but it can be any atom or compound term, e.g. `question(parameters)`.

Using a compound term allows easy partitioning of messages of the same kind in different groups. The `Check` argument is a closure that is converted into a checking goal taking as argument the user answer. The `ask_question/5` implements a read loop that terminates when this checking predicate is true. The question itself is a term that is translated into printing tokens using the `message_tokens/2` multifile predicate described above.

There is also a user-defined multifile predicate for setting default prompt and input streams, `question_prompt_stream(Kind, Component, Prompt, Stream)` *question_prompt_stream/4*.

A usage example of this mechanism can be found in the `debugger` tool where it's used to abstract the user interaction when tracing a goal execution in debug mode.

1.7.5 Predicate properties

We can find the properties of visible predicates by calling the *predicate_property/2* built-in method. For example:

```
| ?- bar::predicate_property(foo(_), Property).
```

Note that this method respects the predicate's scope declarations. For instance, the above call will only return properties for public predicates.

An object's set of visible predicates is the union of all the predicates declared for the object with all the built-in methods and all the Logtalk and Prolog built-in predicates.

The following predicate properties are supported:

scope(Scope) The predicate scope (useful for finding the predicate scope with a single call to `predicate_property/2`)

public, protected, private The predicate scope (useful for testing if a predicate have a specific scope)

static, dynamic All predicates are either static or dynamic (note, however, that a dynamic predicate can only be abolished if it was dynamically declared)

logtalk, prolog, foreign A predicate can be defined in Logtalk source code, Prolog code, or in foreign code (e.g. in C)

built_in The predicate is a built-in predicate

multifile The predicate is declared multifile (i.e. it can have clauses defined in several entities)

meta_predicate(Template) The predicate is declared as a meta-predicate with the specified template

coinductive(Template) The predicate is declared as a coinductive predicate with the specified template

declared_in(Entity) The predicate is declared (using a scope directive) in the specified entity

defined_in(Entity) The predicate definition is looked up in the specified entity (note that this property does not necessarily imply that clauses for the predicate exist in `Entity`; the predicate can simply be false as per the closed-world assumption)

redefined_from(Entity) The predicate is a redefinition of a predicate definition inherited from the specified entity

non_terminal(NonTerminal//Arity) The predicate resulted from the compilation of the specified grammar rule non-terminal

alias_of(Predicate) The predicate (name) is an alias for the specified predicate

alias_declared_in(Entity) The predicate alias is declared in the specified entity

synchronized The predicate is declared as synchronized (i.e. it's a deterministic predicate synchronized using a mutex when using a backend Prolog compiler supporting a compatible multi-threading implementation)

Some properties are only available when the entities are defined in source files and when those source files are compiled with the `source_data` flag turned on:

inline The predicate definition is inlined

auxiliary The predicate is not user-defined but rather automatically generated by the compiler or the term-expansion mechanism

mode(Mode, Solutions) Instantiation, type, and determinism mode for the predicate (which can have multiple modes)

info(ListOfPairs) Documentation key-value pairs as specified in the user-defined `info/2` directive

number_of_clauses(N) The number of clauses for the predicate existing at compilation time (note that this property is not updated at runtime when asserting and retracting clauses for dynamic predicates)

number_of_rules(N) The number of rules for the predicate existing at compilation time (note that this property is not updated at runtime when asserting and retracting clauses for dynamic predicates)

declared_in(Entity, Line) The predicate is declared (using a scope directive) in the specified entity in a source file at the specified line (if applicable)

defined_in(Entity, Line) The predicate is defined in the specified entity in a source file at the specified line (if applicable)

redefined_from(Entity, Line) The predicate is a redefinition of a predicate definition inherited from the specified entity, which is defined in a source file at the specified line (if applicable)

alias_declared_in(Entity, Line) The predicate alias is declared in the specified entity in a source file at the specified line (if applicable)

The properties `declared_in/1-2`, `defined_in/1-2`, and `redefined_from/1-2` do not apply to built-in methods and Logtalk or Prolog built-in predicates. Note that if a predicate is declared in a category imported by the object, it will be the category name — not the object name — that will be returned by the property `declared_in/1`. The same is true for protocol declared predicates.

1.7.6 Finding declared predicates

We can find, by backtracking, all visible user predicates by calling the *current_predicate/1* built-in method. This method respects the predicate's scope declarations. For instance, the following call:

```
| ?- some_object::current_predicate(Name/Arity).
```

will only return user predicates that are declared public. The predicate property `non_terminal/1` may be used to retrieve all grammar rule non-terminals declared for an object. For example:

```
current_non_terminal(Object, NonTerminal//Args) :-
    Object::current_predicate(Name/Arity),
    functor(Predicate, Functor, Arity),
    Object::predicate_property(Predicate, non_terminal(NonTerminal//Args)).
```

Usually, the non-terminal and the corresponding predicate share the same functor but users should not rely on this always being true.

1.7.7 Calling Prolog built-in predicates

In predicate definitions, predicate calls which are not prefixed with a message sending operator (either `::` or `^^`), are compiled to either calls to local predicates or as calls to Logtalk/Prolog built-in predicates. A predicate call is

compiled as a call to a local predicate if the object (or category) contains a scope directive, a definition for the called predicate, or a dynamic declaration for it. When the object (or category) does not contain either a definition of the called predicate or a corresponding dynamic declaration, Logtalk tests if the call corresponds to a Logtalk or Prolog built-in predicate. Calling a predicate which is neither a local predicate nor a Logtalk/Prolog built-in predicate results in a compile time warning. This means that, in the following example:

```
foo :-
    ...,
    write(bar),
    ...
```

the call to the predicate `write/1` will be compiled as a call to the corresponding Prolog built-in predicate unless the object (or category) encapsulating the above definition also contains a predicate named `write/1` or a dynamic declaration for the predicate.

When calling non-standard Prolog built-in predicates or using non-standard Prolog arithmetic functions, you may run into portability problems while trying your applications with different back-end Prolog compilers (non-standard predicates and non-standard arithmetic functions are often specific to a Prolog compiler). You may use the Logtalk compiler flag `portability/1` at *Compiler flags* to help check for problematic calls in your code.

Calling Prolog non-standard meta-predicates

Prolog built-in meta-predicates may only be called locally within objects or categories, i.e. they cannot be used as messages. Compiling calls to non-standard, Prolog built-in meta-predicates can be tricky, however, as there is no standard way of checking if a built-in predicate is also a meta-predicate and finding out which are its meta-arguments. But Logtalk supports override the original meta-predicate template if not programmatically available or usable. For example, assume a `det_call/1` Prolog built-in meta-predicate that takes a goal as argument. We can add to the object (or category) calling it the directive:

```
:- meta_predicate(user:det_call(0)).
```

Another solution is to explicitly declare all non-standard Prolog meta-predicates in the corresponding adapter file using the internal predicate `'$lgt_prolog_meta_predicate'/3`. For example:

```
'$lgt_prolog_meta_predicate'(det_call(_), det_call(0), predicate).
```

The third argument can be either the atom `predicate` or the atom `control_construct`, a distinction that is useful when compiling in debug mode.

1.7.8 Calling Prolog user-defined predicates

Prolog user-defined predicates can be called from within objects or categories by using the *{}/1* compiler bypass control construct. For example:

```
foo :-
    ...,
    {bar},
    ...
```

In alternative, you can also use the *uses/2* directive and write:

```
:- uses(user, [bar/0]).

foo :-
```

```
    ...,
    bar,
    ...
```

Note that `user` is a pseudo-object in Logtalk containing all predicate definitions that are not encapsulated (either in a Logtalk entity or a Prolog module).

Calling Prolog module predicates

To call Prolog module predicates from within objects or categories you can use simply write:

```
foo :-
    ...,
    module:bar,
    ...
```

You can also use in alternative the *use_module/2* directive:

```
:- use_module(module, [bar/0]).

foo :-
    ...,
    bar,
    ...
```

Note that the first argument of the `use_module/2`, when used within an object or a category, is a module name, not a file name. The actual module code should be loaded prior to compilation of Logtalk that uses it. In particular, programmers should not expect that the module be auto-loaded (when using back-end Prolog compilers supporting an autoloading mechanism).

When the module predicate is a meta-predicate but for some reason you don't want its calls to be compiled as such, you can use the *{}/1* compiler bypass control construct as before:

```
foo :-
    ...,
    {module:bar},
    ...
```

This workaround is sometimes necessary when calling module meta-predicates whose meta-predicate templates are ambiguous and cannot be processed by the Logtalk compiler (note, however, that it's often possible to specify an overriding meta-predicate directive within the object or category making the call as explained above).

1.8 Inheritance

The inheritance mechanisms found on object-oriented programming languages allow us the specialization of previously defined objects, avoiding the unnecessary repetition of code. In the context of logic programming, we can interpret inheritance as a form of theory extension: an object will virtually contain, besides its own predicates, all the predicates inherited from other objects that are not redefined by itself.

Logtalk uses a depth-first lookup procedure for finding predicate declarations and predicate definitions, as explained below. The lookup procedures locate the entities holding the predicate declaration and the predicate definition using, respectively, the predicate indicator and the predicate template (constructed from the predicate indicator).

The *alias/2* predicate directive may be used to defining alternative names for inherited predicates, for solving inheritance conflicts, and for giving access to all inherited definitions.

1.8.1 Protocol inheritance

Protocol inheritance refers to the inheritance of predicate declarations (scope directives). These can be contained in objects, in protocols, or in categories. Logtalk supports single and multi-inheritance of protocols: an object or a category may implement several protocols and a protocol may extend several protocols.

Search order for prototype hierarchies

The search order for predicate declarations is first the object, second the implemented protocols (and the protocols that these may extend), third the imported categories (and the protocols that they may implement), and last the objects that the object extends. This search is performed in a depth-first way. When an object inherits two different declarations for the same predicate, by default, only the first one will be considered.

Search order for class hierarchies

The search order for predicate declarations starts in the object classes. Following the classes declaration order, the search starts in the classes implemented protocols (and the protocols that these may extend), third the classes imported categories (and the protocols that they may implement), and last the superclasses of the object classes. This search is performed in a depth-first way. If the object inherits two different declarations for the same predicate, by default only the first one will be considered.

1.8.2 Implementation inheritance

Implementation inheritance refers to the inheritance of predicate definitions. These can be contained in objects or in categories. Logtalk supports multi-inheritance of implementation: an object may import several categories or extend, specialize, or instantiate several objects.

Search order for prototype hierarchies

The search order for predicate definitions is similar to the search for predicate declarations except that implemented protocols are ignored (as they can only contain predicate directives).

Search order for class hierarchies

The search order for predicate definitions is similar to the search for predicate declarations except that implemented protocols are ignored (as they can only contain predicate directives) and that the search starts at the instance itself (that received the message) before proceeding, if no predicate definition is found there, to the instance classes and then to the class superclasses.

Inheritance versus predicate redefinition

When we define a predicate that is already inherited from other object, the inherited definitions are hidden by the new definitions. This is called overriding inheritance: a local definition overrides any inherited ones. For example, assume that we have the following two objects:

```logtalk
:- object(root).

    :- public(bar/1).
    :- public(foo/1).

    bar(root).

    foo(root).
:- end_object.


:- object(descendant,
    extends(root)).

    foo(descendant).
:- end_object.
```

After compiling and loading these objects, we can check the overriding behavior by trying the following queries:

```logtalk
| ?- root::(bar(Bar), foo(Foo)).

Bar = root
Foo = root
yes


| ?- descendant::(bar(Bar), foo(Foo)).

Bar = root
Foo = descendant
yes
```

However, we can explicitly program other behaviors. Let us see a few examples.

Specialization inheritance

Specialization of inherited definitions: the new definition uses the inherited definitions, adding to this new code. This is accomplished by calling the ^^/1 operator in the new definition.

```logtalk
:- object(root).

    :- public(init/0).

    init :-
        write('root init'), nl.

:- end_object.


:- object(descendant,
    extends(root)).

    init :-
        write('descendant init'), nl,
```

(continues on next page)

```
        ^^init.

:- end_object.


| ?- descendant::init.

descendant init
root init

yes
```

Union inheritance

Union of the new with the inherited definitions: all the definitions are taken into account, the calling order being
defined by the inheritance mechanisms. This can be accomplished by writing a clause that just calls, using the *^^/1*
operator, the inherited definitions. The relative position of this clause among the other definition clauses sets the
calling order for the local and inherited definitions.

```
:- object(root).

    :- public(foo/1).

    foo(1).
    foo(2).

:- end_object.


:- object(descendant,
    extends(root)).

    foo(3).
    foo(Foo) :-
        ^^foo(Foo).

:- end_object.


| ?- descendant::foo(Foo).

Foo = 3 ;
Foo = 1 ;
Foo = 2 ;
no
```

Selective inheritance

Hiding of some of the inherited definitions, or differential inheritance: this form of inheritance is normally used in
the representation of exceptions to generic definitions. Here we will need to use the *^^/1* operator to test and possibly
reject some of the inherited definitions.

```
:- object(bird).

    :- public(mode/1).

    mode(walks).
    mode(flies).

:- end_object.


:- object(penguin,
    extends(bird)).

    mode(swims).
    mode(Mode) :-
        ^^mode(Mode),
        Mode \= flies.

:- end_object.


| ?- penguin::mode(Mode).

Mode = swims ;
Mode = walks ;

no
```

1.8.3 Public, protected, and private inheritance

To make all *public predicates* declared via implemented protocols, imported categories, or ancestor objects *protected predicates* or to make all public and protected predicates *private predicates* we prefix the entity's name with the corresponding keyword. For instance:

```
:- object(Object,
    implements(private::Protocol)).

    % all the Protocol public and protected
    % predicates become private predicates
    % for the Object clients

    ...

:- end_object.
```

or:

```
:- object(Class,
    specializes(protected::Superclass)).

    % all the Superclass public predicates become
    % protected predicates for the Class clients

    ...

:- end_object.
```

Omitting the scope keyword is equivalent to using the public scope keyword. For example:

```
:- object(Object,
    imports(public::Category)).
    ...
:- end_object.
```

This is the same as:

```
:- object(Object,
    imports(Category)).
    ...
:- end_object.
```

This way we ensure backward compatibility with older Logtalk versions and a simplified syntax when protected or private inheritance are not used.

1.8.4 Composition versus multiple inheritance

It is not possible to discuss inheritance mechanisms without referring to the long and probably endless debate on single versus multiple inheritance. The single inheritance mechanism can be implemented efficiently but it imposes several limitations on reusing, even if the multiple characteristics we intend to inherit are orthogonal. On the other hand, the multiple inheritance mechanisms are attractive in their apparent capability of modeling complex situations. However, they include a potential for conflict between inherited definitions whose variety does not allow a single and satisfactory solution for all the cases.

Until now, no solution that we might consider satisfactory for all the problems presented by the multiple inheritance mechanisms has been found. From the simplicity of some extensions that use the Prolog search strategy like [McCabe92] or [Moss94] and to the sophisticated algorithms of CLOS [Bobrow_et_al_88], there is no adequate solution for all the situations. Besides, the use of multiple inheritance carries some complex problems in the domain of software engineering, particularly in the reuse and maintenance of the applications. All these problems are substantially reduced if we preferably use in our software development composition mechanisms instead of specialization mechanisms [Taenzer89]. Multiple inheritance is best used as an analysis and project abstraction, rather than as an implementation technique [Shan_et_al_93]. Logtalk provides first-class support for software composition using *Categories*.

Nevertheless, Logtalk supports multi-inheritance by enabling an object to extend, instantiate, or specialize more than one object. The current Logtalk release provides a predicate directive, *alias/2*, which may be used to solve some multi-inheritance conflicts. Lastly, it should be noted that the multi-inheritance support does not affect performance when we use single-inheritance.

1.9 Event-driven programming

The addition of event-driven programming capacities to the Logtalk language [Moura94] is based on a simple but powerful idea:

> The computations must result, not only from message sending, but also from the **observation** of message sending.

The need to associate computations to the occurrence of events was very early recognized in knowledge representation languages, programming languages [Stefik_et_al_86], [Moon86], operative systems [Tanenbaum87], and graphical user interfaces.

With the integration between object-oriented and event-driven programming, we intend to achieve the following goals:

- Minimize the coupling between objects. An object should only contain what is intrinsic to it. If an object observes another object, that means that it should depend only on the (public) protocol of the object observed, and not on the implementation of that same protocol.

- Provide a mechanism for building reflexive systems in Logtalk based on the dynamic behavior of objects in complement to the reflective information of the object's contents and relations.

- Provide a mechanism for easily defining method pre- and post-conditions that can be toggled using the `events` compiler flag. The pre- and post-conditions may be defined in the same object containing the methods or distributed between several objects acting as method monitors.

1.9.1 Definitions

The words *event* and *monitor* have multiple meanings in computer science, so, to avoid misunderstandings, it is advisable that we start by defining them in the Logtalk context.

Event

In an object-oriented system, all computations start through message sending. It thus becomes quite natural to declare that the only event that can occur in this kind of system is precisely the sending of a message. An event can thus be represented by the ordered tuple `(Object, Message, Sender)`.

If we consider message processing an indivisible activity, we can interpret the sending of a message and the return of the control to the object that has sent the message as two distinct events. This distinction allows us to have a more precise control over a system dynamics. In Logtalk, these two types of events have been named `before` and `after`, respectively for message sending and returning. Therefore, we end up by representing an event by the ordered tuple `(Event, Object, Message, Sender)`.

The implementation of the event notion in Logtalk enjoys the following properties:

Independence between the two types of events We can choose to watch only one event type or to process each one of the events associated to a message sending in an independent way.

All events are automatically generated by the message sending mechanism The task of generating events is transparently accomplished by the message sending mechanism. The user just defines which are the events in which he is interested in.

The events watched at any moment can be dynamically changed during program execution The notion of event allows the user not only to have the possibility of observing, but also of controlling and modifying an application behavior, namely by dynamically changing the observed events during program execution. It is our goal to provide the user with the possibility of modeling the largest number of situations.

Monitor

Complementary to the notion of event is the notion of monitor. A monitor is an object that is automatically notified by the message sending mechanisms whenever certain events occur. A monitor should naturally define the actions to be carried out whenever a monitored event occurs.

The implementation of the monitor notion in Logtalk enjoys the following properties:

Any object can act as a monitor The monitor status is a role that any object can perform during its existence. The minimum protocol necessary is declared in the built-in protocol `monitoring`. Strictly speaking, the reference to this protocol is only needed when specializing event handlers. Nevertheless, it is considered good programming practice to always refer the protocol when defining event handlers.

Unlimited number of monitors for each event Several monitors can observe the same event because of distinct reasons. Therefore, the number of monitors per event is bounded only by the available computing resources.

The monitor status of an object can be dynamically changed in runtime This property does not imply that an object
must be dynamic to act as a monitor (the monitor status of an object is not stored in the object).

The execution of actions, defined in a monitor, associated to each event, never affects the term that denotes the message involved
In other words, if the message contains uninstantiated variables, these are not affected by the acting of monitors
associated to the event.

1.9.2 Event generation

For each message that is sent (using the *::/2* message sending mechanism) the runtime system automatically generates
two events. The first — `before event` — is generated when the message is sent. The second — `after event`
— is generated after the message has successfully been executed.

1.9.3 Communicating events to monitors

Whenever a spied event occurs, the message sending mechanisms call the corresponding event handlers directly for all
registered monitors. These calls are made bypassing the message sending primitives in order to avoid potential endless
loops. The event handlers consist in user definitions for the public predicates declared in the `monitoring` built-in
protocol (one for each event kind; see below for more details).

1.9.4 Performance concerns

Ideally, the existence of monitored messages should not affect the processing of the remaining messages. On the other
hand, for each message that has been sent, the system must verify if its respective event is monitored. Whenever
possible, this verification should be performed in constant time and independently of the number of monitored events.
The events representation takes advantage of the first argument indexing performed by most Prolog compilers, which
ensure — in the general case — access in constant time.

Event-support can be turned off on a per-object (or per-category) basis using the *events/1* compiler flag. With event-
support turned off, Logtalk uses optimized code for processing message sending calls that skips the checking of
monitored events, resulting in a small but measurable performance improvement.

1.9.5 Monitor semantics

The established semantics for monitors actions consists on considering its success as a necessary condition so that a
message can succeed:

- All actions associated to events of type `before` must succeed, so that the message processing can start.

- All actions associated to events of type `after` also have to succeed so that the message itself succeeds. The
  failure of any action associated to an event of type `after` forces backtracking over the message execution (the
  failure of a monitor never causes backtracking over the preceding monitor actions).

Note that this is the most general choice. If we wish a transparent presence of monitors in a message processing, we
just have to define the monitor actions in such a way that they never fail (which is very simple to accomplish).

1.9.6 Activation order of monitors

Ideally, whenever there are several monitors defined for the same event, the calling order should not interfere with the
result. However, this is not always possible. In the case of an event of type `before`, the failure of a monitor prevents
a message from being sent and prevents the execution of the remaining monitors. In case of an event of type `after`,

a monitor failure will force backtracking over message execution. Different orders of monitor activation can therefore lead to different results if the monitor actions imply object modifications unrecoverable in case of backtracking. Therefore, the order for monitor activation should be assumed as arbitrary. In effect, to assume or to try to impose a specific sequence requires a global knowledge of an application dynamics, which is not always possible. Furthermore, that knowledge can reveal itself as incorrect if there is any changing in the execution conditions. Note that, given the independence between monitors, it does not make sense that a failure forces backtracking over the actions previously executed.

1.9.7 Event handling

Logtalk provides three built-in predicates for event handling. These predicates enable you to find what events are defined, to define new events and to abolish events when they are no longer needed. If you plan to use events extensively in your application, then you should probably define a set of objects that use the built-in predicates described below to implement more sophisticated and high-level behavior.

Defining new events

New events can be defined using the Logtalk built-in predicate *define_events/5*:

```
| ?- define_events(Event, Object, Message, Sender, Monitor).
```

Note that if any of the `Event`, `Object`, `Message`, and `Sender` arguments is a free variable or contains free variables, this call will define a **set** of matching events.

Abolishing defined events

Events that are no longer needed may be abolished using the *abolish_events/5* built-in predicate:

```
| ?- abolish_events(Event, Object, Message, Sender, Monitor).
```

If called with free variables, this goal will remove all matching events.

Finding defined events

The events that are currently defined can be retrieved using the Logtalk built-in predicate *current_event/5*:

```
| ?- current_event(Event, Object, Message, Sender, Monitor).
```

Note that this predicate will return **sets** of matching events if some of the returned arguments are free variables or contain free variables.

Defining event handlers

The `monitoring` built-in protocol declares two public predicates, *before/3* and *after/3*, that are automatically called to handle `before` and `after` events. Any object that plays the role of monitor must define one or both of these event handler methods:

```
before(Object, Message, Sender) :-
    ... .

after(Object, Message, Sender) :-
    ... .
```

The arguments in both methods are instantiated by the message sending mechanisms when a monitored event occurs. For example, assume that we want to define a monitor called `tracer` that will track any message sent to an object by printing a describing text to the standard output. Its definition could be something like:

```logtalk
:- object(tracer,
    % built-in protocol for event handler methods
    implements(monitoring)).

    before(Object, Message, Sender) :-
        write('call: '), writeq(Object),
        write(' <-- '), writeq(Message),
        write(' from '), writeq(Sender), nl.

    after(Object, Message, Sender) :-
        write('exit: '), writeq(Object),
        write(' <-- '), writeq(Message),
        write(' from '), writeq(Sender), nl.

:- end_object.
```

Assume that we also have the following object:

```logtalk
:- object(any).

    :- public(bar/1) .
    :- public(foo/1) .

    bar(bar).

    foo(foo).

:- end_object.
```

After compiling and loading both objects (note that the object `any` must be compiled with the flag `events(allow)`), we can start tracing every message sent to any object by calling the `define_events/5` built-in predicate:

```
| ?- define_events(_, _, _, _, tracer).

yes
```

From now on, every message sent to any object will be traced to the standard output stream:

```
| ?- any::bar(X).

call: any <-- bar(X) from user
exit: any <-- bar(bar) from user
X = bar

yes
```

To stop tracing, we can use the `abolish_events/5` built-in predicate:

```
| ?- abolish_events(_, _, _, _, tracer).

yes
```

The `monitoring` protocol declares the event handlers as public predicates. If necessary, protected or private implementation of the protocol may be used in order to change the scope of the event handler predicates. Note that the

message sending processing mechanisms are able to call the event handlers irrespective of their scope. Nevertheless, the scope of the event handlers may be restricted in order to prevent other objects from calling them.

1.10 Multi-threading programming

Logtalk provides **experimental** support for multi-threading programming on selected Prolog compilers. Logtalk makes use of the low-level Prolog built-in predicates that implement message queues and interface with POSIX threads and mutexes (or a suitable emulation), providing a small set of high-level predicates and directives that allows programmers to easily take advantage of modern multi-processor and multi-core computers without worrying about the details of creating, synchronizing, or communicating with threads. Logtalk multi-threading programming integrates with object-oriented programming providing a threaded engines API, enabling objects and categories to prove goals concurrently, and supporting synchronous and asynchronous messages.

1.10.1 Enabling multi-threading support

Multi-threading support may be disabled by default. It can be enabled on the Prolog adapter files of supported compilers by setting the read-only compiler flag `threads` to `supported`.

1.10.2 Enabling objects to make multi-threading calls

The *threaded/0* object directive is used to enable an object to make multi-threading calls:

```
:- threaded.
```

This directive results in the automatic creation and set up of an object message queue when the object is loaded or created at runtime. Object message queues are used for exchanging thread notifications and for storing concurrent goal solutions and replies to the *multi-threading calls* made within the object. The message queue for the pseudo-object `user` is automatically created at Logtalk startup (provided that multi-threading programming is supported and enabled for the chosen Prolog compiler).

1.10.3 Multi-threading built-in predicates

Logtalk provides a small set of built-in predicates for multi-threading programming. For simple tasks where you simply want to prove a set of goals, each one in its own thread, Logtalk provides a *threaded/1* built-in predicate. The remaining predicates allow for fine-grained control, including postponing retrieving of thread goal results at a later time, supporting non-deterministic thread goals, and making *one-way* asynchronous calls. Together, these predicates provide high-level support for multi-threading programming, covering most common use cases.

Proving goals concurrently using threads

A set of goals may be proved concurrently by calling the Logtalk built-in predicate *threaded/1*. Each goal in the set runs in its own thread.

When the `threaded/1` predicate argument is a *conjunction* of goals, the predicate call is akin to *and-parallelism*. For example, assume that we want to find all the prime numbers in a given interval, `[N, M]`. We can split the interval in two parts and then span two threads to compute the prime numbers in each sub-interval:

```
prime_numbers(N, M, Primes) :-
    M > N,
    N1 is N + (M - N) // 2,
    N2 is N1 + 1,
    threaded((
        prime_numbers(N2, M, [], Acc),
        prime_numbers(N, N1, Acc, Primes)
    )).

prime_numbers(N, M, Acc, Primes) :-
    ...
```

The `threaded/1` call terminates when the two implicit threads terminate. In a computer with two or more processors (or with a processor with two or more cores) the code above can be expected to provide better computation times when compared with single-threaded code for sufficiently large intervals.

When the `threaded/1` predicate argument is a *disjunction* of goals, the predicate call is akin to *or-parallelism*, here reinterpreted as a set of goals *competing* to find a solution. For example, consider the different methods that we can use to find the roots of real functions. Depending on the function, some methods will faster than others. Some methods will converge into the solution while others may diverge and never find it. We can try all the methods simultaneously by writing:

```
find_root(Function, A, B, Error, Zero, Algorithm) :-
    threaded((
        (bisection::find_root(Function, A, B, Error, Zero), Algorithm = bisection)
    ;   (newton::find_root(Function, A, B, Error, Zero), Algorithm = newton)
    ;   (muller::find_root(Function, A, B, Error, Zero), Algorithm = muller)
    )).
```

The above `threaded/1` goal succeeds when one of the implicit threads succeeds in finding the function root, leading to the termination of all the remaining competing threads.

The `threaded/1` built-in predicate is most useful for lengthy, independent deterministic computations where the computational costs of each goal outweigh the overhead of the implicit thread creation and management.

Proving goals asynchronously using threads

A goal may be proved asynchronously using a new thread by calling the Logtalk built-in predicate *threaded_call/1-2*. Calls to this predicate are always true and return immediately (assuming a callable argument). The term representing the goal is copied, not shared with the thread. The thread computes the first solution to the goal, posts it to the message queue of the object from where the `threaded_call/1` predicate was called, and suspends waiting for either a request for an alternative solution or for the program to commit to the current solution.

The results of proving a goal asynchronously in a new thread may be later retrieved by calling the Logtalk built-in predicate *threaded_exit/1-2* within the same object where the call to the `threaded_call/1` predicate was made. The `threaded_exit/1` calls suspend execution until the results of the `threaded_call/1` calls are sent back to the object message queue.

The `threaded_exit/1` predicate allow us to retrieve alternative solutions through backtracking (if you want to commit to the first solution, you may use the *threaded_once/1-2* predicate instead of the `threaded_call/1` predicate). For example, assuming a `lists` object implementing the usual `member/2` predicate, we could write:

```
| ?- threaded_call(lists::member(X, [1,2,3])).

X = _G189
yes
```

```
| ?- threaded_exit(lists::member(X, [1,2,3])).

X = 1 ;
X = 2 ;
X = 3 ;
no
```

In this case, the `threaded_call/1` and the `threaded_exit/1` calls are made within the pseudo-object *user*. The implicit thread running the `lists::member/2` goal suspends itself after providing a solution, waiting for a request to an alternative solution; the thread is automatically terminated when the runtime engine detects that back-tracking to the `threaded_exit/1` call is no longer possible.

Calls to the `threaded_exit/1` predicate block the caller until the object message queue receives the reply to the asynchronous call. The predicate *threaded_peek/1-2* may be used to check if a reply is already available without removing it from the thread queue. The `threaded_peek/1` predicate call succeeds or fails immediately without blocking the caller. However, keep in mind that repeated use of this predicate is equivalent to polling a message queue, which may severely hurt performance.

Be careful when using the `threaded_exit/1` predicate inside failure-driven loops. When all the solutions have been found (and the thread generating them is therefore terminated), re-calling the predicate will generate an exception. Note that failing instead of throwing an exception is not an acceptable solution as it could be misinterpreted as a failure of the `threaded_exit/1` argument.

The example on the previous section with prime numbers could be rewritten using the `threaded_call/1` and `threaded_exit/1` predicates:

```
prime_numbers(N, M, Primes) :-
    M > N,
    N1 is N + (M - N) // 2,
    N2 is N1 + 1,
    threaded_call(prime_numbers(N2, M, [], Acc)),
    threaded_call(prime_numbers(N, N1, Acc, Primes)),
    threaded_exit(prime_numbers(N2, M, [], Acc)),
    threaded_exit(prime_numbers(N, N1, Acc, Primes)).

prime_numbers(N, M, Acc, Primes) :-
    ...
```

When using asynchronous calls, the link between a `threaded_exit/1` call and the corresponding `threaded_call/1` call is established using unification. If there are multiple `threaded_call/1` calls for a matching `threaded_exit/1` call, the connection can potentially be established with any of them. Nevertheless, you can easily use a tag the calls by using the extended *threaded_call/1-2* and *threaded_exit/1-2* built-in predicates. For example:

```
?- threaded_call(member(X, [1,2,3]), Tag).

Tag = 1
yes

?- threaded_call(member(X, [1,2,3]), Tag).

Tag = 2
yes

?- threaded_exit(member(X, [1,2,3]), 2).
```

```
X = 1 ;
X = 2 ;
X = 3
yes
```

When using these predicates, the tags shall be considered as an opaque term; users shall not rely on its type.

1.10.4 One-way asynchronous calls

Sometimes we want to prove a goal in a new thread without caring about the results. This may be accomplished by using the built-in predicate *threaded_ignore/1*. For example, assume that we are developing a multi-agent application where an agent may send an "happy birthday" message to another agent. We could write:

```
..., threaded_ignore(agent::happy_birthday), ...
```

The call succeeds with no reply of the goal success, failure, or even exception ever being sent back to the object making the call. Note that this predicate implicitly performs a deterministic call of its argument.

1.10.5 Asynchronous calls and synchronized predicates

Proving a goal asynchronously using a new thread may lead to problems when the goal results in side effects such as input/output operations or modifications to an *object database*. For example, if a new thread is started with the same goal before the first one finished its job, we may end up with mixed output, a corrupted database, or unexpected goal failures. In order to solve this problem, predicates (and grammar rule non-terminals) with side effects can be declared as *synchronized* by using the *synchronized/1* predicate directive. Proving a query to a synchronized predicate (or synchronized non-terminal) is internally protected by a mutex, thus allowing for easy thread synchronization. For example:

```
:- synchronized(db_update/1).    % ensure thread synchronization

db_update(Update) :-
    % predicate with side-effects
    ...
```

A second example: assume an object defining two predicates for writing, respectively, even and odd numbers in a given interval to the standard output. Given a large interval, a goal such as:

```
| ?- threaded_call(obj::odd_numbers(1,100)),
     threaded_call(obj::even_numbers(1,100)).

1 3 2 4 6 8 5 7 10 ...
...
```

will most likely result in a mixed up output. By declaring the odd_numbers/2 and even_numbers/2 predicates synchronized:

```
:- synchronized([
    odd_numbers/2,
    even_numbers/2]).
```

one goal will only start after the other one finished:

```
| ?- threaded_ignore(obj::odd_numbers(1,99)),
    threaded_ignore(obj::even_numbers(1,99)).

1 3 5 7 9 11 ...
...
2 4 6 8 10 12 ...
...
```

Note that, in a more realistic scenario, the two `threaded_ignore/1` calls would be made concurrently from different objects. Using the same synchronized directive for a set of predicates imply that they all use the same mutex, as required for this example.

As each Logtalk entity is independently compiled, this directive must be included in every object or category that contains a definition for the described predicate, even if the predicate declaration is inherited from another entity, in order to ensure proper compilation. Note that a synchronized predicate cannot be declared dynamic. To ensure atomic updates of a dynamic predicate, declare as synchronized the predicate performing the update.

Synchronized predicates may be used as wrappers to messages sent to objects that are not multi-threading aware. For example, assume a `random` object defining a `random/1` predicate that generates random numbers, using side effects on its implementation (e.g. for storing the generator seed). We can specify and define e.g. a `sync_random/1` predicate as follows:

```
:- synchronized(sync_random/1).

sync_random(Random) :-
    random::random(Random).
```

and then always use the `sync_random/1` predicate instead of the predicate `random/1` from multi-threaded code.

The synchronization entity and predicate directives may be used when defining objects that may be reused in both single-threaded and multi-threaded Logtalk applications. The directives are simply ignored (i.e. the synchronized predicates are interpreted as normal predicates) when the objects are used in a single-threaded application.

1.10.6 Synchronizing threads through notifications

Declaring a set of predicates as synchronized can only ensure that they are not executed at the same time by different threads. Sometimes we need to suspend a thread not on a synchronization lock but on some condition that must hold true for a thread goal to proceed. I.e. we want a thread goal to be suspended until a condition becomes true instead of simply failing. The built-in predicate *threaded_wait/1* allows us to suspend a predicate execution (running in its own thread) until a notification is received. Notifications are posted using the built-in predicate *threaded_notify/1*. A notification is a Prolog term that a programmer chooses to represent some condition becoming true. Any Prolog term can be used as a notification argument for these predicates. Related calls to the `threaded_wait/1` and `threaded_notify/1` must be made within the same object, *this*, as the object message queue is used internally for posting and retrieving notifications.

Each notification posted by a call to the `threaded_notify/1` predicate is consumed by a single `threaded_wait/1` predicate call (i.e. these predicates implement a peer-to-peer mechanism). Care should be taken to avoid deadlocks when two (or more) threads both wait and post notifications to each other.

1.10.7 Engines

Threaded *engines* provide an alternative to the multi-threading predicates described in the previous sections. An engine is a computing thread whose solutions can be lazily computed and retrieved. In addition, an engine also supports a term queue that allows passing arbitrary terms to the engine.

An engine is created by calling the *threaded_engine_create/3* built-in predicates. For example:

```
| ?- threaded_engine_create(X, member(X, [1,2,3]), worker).
yes
```

The first argument is an *answer template* to be used for retrieving solution bindings. The user can name the engine, as in this example where the atom `worker` is used, or have the runtime generate a name, which should be treated as an opaque term.

Engines are scoped by the object within which the `threaded_engine_create/3` call takes place. Thus, different objects can create engines with the same names with no conflicts. Moreover, engines share the visible predicates of the object creating them.

The engine computes the first solution of its goal argument and suspends waiting for it to be retrieved. Solutions can be retrieved one at a time using the *threaded_engine_next/2* built-in predicate:

```
| ?- threaded_engine_next(worker, X).
X = 1
yes
```

The call blocks until a solution is available and fails if there are no solutions left. After returning a solution, this predicate signals the engine to start computing the next one. Note that this predicate is deterministic. In contrast with the `threaded_exit/1-2` built-in predicates, retrieving the next solution requires calling the predicate again instead of by backtracking into its call.

There is also an alternative reified version of the predicate, *threaded_engine_next_reified/2*, which returns the `the(Answer)`, `no`, and `exception(Error)` terms as answers.

Engines must be explicitly terminated using the *threaded_engine_destroy/1* built-in predicate:

```
| ?- threaded_engine_destroy(worker).
yes
```

A common usage pattern for engines is to define a recursive predicate that uses the engine term queue to retrieve a task to be performed. For example, assume we define the following predicate:

```
loop :-
    threaded_engine_fetch(Task),
    handle(Task),
    loop.
```

The *threaded_engine_fetch/1* built-in predicate fetches a task for the engine term queue. The engine clients would use the *threaded_engine_post/2* built-in predicate to post tasks into the engine term queue. The engine would be created using the call:

```
| ?- threaded_engine_create(none, loop, worker).

yes
```

The `handle/1` predicate, after performing a task, can use the *threaded_engine_yield/1* built-in predicate to make the task results available for consumption using the `threaded_engine_next/2` built-in predicate. Blocking semantics are used by these two predicates: the `threaded_engine_yield/1` predicate blocks until the returned solution is consumed while the `threaded_engine_next/2` predicate blocks until a solution becomes available.

1.10.8 Multi-threading performance

The performance of multi-threading applications is highly dependent on the back-end Prolog compiler, on the operating-system, and on the use of *dynamic binding* and dynamic predicates. All compatible back-end Prolog com-

pilers that support multi-threading features make use of POSIX threads or *pthreads*. The performance of the underlying pthreads implementation can exhibit significant differences between operating systems. An important point is synchronized access to dynamic predicates. As different threads may try to simultaneously access and update dynamic predicates, these operations must be protected by a lock, usually implemented using a mutex. Poor mutex lock operating-system performance, combined with a large number of collisions by several threads trying to acquire the same lock, often result in severe performance penalties. Thus, whenever possible, avoid using dynamic predicates and dynamic binding.

1.11 Error handling

All error handling is done in Logtalk by using the ISO defined `catch/3` and `throw/1` predicates [ISO95]. Errors thrown by Logtalk have the following format:

```
error(Error, logtalk(Goal, ExecutionContext))
```

In this exception term, `Goal` is the goal that triggered the error `Error` and `ExecutionContext` is the context in which `Goal` is called. For example:

```
error(
    permission_error(modify,private_predicate,p),
    logtalk(foo::abolish(p/0), _)
)
```

Note, however, that `Goal` and `ExecutionContext` can be unbound or only partially instantiated when the corresponding information is not available (e.g. due to compiler optimizations that throw away the necessary error context information). The `ExecutionContext` argument is an opaque term that can be decoded using the `logtalk::execution_context/7` predicate.

1.11.1 Compiler warnings and errors

The current Logtalk compiler uses the `read_term/3` ISO Prolog defined built-in predicate to read and compile a Logtalk source file. One consequence of this is that invalid Prolog terms or syntax errors may abort the compilation process with limited information given to the user (due to the inherent limitations of the `read_term/3` predicate).

If all the terms in a source file are valid, then there is a set of errors or potential errors, described below, that the compiler will try to detect and report, depending on the used compiler flags (see the *Compiler flags* section of this manual for details).

Unknown entities

The Logtalk compiler warns about any referenced entity that is not currently loaded. The warning may reveal a misspell entity name or just an entity that it will be loaded later. Out-of-oder loading should be avoided when possible as it prevents some code optimizations such as *static binding* of messages to methods.

Singleton variables

Singleton variables in a clause are often misspell variables and, as such, one of the most common errors when programming in Prolog. When the backend Prolog compiler complies with the Prolog ISO standard or at least supports the ISO predicate `read_term/3` called with the option `singletons(S)`, the Logtalk compiler warns about any singleton variable found while compiling a source file.

Redefinition of Prolog built-in predicates

The Logtalk compiler will warn us of any redefinition of a Prolog built-in predicate inside an object or category. Sometimes the redefinition is intended. In other cases, the user may not be aware that the used backend Prolog compiler may already provide the predicate as a built-in or may want to ensure code portability among several Prolog compilers with different sets of built-in predicates.

Redefinition of Logtalk built-in predicates

Similar to the redefinition of Prolog built-in predicates, the Logtalk compiler will warn us if we try to redefine a Logtalk built-in. The redefinition will probably be an error in almost all (if not all) cases.

Redefinition of Logtalk built-in methods

An error will be thrown if we attempt to redefine a Logtalk built-in method inside an entity. The default behavior is to report the error and abort the compilation of the offending entity.

Misspell calls of local predicates

A warning will be reported if Logtalk finds (in the body of a predicate definition) a call to a local predicate that is not defined, built-in (either in Prolog or in Logtalk) or declared dynamic. In most cases these calls are simple misspell errors.

Portability warnings

A warning will be reported if a predicate clause contains a call to a non-ISO specified built-in predicate or arithmetic function, Portability warnings are also reported for non-standard flags or flag values. These warnings often cannot be avoided due to the limited scope of the ISO Prolog standard.

Deprecated elements

A warning will be reported if a deprecated directive or control construct is used.

Missing directives

A warning will be reported for any missing dynamic, discontiguous, meta-predicate, and public predicate directive.

Duplicated directives

A warning will be reported for any duplicated scope, multifile, dynamic, discontiguous, meta-predicate, and meta-non-terminal directives. Note that conflicting directives for the same predicate are handled as errors, not as duplicated directive warnings.

Goals that are always true or false

A warning will be reported for any goal that is always true or false. This is usually caused by typos in the code. For example, writing `X == y` instead of `X == Y`.

Trivial fails

A warning will be reported for any call to a local static predicate with no matching clause.

Suspicious calls

A warning will be reported for calls that are syntactically correct but most likely a semantic error. An example is `::/1` calls in clauses that apparently are meant to implement recursive predicate definitions where the user intention is to call the local predicate.

Lambda variables

A warning will be reported for *lambda expressions* with unclassified variables (not listed as either *lambda free* or *lambda parameter* variables) or where variables play a dual role (as both lambda free and lambda parameter variables).

Redefinition of predicates declared in `uses/2` and `use_module/2` directives

A error will be reported for any attempt to define locally a predicate that is already listed in an *uses/2* or in an *use_module/2* directive.

Other warnings and errors

The Logtalk compiler will throw an error if it finds a predicate clause or a directive that cannot be parsed. The default behavior is to report the error and abort the compilation of the offending entity.

1.11.2 Runtime errors

This section briefly describes runtime errors that result from misuse of Logtalk built-in predicates, built-in methods or from message sending. For a complete and detailed description of runtime errors please consult the Reference Manual.

Logtalk built-in predicates

Most Logtalk built-in predicates checks the type and mode of the calling arguments, throwing an exception in case of misuse.

Logtalk built-in methods

Most Logtalk built-in method checks the type and mode of the calling arguments, throwing an exception in case of misuse.

Message sending

The message sending mechanisms always check if the receiver of a message is a defined object and if the message corresponds to a declared predicate within the scope of the sender. The built-in protocol `forwarding` declares a predicate, *forward/1*, which is automatically called (if defined) by the runtime for any message that the receiving object does not understand. The usual definition for this error handler is to delegate or forward the message to another object that might be able to answer it:

```
forward(Message) :-
    % forward the message while preserving the sender
    [Object::Message].
```

If preserving the original sender is not required, this definition can be simplified to:

```
forward(Message) :-
    Object::Message.
```

More sophisticated definitions are, of course, possible.

1.12 Documenting applications

By setting the compiler flag source_data, Logtalk saves all relevant documenting information collected when compiling a source file. The provided lgtdoc tool can access this information by using Logtalk's reflection support and generate a documentation file for each compiled entity (object, protocol, or category) in XML format. Contents of the XML file include the entity name, type, and compilation mode (static or dynamic), the entity relations with other entities, and a description of any declared predicates (name, compilation mode, scope, . . .).

The XML documentation files can be enriched with arbitrary user-defined information, either about an entity or about its predicates, by using the two directives described below.

1.12.1 Documenting directives

Logtalk supports two documentation directives for providing arbitrary user-defined information about an entity or a predicate. These two directives complement other Logtalk directives that also provide important documentation information like *mode/2*.

Entity directives

Arbitrary user-defined entity information can be represented using the *info/1* directive:

```
:- info([
    Key1 is Value1,
    Key2 is Value2,
    ...
]).
```

In this pattern, keys should be atoms and values should be ground terms. The following keys are predefined and may be processed specially by Logtalk:

comment Comment describing entity purpose (an atom).

author Entity author(s) (an atom or a compound term {entity} where entity is the name of an XML entity defined in the custom.ent file).

version Version number (a number).

date Date of last modification (formatted as Year/Month/Day).

parameters Parameter names and descriptions for parametric entities (a list of key-values where both keys and values are atoms).

parnames Parameter names for parametric entities (a list of atoms; a simpler version of the previous key, used when parameter descriptions are deemed unnecessary).

copyright Copyright notice for the entity source code (an atom or a compound term `{entity}` where `entity` is the name of an XML entity defined in the `custom.ent` file).

license License terms for the entity source code; usually, just the license name (an atom or a compound term `{entity}` where `entity` is the name of an XML entity defined in the `custom.ent` file).

remarks List of general remarks about the entity using the format *Topic – Text*. Both the topic and the text must be atoms.

see_also List of related entities.

For example:

```
:- info([
    version is 2.1,
    author is 'Paulo Moura',
    date is 2000/4/20,
    comment is 'Building representation.',
    diagram is 'UML Class Diagram #312'
]).
```

Use only the keywords that make sense for your application and remember that you are free to invent your own keywords. All key-value pairs can be retrieved programmatically using the reflection API and are visible to `lgtdoc` tool.

Predicate directives

Arbitrary user-defined predicate information can be represented using the *info/2* directive:

```
:- info(Name/Arity, [
    Key1 is Value1,
    Key2 is Value2,
    ...
]).
```

The first argument can also a grammar rule non-terminal indicator, `Name//Arity`. Keys should be atoms and values should be bound terms. The following keys are predefined and may be processed specially by Logtalk:

comment Comment describing predicate purpose (an atom).

arguments Names and descriptions of predicate arguments for pretty print output (a list of key-values where both keys and values are atoms).

argnames Names of predicate arguments for pretty print output (a list of atoms; a simpler version of the previous key, used when argument descriptions are deemed unnecessary).

allocation Objects where we should define the predicate. Some possible values are `container`, `descendants`, `instances`, `classes`, `subclasses`, and `any`.

redefinition Describes if predicate is expected to be redefined and, if so, in what way. Some possible values are `never`, `free`, `specialize`, `call_super_first`, `call_super_last`.

exceptions List of possible exceptions throw by the predicate using the format *Description – Exception term*. The description must be an atom. The exception term must be a non-variable term.

examples List of typical predicate call examples using the format *Description – Predicate call – Variable bindings*. The description must be an atom. The predicate call term must be a non-variable term. The variable bindings term uses the format {*Variable=Term*, ...}. When there are no variable bindings, the success or failure of the predicate call should be represented by the terms `{yes}` or `{no}`, respectively.

remarks List of general remarks about the predicate using the format *Topic – Text*. Both the topic and the text must be atoms.

For example:

```
:- info(color/1, [
    comment is 'Table of defined colors.',
    argnames is ['Color'],
    constraint is 'Up to four visible colors allowed.'
]).
```

As with the `info/1` directive, use only the keywords that make sense for your application and remember that you are free to invent your own keywords. All key-value pairs can also be retrieved programmatically using the reflection API and are visible to `lgtdoc` tool.

1.12.2 Processing and viewing documenting files

The `lgtdoc` tool generates an XML documenting file per entity. It can also generate library, directory, entity, and predicate indexes when documenting libraries and directories. For example, assuming the default filename extensions, a `trace` object and a `sort(_)` parametric object will result in `trace_0.xml` and `sort_1.xml` XML files.

Each entity XML file contains references to two other files, an XML specification file and a XSL style-sheet file. The XML specification file can be either a DTD file (`logtalk_entity.dtd`) or an XML Scheme file (`logtalk_entity.xsd`). The [XSL] style-sheet file is responsible for converting the XML files to some desired format such as HTML or PDF. The default names for the XML specification file and the XSL style-sheet file are defined by the `lgtdoc` tool but can be overridden by passing a list of options to the tool predicates. The `lgtdoc/xml` sub-directory in the Logtalk installation directory contains the XML specification files described above, along with several sample XSL style-sheet files and sample scripts for converting XML documenting files to several formats (e.g. reStructuredText, Markdown, HTML, and PDF). Please read the `NOTES` file included in the directory for details. You may use the supplied sample files as a starting point for generating the documentation of your Logtalk applications.

The Logtalk DTD file, `logtalk_entity.dtd`, contains a reference to a user-customizable file, `custom.ent`, which declares XML entities for source code author names, license terms, and copyright string. After editing the `custom.ent` file to reflect your personal data, you may use the XML entities on `info/1` documenting directives. For example, assuming that the XML entities are named *author*, *license*, and *copyright* we may write:

```
:- info([
    version is 1.1,
    author is {author},
    license is {license},
    copyright is {copyright}
]).
```

The entity references are replaced by the value of the corresponding XML entity when the XML documenting files are processed (**not** when they are generated; this notation is just a shortcut to take advantage of XML entities).

The `lgtdoc` tool supports a set of options that can be used to control the generation of the XML documentation files. Please see the tool documentation for details.

1.12.3 Inline formatting in comments text

Inline formatting in comments text can be accomplished by using Markdown (or reStructuredText) syntax and converting XML documenting files to Markdown (or reStructuredText) files (and these, if required, to e.g. HTML, ePub, or PDF formats).

1.13 Installing Logtalk

This page provides an overview of Logtalk installation requirements and instructions and a description of the files contained on the Logtalk distribution. For detailed, up-to-date installation and configuration instructions, please see the `README.md`, `INSTALL.md`, and `CUSTOMIZE.md` files distributed with Logtalk. The broad compatibility of Logtalk, both with Prolog compilers and operating-systems, together with all the possible user scenarios, means that installation can vary from very simple by running an installer or a couple of scripts to the need of patching both Logtalk and Prolog compilers to workaround the lack of strong Prolog standards or to cope with the requirements of less common operating-systems.

The preferred installation scenario is to have Logtalk installed in a system-wide location, thus available for all users, and a local copy of user-modifiable files on each user home directory (even when you are the single user of your computer). This scenario allows each user to independently customize Logtalk and to freely modify the provided libraries and programming examples. Logtalk installers, installation shell scripts, and Prolog integration scripts favor this installation scenario, although alternative installation scenarios are always possible. The installers set two environment variables, `LOGTALKHOME` and `LOGTALKUSER`, pointing, respectively, to the Logtalk installation folder and to the Logtalk user folder.

User applications should preferable be kept outside of the Logtalk user folder created by the installation process, however, as updating Logtalk often results in updating the contents of this folder. If your applications depend on customizations to the distribution files, backup those changes before updating Logtalk.

1.13.1 Hardware and software requirements

Computer and operating system

Logtalk is compatible with almost any computer/operating-system with a modern, standards compliant, Prolog compiler available.

Prolog compiler

Logtalk requires a backend Prolog compiler supporting official and de facto standards. Capabilities needed by Logtalk that are not defined in the official ISO Prolog Core standard include:

- access to predicate properties
- operating-system access predicates
- de facto standard predicates not (yet) specified in the official standard

Logtalk needs access to the predicate property `built_in` to properly compile objects and categories that contain Prolog built-in predicates calls. In addition, some Logtalk built-ins need to know the dynamic/static status of predicates to ensure correct application. The ISO standard for Prolog modules defines a `predicate_property/2` predicate that is already implemented by most Prolog compilers. Note that if these capabilities are not built-in the user cannot easily define them.

For optimal performance, Logtalk requires that the Prolog compiler supports **first-argument indexing** for both static and dynamic code (most modern compilers support this feature).

Since most Prolog compilers are moving closer to the ISO Prolog standard [ISO95], it is advisable that you try to use the most recent version of your favorite Prolog compiler.

1.13.2 Logtalk installers

Logtalk installers are available for macOS, Linux, and Microsoft Windows. Depending on the chosen installer, some tasks (e.g. setting environment variables or integrating Logtalk with some Prolog compilers) may need to be performed manually.

1.13.3 Source distribution

Logtalk sources are available in a `tar` archive compressed with `bzip2`, `lgt3xxx.tar.bz2`. You may expand the archive by using a decompressing utility or by typing the following commands at the command-line:

```
% tar -jxvf lgt3xxx.tar.bz2
```

This will create a sub-directory named `lgt3xxx` in your current directory. Almost all files in the Logtalk distribution are text files. Different operating-systems use different end-of-line codes for text files. Ensure that your decompressing utility converts the end-of-lines of all text files to match your operating system.

1.13.4 Directories and files organization

In the Logtalk installation directory, you will find the following files and directories:

`BIBLIOGRAPHY.bib` – Logtalk bibliography in BibTeX format

`CUSTOMIZE.md` – Logtalk end-user customization instructions

`INSTALL.md` – Logtalk installation instructions

`LICENSE.txt` – Logtalk user license

`NOTICE.txt` – Logtalk copyright notice

`QUICK_START.md` – Quick start instructions for those that do not like to read manuals

`README.md` – several useful information

`RELEASE_NOTES.md` – release notes for this version

`UPGRADING.md` – instructions on how to upgrade your programs to the current Logtalk version

`VERSION.txt` – file containing the current Logtalk version number (used for compatibility checking when upgrading Logtalk)

`loader-sample.lgt` – sample loader file for user applications

`settings-sample.lgt` – sample file for user-defined Logtalk settings

`tester-sample.lgt` – sample file for helping to automate running user application unit tests

adapters `NOTES.md` – notes on the provided adapter files `template.pl` – template adapter file `...` – specific adapter files

coding `NOTES.md` – notes on syntax highlighter and text editor support files providing syntax coloring for publishing and editing Logtalk source code `...` – syntax coloring support files

contributions `NOTES.md` – notes on the user-contributed code `...` – user-contributed code files

core `NOTES.md` – notes on the current status of the compiler and runtime `...` – core source files

docs `NOTES.md` – notes on the provided documentation for core, library, tools, and contributions entities `index.html` – root document for all entities documentation `...` – other entity documentation files

examples `NOTES.md` – short description of the provided examples `bricks NOTES.md` – example description and other notes `SCRIPT.txt` – step by step example tutorial `loader.lgt` – loader utility file for the example objects `...` – bricks example source files

`...` – other examples

integration `NOTES.md` – notes on scripts for Logtalk integration with Prolog compilers `...` – Prolog integration scripts

library `NOTES.md` – short description of the library contents `all_loader.lgt` – loader utility file for all library entities `...` – library source files

man `...` – POSIX man pages for the shell scripts

manuals `NOTES.md` – notes on the provided documentation `bibliography.html` – bibliography `glossary.html` – glossary `index.html` – root document for all documentation `...` – other documentation files

paths `NOTES.md` – description on how to setup library and examples paths `paths.pl` – default library and example paths

scratch `NOTES.md` – notes on the scratch directory

scripts `NOTES.md` – notes on scripts for Logtalk user setup, packaging, and installation `...` – packaging, installation, and setup scripts

tests `NOTES.md` – notes on the current status of the unit tests `...` – unit tests for built-in features

tools `NOTES.md` – notes on the provided programming tools `...` – programming tools

Adapter files

Adapter files provide the glue code between the Logtalk compiler/runtime and a Prolog compiler. Each adapter file contains two sets of predicates: ISO Prolog standard predicates and directives not built-in in the target Prolog compiler and Logtalk specific predicates.

Logtalk already includes ready to use adapter files for most academic and commercial Prolog compilers. If an adapter file is not available for the compiler that you intend to use, then you need to build a new one, starting from the included `template.pl` file. Start by making a copy of the template file. Carefully check (or complete if needed) each listed definition. If your Prolog compiler conforms to the ISO standard, this task should only take you a few minutes. In most cases, you can borrow code from the predefined adapter files. If you are unsure that your Prolog compiler provides all the ISO predicates needed by Logtalk, try to run the system by setting the unknown predicate error handler to report as an error any call to a missing predicate. Better yet, switch to a modern, ISO compliant, Prolog compiler. If you send me your adapter file, with a reference to the target Prolog compiler, maybe I can include it in the next release of Logtalk.

The adapter files specifies *default* values for most of the Logtalk compiler flags. Most of these compiler flags are described in the next section. A few of these flags have read-only values and cannot be changed at runtime. These are:

settings_file Allows or disables loading of *settings files* at startup. Possible values are `allow`, `restrict`, and `deny`. The usual default value is `allow` but it can be changed by editing the adapter file when e.g. embedding Logtalk in a compiled application. With a value of `allow`, settings files are searched in the startup directory, in the Logtalk user directory, and in the user home directory. With a value of `restrict`, settings files are only searched in the Logtalk user directory and in the user home directory.

prolog_dialect Name of the back-end Prolog compiler (an atom). This flag can be used for conditional compilation of Prolog specific code.

prolog_version Version of the back-end Prolog compiler (a compound term, `v(Major, Minor, Patch)`, whose arguments are integers). This flag availability depends on the Prolog compiler. Checking the value of this flag fails for any Prolog compiler that does not provide access to version data.

prolog_compatible_version Compatible version of the back-end Prolog compiler (a compound term, usually with the format `@>=(v(Major, Minor, Patch))`, whose arguments are integers). This flag availability depends on the Prolog compiler. Checking the value of this flag fails for any Prolog compiler that does not provide access to version data.

prolog_conformance Level of conformance of the back-end Prolog compiler with the ISO Prolog Core standard. The possible values are `strict` for compilers claiming strict conformance and `lax` for compilers claiming only broad conformance.

unicode Informs Logtalk if the back-end Prolog compiler supports the Unicode standard. Possible flag values are `unsupported`, `full` (all Unicode planes supported), and `bmp` (supports only the Basic Multilingual Plane).

encoding_directive Informs Logtalk if the back-end Prolog compiler supports the *encoding/1* directive. This directive is used for declaring the text encoding of source files. Possible flag values are `unsupported`, `full` (can be used in both Logtalk source files and compiler generated Prolog files), and `source` (can be used only in Logtalk source files).

tabling Informs Logtalk if the back-end Prolog compiler provides tabling programming support. Possible flag values are `unsupported` and `supported`.

engines Informs if the back-end Prolog compiler provides the required low level multi-threading programming support for Logtalk threaded engines. Possible flag values are `unsupported` and `supported`.

threads Informs if the back-end Prolog compiler provides the required low level multi-threading programming support for all high-level Logtalk multi-threading features. Possible flag values are `unsupported` and `supported`.

modules Informs Logtalk if the back-end Prolog compiler provides suitable module support. Possible flag values are `unsupported` and `supported` (Logtalk provides limited support for compiling Prolog modules as objects).

coinduction Informs Logtalk if the back-end Prolog compiler provides the minimal support for cyclic terms necessary for working with coinductive predicates. Possible flag values are `unsupported` and `supported`.

Settings files

Although is always possible to edit the back-end Prolog compiler adapter files, the recommended solution to customize compiler flags is to edit the `settings.lgt` file in the Logtalk user folder or in the user home folder. Depending on the back-end Prolog compiler and on the operating-system, is also possible to define per-project settings files by creating a `settings.lgt` file in the project directory and by starting Logtalk from this directory. At startup, Logtalk tries to load a `settings.lgt` file from the startup directory (assuming that the read-only `settings` flag is set to `allow`). If not found, Logtalk tries to load a `settings.lgt` file from the Logtalk user folder. If still not found, Logtalk tries to load a `settings.lgt` file from the user home folder. When no settings files are found, Logtalk will use the default compiler flag values set on the back-end Prolog compiler adapter files. When limitations of the back-end Prolog compiler or on the operating-system prevent Logtalk from finding the settings files, these can always be loaded manually after Logtalk startup.

Settings files are normal Logtalk source files (although when automatically loaded by Logtalk they are compiled silently with any errors being simply ignored). The usual contents is an `initialization/1` Prolog directive containing calls to the *set_logtalk_flag/2* Logtalk built-in predicate and asserting clauses for the *logtalk_library_path/2* multifile dynamic predicate. Note that the *set_logtalk_flag/2* directive cannot be used as its scope is local to the source file being compiled. For example, one of the troubles of writing portable applications is the different feature sets of Prolog compilers. A typical issue is the lack of support for tabling. Using the Logtalk support for conditional compilation you could write:

```
:- if(current_logtalk_flag(tabling, supported)).

    % add tabling directives to the source code
```

```
    :- table(foo/1).
    :- table(bar/2).

:- endif.
```

The Logtalk flag `prolog_dialect` may also be used with the conditional compilation directives in order to define a single settings file that can be used with several back-end Prolog compilers. For example:

```
:- if(current_logtalk_flag(prolog_dialect, yap)).

    % YAP specific settings
    ...

:- elif(current_logtalk_flag(prolog_dialect, gnu)).

    % GNU Prolog specific settings
    ...

:- else.

    % generic Prolog settings

:- endif.
```

Logtalk compiler and runtime

The `compiler` sub-directory contains the Prolog source file(s) that implement the Logtalk compiler and the Logtalk runtime. The compiler and the runtime may be split in two (or more) separate files or combined in a single file, depending on the Logtalk release that you are installing.

Library

Starting from version 2.7.0, Logtalk contains a standard library of useful objects, categories, and protocols. Read the corresponding `NOTES.md` file for details about the library contents.

Examples

Logtalk 2.x and 3.x contain new implementations of some of the examples provided with previous 1.x versions. The sources of each one of these examples can be found included in a subdirectory with the same name, inside the directory examples. The majority of these examples include a file named `SCRIPT.txt` that contains cases of simple utilization. Some examples may depend on other examples and library objects to work properly. Read the corresponding `NOTES.md` file for details before running an example.

Logtalk source files

Logtalk source files are text files containing one or more entity definitions (objects, categories, or protocols). The Logtalk source files may also contain plain Prolog code. The extension `.lgt` is normally used. Logtalk compiles these files to plain Prolog by appending to the file name a suffix derived from the extension and by replacing the `.lgt` extension with `.pl` (`.pl` is the default Prolog extension; if your Prolog compiler expects the Prolog source filenames to end with a specific, different extension, you can set it in the corresponding adapter file).

1.14 Writing, running, and debugging applications

1.14.1 Writing applications

For a successful programming in Logtalk, you need a good working knowledge of Prolog and an understanding of the principles of object-oriented programming. Most guidelines for writing good Prolog code apply as well to Logtalk programming. To those guidelines, you should add the basics of good object-oriented design.

One of the advantages of a system like Logtalk is that it enable us to use the currently available object-oriented methodologies, tools, and metrics [Champaux92] in Prolog programming. That said, writing applications in Logtalk is similar to writing applications in Prolog: we define new predicates describing what is true about our domain objects, about our problem solution. We encapsulate our predicate directives and definitions inside new objects, categories and protocols that we create by hand with a text editor or by using the Logtalk built-in predicates. Some of the information collected during the analysis and design phases can be integrated in the objects, categories and protocols that we define by using the available entity and predicate documenting directives.

Source files

Logtalk source files may define any number of entities (objects, categories, or protocols) and Prolog code. If you prefer to define each entity in its own source file, then it is recommended that the source file be named after the entity identifier. For parametric objects, the identifier arity can be appended to the identifier functor. By default, all Logtalk source files use the extension `.lgt` but this is optional and can be set in the adapter files. Intermediate Prolog source files (generated by the Logtalk compiler) have, by default, a `_lgt` suffix and a `.pl` extension. Again, this can be set to match the needs of a particular Prolog compiler in the corresponding adapter file. For example, we may define an object named `vehicle` and save it in a `vehicle.lgt` source file that will be compiled to a `vehicle_lgt.pl` Prolog file. If we have a `sort(_)` parametric object we can save it on a `sort_1.lgt` source file that will be compiled to a `sort_1_lgt.pl` Prolog file. This name scheme helps avoid file name conflicts (remember that all Logtalk entities share the same name space). To further prevent file name conflicts, depending on the backend compiler, the names of the intermediate Prolog files may include a directory hash.

Logtalk source files may contain Prolog code interleaved with Logtalk entity definitions. Plain Prolog code is usually copied as-is to the corresponding Prolog output file (except, of course, if subject to the term-expansion mechanism). Prolog modules are compiled as objects. The following Prolog directives are processed when read (thus affecting the compilation of the source code that follows): `ensure_loaded/1`, `use_module/1-2`, `op/3`, and `set_prolog_flag/2`. The `initialization/1` Prolog directive may be used for defining an initialization goal to be executed when loading a source file. Most calls to Logtalk built-in predicates from file `initialization/1` directives are compiled for better performance.

The text encoding used in a source file may be declared using the *encoding/1* directive when running Logtalk with back-end Prolog compilers that support multiple encodings (check the `encoding_directive` flag in the adapter file of your Prolog compiler). The encoding used (and, in the case of a Unicode encoding, any BOM present) in a source file will be used for the intermediate Prolog file generated by the compiler. Logtalk uses the encoding names specified by IANA (in those cases where a preferred MIME name alias is specified, the alias is used instead).

Logtalk source files can include the text of other files by using the *include/1* directive. Although there's also a standard Prolog `include/1` directive, any occurrences of this directive in a Logtalk source file is handled by the Logtalk compiler, not by the backend Prolog compiler.

Portable applications

Logtalk is compatible with almost all modern Prolog compilers. However, this does not necessarily imply that your Logtalk applications will have the same level of portability. If possible, you should only use in your applications Logtalk built-in predicates and ISO Prolog specified built-in predicates and arithmetic functions. If you need to use built-in predicates (or built-in arithmetic functions) that may not be available in other Prolog compilers, you should

try to encapsulate the non-portable code in a small number of objects and provide a portable **interface** for that code through the use of Logtalk protocols. An example will be code that access operating-system specific features. The Logtalk compiler can warn you of the use of non-ISO specified built-in predicates and arithmetic functions by using the `portability/1` compiler flag.

Conditional compilation

Logtalk supports conditional compilation within source files using the *if/1*, *elif/1*, *else/0*, and *endif/0* directives. This support is similar to the support found in several Prolog systems such as ECLiPSe, SWI-Prolog, or YAP.

Avoiding common errors

Try to write objects and protocol documentation **before** writing any other code; if you are having trouble documenting a predicate perhaps we need to go back to the design stage.

Try to avoid lengthy hierarchies. Composition is often a better choice over inheritance for defining new objects (Logtalk supports component-based programming through the use of categories). In addition, prototype-based hierarchies are semantically simpler than class-based hierarchies.

Dynamic predicates or dynamic entities are sometimes needed, but we should always try to minimize the use of non-logical features such as asserts and retracts.

Since each Logtalk entity is independently compiled, if an object inherits a dynamic or a meta-predicate predicate, then the respective directives must be repeated to ensure a correct compilation.

In general, Logtalk does not verify if a user predicate call/return arguments comply with the declared modes. On the other hand, Logtalk built-in predicates, built-in methods, and message sending control structures are fully checked for calling mode errors.

Logtalk error handling strongly depends on the ISO compliance of the chosen Prolog compiler. For instance, the error terms that are generated by some Logtalk built-in predicates assume that the Prolog built-in predicates behave as defined in the ISO standard regarding error conditions. In particular, if your Prolog compiler does not support a `read_term/3` built-in predicate compliant with the ISO Prolog Standard definition, then the current version of the Logtalk compiler may not be able to detect misspell variables in your source code.

Coding style guidelines

It is suggested that all code between an entity opening and closing directives be indented by one tab stop. When defining entity code, both directives and predicates, Prolog coding style guidelines may be applied. All Logtalk source files, examples, and standard library entities use tabs (the recommended setting is a tab width equivalent to 4 spaces) for laying out code. Closed related entities can be defined in the same source file. However, for best performance, is often necessary to have an entity per source file. Entities that might be useful in different contexts (such as library entities) are best defined in their own source files.

1.14.2 Running a Logtalk session

We run Logtalk inside a normal Prolog session, after loading the necessary files. Logtalk extends but does not modify your Prolog compiler. We can freely mix Prolog queries with the sending of messages and our applications can be made of both normal Prolog clauses and object definitions.

Starting Logtalk

Depending on your Logtalk installation, you may use a script or a shortcut to start Logtalk with your chosen Prolog compiler. On POSIX operating systems, the scripts should be available from the command-line; scripts are named upon the used Prolog compilers. On Windows, the shortcuts should be available from the Start Menu. If no scripts or shortcuts are available for your installation, operating-system, or Prolog compiler, you can always start a Logtalk session by performing the following steps:

1. Start your Prolog compiler.

2. Load the appropriate adapter file for your compiler. Adapter files for most common Prolog compilers can be found in the `adapters` subdirectory.

3. Load the library paths file corresponding to your Logtalk installation contained in the `paths` subdirectory.

4. Load the Logtalk compiler/runtime files contained in the `compiler` subdirectory.

Note that the adapter files, compiler/runtime files, and library paths file are Prolog source files. The predicate called to load (and compile) them depends on your Prolog compiler. In case of doubt, consult your Prolog compiler reference manual or take a look at the definition of the predicate `'$lgt_load_prolog_code'/3` in the corresponding adapter file.

Most Prolog compilers support automatic loading of an initialization file, which can include the necessary directives to load both the Prolog adapter file and the Logtalk compiler. This feature, when available, allows automatic loading of Logtalk when you start your Prolog compiler.

Compiling and loading your applications

Your applications will be made of source files containing your objects, protocols, and categories. The source files can be compiled to disk by calling the Logtalk built-in predicate *logtalk_compile/1*:

```
| ?- logtalk_compile([source_file1, source_file2, ...]).
```

This predicate runs the compiler on each file and, if no fatal errors are found, outputs Prolog source files that can then be consulted or compiled in the usual way by your Prolog compiler.

To compile to disk and also load into memory the source files we can use the Logtalk built-in predicate *logtalk_load/1*:

```
| ?- logtalk_load([source_file1, source_file2, ...]).
```

This predicate works in the same way of the predicate `logtalk_compile/1` but also loads the compiled files into memory.

Both predicates expect a source file name or a list of source file names as an argument. The Logtalk source file name extension, as defined in the adapter file (by default, `.lgt`), can be omitted.

If you have more than a few source files then you may want to use a loader helper file containing the calls to the `logtalk_load/1-2` predicates. Consulting or compiling the loader file will then compile and load all your Logtalk entities into memory (see below for details).

With most Prolog back-end compilers, you can use the shorthands `{File}` for `logtalk_load(File)` and `{File1, File2, ...}` for `logtalk_load([File1, File2, ...])`. The use these shorthands should be restricted to the Logtalk/Prolog top-level interpreter as they are not part of the language specification and may be commented out in case of conflicts with backend Prolog compiler features.

The built-in predicate *logtalk_make/0* can be used to reload all modified source files. Files are also reloaded when the compilation mode changes. For example, assume that you have loaded your application files and found a bug. You can easily recompile the files in debug mode by using the queries:

```
| ?- set_logtalk_flag(debug, on).
...

| ?- logtalk_make.
...
```

After debugging and fixing the bugs, you can reload the files in normal (or optimized) mode by turning the debug flag off and calling the `logtalk_make/0` predicate again.

An extended version of this predicate, *logtalk_make/1*, accepts `all`, `clean`, `check`, `circular`, `documentation`, and `caches` arguments for, respectively, reloading modified Logtalk source files, deleting any intermediate files generated by the compilation of Logtalk source files, checking for code issues, listing of circular dependencies, generating documentation, and deleting dynamic binding caches. With most Prolog backend compilers, you can use the shorthands `{*}` for `logtalk_make(all)`, `{!}` for `logtalk_make(clean)`, `{?}` for `logtalk_make(missing)`, `{@}` for `logtalk_make(circular)`, `{#}` for `logtalk_make(documentation)`, and `{$}` for `logtalk_make(caches)`. The `logtalk_make(clean)` goal can be specially useful before switching backend Prolog compilers as the generated intermediate files may not be compatible. The `logtalk_make(caches)` goal is usually used when benchmarking compiler performance improvements.

Loader utility files

Most examples directories contain a Logtalk utility file that can be used to load all included source files. These loader utility files are usually named `loader.lgt` or contain the word "loader" in their name. Loader files are ordinary source file and thus compiled and loaded like any source file. For an example loader file named `loader.lgt` we would type:

```
| ?- logtalk_load(loader).
```

Usually these files contain a call to the Logtalk built-in predicates *set_logtalk_flag/2* (e.g. for setting global, *project-specific*, flag values) and *logtalk_load/1* or *logtalk_load/2* (for loading project files), wrapped inside a Prolog `initialization/1` directive. For instance, if your code is split in three Logtalk source files named `source1.lgt`, `source2.lgt`, and `source3.lgt`, then the contents of your loader file could be:

```
:- initialization((
    % set project-specific global flags
    set_logtalk_flag(events, allow),
    % load the project source files
    logtalk_load([source1, source2, source3])
)).
```

Another example of directives that are often used in a loader file would be `op/3` directives declaring global operators needed by your application. Loader files are also often used for setting source file-specific compiler flags (this is useful even when you only have a single source file if you always load it with using the same set of compiler flags). For example:

```
:- initialization((
    % set project-specific global flags
    set_logtalk_flag(underscore_variables, dont_care),
    set_logtalk_flag(source_data, off),
    % load the project source files
    logtalk_load(
        [source1, source2, source3],
        % source file-specific flags
        [portability(warning)]),
```

(continues on next page)

```
    logtalk_load(
        [source4, source5],
        % source file-specific flags
        [portability(silent)])
)).
```

To take the best advantage of loader files, define a clause for the multifile and dynamic `logtalk_library_path/`
`2` predicate for the directory containing your source files as explained in the next section.

A common mistake is to try to set compiler flags using `logtalk_load/2` with a loader file. For example, by
writing:

```
| ?- logtalk_load(loader, [underscore_variables(dont_care), source_data(off)]).
```

This will not work as you might expect as the compiler flags will only be used in the compilation of the `loader.lgt`
file itself and will not affect the compilation of files loaded through the `initialization/1` directive contained on
the loader file.

Libraries of source files

Logtalk defines a *library* simply as a directory containing source files. Library locations can be specified by defining
or asserting clauses for the dynamic and multifile predicate *logtalk_library_path/2*. For example:

```
:- multifile(logtalk_library_path/2).
:- dynamic(logtalk_library_path/2).

logtalk_library_path(shapes, '$LOGTALKUSER/examples/shapes/').
```

The first argument of the predicate is used as an alias for the path on the second argument. Library aliases may also
be used on the second argument. For example:

```
:- multifile(logtalk_library_path/2).
:- dynamic(logtalk_library_path/2).

logtalk_library_path(lgtuser, '$LOGTALKUSER/').
logtalk_library_path(examples, lgtuser('examples/')).
logtalk_library_path(viewpoints, examples('viewpoints/')).
```

This allows us to load a library source file without the need to first change the current working directory to the library
directory and then back to the original directory. For example, in order to load a `loader.lgt` file, contained in a
library named `viewpoints`, we just need to type:

```
| ?- logtalk_load(viewpoints(loader)).
```

The best way to take advantage of this feature is to load at startup a source file containing clauses for the
`logtalk_library_path/2` predicate needed for all available libraries. This allows us to load library source
files or entire libraries without worrying about libraries paths, improving code portability. The directory paths on the
second argument should always end with the path directory separator character. Most back-end Prolog compilers al-
lows the use of environment variables in the second argument of the `logtalk_library_path/2` predicate. Use
of POSIX relative paths (e.g. `'../'` or `'./'`) for top-level library directories (e.g. `lgtuser` in the example above)
is not advised as different back-end Prolog compilers may start with different initial working directories, which may
result in portability problems of your loader files.

The library notation provides functionality inspired by the `file_search_path/2` mechanism introduced by Quin-
tus Prolog and later adopted by some other Prolog compilers.

Compiler linter

The compiler includes a linter that checks for a wide range of possible problems in source files. Notably, the compiler checks for unknown entities, unknown predicates, undefined predicates (i.e. predicates that are declared but not defined), missing directives (including missing `dynamic/1` and `meta_predicate/1` directives), redefined built-in predicates, calls to non-portable predicates, singleton variables, tautology and falsehood goals (i.e. goals that are can be replaced by `true` or `fail`), and trivial fails (i.e. calls to predicates with no match clauses). Some of the linter warnings are controlled by compiler flags. See the next section for details.

Compiler flags

The *logtalk_load/1* and *logtalk_compile/1* always use the current set of default compiler flags as specified in your settings file and the Logtalk adapter files or changed for the current session using the built-in predicate *set_logtalk_flag/2*. Although the default flag values cover the usual cases, you may want to use a different set of flag values while compiling or loading some of your Logtalk source files. This can be accomplished by using the *logtalk_load/2* or the *logtalk_compile/2* built-in predicates. These two predicates accept a list of options affecting how a Logtalk source file is compiled and loaded:

```
| ?- logtalk_compile(Files, Options).
```

or:

```
| ?- logtalk_load(Files, Options).
```

In fact, the `logtalk_load/1` and `logtalk_compile/1` predicates are just shortcuts to the extended versions called with the default compiler flag values. The options are represented by a compound term where the functor is the flag name and the sole argument is the flag value.

We may also change the default flag values from the ones loaded from the adapter file by using the *set_logtalk_flag/2* built-in predicate. For example:

```
| ?- set_logtalk_flag(unknown_entities, silent).
```

The current default flags values can be enumerated using the *current_logtalk_flag/2* built-in predicate:

```
| ?- current_logtalk_flag(unknown_entities, Value).

Value = silent
yes
```

Logtalk also implements a *set_logtalk_flag/2* directive, which can be used to set flags within a source file or within an entity. For example:

```
% compile objects in this source file with event support
:- set_logtalk_flag(events, allow).

:- object(foo).

    % compile this object with support
    % for dynamic predicate declarations
    :- set_logtalk_flag(dynamic_declarations, allow).
    ...

:- end_object.

...
```

Note that the scope of the `set_logtalk_flag/2` directive is local to the entity or to the source file containing it.

Version flags

version_data(Value) Read-only flag whose value is the compound term `logtalk(Major,Minor, Patch,Status)`. The first three arguments are integers and the last argument is an atom, possibly empty, representing version status: `aN` for alpha versions, `bN` for beta versions, `rcN` for release candidates (with `N` being a natural number), and `stable` for stable versions. The `version_data` flag is also a de facto standard for Prolog compilers.

Lint flags

unknown_entities(Option) Controls the unknown entity warnings, resulting from loading an entity that references some other entity that is not currently loaded. Possible option values are `warning` (the usual default) and `silent`. Note that these warnings are not always avoidable, specially when using reflective designs of class-based hierarchies.

unknown_predicates(Option) Defines the compiler behavior when calls to unknown predicates (or non-terminals) are found. An unknown predicate is a called predicate that is neither locally declared or defined. Possible option values are `error`, `warning` (the usual default), and `silent` (not recommended).

undefined_predicates(Option) Defines the compiler behavior when calls to declared but undefined predicates (or non-terminals) are found. Note that calls to declared but undefined predicates (or non-terminals) fail as per closed-world assumption. Possible option values are `error`, `warning` (the usual default), and `silent` (not recommended).

portability(Option) Controls the non-ISO specified Prolog built-in predicate and non-ISO specified Prolog built-in arithmetic function calls warnings plus use of non-standard Prolog flags and/or flag values. Possible option values are `warning` and `silent` (the usual default).

missing_directives(Option) Controls the missing predicate directive warnings. Possible option values are `warning` (the usual default) and `silent` (not recommended).

duplicated_directives(Option) Controls the duplicated predicate directive warnings. Possible option values are `warning` (the usual default) and `silent` (not recommended). Note that conflicting directives for the same predicate are handled as errors, not as duplicated directive warnings.

trivial_goal_fails(Option) Controls the printing of warnings warnings for calls to local static predicates with no matching clauses. Possible option values are `warning` (the usual default) and `silent` (not recommended).

always_true_or_false_goals(Option) Controls the printing of warnings for goals that are always true or false. Possible option values are `warning` (the usual default) and `silent` (not recommended).

lambda_variables(Option) Controls the printing of lambda variable related warnings. Possible option values are `warning` (the usual default) and `silent` (not recommended).

suspicious_calls(Option) Controls the printing of suspicious call warnings. Possible option values are `warning` (the usual default) and `silent` (not recommended).

redefined_built_ins(Option) Controls the Logtalk and Prolog built-in predicate redefinition warnings. Possible option values are `warning` (the usual default) and `silent`. Warnings about redefined Prolog built-in predicates are often the result of running a Logtalk application on several Prolog compilers as each Prolog compiler defines its set of built-in predicates.

singleton_variables(Option) Controls the singleton variable warnings. Possible option values are `warning` (the usual default) and `silent` (not recommended).

underscore_variables(Option) Controls the interpretation of variables that start with an underscore (excluding the anonymous variable) that occur once in a term as either don't care variables or singleton variables. Possible option values are `dont_care` and `singletons` (the usual default). Note that, depending on your Prolog compiler, the `read_term/3` built-in predicate may report variables that start with an underscore as singleton variables. There is no standard behavior, hence this option.

Optional features compilation flags

complements(Option) Allows objects to be compiled with support for complementing categories turned off in order to improve performance and security. Possible option values are `allow` (allow complementing categories to override local object predicate declarations and definitions), `restrict` (allow complementing categories to add predicate declarations and definitions to an object but not to override them), and `deny` (ignore complementing categories; the usual default). This option can be used on a per-object basis. Note that changing this option is of no consequence for objects already compiled and loaded.

dynamic_declarations(Option) Allows objects to be compiled with support for dynamic declaration of new predicates turned off in order to improve performance and security. Possible option values are `allow` and `deny` (the usual default). This option can be used on a per-object basis. Note that changing this option is of no consequence for objects already compiled and loaded. This option is only checked when sending an `asserta/1` or `assertz/1` message to an object. Local asserting of new predicates is always allowed.

events(Option) Allows message sending calls to be compiled with event-driven programming support disable in order to improve performance. Possible option values are `allow` and `deny` (the usual default). Objects (and categories) compiled with this option set to `deny` use optimized code for message-sending calls that does not trigger events. As such, this option can be used on a per-object (or per-category) basis. Note that changing this option is of no consequence for objects already compiled and loaded.

context_switching_calls(Option) Allows context switching calls (`<</2`) to be either allowed or denied. Possible option values are `allow` and `deny`. The default flag vale is `allow`. Note that changing this option is of no consequence for objects already compiled and loaded.

Back-end Prolog compiler and loader flags

prolog_compiler(Flags) List of compiler flags for the generated Prolog files. The valid flags are specific to the used Prolog backend compiler. The usual default is the empty list. These flags are passed to the backend Prolog compiler built-in predicate that is responsible for compiling to disk a Prolog file. For Prolog compilers that don't provide separate predicates for compiling and loading a file, use instead the `prolog_loader/1` flag.

prolog_loader(Flags) List of loader flags for the generated Prolog files. The valid flags are specific to the used Prolog backend compiler. The usual default is the empty list. These flags are passed to the backend Prolog compiler built-in predicate that is responsible for loading a (compiled) Prolog file.

Other flags

scratch_directory(Directory) Sets the directory to be used to store the temporary files generated when compiling Logtalk source files. This directory can be specified using an atom or using library notation. The directory must always end with a slash. The default value is a sub-directory of the source files directory, either `'./lgt_tmp/'` or `'./.lgt_tmp/'` (depending on the back-end Prolog compiler and operating-system). Relative directories must always start with `'./'` due to the lack of a portable solution to check if a path is relative or absolute.

report(Option) Controls the default printing of messages. Possible option values are `on` (by usual default, print all messages that are not intercepted by the user), `warnings` (only print warning and error messages that are not intercepted by the user), and `off` (do not print any messages that are not intercepted by the user).

code_prefix(Character) Enables the definition of prefix for all functors of Prolog code generated by the Logtalk compiler. The option value must be a single character atom. Its default value is `'$'`. Specifying a code prefix provides a way to solve possible conflicts between Logtalk compiled code and other Prolog code. In addition, some Prolog compilers automatically hide predicates whose functor start with a specific prefix such as the character `$`. Although this is not a read-only flag, it should only be changed at startup time and before loading any source files.

optimize(Option) Controls the compiler optimizations. Possible option values are `on` (used by default for deployment) and `off` (used by default for development). Compiler optimizations include the use of static binding whenever possible, the removal of redundant calls to `true/0` from predicate clauses, the removal of redundant unifications when compiling grammar rules, and inlining of predicate definitions with a single clause that links to a local predicate, to a plain Prolog built-in (or foreign) predicate, or to a Prolog module predicate with the same arguments. Care should be taken when developing applications with this flag turned on as changing and reloading a file may render *static binding* optimizations invalid for code defining in other loaded files. Turning on this flag automatically turns off the `debug` flag.

source_data(Option) Defines how much information is retained when compiling a source file. Possible option values are `on` (the usual default for development) and `off`. With this flag set to `on`, Logtalk will keep the information represented using documenting directives plus source location data (including source file names and line numbers). This information can be retrieved using reflection and is useful for documenting, debugging, and integration with third-party development tools. This flag can be turned off in order to generate more compact code.

debug(Option) Controls the compilation of source files in debug mode (the Logtalk default debugger can only be used with files compiled in this mode). Also controls, by default, printing of `debug>` and `debug(Topic)` messages. Possible option values are `on` and `off` (the usual default). Turning on this flag automatically turns off the `optimize` flag.

reload(Option) Defines the reloading behavior for source files. Possible option values are `skip` (skip loading of already loaded files; this value can be used to get similar functionality to the Prolog directive `ensure_loaded/1` but should be used only with fully debugged code), `changed` (the usual default; reload files only when they are changed since last loaded provided that the any explicit flags and the compilation mode are the same as before), and `always` (always reload files).

relative_to(Directory) Defines a base directory for resolving relative source file paths. The default value is the directory of the source file being compiled.

hook(Object) Allows the definition of compiler hooks that are called for each term read form a source file and for each compiled goal. This option specifies an object (which can be the pseudo-object `user`) implementing the `expanding` built-in protocol. The hook object must be compiled and loaded when this option is used. It's also possible to specify a Prolog module instead of a Logtalk object but the module must be pre-loaded and its identifier must be different from any object identifier. The object is expected to define clauses for the *term_expansion/2* and *goal_expansion/2* predicates. In the case of the `term_expansion/2` predicate, the first argument is the term read form the source file while the second argument returns a list of terms corresponding to the expansion of the first argument. In the case of the `goal_expansion/2` predicate, the second argument should be a goal resulting from the expansion of the goal in the first argument. The predicate `goal_expansion/2` is recursively called on the expanded goal until a fixed point is reached. Care must be taken to avoid compilation loops.

clean(Option) Controls cleaning of the intermediate Prolog files generated when compiling Logtalk source files. Possible option values are `off` and `on` (the usual default). When turned on, this flag also forces recompilation of all source files, disregarding any existing intermediate files. Thus, it is strong advisable to turn on this flag when switching backend Prolog compilers as the intermediate files generated by the compilation of source files may not be portable (due to differences in the implementation of the standard `write_canonical/2` predicate).

1.14. Writing, running, and debugging applications   **91**

User-defined flags

Logtalk provides a *create_logtalk_flag/3* predicate that can be used for defining new flags.

Reloading and smart compilation of source files

As a general rule, reloading source files should never occur in production code and should be handled with care in development code. Reloading a Logtalk source file usually requires reloading the intermediate Prolog file that is generated by the Logtalk compiler. The problem is that there is no standard behavior for reloading Prolog files. For static predicates, almost all Prolog compilers replace the old definitions with the new ones. However, for dynamic predicates, the behavior depends on the Prolog compiler. Most compilers replace the old definitions but some of them simply append the new ones, which usually leads to trouble. See the compatibility notes for the back-end Prolog compiler you intend to use for more information. There is an additional potential problem when using multi-threading programming. Reloading a threaded object does not recreate from scratch its old message queue, which may still be in use (e.g. threads may be waiting on it).

When using library entities and stable code, you can avoid reloading the corresponding source files (and, therefore, recompiling them) by setting the compiler flag `reload` to `skip`. For code under development, you can turn off the `clean` flag to avoid recompiling files that have not been modified since last compilation (assuming that back-end Prolog compiler that you are using supports retrieving of file modification dates). You can disable deleting the intermediate files generated when compiling source files by changing the default flag value in your settings file, by using the corresponding compiler flag with the compiling and loading built-in predicates, or, for the remaining of a working session, by using the call:

```
| ?- set_logtalk_flag(clean, off).
```

Some caveats that you should be aware. First, some warnings that might be produced when compiling a source file will not show up if the corresponding object file is up-to-date because the source file is not being (re)compiled. Second, if you are using several Prolog compilers with Logtalk, be sure to perform the first compilation of your source files with smart compilation turned off: the intermediate Prolog files generated by the Logtalk compiler may be not compatible across Prolog compilers or even for the same Prolog compiler across operating systems (e.g. due to the use of different character encodings or end-of-line characters).

Using Logtalk for batch processing

If you use Logtalk for batch processing, you probably want to turn off the option `report` to suppress all messages of type `banner`, `comment`, `comment(_)`, `warning`, and `warning(_)` that are normally printed. Note that error messages and messages providing information requested by the user will still be printed.

1.14.3 Optimizing performance

The default compiler flag settings are appropriated for the **development** but not necessarily for the **deployment** of applications. To minimize the generated code size, turn the `source_data` flag off. To optimize runtime performance, turn on the `optimize` flag.

Pay special attention to file compilation/loading order. Whenever possible, compile/load your files taking into account file dependencies to enable *static binding* optimizations. The easiest way to find the dependencies and thus the best compilation/loading order is to use the `diagrams` tool to generate a file dependency diagram for your application.

Minimize the use of dynamic predicates. Parametric objects can often be used in alternative. When dynamic predicates cannot be avoided, try to make them private. Declaring a dynamic predicate also as a private predicate allows the compiler to optimize local calls to the database methods (e.g. *assertz/1* and *retract/1*) that handle the predicate.

Sending a *message to self* implies *dynamic binding* but there are often cases where *::/1* is misused to call an imported or inherited predicate that is never going to be redefined in a descendant. In these cases, a *super call*, *^^/1*, can be used instead with the benefit of often enabling static binding. Most of the guidelines for writing efficient Prolog code also apply to Logtalk code. In particular, define your predicates to take advantage of first-argument indexing. In the case of recursive predicates, define them as tail-recursive predicates whenever possible.

1.14.4 Debugging Logtalk applications

The Logtalk distribution includes in its `tools` directory a command-line debugger, implemented as a Logtalk application. It can be loaded by typing:

```
| ?- logtalk_load(debugger(loader)).
```

This tool implements debugging features similar to those found on most Prolog systems. There are some differences, however, between the usual implementation of Prolog debuggers and the current implementation of the Logtalk debugger that you should be aware. First, unlike some Prolog debuggers, the Logtalk debugger is not implemented as a meta-interpreter. This translates to a different, although similar, set of debugging features when compared with some of the more sophisticated Prolog debuggers. Second, debugging is only possible for entities compiled in debug mode. When compiling an entity in debug mode, Logtalk decorates clauses with source information to allow tracing of the goal execution. Third, implementation of spy points allows the user to specify the execution context for entering the debugger. This feature is a consequence of the encapsulation of predicates inside objects.

Compiling source files and entities in debug mode

Compilation of source files in debug mode is controlled by the compiler flag `debug`. The default value for this flag, usually `off`, is defined in the adapter files. Its value may be changed at runtime by writing:

```
| ?- set_logtalk_flag(debug, on).
```

In alternative, if we want to compile only some source files in debug mode, we may instead write:

```
| ?- logtalk_load([file1, file2, ...], [debug(on)]).
```

The compiler flag `clean` should be turned on whenever the debug flag is turned on at runtime. This is necessary because debug code would not be generated for files previously compiled in normal mode if there are no changes to the source files.

After loading the debugger, we may check or enumerate, by backtracking, all loaded entities compiled in debug mode as follows:

```
| ?- debugger::debugging(Entity).
```

To compile only a specific entity in debug mode, use the *set_logtalk_flag/2* directive inside the entity.

Logtalk Procedure Box model

Logtalk uses a *Procedure Box model* similar to those found on most Prolog compilers. The traditional Prolog procedure box model uses four ports (*call*, *exit*, *redo*, and *fail*) for describing control flow when a predicate clause is used during program execution:

call predicate call

exit success of a predicate call

redo backtracking into a predicate

fail failure of a predicate call

Logtalk, as found on some recent Prolog compilers, adds a port for dealing with exceptions thrown when calling a predicate:

exception predicate call throws an exception

In addition to the ports described above, Logtalk adds two more ports, *fact* and *rule*, which show the result of the unification of a goal with, respectively, a fact and a rule head:

fact unification success between a goal and a fact

rule unification success between a goal and a rule head

The user may define for which ports the debugger should pause for user interaction by specifying a list of leashed ports. For example:

```
| ?- debugger::leash([call, exit, fail]).
```

Alternatively, the user may use an atom abbreviation for a pre-defined set of ports. For example:

```
| ?- debugger::leash(loose).
```

The abbreviations defined in Logtalk are similar to those defined on some Prolog compilers:

none []

loose [fact, rule, call]

half [fact, rule, call, redo]

tight [fact, rule, call, redo, fail, exception]

full [fact, rule, call, exit, redo, fail, exception]

By default, the debugger pauses at every port for user interaction.

Defining spy points

Logtalk spy points can be defined by simply stating which file line numbers or predicates should be spied, as in most Prolog debuggers, or by fully specifying the context for activating a spy point. In the case of line number spy points, the line number must correspond to the first line of an entity clause. To simplify the definition of line number spy points, these are specified using the entity identifier instead of the file name (as all entities share a single namespace, an entity can only be defined in a single file).

Defining line number and predicate spy points

Line number and predicate spy points are specified using the method `spy/1`. The argument can be either a pair entity identifier - line number (`Entity-Line`) or a predicate indicator (`Name/Arity`) or a list of spy points. For example:

```
| ?- debugger::spy(person-42).

Spy points set.
yes

| ?- debugger::spy(foo/2).

Spy points set.
yes
```

```
| ?- debugger::spy([foo/4, bar/1]).

Spy points set.
yes
```

Line numbers and predicate spy points can be removed by using the method `nospy/1`. The argument can be a spy point, a list of spy points, or a non-instantiated variable in which case all spy points will be removed. For example:

```
| ?- debugger::nospy(_).

All matching predicate spy points removed.
yes
```

Defining context spy points

A context spy point is a term describing a message execution context and a goal:

```
(Sender, This, Self, Goal)
```

The debugger is evoked whenever the execution context is true and when the spy point goal unifies with the goal currently being executed. Variable bindings resulting from the unification between the current goal and the goal argument are discarded. The user may establish any number of context spy points as necessary. For example, in order to call the debugger whenever a predicate defined on an object named `foo` is called we may define the following spy point:

```
| ?- debugger::spy(_, foo, _, _).

Spy point set.
yes
```

For example, we can spy all calls to a `foo/2` predicate by setting the condition:

```
| ?- debugger::spy(_, _, _, foo(_, _)).

Spy point set.
yes
```

The method `nospy/4` may be used to remove all matching spy points. For example, the call:

```
| ?- debugger::nospy(_, _, foo, _).

All matching context spy points removed.
yes
```

will remove all context spy points where the value of *self* matches the atom `foo`.

Removing all spy points

We may remove all line number, predicate, and context spy points by using the `nospyall/0` predicate:

```
| ?- debugger::nospyall.

All line number spy points removed.
All predicate spy points removed.
All context spy points removed.
yes
```

Tracing program execution

Logtalk allows tracing of execution for all objects compiled in debug mode. To start the debugger in trace mode, write:

```
| ?- debugger::trace.

yes
```

Note that, when tracing, spy points will be ignored. While tracing, the debugger will pause for user input at each leashed port, printing an informative message with the port name and the current goal. Before the port number, when a spy point is set for the current clause or goal, the debugger will print a # character for line number spy points, a + character for predicate spy points, and a * character for context spy points. The debugger also provides determinism information by prefixing the exit port with a * character when a call succeeds with choice-points pending. After the port name, the debugger prints the goal invocation number. This invocation number is unique and can be used to correlate the port trace messages.

To stop tracing and turning off the debugger, write:

```
| ?- debugger::notrace.

yes
```

Debugging using spy points

Tracing a program execution may generate large amounts of debugging data. Debugging using spy points allows the user to concentrate its attention in specific points of its code. To start a debugging session using spy points, write:

```
| ?- debugger::debug.

yes
```

At the beginning of a port description, the debugger will print a #, +, or * character before the current goal if there is, respectively, a line number, a predicate, or a context spy point defined.

To stop the debugger, write:

```
| ?- debugger::nodebug.

yes
```

Note that stopping the debugger does not remove any defined spy points.

Debugging commands

The debugger pauses at leashed ports when tracing or when finding a spy point for user interaction. The commands available are as follows:

c — creep go on; you may use the spacebar, return, or enter keys in alternative

l — leap continues execution until the next spy point is found

s — skip skips debugging for the current goal; valid at call, redo, and unification ports

q — quasi-skip skips debugging until returning to the current goal or reaching a spy point; valid at call and redo ports

r — retry retries the current goal but side-effects are not undone; valid at the fail port

j — jump reads invocation number and continues execution until a port is reached for that number

z — zap reads port name and continues execution until that port is reached reads negated port name and continues execution until a port other than the negated port is reached

i — ignore ignores goal, assumes that it succeeded; valid at call and redo ports

f — fail forces backtracking; may also be used to convert an exception into a failure

n — nodebug turns off debugging

@ — command; ! can be used in alternative reads and executes a query

b — break suspends execution and starts new interpreter; type `end_of_file` to terminate

a — abort returns to top level interpreter

Q — quit quits Logtalk

p — print writes current goal using the print/1 predicate if available

d — display writes current goal without using operator notation

w — write writes current goal quoting atoms if necessary

$ — dollar outputs the compiled form of the current goal (for low-level debugging)

x — context prints execution context

. — file prints file, entity, predicate, and line number information at an unification port

e — exception prints exception term thrown by the current goal

= — debugging prints debugging information

< — write depth sets the write term depth (set to 0 to reset)

* — add adds a context spy point for the current goal

/ — remove removes a context spy point for the current goal

+ — add adds a predicate spy point for the current goal

- — remove removes a predicate spy point for the current goal

 # — add adds a line number spy point for the current clause

| — remove removes a line number spy point for the current clause

h — condensed help prints list of command options

? — extended help prints list of command options

Context-switching calls

Logtalk provides a control construct, `<</2`, which allows the execution of a query within the context of an object. Common debugging uses include checking an object local predicates (e.g. predicates representing internal dynamic state) and sending a message from within an object. This control construct may also be used to write unit tests.

Consider the following toy example:

```
:- object(broken).

    :- public(a/1).

    a(A) :- b(A, B), c(B).
    b(1, 2). b(2, 4). b(3, 6).
    c(3).

:- end_object.
```

Something is wrong when we try the object public predicate, `a/1`:

```
| ?- broken::a(A).

no
```

For helping diagnosing the problem, instead of compiling the object in debug mode and doing a *trace* of the query to check the clauses for the non-public predicates, we can instead simply type:

```
| ?- broken << c(C).

C = 3
yes
```

The `<</2` control construct works by switching the execution context to the object in the first argument and then compiling and executing the second argument within that context:

```
| ?- broken << (self(Self), sender(Sender), this(This)).

Self = broken
Sender = broken
This = broken

yes
```

As exemplified above, the `<</2` control construct allows you to call an object local and private predicates. However, it is important to stress that we are not bypassing or defeating an object predicate scope directives. The calls take place within the context of the specified object, not within the context of the object making the `<</2` call. Thus, the `<</2` control construct implements a form of *execution-context switching*.

The availability of the `<</2` control construct is controlled by the compiler flag `context_switching_calls` (its default value is defined in the adapter files of the back-end Prolog compilers).

Using compilation hooks and term expansion for debugging

It is possible to use compilation hooks and the term expansion mechanism for conditional compilation of debugging goals. Assume that we chose the predicate `debug/1` to represent debug goals. For example:

```
append([], List, List) :-
    debug((write('Base case: '), writeq(append([], List, List)), nl)).
append([Head| Tail], List, [Head| Tail2]) :-
    debug((write('Recursive case: '), writeq(append(Tail, List, Tail2)), nl)),
    append(Tail, List, Tail2).
```

When debugging, we want to call the argument of the predicate debug/1. This can be easily accomplished by defining a hook object containing the following definition for goal_expansion/2:

```
goal_expansion(debug(Goal), Goal).
```

When not debugging, we can use a second hook object to discard the debug/1 calls by defining the predicate goal_expansion/2 as follows:

```
goal_expansion(debug(_), true).
```

The Logtalk compiler automatically removes any redundant calls to the built-in predicate true/0 when compiling object predicates.

Debugging messages

Calls to the logtalk::print_message/3 predicate where the message kind is either debug or debug(_) are only printed, by default, when the debug flag is turned on. Note that using these messages does not require compiling the code in debug mode, only turning on the flag. To avoid having to define message_tokens//2 grammar rules for translating each debug message, Logtalk provides default tokenization for four *meta-messages* that cover the most common cases:

@Message By default, the message is printed as passed to the write/1 predicate followed by a newline.

Key-Value By default, the message is printed as Key: Value followed by a newline. The value is printed as passed to the writeq/1 predicate.

List By default, the list items are printed indented one per line. The items are preceded by a dash and printed as passed to the writeq/1 predicate.

Title::List By default, the title is printed followed by a newline and the indented list items, one per line. The items are preceded by a dash and printed as passed to the writeq/1 predicate.

These print messages goals can always be combined with hooks as described in the previous section to remove them in production ready code.

Some simple examples of using these meta-messages:

```
| ?- logtalk::print_message(debug, core, @'Phase 1 completed').
yes

| ?- set_logtalk_flag(debug, on).
yes

| ?- logtalk::print_message(debug, core, @'Phase 1 completed').
>>> Phase 1 completed
yes

| ?- logtalk::print_message(debug, core, answer-42).
>>> answer: 42
yes
```

```
| ?- logtalk::print_message(debug, core, ['Arthur Dent','Ford Prefect','Marvin']).
>>> - 'Arthur Dent'
>>> - 'Ford Prefect'
>>> - 'Marvin'
yes

| ?- logtalk::print_message(debug, core, 'Names'::['Arthur Dent','Ford Prefect',
↪'Marvin']).
>>> Names:
>>> - 'Arthur Dent'
>>> - 'Ford Prefect'
>>> - 'Marvin'
yes
```

1.15 Prolog Integration and Migration Guide

An application may include plain Prolog files, Prolog modules, and Logtalk objects. This is a perfectly valid way of developing a complex application and, in some cases, it might be the most appropriated solution. Modules may be used for legacy code or when a simple encapsulation mechanism is adequate. Logtalk objects may be used when more powerful encapsulation, abstraction, and reuse features are necessary. Logtalk supports the compilation of source files containing both plain Prolog and Prolog modules. This guide provides tips for integrating and migrating plain Prolog code and Prolog module code to Logtalk. Step-by-step instructions are provided for encapsulating plain Prolog code in objects, converting Prolog modules into objects, and compiling and reusing Prolog modules as objects from inside Logtalk. An interesting application of the techniques described in this guide is a solution for running a Prolog application which uses modules on a Prolog compiler with no module system.

1.15.1 Source files with both Prolog code and Logtalk code

Logtalk source files may contain plain Prolog code intermixed with Logtalk code. The Logtalk compiler simply copies the plain Prolog code as-is to the generated Prolog file. With Prolog modules, it is assumed that the module code starts with a `module/1-2` directive and ends at the end of the file. There is no module ending directive which would allowed us to define more than one module per file. In fact, most if not all Prolog module systems always define a single module per file. Some of them mandate that the `module/1-2` directive be the first term on a source file. As such, when the Logtalk compiler finds a `module/1-2` directive, it assumes that all code that follows until the end of the file belongs to the module.

1.15.2 Encapsulating plain Prolog code in objects

Most applications consist of several plain Prolog source files, each one defining a few top-level predicates and auxiliary predicates that are not meant to be directly called by the user. Encapsulating plain Prolog code in objects allows us to make clear the different roles of each predicate, to hide implementation details, to prevent auxiliary predicates from being called outside the object, and to take advantage of Logtalk advanced code encapsulating and reusing features.

Encapsulating Prolog code using Logtalk objects is simple. First, for each source file, add an opening object directive, *object/1-5*, to the beginning of the file and an ending object directive, *end_object/0*, to end of the file. Choose an object name that reflects the purpose of source file code (this is a good opportunity for code refactoring if necessary). Second, add public predicate directives, *public/1*, for the top-level predicates that are used directly by the user or called from other source files. Third, we need to be able to call from inside an object predicates defined in other source files/objects. The easiest solution, which has the advantage of not requiring any changes to the predicate definitions, is to use the *uses/2* directive. If your Prolog compiler supports cross-referencing tools, you may use them to help you make sure

that all calls to predicates on other source files/objects are listed in the *uses/2* directives. In alternative, compiling the resulting objects with the Logtalk `unknown_predicates` and `portability` flags set to `warning` will help you identify calls to predicates defined on other converted source files and possible portability issues.

Prolog multifile predicates

Prolog *multifile* predicates are used when clauses for the same predicate are spread among several source files. When encapsulating plain Prolog code that uses multifile predicates, is often the case that the clauses of the multifile predicates get spread between different objects and categories but conversion is straight-forward. In the Logtalk object (or category) holding the multifile predicate *primary declaration*, add a *predicate scope directive* and a *multifile/1* directive. In all other objects (or categories) defining clauses for the multifile predicate, add a `multifile/1` directive and predicate clauses using the format:

```
:- multifile(Entity::Name/Arity).

Entity::Functor(...) :-
    ...
```

See the User Manual section on the `multifile/1` predicate directive for more information. An alternative solution is to simply keep the clauses for the multifile predicates as plain Prolog code and define, if necessary, a parametric object to encapsulate all predicates working with the multifile predicate clauses. For example, assume the following `multifile/1` directive:

```
% city(Name, District, Population, Neighbors)
:- multifile(city/4).
```

We can define a parametric object with `city/4` as its identifier:

```
:- object(city(_Name, _District, _Population, _Neighbors)).

    % predicates for working with city/4 clauses

:- end_object.
```

This solution is preferred when the multifile predicates are used to represent large tables of data. See the section on *Parametric objects* for more details.

1.15.3 Converting Prolog modules into objects

Converting Prolog modules into objects may allow an application to run on a wider range of Prolog compilers, overcoming compatibility problems. Some Prolog compilers don't support a module system. Among those Prolog compilers which support a module system, the lack of standardization leads to several issues, specially with semantics, operators, and meta-predicates. In addition, the conversion allows you to take advantage of Logtalk more powerful abstraction and reuse mechanisms such as separation between interface from implementation, inheritance, parametric objects, and categories.

Converting a Prolog module into an object is easy as long as the directives used in the module are supported by Logtalk (see below). Assuming that this is the case, apply the following steps:

1. Convert the module `module/1` directive into an opening object directive, *object/1-5*, using the module name as the object name. For `module/2` directives apply the same conversion and convert the list of exported predicates into Logtalk *public/1* predicate directives.

2. Add a closing object directive, *end_object/0*, at the end of the module code.

3. Convert any `export/1` directives into `public/1` predicate directives.

4. Convert any `use_module/1` directives into `use_module/2` directives (see next section).

5. Convert any `use_module/2` directives referencing other modules also being converted to objects into Logtalk *uses/2* directives. If the referenced modules are not being converted into objects, simply keep the `use_module/2` directives unchanged.

6. Convert any `meta_predicate/1` directives into Logtalk *meta_predicate/1* directives by replacing the module meta-argument indicator, `:`, with the Logtalk meta-argument indicator, `0`. Closures must be represented using an integer denoting the number of additional arguments that will be appended to construct a goal. Arguments which are not meta-arguments are represented by the `*` character.

7. Convert any explicit qualified calls to module predicates to messages by replacing the `:/2` operator with the *::/2* message sending operator, assuming that the referenced modules are also being converted into objects. Calls in the pseudo-module `user` can simply be encapsulated using the *{}/1* Logtalk external call control construct. You can also use instead an *uses/2* directive where the first argument would be the atom `user` and the second argument a list of all external predicates. This alternative has the advantage of not requiring changes to the code making the predicate calls.

8. If your module uses the database built-in predicates to implement module local mutable state using dynamic predicates, add both *private/1* and *dynamic/1* directives for each dynamic predicate.

9. If your module declares or defines clauses for multifile module predicates, replace the `:/2` functor by `::/2` in the `multifile/1` directives and in the clause heads (assuming that all modules defining the multifile predicates are converted into objects; if that is not the case, just keep the `multifile/1` directives and the clause heads as-is).

10. Compile the resulting objects with the Logtalk `unknown_predicates`, and `portability` flags set to `warning` to help you locate possible issues and calls to proprietary Prolog built-in predicates and to predicates defined on other converted modules. In order to improve code portability, check the Logtalk library for possible alternatives to the use of proprietary Prolog built-in predicates.

Before converting your modules to objects, you may try to compile them first as objects (using the *logtalk_compile/1* Logtalk built-in predicates) to help identify any issues that must be dealt with when doing the conversion to objects. Note that Logtalk supports compiling Prolog files as Logtalk source code without requiring changes to the file name extensions.

1.15.4 Compiling Prolog modules as objects

An alternative to convert Prolog modules into objects is to just compile the Prolog source files using the `logtalk_load/1-2` and `logtalk_compile/1-2` predicates (set the Logtalk `portability` flag set to `warning` to help you catch any unnoticed cross-module predicate calls). This allows you to reuse existing module code as objects. This has the advantage of requiring little if any code changes. There are, however, some limitations that you must be aware. These limitations are a consequence of the lack of standardization of Prolog module systems.

Supported module directives

Currently, Logtalk supports the following module directives:

module/1 The module name becomes the object name.

module/2 The module name becomes the object name. The exported predicates become public object predicates. The exported grammar rule non-terminals become public grammar rule non-terminals. The exported operators become public object operators but are not active elsewhere when loading the code.

use_module/2 This directive is compiled as a Logtalk *uses/2* directive in order to ensure correct compilation of the module predicate clauses. The first argument of this directive must be the module **name** (an atom), not a module file specification (the adapter files attempt to use the Prolog dialect level term-expansion mechanism

to find the module name from the module file specification). Note that the module is not automatically loaded by Logtalk (as it would be when compiling the directive using Prolog instead of Logtalk; the programmer may also want the specified module to be compiled as an object). The second argument must be a predicate indicator (`Name/Arity`), a grammar rule non-terminal indicator (`Name//Arity`), a operator declaration, or a list of predicate indicators, grammar rule non-terminal indicators, and operator declarations.

export/1 Exported predicates are compiled as public object predicates. The argument must be a predicate indicator (`Name/Arity`), a grammar rule non-terminal indicator (`Name//Arity`), an operator declaration, or a list of predicate indicators, grammar rule non-terminal indicators, and operator declarations.

reexport/2 Reexported predicates are compiled as public object predicates. The first argument is the module name. The second argument must be a predicate indicator (`Name/Arity`), a grammar rule non-terminal indicator (`Name//Arity`), an operator declaration, or a list of predicate indicators, grammar rule non-terminal indicators, and operator declarations.

meta_predicate/1 Module meta-predicates become object meta-predicates. Only predicate arguments marked as goals or closures (using an integer) are interpreted as meta-arguments. In addition, Prolog module meta-predicates and Logtalk meta-predicates don't share the same explicit-qualification calling semantics: in Logtalk, meta-arguments are always called in the context of the *sender*.

A common issue when compiling modules as objects is the use of the atoms `dynamic`, `discontiguous`, and `multifile` as operators in directives. For better portability avoid this usage. For example, write:

```
:- dynamic([foo/1, bar/2]).
```

instead of:

```
:- dynamic foo/1, bar/2.
```

Another common issue is missing `meta_predicate/1`, `dynamic/1`, `discontiguous/1`, and `multifile/1` predicate directives. Logtalk allows detection of most missing directives (by setting its `missing_directives` flag to `warning`).

When compiling modules as objects, you probably don't need event support turned on. You may use the compiler flag `events(deny)` with the Logtalk compiling and loading built-in methods for a small performance gain for the compiled code.

Current limitations and workarounds

The `reexport/1` and `use_module/1` directives are not directly supported by the Logtalk compiler. But most Prolog adapter files provide support for compiling these directives using Logtalk's first stage of its term-expansion mechanism. Nevertheless, these directives can be converted, respectively, into `reexport/2` and `use_module/2` directives by finding which predicates exported by the specified modules are reexported or imported into the module containing the directive. Finding the names of the imported predicates that are actually used is easy. First, comment out the `use_module/1` directives and compile the file (making sure that the compiler flag `unknown_predicates` is set to `warning`). Logtalk will print a warning with a list of predicates that are called but never defined. Second, use these list to replace the `reexport/1` and `use_module/1` directives by, respectively, `reexport/2` and `use_module/2` directives. You should then be able to compile the modified Prolog module as an object.

Although Logtalk supports term and goal expansion mechanisms, the semantics are different from similar mechanisms found in some Prolog compilers. In particular, Logtalk does not support defining term and goal expansions clauses in a source file for expanding the source file itself. Logtalk forces a clean separation between expansions clauses and the source files that will be subject to source-to-source expansions by using *hook objects*.

1.15.5 Dealing with proprietary Prolog directives and predicates

Most Prolog compilers define proprietary, non-standard, directives and predicates that may be used in both plain code and module code. Non-standard Prolog built-in predicates are usually not problematic, as Logtalk is usually able to identify and compile them correctly (but see the notes on built-in meta-predicates for possible caveats). However, Logtalk will generate compilation errors on source files containing proprietary directives unless you first specify how the directives should be handled. Several actions are possible on a per-directive basis: ignoring the directive (i.e. do not copy the directive, although a goal can be proved as a consequence), rewriting and copy the directive to the generated Prolog files, or rewriting and recompiling the resulting directive. To specify these actions, the adapter files contain clauses for the `'$lgt_prolog_term_expansion'`/2 predicate. For example, assume that a given Prolog compiler defines a `comment/2` directive for predicates using the format:

```
:- comment(foo/2, "Brief description of the predicate").
```

We can rewrite this predicate into a Logtalk `info/2` directive by defining a suitable clause for the `'$lgt_prolog_term_expansion'`/2 predicate:

```
'$lgt_prolog_term_expansion'(comment(F/A, String), info(F/A, [comment is Atom])) :-
    atom_codes(Atom, String).
```

This Logtalk feature can be used to allow compilation of legacy Prolog code without the need of changing the sources. When used, is advisable to set the `portability/1` compiler flag to `warning` in order to more easily identify source files that are likely non-portable across Prolog compilers.

A second example, where a proprietary Prolog directive is discarded after triggering a side effect:

```
'$lgt_prolog_term_expansion'(load_foreign_files(Files,Libs,InitRoutine), []) :-
    load_foreign_files(Files,Libs,InitRoutine).
```

In this case, although the directive is not copied to the generated Prolog file, the foreign library files are loaded as a side effect of the Logtalk compiler calling the `'$lgt_prolog_term_expansion'`/2 hook predicate.

1.15.6 Calling Prolog module predicates

Prolog module predicates can be called from within objects or categories by simply using explicit module qualification, i.e. by writing `Module:Goal` or `Goal@Module` (depending on the module system). Logtalk also supports the use of `use_module/2` directives in object and categories (with the restriction that the first argument of the directive must be the actual module name and not the module file name or the module file path). In this case, these directives are parsed in a similar way to Logtalk *uses/2* directives, with calls to the specified module predicates being automatically translated to `Module:Goal` calls. For example, assume a `clpfd` Prolog module implementing a finite domain constraint solver. You could write:

```
:- object(puzzle).

    :- public(puzzle/1).

    :- use_module(clpfd, [
        all_different/1, ins/2, label/1,
        (#=)/2, (#\=)/2,
        op(700, xfx, #=), op(700, xfx, #\=)
    ]).

    puzzle([S,E,N,D] + [M,O,R,E] = [M,O,N,E,Y]) :-
        Vars = [S,E,N,D,M,O,R,Y],
        Vars ins 0..9,
```

```
        all_different(Vars),
                S*1000 + E*100 + N*10 + D +
                M*1000 + O*100 + R*10 + E #=
        M*10000 + O*1000 + N*100 + E*10 + Y,
        M #\= 0, S #\= 0,
        label([M,O,N,E,Y]).

:- end_object.
```

As a general rule, the Prolog modules should be loaded (e.g. in the auxiliary Logtalk loader files) *before* compiling objects that make use of module predicates. Moreover, the Logtalk compiler does not generate code for the automatic loading of modules referenced in `use_module/1-2` directives. This is a consequence of the lack of standardization of these directives, whose first argument can be a module name, a straight file name, or a file name using some kind of library notation, depending on the back-end Prolog compiler. Worse, modules are sometimes defined in files with names different from the module names requiring finding, opening, and reading the file in order to find the actual module name.

Logtalk supports the declaration of predicate aliases in `use_module/2` directives used within object and categories. For example, the ECLiPSe IC Constraint Solvers define a `::/2` variable domain operator that clashes with the Logtalk `::/2` message sending operator. We can solve the conflict by writing:

```
:- use_module(ic, [(::)/2 as ins/2]).
```

With this directive, calls to the `ins/2` predicate alias will be automatically compiled by Logtalk to calls to the `::/2` predicate in the `ic` module.

When calling Prolog module meta-predicates, the Logtalk compiler may need help to understand the corresponding meta-predicate template. Despite some recent progress in standardization of the syntax of `meta_predicate/1` directives and of the `meta_predicate/1` property returned by the `predicate_property/2` reflection predicate, portability is still a problem. Thus, Logtalk allows the original `meta_predicate/1` directive to be overridden with a local one that Logtalk can make sense of. Note that the Logtalk library provides implementations of common meta-predicates, which can be used in place of module meta-predicates.

Logtalk allows you to send a message to a module in order to call one of its predicates. This is usually not advised as it implies a performance penalty when compared to just using the `Module:Call` notation. Moreover, this works only if there is no object with the same name as the module you are targeting. This feature is necessary, however, in order to properly support compilation of modules containing `use_module/2` directives as objects. If the modules specified in the `use_module/2` directives are not compiled as objects but are instead loaded as-is by Prolog, the exported predicates would need to be called using the `Module:Call` notation but the converted module will be calling them through message sending. Thus, this feature ensures that, on a module compiled as an object, any predicate calling other module predicates will work as expected either these other modules are loaded as-is or also compiled as objects.

1.15.7 Compiling Prolog module multifile predicates

Some Prolog module libraries, e.g. constraint packages, expect clauses for some library predicates to be defined in other modules. This is accomplished by declaring the library predicate *multifile* and by explicitly prefixing predicate clause heads with the library module identifier. For example:

```
:- multifile(clpfd:run_propagator/2).
clpfd:run_propagator(..., ...) :-
    ...
```

Logtalk supports the compilation of such clauses within objects and categories. While the clause head is compiled as-is, the clause body is compiled in the same way as a regular object or category predicate, thus allowing calls to local object or category predicates. For example:

```
:- object(...).

    :- multifile(clpfd:run_propagator/2).
    clpfd:run_propagator(..., ...) :-
        ...      % calls to local object predicates

:- end_object.
```

The Logtalk compiler will print a warning if the `multifile/1` directive is missing. These multifile predicates may also be declared dynamic using the same `Module:Name/Arity` notation.

REFERENCE MANUAL

2.1 Grammar

The Logtalk grammar is here described using Backus-Naur Form syntax. Non-terminal symbols in *italics* have the definition found in the ISO Prolog Core standard. Terminal symbols are represented in a `fixed width font` and between double-quotes.

2.1.1 Entities

entity ::=
>   object |
>   category |
>   protocol

2.1.2 Object definition

object ::=
>   begin_object_directive [object_terms] end_object_directive.

begin_object_directive ::=
>   ":- `object(`" object_identifier ["`,`" object_relations] "`) .`"

end_object_directive ::=
>   ":- `end_object.`"

object_relations ::=
>   prototype_relations |
>   non_prototype_relations

prototype_relations ::=
>   prototype_relation |
>   prototype_relation "`,`" prototype_relations

prototype_relation ::=
      implements_protocols |
      imports_categories |
      extends_objects

non_prototype_relations ::=
      non_prototype_relation |
      non_prototype_relation ", " non_prototype_relations

non_prototype_relation ::=
      implements_protocols |
      imports_categories |
      instantiates_classes |
      specializes_classes

2.1.3 Category definition

category ::=
      begin_category_directive [category_terms] end_category_directive.

begin_category_directive ::=
      ":- category(" category_identifier ["," category_relations] ") ."

end_category_directive ::=
      ":- end_category."

category_relations ::=
      category_relation |
      category_relation ", " category_relations

category_relation ::=
      implements_protocols |
      extends_categories |
      complements_objects

2.1.4 Protocol definition

protocol ::=
      begin_protocol_directive [protocol_directives]
      end_protocol_directive.

begin_protocol_directive ::=

> ```
> ":- protocol(" protocol_identifier [ "," extends_protocols ] ")."
> ```

end_protocol_directive ::=
> ```
> ":- end_protocol."
> ```

2.1.5 Entity relations

extends_protocols ::=
> "`extends(`" extended_protocols "`)`"

extends_objects ::=
> "`extends(`" extended_objects "`)`"

extends_categories ::=
> "`extends(`" extended_categories "`)`"

implements_protocols ::=
> "`implements(`" implemented_protocols "`)`"

imports_categories ::=
> "`imports(`" imported_categories "`)`"

instantiates_classes ::=
> "`instantiates(`" instantiated_objects "`)`"

specializes_classes ::=
> "`specializes(`" specialized_objects "`)`"

complements_objects ::=
> "`complements(`" complemented_objects "`)`"

Implemented protocols

implemented_protocols ::=
> implemented_protocol |
> implemented_protocol_sequence |
> implemented_protocol_list

implemented_protocol ::=
> protocol_identifier |
> scope "`::`" protocol_identifier

implemented_protocol_sequence ::=
      implemented_protocol |
      implemented_protocol "," implemented_protocol_sequence

implemented_protocol_list ::=
      "[" implemented_protocol_sequence "]"

Extended protocols

extended_protocols ::=
      extended_protocol |
      extended_protocol_sequence |
      extended_protocol_list

extended_protocol ::=
      protocol_identifier |
      scope "::" protocol_identifier

extended_protocol_sequence ::=
      extended_protocol |
      extended_protocol "," extended_protocol_sequence

extended_protocol_list ::=
      "[" extended_protocol_sequence "]"

Imported categories

imported_categories ::=
      imported_category |
      imported_category_sequence |
      imported_category_list

imported_category ::=
      category_identifier |
      scope "::" category_identifier

imported_category_sequence ::=
      imported_category |
      imported_category "," imported_category_sequence

imported_category_list ::=
      "[" imported_category_sequence "]"

Extended objects

extended_objects ::=
> extended_object |
> extended_object_sequence |
> extended_object_list

extended_object ::=
> object_identifier |
> scope "::" object_identifier

extended_object_sequence ::=
> extended_object |
> extended_object "," extended_object_sequence

extended_object_list ::=
> "[" extended_object_sequence "]"

Extended categories

extended_categories ::=
> extended_category |
> extended_category_sequence |
> extended_category_list

extended_category ::=
> category_identifier |
> scope "::" category_identifier

extended_category_sequence ::=
> extended_category |
> extended_category "," extended_category_sequence

extended_category_list ::=
> "[" extended_category_sequence "]"

Instantiated objects

instantiated_objects ::=
> instantiated_object |
> instantiated_object_sequence |
> instantiated_object_list

instantiated_object ::=
      object_identifier |
      scope ":" object_identifier

instantiated_object_sequence ::=
      instantiated_object
      instantiated_object "," instantiated_object_sequence |

instantiated_object_list ::=
      "[" instantiated_object_sequence "]"

Specialized objects

specialized_objects ::=
      specialized_object |
      specialized_object_sequence |
      specialized_object_list

specialized_object ::=
      object_identifier |
      scope ":" object_identifier

specialized_object_sequence ::=
      specialized_object |
      specialized_object "," specialized_object_sequence

specialized_object_list ::=
      "[" specialized_object_sequence "]"

Complemented objects

complemented_objects ::=
      object_identifier |
      complemented_object_sequence |
      complemented_object_list

complemented_object_sequence ::=
      object_identifier |
      object_identifier "," complemented_object_sequence

complemented_object_list ::=
      "[" complemented_object_sequence "]"

Entity and predicate scope

scope ::=
    "public" |
    "protected" |
    "private"

2.1.6 Entity identifiers

entity_identifiers ::=
    entity_identifier |
    entity_identifier_sequence |
    entity_identifier_list

entity_identifier ::=
    object_identifier |
    protocol_identifier |
    category_identifier

entity_identifier_sequence ::=
    entity_identifier |
    entity_identifier "," entity_identifier_sequence

entity_identifier_list ::=
    "[" entity_identifier_sequence "]"

Object identifiers

object_identifiers ::=
    object_identifier |
    object_identifier_sequence |
    object_identifier_list

object_identifier ::=
    atom |
    compound

object_identifier_sequence ::=
    object_identifier |
    object_identifier "," object_identifier_sequence

object_identifier_list ::=
    "[" object_identifier_sequence "]"

Category identifiers

category_identifiers ::=
      category_identifier |
      category_identifier_sequence |
      category_identifier_list

category_identifier ::=
      atom |
      compound

category_identifier_sequence ::=
      category_identifier |
      category_identifier "," category_identifier_sequence

category_identifier_list ::=
      "[" category_identifier_sequence "]"

Protocol identifiers

protocol_identifiers ::=
      protocol_identifier |
      protocol_identifier_sequence |
      protocol_identifier_list

protocol_identifier ::=
      atom

protocol_identifier_sequence ::=
      protocol_identifier |
      protocol_identifier "," protocol_identifier_sequence

protocol_identifier_list ::=
      "[" protocol_identifier_sequence "]"

Module identifiers

module_identifier ::=
      atom

2.1.7 Source file names

source_file_names ::=

source_file_name |
source_file_name_list

source_file_name ::=
    atom |
    library_source_file_name

library_source_file_name ::=
    library_name " (" *atom* ") "

library_name ::=
    atom

source_file_name_sequence ::=
    source_file_name |
    source_file_name "," source_file_name_sequence

source_file_name_list ::=
    "[" source_file_name_sequence "]"

2.1.8 Terms

Object terms

object_terms ::=
    object_term |
    object_term object_terms

object_term ::=
    object_directive |
    clause |
    grammar_rule

Category terms

category_terms ::=
    category_term |
    category_term category_terms

category_term ::=
    category_directive |
    clause |

grammar_rule

2.1.9 Directives

Source file directives

source_file_directives ::=
    source_file_directive |
    source_file_directive source_file_directives

source_file_directive ::=
    ":- `encoding(`" *atom* "`).`" |
    ":- `set_logtalk_flag(`" *atom* "`,`" *nonvar* "`).`" |
    ":- `include(`" source_file_name "`).`"
    Prolog directives

Conditional compilation directives

conditional_compilation_directives ::=
    conditional_compilation_directive |
    conditional_compilation_directive conditional_compilation_directives

conditional_compilation_directive ::=
    ":- `if(`" *callable* "`).`" |
    ":- `elif(`" *callable* "`).`" |
    ":- `else.`" |
    ":- `endif.`"

Object directives

object_directives ::=
    object_directive |
    object_directive object_directives

object_directive ::=
    ":- `initialization(`" *callable* "`).`" |
    ":- `built_in.`" |
    ":- `threaded.`" |
    ":- `dynamic.`" |
    ":- `info(`" entity_info_list "`).`" |
    ":- `set_logtalk_flag(`" *atom* "`,`" *nonvar* "`).`" |
    ":- `include(`" source_file_name "`).`" |
    predicate_directives

Category directives

category_directives ::=
>   category_directive |
>   category_directive category_directives

category_directive ::=
>   ":- built_in."|
>   ":- dynamic."|
>   ":- info(" entity_info_list ")."|
>   ":- set_logtalk_flag(" *atom* "," *nonvar* ")."|
>   ":- include(" source_file_name ")."|
>   predicate_directives

Protocol directives

protocol_directives ::=
>   protocol_directive |
>   protocol_directive protocol_directives

protocol_directive ::=
>   ":- built_in."|
>   ":- dynamic."|
>   ":- info(" entity_info_list ")."|
>   ":- set_logtalk_flag(" *atom* "," *nonvar* ")."|
>   ":- include(" source_file_name ")."|
>   predicate_directives

Predicate directives

predicate_directives ::=
>   predicate_directive |
>   predicate_directive predicate_directives

predicate_directive ::=
>   alias_directive |
>   synchronized_directive |
>   uses_directive |
>   use_module_directive |
>   scope_directive |
>   mode_directive |
>   meta_predicate_directive |
>   meta_non_terminal_directive |
>   info_directive |
>   dynamic_directive |

discontiguous_directive |
multifile_directive |
coinductive_directive |
operator_directive

alias_directive ::=
    ":- `alias`(" entity_identifier "," predicate_indicator_alias_list ") ." |
    ":- `alias`(" entity_identifier "," non_terminal_indicator_alias_list ") ."

synchronized_directive ::=
    ":- `synchronized`(" predicate_indicator_term | non_terminal_indicator_term ") ."

uses_directive ::=
    ":- `uses`(" object_identifier "," predicate_indicator_alias_list ") ."

use_module_directive ::=
    ":- `use_module`(" module_identifier "," module_predicate_indicator_alias_list ") ." |

scope_directive ::=
    ":- `public`(" predicate_indicator_term | non_terminal_indicator_term ") ." |
    ":- `protected`(" predicate_indicator_term | non_terminal_indicator_term ") ." |
    ":- `private`(" predicate_indicator_term | non_terminal_indicator_term ") ."

mode_directive ::=
    ":- `mode`(" predicate_mode_term | non_terminal_mode_term "," number_of_proofs ") ."

meta_predicate_directive ::=
    ":- `meta_predicate`(" meta_predicate_template_term ") ."

meta_non_terminal_directive ::=
    ":- `meta_non_terminal`(" meta_non_terminal_template_term ") ."

info_directive ::=
    ":- `info`(" predicate_indicator | non_terminal_indicator "," predicate_info_list ") ."

dynamic_directive ::=
    ":- `dynamic`(" qualified_predicate_indicator_term | qualified_non_terminal_indicator_term ") ."

discontiguous_directive ::=
    ":- `discontiguous`(" predicate_indicator_term |

non_terminal_indicator_term ") ."

multifile_directive ::=
    ":- multifile(" qualified_predicate_indicator_term |
    qualified_non_terminal_indicator_term ") ."

coinductive_directive ::=
    ":- coinductive(" predicate_indicator_term |
    coinductive_predicate_template_term ") ."

predicate_indicator_term ::=
    predicate_indicator |
    predicate_indicator_sequence |
    predicate_indicator_list

predicate_indicator_sequence ::=
    predicate_indicator |
    predicate_indicator "," predicate_indicator_sequence

predicate_indicator_list ::=
    "[" predicate_indicator_sequence "]"

qualified_predicate_indicator_term ::=
    qualified_predicate_indicator |
    qualified_predicate_indicator_sequence |
    qualified_predicate_indicator_list

qualified_predicate_indicator_sequence ::=
    qualified_predicate_indicator |
    qualified_predicate_indicator "," qualified_predicate_indicator_sequence

qualified_predicate_indicator_list ::=
    "[" qualified_predicate_indicator_sequence "]"

qualified_predicate_indicator ::=
    predicate_indicator |
    object_identifier "::" predicate_indicator |
    category_identifier "::" predicate_indicator |
    module_identifier ":" predicate_indicator

predicate_indicator_alias ::=

> *predicate_indicator* |
> *predicate_indicator* "`as`" *predicate_indicator* |
> *predicate_indicator* "`::`" *predicate_indicator* |
> *predicate_indicator* "`:`" *predicate_indicator*

predicate_indicator_alias_sequence ::=
> predicate_indicator_alias |
> predicate_indicator_alias "`,`" predicate_indicator_alias_sequence

predicate_indicator_alias_list ::=
> "`[`" predicate_indicator_alias_sequence "`]`"

module_predicate_indicator_alias ::=
> *predicate_indicator* |
> *predicate_indicator* "`as`" *predicate_indicator* |
> *predicate_indicator* "`:`" *predicate_indicator*

module_predicate_indicator_alias_sequence ::=
> module_predicate_indicator_alias |
> module_predicate_indicator_alias "`,`" module_predicate_indicator_alias_sequence

module_predicate_indicator_alias_list ::=
> "`[`" module_predicate_indicator_alias_sequence "`]`"

non_terminal_indicator_term ::=
> non_terminal_indicator |
> non_terminal_indicator_sequence |
> non_terminal_indicator_list

non_terminal_indicator_sequence ::=
> non_terminal_indicator |
> non_terminal_indicator "`,`" non_terminal_indicator_sequence

non_terminal_indicator_list ::=
> "`[`" non_terminal_indicator_sequence "`]`"

non_terminal_indicator ::=
> *functor* "`//`" *arity*

qualified_non_terminal_indicator_term ::=
> qualified_non_terminal_indicator |

      qualified_non_terminal_indicator_sequence |
      qualified_non_terminal_indicator_list

qualified_non_terminal_indicator_sequence ::=
      qualified_non_terminal_indicator |
      qualified_non_terminal_indicator "`,`" qualified_non_terminal_indicator_sequence

qualified_non_terminal_indicator_list ::=
      "`[`" qualified_non_terminal_indicator_sequence "`]`"

qualified_non_terminal_indicator ::=
      non_terminal_indicator |
      object_identifier "`::`" non_terminal_indicator |
      category_identifier "`::`" non_terminal_indicator |
      module_identifier "`:`" non_terminal_indicator

non_terminal_indicator_alias ::=
      non_terminal_indicator |
      non_terminal_indicator "`as`" *non_terminal_indicator*
      non_terminal_indicator "`::`" *non_terminal_indicator*

non_terminal_indicator_alias_sequence ::=
      non_terminal_indicator_alias |
      non_terminal_indicator_alias "`,`"
      non_terminal_indicator_alias_sequence

non_terminal_indicator_alias_list ::=
      "`[`" non_terminal_indicator_alias_sequence "`]`"

coinductive_predicate_template_term ::=
      coinductive_predicate_template |
      coinductive_predicate_template_sequence |
      coinductive_predicate_template_list

coinductive_predicate_template_sequence ::=
      coinductive_predicate_template |
      coinductive_predicate_template "`,`"
      coinductive_predicate_template_sequence

coinductive_predicate_template_list ::=
      "`[`" coinductive_predicate_template_sequence "`]`"

coinductive_predicate_template ::=
    atom " (" coinductive_mode_terms ") "

coinductive_mode_terms ::=
    coinductive_mode_term |
    coinductive_mode_terms ", " coinductive_mode_terms

coinductive_mode_term ::=
    "+" | "-"

predicate_mode_term ::=
    atom " (" mode_terms ") "

non_terminal_mode_term ::=
    atom " (" mode_terms ") "

mode_terms ::=
    mode_term |
    mode_term ", " mode_terms

mode_term ::=
    "@" [type] | "+" [type] | "-" [type] | "?" [
    type] |
    "++" [type] | "--" [type]

type ::=
    prolog_type | logtalk_type | user_defined_type

prolog_type ::=
    "term" | "nonvar" | "var" |
    "compound" | "ground" | "callable" | "list" |
    "atomic" | "atom" |
    "number" | "integer" | "float"

logtalk_type ::=
    "object" | "category" | "protocol" |
    "event"

user_defined_type ::=
    atom |
    compound

number_of_proofs ::=
    "`zero`"|"`zero_or_one`"|"`zero_or_more`"|"`one`"|
    "`one_or_more`"|"`one_or_error`"|"`error`"

meta_predicate_template_term ::=
    meta_predicate_template |
    meta_predicate_template_sequence |
    meta_predicate_template_list

meta_predicate_template_sequence ::=
    meta_predicate_template |
    meta_predicate_template "`,`" meta_predicate_template_sequence

meta_predicate_template_list ::=
    "`[`" meta_predicate_template_sequence "`]`"

meta_predicate_template ::=
    object_identifier "`::`" *atom* " (" meta_predicate_specifiers ") " |
    category_identifier "`::`" *atom* " (" meta_predicate_specifiers ") " |
    atom " (" meta_predicate_specifiers ") "

meta_predicate_specifiers ::=
    meta_predicate_specifier |
    meta_predicate_specifier "`,`" meta_predicate_specifiers

meta_predicate_specifier ::=
    non-negative integer | "`::`" | "^" |
    "`*`"

meta_non_terminal_template_term ::=
    meta_predicate_template_term

entity_info_list ::=
    "`[]`" |
    "`[`" entity_info_item "`is`" *nonvar* "|" entity_info_list
    "`]`"

entity_info_item ::=
    "`comment`"|"`remarks`"|
    "`author`"|"`version`"|"`date`"|
    "`copyright`"|"`license`"|

"parameters"|"parnames"|
"see_also"|
atom

predicate_info_list ::=
"[]"|
"[" predicate_info_item "is" *nonvar* "|" predicate_info_list "]"

predicate_info_item ::=
"comment"|"remarks"|
"arguments"|"argnames"|
"redefinition"|"allocation"|
"examples"|"exceptions"|
atom

2.1.10 Clauses and goals

clause ::=
object_identifier "::" *head* ":-" *body* |
module_identifier ":" *head* ":-" *body* |
head :- *body* |
fact

goal ::=
message_sending |
super_call |
external_call |
context_switching_call |
callable

message_sending ::=
message_to_object |
message_delegation |
message_to_self

message_to_object ::=
receiver "::" messages

message_delegation ::=
"[" message_to_object "]"

message_to_self ::=

"::" messages

super_call ::=
    "^^" message

messages ::=
    message |
    "(" message "," messages ")" |
    "(" message ";" messages ")" |
    "(" message "->" messages ")"

message ::=
    callable |
    variable

receiver ::=
    "{" *callable* "}" |
    object_identifier |
    variable

external_call ::=
    "{" *callable* "}"

context_switching_call ::=
    object_identifier "<<" *goal*

2.1.11 Lambda expressions

lambda_expression ::=
    lambda_free_variables "/" lambda_parameters ">>" *callable* |
    lambda_free_variables "/" *callable* |
    lambda_parameters ">>" *callable*

lambda_free_variables ::=
    "{" *conjunction of variables* "}" |
    "{" *variable* "}" |
    "{}"

lambda_parameters ::=
    list of terms |
    "[]"

2.1.12 Entity properties

category_property ::=

    "static"|

    "dynamic"|

    "built_in"|

    "file(" *atom* ")"|

    "file(" *atom* "," *atom* ")"|

    "lines(" *integer* "," *integer* ")"|

    "events"|

    "source_data"|

    "public(" predicate_indicator_list ")"|

    "protected(" predicate_indicator_list ")"|

    "private(" predicate_indicator_list ")"|

    "declares(" predicate_indicator "," predicate_declaration_property_list ")"|

    "defines(" predicate_indicator "," predicate_definition_property_list ")"|

    "includes(" predicate_indicator "," object_identifier|category_identifier ","
predicate_definition_property_list ")"|

    "provides(" predicate_indicator "," object_identifier|category_identifier ","
predicate_definition_property_list ")"|

    "alias(" predicate_indicator "," predicate_alias_property_list ")"|

    "calls(" predicate "," predicate_call_update_property_list ")"|

    "updates(" predicate "," predicate_call_update_property_list ")"|

    "number_of_clauses(" *integer* ")"|

    "number_of_rules(" *integer* ")"|

    "number_of_user_clauses(" *integer* ")"|

    "number_of_user_rules(" *integer* ")"|

    "debugging"

object_property ::=

    "static"|

    "dynamic"|

    "built_in"|

    "threaded"|

    "file(" *atom* ")"|

    "file(" *atom* "," *atom* ")"|

    "lines(" *integer* "," *integer* ")"|

    "context_switching_calls"|

    "dynamic_declarations"|

    "events"|

    "source_data"|

    "complements(" "allow"|"restrict" ")"|

    "complements"|

    "public(" predicate_indicator_list ")"|

    "protected(" predicate_indicator_list ")"|

    "private(" predicate_indicator_list ")"|

    "declares(" predicate_indicator "," predicate_declaration_property_list ")"|

"`defines`(" predicate_indicator "`,`" predicate_definition_property_list ")" |

"`includes`(" predicate_indicator "`,`" object_identifier | category_identifier "`,`"
predicate_definition_property_list ")" |

"`provides`(" predicate_indicator "`,`" object_identifier | category_identifier "`,`"
predicate_definition_property_list ")"

"`alias`(" predicate_indicator "`,`" predicate_alias_property_list ")" |

"`calls`(" predicate "`,`" predicate_call_update_property_list ")" |

"`updates`(" predicate "`,`" predicate_call_update_property_list ")" |

"`number_of_clauses`(" *integer* ")" |

"`number_of_rules`(" *integer* ")" |

"`number_of_user_clauses`(" *integer* ")"

"`number_of_user_rules`(" *integer* ")" |

"`module`|"

"`debugging`"

protocol_property ::=

"`static`" |

"`dynamic`" |

"`built_in`" |

"`source_data`" |

"`file`(" *atom* ")" |

"`file`(" *atom* "`,`" *atom* ")" |

"`lines`(" *integer* "`,`" *integer* ")" |

"`public`(" predicate_indicator_list ")" |

"`protected`(" predicate_indicator_list ")" |

"`private`(" predicate_indicator_list ")" |

"`declares`(" predicate_indicator "`,`" predicate_declaration_property_list ")" |

"`alias`(" predicate_indicator "`,`" predicate_alias_property_list ")" |

"`debugging`"

predicate_declaration_property_list ::=

"`[`" predicate_declaration_property_sequence "`]`"

predicate_declaration_property_sequence ::=

predicate_declaration_property |

predicate_declaration_property "`,`"
predicate_declaration_property_sequence

predicate_declaration_property ::=

"`static`" | "`dynamic`" |

"`scope`(" scope ")" |

"`private`" | "`protected`" | "`public`" |

"`coinductive`" |

"`multifile`" |

"`synchronized`" |

`"meta_predicate("` meta_predicate_template `")"` |
`"coinductive("` coinductive_predicate_template `")"` |
`"non_terminal("` non_terminal_indicator `")"` |
`"include("` *atom* `")"` |
`"line_count("` *integer* `")"` |
`"mode("` predicate_mode_term | non_terminal_mode_term `","` number_of_proofs `")"` |
`"info("` *list* `")"`

predicate_definition_property_list ::=
    `"["` predicate_definition_property_sequence `"]"`

predicate_definition_property_sequence ::=
    predicate_definition_property |
    predicate_definition_property `","`
    predicate_definition_property_sequence

predicate_definition_property ::=
    `"inline"` | `"auxiliary"` |
    `"non_terminal("` non_terminal_indicator `")"` |
    `"include("` *atom* `")"` |
    `"line_count("` *integer* `")"` |
    `"number_of_clauses("` *integer* `")"` |
    `"number_of_rules("` *integer* `")"`

predicate_alias_property_list ::=
    `"["` predicate_alias_property_sequence `"]"`

predicate_alias_property_sequence ::=
    predicate_alias_property |
    predicate_alias_property `","` predicate_alias_property_sequence

predicate_alias_property ::=
    `"for("` predicate_indicator `")"` |
    `"from("` entity_identifier `")"` |
    `"non_terminal("` non_terminal_indicator `")"` |
    `"include("` *atom* `")"` |
    `"line_count("` *integer* `")"`

predicate ::=
    predicate_indicator |
    `"^^"` predicate_indicator |
    `"::"` predicate_indicator |
    variable `"::"` predicate_indicator |

object_identifier "`::`" predicate_indicator |
variable "`:`" predicate_indicator |
module_identifier "`:`" predicate_indicator

predicate_call_update_property_list ::=
   "`[`" predicate_call_update_property_sequence "`]`"

predicate_call_update_property_sequence ::=
   predicate_call_update_property |
   predicate_call_update_property "`,`"
   predicate_call_update_property_sequence

predicate_call_update_property ::=
   "`caller(`" predicate_indicator "`)`" |
   "`include(`" *atom* "`)`" |
   "`line_count(`" *integer* "`)`" |
   "`as(`" predicate_indicator "`)`"

2.1.13 Predicate properties

predicate_property ::=
   "`static`" | "`dynamic`" |
   "`scope(`" scope "`)`" |
   "`private`" | "`protected`" | "`public`" |
   "`logtalk`" | "`prolog`" | "`foreign`" |
   "`coinductive(`" coinductive_predicate_template "`)`" |
   "`multifile`" |
   "`synchronized`" |
   "`built_in`" |
   "`inline`" |
   "`declared_in(`" entity_identifier "`)`" |
   "`defined_in(`" object_identifier | category_identifier "`)`" |
   "`redefined_from(`" object_identifier | category_identifier "`)`" |
   "`meta_predicate(`" meta_predicate_template "`)`" |
   "`alias_of(`" callable "`)`" |
   "`alias_declared_in(`" entity_identifier "`)`" |
   "`non_terminal(`" non_terminal_indicator "`)`" |
   "`mode(`" predicate_mode_term | non_terminal_mode_term "`,`" number_of_proofs "`)`" |
   "`info(`" *list* "`)`" |
   "`number_of_clauses(`" *integer* "`)`" |
   "`number_of_rules(`" *integer* "`)`" |
   "`declared_in(`" entity_identifier "`,`" *integer* "`)`" |
   "`defined_in(`" object_identifier | category_identifier "`,`" *integer*"`)`" |
   "`redefined_from(`" object_identifier | category_identifier "`,`" *integer*"`)`" |
   "`alias_declared_in(`" entity_identifier "`,`" *integer*"`)`"

2.1.14 Compiler flags

compiler_flag ::=
    flag(flag_value)

2.2 Directives

2.2.1 Source file directives

encoding/1

Description

```
encoding(Encoding)
```

Declares the source file text encoding. Only supported on some back-end Prolog compilers. When used, this directive must be the first term in the source file in the first line. Currently recognized encodings values include `'US-ASCII'`, `'ISO-8859-1'`, `'ISO-8859-2'`, `'ISO-8859-15'`, `'UCS-2'`, `'UCS-2LE'`, `'UCS-2BE'`, `'UTF-8'`, `'UTF-16'`, `'UTF-16LE'`, `'UTF-16BE'`, `'UTF-32'`, `'UTF-32LE'`, `'UTF-32BE'`, `'Shift_JIS'`, and `'EUC-JP'`. Be sure to use an encoding supported by the chosen back-end Prolog compiler (whose adapter file must define a table that translates between the Logtalk and Prolog specific atoms that represent each supported encoding). When writing portable code that cannot be expressed using ASCII, `'UTF-8'` is usually the best choice.

Template and modes

```
encoding(+atom)
```

Examples

```
:- encoding('UTF-8').
```

include/1

Description

```
include(File)
```

Includes a file contents, which must be valid terms, at the place of occurrence of the directive. The file can be specified as a relative path, an absolute path, or using library notation and is expanded as a source file name. Relative paths are interpreted as relative to the path of the file containing the directive.

When using the reflection API, predicates from an included file can be distinguished from predicates from the main file by looking for the `include/1` predicate declaration or definition property. For the included predicates, the `line_count/1` property stores the term line number in the included file.

This directive can be used as either a source file directive or an entity directive. As an entity directive, it can be used both in entities defined in source files and with the entity creation built-in predicates.

Note that when using this directive as an argument in a call to the *create_object/4* and *create_category/4* predicates, the objects will not be recreated or redefined when the included file(s) are modified and the *logtalk_make/0* predicate or the *logtalk_make/1* predicate (with target `all`) are called.

Template and modes

```
include(@source_file_name)
```

Examples

```
:- include(data('raw_1.txt')).

:- include('factbase.pl').

:- include('/home/me/databases/cities.pl').

| ?- create_object(cities, [], [public(city/4), include('/home/me/dbs/cities.pl')],␣
→[]).
```

initialization/1

Description

```
initialization(Goal)
```

When used within an object, this directive defines a goal to be called after the object has been loaded into memory. When used at a global level within a source file, this directive defines a goal to be called after the compiled source file is loaded into memory.

Multiple initialization directives can be used in a source file or in an object. Their goals will be called in order at loading time.

Categories and protocols cannot contain initialization directives as the initialization goals would lack a complete execution context which is only available for objects.

Although technically a global `initialization/1` directive in a source file is a Prolog directive, calls to Logtalk built-in predicates from it are usually compiled to improve performance and providing better support for embedded applications.

Template and modes

```
initialization(@callable)
```

Examples

```
:- initialization(init).
```

op/3

Description

```
op(Precedence, Associativity, Operator)
```

Declares operators. Operators declared inside entities have local scope. Global operators can be declared inside a source file by writing the respective directives before the entity opening directives.

Template and modes

```
op(+integer, +associativity, +atom_or_atom_list)
```

Examples

```
:- op(200, fy, +).
:- op(200, fy, ?).
:- op(200, fy, @).
:- op(200, fy, -).
```

See also:

current_op/3

set_logtalk_flag/2

Description

```
set_logtalk_flag(Flag, Value)
```

Sets Logtalk flag values. The scope of this directive is the entity or the source file containing it. For global scope, use the corresponding *set_logtalk_flag/2* built-in predicate called from an *initialization/1* directive.

Template and modes

```
set_logtalk_flag(+atom, +nonvar)
```

Errors

Flag is a variable: `instantiation_error`

Value is a variable: `instantiation_error`

Flag is not an atom: `type_error(atom, Flag)`

Flag is neither a variable nor a valid flag: `domain_error(flag, Flag)`

Value is not a valid value for flag Flag: `domain_error(flag_value, Flag + Value)`

Flag is a read-only flag: `permission_error(modify, flag, Flag)`

Examples

```
:- set_logtalk_flag(unknown_entities, silent).
```

2.2.2 Conditional compilation directives

if/1

Description

```
if(Goal)
```

Starts conditional compilation. The code following the directive is compiled if `Goal` is true. The goal is subjected to goal expansion when the directive occurs in a source file.

Conditional compilation goals cannot depend on predicate definitions contained in the same source file that contains the conditional compilation directives (as those predicates only become available after the file is fully compiled and loaded).

Template and modes

```
if(@callable)
```

Examples

```
:- if(current_prolog_flag(double_quotes, atom)).
```

See also:

elif/1, *else/0*, *endif/0*

elif/1

Description

```
elif(Goal)
```

Supports embedded conditionals when performing conditional compilation. The code following the directive is compiled if `Goal` is true. The goal is subjected to goal expansion when the directive occurs in a source file.

Conditional compilation goals cannot depend on predicate definitions contained in the same source file that contains the conditional compilation directives (as those predicates only become available after the file is fully compiled and loaded).

Template and modes

```
elif(@callable)
```

Examples

```
:- elif(predicate_property(callable(_), built_in)).
```

See also:

else/0, *endif/0*, *if/1*

else/0

Description

```
else
```

Starts an *else* branch when performing conditional compilation.

Template and modes

```
else
```

Examples

```
:- else.
```

See also:

elif/1, *endif/0*, *if/1*

endif/0

Description

```
endif
```

Ends conditional compilation for the matching `if/1` directive.

Template and modes

```
endif
```

Examples

```
:- endif.
```

See also:

elif/1, *else/0*, *if/1*

2.2.3 Entity directives

built_in/0

Description

```
built_in
```

Declares an entity as built-in. Built-in entities cannot be redefined once loaded.

Template and modes

```
built_in
```

Examples

```
:- built_in.
```

category/1-3

Description

```
category(Category)

category(Category,
    implements(Protocols))

category(Category,
    extends(Categories))

category(Category,
    complements(Objects))

category(Category,
```

```
    implements(Protocols),
    extends(Categories))

category(Category,
    implements(Protocols),
    complements(Objects))

category(Category,
    extends(Categories),
    complements(Objects))

category(Category,
    implements(Protocols),
    extends(Categories),
    complements(Objects))
```

Starting category directive.

Template and modes

```
category(+category_identifier)

category(+category_identifier,
    implements(+implemented_protocols))

category(+category_identifier,
    extends(+extended_categories))

category(+category_identifier,
    complements(+complemented_objects))

category(+category_identifier,
    implements(+implemented_protocols),
    extends(+extended_categories))

category(+category_identifier,
    implements(+implemented_protocols),
    complements(+complemented_objects))

category(+category_identifier,
    extends(+extended_categories),
    complements(+complemented_objects))

category(+category_identifier,
    implements(+implemented_protocols),
    extends(+extended_categories),
    complements(+complemented_objects))
```

Examples

```
:- category(monitoring).

:- category(monitoring,
```

```
        implements(monitoringp)).

:- category(attributes,
    implements(protected::variables)).

:- category(extended,
    extends(minimal)).

:- category(logging,
    implements(monitoring),
    complements(employee)).
```

See also:

end_category/0

dynamic/0

Description

```
dynamic
```

Declares an entity and its contents as dynamic. Dynamic entities can be abolished at runtime.

Template and modes

```
dynamic
```

Examples

```
:- dynamic.
```

See also:

dynamic/1, *object_property/2*, *protocol_property/2*, *category_property/2*

end_category/0

Description

```
end_category
```

Ending category directive.

Template and modes

```
end_category
```

Examples

```
:- end_category.
```

See also:

category/1-3

end_object/0

Description

```
end_object
```

Ending object directive.

Template and modes

```
end_object
```

Examples

```
:- end_object.
```

See also:

object/1-5

end_protocol/0

Description

```
end_protocol
```

Ending protocol directive.

Template and modes

```
end_protocol
```

Examples

```
:- end_protocol.
```

See also:

protocol/1-2

info/1

Description

```
info(List)
```

Documentation directive for objects, protocols, and categories. The directive argument is a list of pairs using the format *Key is Value*. See the *Entity directives* section for a description of the default keys.

Template and modes

```
info(+entity_info_list)
```

Examples

```
:- info([
    version is 1.0,
    author is 'Paulo Moura',
    date is 2000/4/20,
    comment is 'List protocol.'
]).
```

See also:

info/2, *object_property/2*, *protocol_property/2*, *category_property/2*

object/1-5

Description

Stand-alone objects (prototypes)

```
object(Object)

object(Object,
    implements(Protocols))

object(Object,
    imports(Categories))

object(Object,
    implements(Protocols),
    imports(Categories))
```

Prototype extensions

```
object(Object,
    extends(Objects))

object(Object,
    implements(Protocols),
```

```
    extends(Objects))

object(Object,
    imports(Categories),
    extends(Objects))

object(Object,
    implements(Protocols),
    imports(Categories),
    extends(Objects))
```

Class instances

```
object(Object,
    instantiates(Classes))

object(Object,
    implements(Protocols),
    instantiates(Classes))

object(Object,
    imports(Categories),
    instantiates(Classes))

object(Object,
    implements(Protocols),
    imports(Categories),
    instantiates(Classes))
```

Classes

```
object(Object,
    specializes(Classes))

object(Object,
    implements(Protocols),
    specializes(Classes))

object(Object,
    imports(Categories),
    specializes(Classes))

object(Object,
    implements(Protocols),
    imports(Categories),
    specializes(Classes))
```

Classes with metaclasses

```
object(Object,
    instantiates(Classes),
    specializes(Classes))

object(Object,
    implements(Protocols),
    instantiates(Classes),
    specializes(Classes))
```

```
object(Object,
    imports(Categories),
    instantiates(Classes),
    specializes(Classes))

object(Object,
    implements(Protocols),
    imports(Categories),
    instantiates(Classes),
    specializes(Classes))
```

Starting object directive.

Template and modes

Stand-alone objects (prototypes)

```
object(+object_identifier)

object(+object_identifier,
    implements(+implemented_protocols))

object(+object_identifier,
    imports(+imported_categories))

object(+object_identifier,
    implements(+implemented_protocols),
    imports(+imported_categories))
```

Prototype extensions

```
object(+object_identifier,
    extends(+extended_objects))

object(+object_identifier,
    implements(+implemented_protocols),
    extends(+extended_objects))

object(+object_identifier,
    imports(+imported_categories),
    extends(+extended_objects))

object(+object_identifier,
    implements(+implemented_protocols),
    imports(+imported_categories),
    extends(+extended_objects))
```

Class instances

```
object(+object_identifier,
    instantiates(+instantiated_objects))

object(+object_identifier,
    implements(+implemented_protocols),
```

```
    instantiates(+instantiated_objects))

object(+object_identifier,
    imports(+imported_categories),
    instantiates(+instantiated_objects))

object(+object_identifier,
    implements(+implemented_protocols),
    imports(+imported_categories),
    instantiates(+instantiated_objects))
```

Classes

```
object(+object_identifier,
    specializes(+specialized_objects))

object(+object_identifier,
    implements(+implemented_protocols),
    specializes(+specialized_objects))

object(+object_identifier,
    imports(+imported_categories),
    specializes(+specialized_objects))

object(+object_identifier,
    implements(+implemented_protocols),
    imports(+imported_categories),
    specializes(+specialized_objects))
```

Class with metaclasses

```
object(+object_identifier,
    instantiates(+instantiated_objects),
    specializes(+specialized_objects))

object(+object_identifier,
    implements(+implemented_protocols),
    instantiates(+instantiated_objects),
    specializes(+specialized_objects))

object(+object_identifier,
    imports(+imported_categories),
    instantiates(+instantiated_objects),
    specializes(+specialized_objects))

object(+object_identifier,
    implements(+implemented_protocols),
    imports(+imported_categories),
    instantiates(+instantiated_objects),
    specializes(+specialized_objects))
```

Examples

```
:- object(list).

:- object(list,
    implements(listp)).

:- object(list,
    extends(compound)).

:- object(list,
    implements(listp),
    extends(compound)).

:- object(object,
    imports(initialization),
    instantiates(class)).

:- object(abstract_class,
    instantiates(class),
    specializes(object)).

:- object(agent,
    imports(private::attributes)).
```

See also:

end_object/0

protocol/1-2

Description

```
protocol(Protocol)

protocol(Protocol,
    extends(Protocols))
```

Starting protocol directive.

Template and modes

```
protocol(+protocol_identifier)

protocol(+protocol_identifier,
    extends(+extended_protocols))
```

Examples

```
:- protocol(listp).

:- protocol(listp,
    extends(compoundp)).
```

(continues on next page)

```
:- protocol(queuep,
    extends(protected::listp)).
```

See also:

end_protocol/0

threaded/0

Description

```
threaded
```

Declares that an object supports concurrent calls and asynchronous messages. Any object containing calls to the built-in multi-threading predicates (or importing a category that contains such calls) must include this directive.

Template and modes

```
threaded
```

Examples

```
:- threaded.
```

See also:

synchronized/1 object_property/2

2.2.4 Predicate directives

alias/2

Description

```
alias(Entity, PredicateAliases)
alias(Entity, NonTerminalAliases)
```

Declares predicate and grammar rule non-terminal aliases. A predicate (non-terminal) alias is an alternative name for a predicate (non-terminal) declared or defined in an extended protocol, an implemented protocol, an extended category, an imported category, an extended prototype, an instantiated class, or a specialized class. Predicate aliases may be used to solve conflicts between imported or inherited predicates. It may also be used to give a predicate (non-terminal) a name more appropriated in its usage context. This directive may be used in objects, protocols, and categories.

Predicate (and non-terminal) aliases are specified using (preferably) the notation `Name/Arity as Alias/Arity` or, in alternative, the notation `Name/Arity::Alias/Arity`.

It is also possible to declare predicate and grammar rule non-terminal aliases in implicit qualification directives for sending messages to objects and calling module predicates.

Template and modes

```
alias(@entity_identifier, +list(predicate_indicator_alias))
alias(@entity_identifier, +list(non_terminal_indicator_alias))
```

Examples

```
:- alias(list, [member/2 as list_member/2]).
:- alias(set,  [member/2 as set_member/2]).

:- alias(words, [singular//0 as peculiar//0]).
```

See also:

uses/2, *use_module/2*

coinductive/1

Description

```
coinductive(Name/Arity)
coinductive((Functor1/Arity1, ...))
coinductive([Functor1/Arity1, ...])

coinductive(Template)
coinductive((Template1, ...))
coinductive([Template1, ...])
```

This is an **experimental** directive, used for declaring coinductive predicates. Requires a back-end Prolog compiler with minimal support for cyclic terms. The current implementation of coinduction allows the generation of only the *basic cycles* but all valid solutions should be recognized. Use a predicate indicator as argument when all the coinductive predicate arguments are relevant for coinductive success. Use a template when only some coinductive predicate arguments (represented by a "+") should be considered when testing for coinductive success (represent the arguments that should be disregarded by a "−"). It's possible to define local *coinductive_success_hook/1-2* predicates that are automatically called with the coinductive predicate term resulting from a successful unification with an ancestor goal as first argument. The second argument, when present, is the coinductive hypothesis (i.e. the ancestor goal) used. These hook predicates can provide an alternative to the use of tabling when defining some coinductive predicates. There is no overhead when these hook predicates are not defined.

This directive must precede any calls to the declared coinductive predicates.

Template and modes

```
coinductive(+predicate_indicator_term)
coinductive(+coinductive_predicate_template_term)
```

Examples

```
:- coinductive(comember/2).
:- coinductive(controller(+,+,+,-,-)).
```

See also:

coinductive_success_hook/1-2, *predicate_property/2*

discontiguous/1

Description

```
discontiguous(Name/Arity)
discontiguous((Functor1/Arity1, ...))
discontiguous([Functor1/Arity1, ...])

discontiguous(Name//Arity)
discontiguous((Functor1//Arity1, ...))
discontiguous([Functor1//Arity1, ...])
```

Declares discontiguous predicates and discontiguous grammar rule non-terminals. The use of this directive should be avoided as not all backend Prolog compilers support discontiguous predicates.

Template and modes

```
discontiguous(+predicate_indicator_term)
discontiguous(+non_terminal_indicator_term)
```

Examples

```
:- discontiguous(counter/1).

:- discontiguous((lives/2, works/2)).

:- discontiguous([db/4, key/2, file/3]).
```

dynamic/1

Description

```
dynamic(Name/Arity)
dynamic((Functor1/Arity1, ...))
dynamic([Functor1/Arity1, ...])

dynamic(Entity::Name/Arity)
dynamic((Entity1::Functor1/Arity1, ...))
dynamic([Entity1::Functor1/Arity1, ...])

dynamic(Module:Name/Arity)
dynamic((Module1:Functor1/Arity1, ...))
```

(continues on next page)

```
dynamic([Module1:Functor1/Arity1, ...])

dynamic(Name//Arity)
dynamic((Functor1//Arity1, ...))
dynamic([Functor1//Arity1, ...])

dynamic(Entity::Name//Arity)
dynamic((Entity1::Functor1//Arity1, ...))
dynamic([Entity1::Functor1//Arity1, ...])

dynamic(Module:Name//Arity)
dynamic((Module1:Functor1//Arity1, ...))
dynamic([Module1:Functor1//Arity1, ...])
```

Declares dynamic predicates and dynamic grammar rule non-terminals. Note that an object can be static and have both static and dynamic predicates/non-terminals. Dynamic predicates cannot be declared as synchronized. When the dynamic predicates are local to an object, declaring them also as private predicates allows the Logtalk compiler to generate optimized code for asserting and retracting predicate clauses. Categories can also contain dynamic predicate directives but cannot contain clauses for dynamic predicates.

The predicate indicators (or non-terminal indicators) can be explicitly qualified with an object, category, or module identifier when the predicates (or non-terminals) are also declared multifile.

Note that dynamic predicates cannot be declared synchronized (when necessary, declare the predicates updating the dynamic predicates as synchronized).

Template and modes

```
dynamic(+qualified_predicate_indicator_term)
dynamic(+qualified_non_terminal_indicator_term)
```

Examples

```
:- dynamic(counter/1).

:- dynamic((lives/2, works/2)).

:- dynamic([db/4, key/2, file/3]).
```

See also:

dynamic/0, *predicate_property/2*

info/2

Description

```
info(Name/Arity, List)
info(Name//Arity, List)
```

Documentation directive for predicates and grammar rule non-terminals. The first argument is either a predicate indicator or a grammar rule non-terminal indicator. The second argument is a list of pairs using the format *Key is Value*. See the *Predicate directives* section for a description of the default keys.

Template and modes

```
info(+predicate_indicator, +predicate_info_list)
info(+non_terminal_indicator, +predicate_info_list)
```

Examples

```
:- info(empty/1, [
    comment is 'True if the argument is an empty list.',
    argnames is ['List']
]).

:- info(sentence//0, [
    comment is 'Rewrites a sentence into a noun phrase and a verb phrase.'
]).
```

See also:

info/1, *mode/2*, *predicate_property/2*

meta_predicate/1

Description

```
meta_predicate(MetaPredicateTemplate)
meta_predicate((MetaPredicateTemplate1, ...))
meta_predicate([MetaPredicateTemplate1, ...])

meta_predicate(Entity::MetaPredicateTemplate)
meta_predicate((Entity1::MetaPredicateTemplate1, ...))
meta_predicate([Entity1::MetaPredicateTemplate1, ...])

meta_predicate(Module:MetaPredicateTemplate)
meta_predicate((Module1:MetaPredicateTemplate1, ...))
meta_predicate([Module1:MetaPredicateTemplate1, ...])
```

Declares meta-predicates, i.e., predicates that have arguments that will be called as goals. An argument may also be a *closure* instead of a goal if the meta-predicate uses the *call/1-N* Logtalk built-in methods to construct and call the actual goal from the closure and the additional arguments.

Meta-arguments which are goals are represented by the integer 0. Meta-arguments which are closures are represented by a positive integer, N, representing the number of additional arguments that will be appended to the closure in order to construct the corresponding meta-call. Normal arguments are represented by the atom *. Meta-arguments are always called in the meta-predicate calling context, not in the meta-predicate definition context.

Logtalk allows the use of this directive to override the original meta-predicate directive. This is sometimes necessary when calling Prolog module meta-predicates due to the lack of standardization of the syntax of the meta-predicate templates.

Template and modes

```
meta_predicate(+meta_predicate_template_term)

meta_predicate(+object_identifier::+meta_predicate_template_term)
meta_predicate(+category_identifier::+meta_predicate_template_term)

meta_predicate(+module_identifier:+meta_predicate_template_term)
```

Examples

```
:- meta_predicate(findall(*, 0, *)).

:- meta_predicate(forall(0, 0)).

:- meta_predicate(maplist(2, *, *)).
```

See also:

meta_non_terminal/1, *predicate_property/2*

meta_non_terminal/1

Description

```
meta_non_terminal(MetaNonTerminalTemplate)
meta_non_terminal((MetaNonTerminalTemplate1, ...))
meta_non_terminal([MetaNonTerminalTemplate1, ...])

meta_non_terminal(Entity::MetaNonTerminalTemplate)
meta_non_terminal((Entity1::MetaNonTerminalTemplate1, ...))
meta_non_terminal([Entity1::MetaNonTerminalTemplate1, ...])

meta_non_terminal(Module:MetaNonTerminalTemplate)
meta_non_terminal((Module1:MetaNonTerminalTemplate1, ...))
meta_non_terminal([Module1:MetaNonTerminalTemplate1, ...])
```

Declares meta-non-terminals, i.e., non-terminals that have arguments that will be called as non-terminals (or grammar rule bodies). An argument may also be a *closure* instead of a goal if the non-terminal uses the *call//1-N* Logtalk built-in methods to construct and call the actual non-terminal from the closure and the additional arguments.

Meta-arguments which are non-terminals are represented by the integer 0. Meta-arguments which are closures are represented by a positive integer, N, representing the number of additional arguments that will be appended to the closure in order to construct the corresponding meta-call. Normal arguments are represented by the atom *. Meta-arguments are always called in the meta-non-terminal calling context, not in the meta-non-terminal definition context.

Template and modes

```
meta_non_terminal(+meta_non_terminal_template_term)

meta_non_terminal(+object_identifier::+meta_non_terminal_template_term)
```

(continues on next page)

```
meta_non_terminal(+category_identifier::+meta_non_terminal_template_term)

meta_non_terminal(+module_identifier:+meta_non_terminal_template_term)
```

Examples

```
:- meta_non_terminal(findall(*, 0, *)).

:- meta_non_terminal(forall(0, 0)).

:- meta_non_terminal(maplist(2, *, *)).
```

See also:

meta_predicate/1, *predicate_property/2*

mode/2

Description

```
mode(Mode, NumberOfProofs)
```

Most predicates can be used with several instantiations modes. This directive enables the specification of each instantiation mode and the corresponding number of proofs (not necessarily distinct solutions). You may also use this directive for documenting grammar rule non-terminals.

Template and modes

```
mode(+predicate_mode_term, +number_of_proofs)
mode(+non_terminal_mode_term, +number_of_proofs)
```

Examples

```
:- mode(atom_concat(-atom, -atom, +atom), one_or_more).
:- mode(atom_concat(+atom, +atom, -atom), one).

:- mode(var(@term), zero_or_one).

:- mode(solve(+callable, -list(atom)), zero_or_one).
```

See also:

info/2, *predicate_property/2*

multifile/1

Description

```
multifile(Name/Arity)
multifile((Functor1/Arity1, ...))
multifile([Functor1/Arity1, ...])

multifile(Entity::Name/Arity)
multifile((Entity1::Functor1/Arity1, ...))
multifile([Entity1::Functor1/Arity1, ...])

multifile(Module:Name/Arity)
multifile((Module1:Functor1/Arity1, ...))
multifile([Module1:Functor1/Arity1, ...])

multifile(Name//Arity)
multifile((Functor1//Arity1, ...))
multifile([Functor1//Arity1, ...])

multifile(Entity::Name//Arity)
multifile((Entity1::Functor1//Arity1, ...))
multifile([Entity1::Functor1//Arity1, ...])

multifile(Module:Name//Arity)
multifile((Module1:Functor1//Arity1, ...))
multifile([Module1:Functor1//Arity1, ...])
```

Declares multifile predicates and multifile grammar rule non-terminals. In the case of object or category multifile predicates, the predicate (or non-terminal) must also have a scope directive in the object or category holding its *primary declaration* (i.e. the declaration without the `Entity::` prefix). Entities holding multifile predicate primary declarations must be compiled and loaded prior to any entities contributing with clauses for the multifile predicates.

Protocols cannot declare multifile predicates as protocols cannot contain predicate definitions.

Template and modes

```
multifile(+qualified_predicate_indicator_term)
multifile(+qualified_non_terminal_indicator_term)
```

Examples

```
:- multifile(table/3).
:- multifile(user::hook/2).
```

See also:

public/1, *protected/1*, *private/1*, *predicate_property/2*

private/1

Description

```
private(Name/Arity)
private((Functor1/Arity1, ...))
private([Functor1/Arity1, ...])

private(Name//Arity)
private((Functor1//Arity1, ...))
private([Functor1//Arity1, ...])

private(op(Precedence, Associativity, Operator))
```

Declares private predicates, private grammar rule non-terminals, and private operators. A private predicate can only be called from the object containing the private directive. A private non-terminal can only be used in a call of the *phrase/2* and *phrase/3* methods from the object containing the private directive.

Template and modes

```
private(+predicate_indicator_term)
private(+non_terminal_indicator_term)
private(+operator_declaration)
```

Examples

```
:- private(counter/1).

:- private((init/1, free/1)).

:- private([data/3, key/1, keys/1]).
```

See also:

protected/1, *public/1*, *predicate_property/2*

protected/1

Description

```
protected(Name/Arity)
protected((Functor1/Arity1, ...))
protected([Functor1/Arity1, ...])

protected(Name//Arity)
protected((Functor1//Arity1, ...))
protected([Functor1//Arity1, ...])

protected(op(Precedence, Associativity, Operator))
```

Declares protected predicates, protected grammar rule non-terminals, and protected operators. A protected predicate can only be called from the object containing the directive or from an object that inherits the directive. A protected non-terminal can only be used as an argument in a *phrase/2* and *phrase/3* calls from the object containing the directive

or from an object that inherits the directive. Protected operators are not inherited but declaring them provides useful information for defining descendant objects.

Template and modes

```
protected(+predicate_indicator_term)
protected(+non_terminal_indicator_term)
protected(+operator_declaration)
```

Examples

```
:- protected(init/1).

:- protected((print/2, convert/4)).

:- protected([load/1, save/3]).
```

See also:

private/1, *public/1*, *predicate_property/2*

public/1

Description

```
public(Name/Arity)
public((Functor1/Arity1, ...))
public([Functor1/Arity1, ...])

public(Name//Arity)
public((Functor1//Arity1, ...))
public([Functor1//Arity1, ...])

public(op(Precedence, Associativity, Operator))
```

Declares public predicates, public grammar rule non-terminals, and public operators. A public predicate can be called from any object. A public non-terminal can be used as an argument in *phrase/2* and *phrase/3* calls from any object. Public operators are not exported but declaring them provides useful information for defining client objects.

Template and modes

```
public(+predicate_indicator_term)
public(+non_terminal_indicator_term)
public(+operator_declaration)
```

Examples

```
:- public(ancestor/1).

:- public((instance/1, instances/1)).

:- public([leaf/1, leaves/1]).
```

See also:

private/1, *protected/1*, *predicate_property/2*

synchronized/1

Description

```
synchronized(Name/Arity)
synchronized((Functor1/Arity1, ...))
synchronized([Functor1/Arity1, ...])

synchronized(Name//Arity)
synchronized((Functor1//Arity1, ...))
synchronized([Functor1//Arity1, ...])
```

Declares synchronized predicates and synchronized grammar rule non-terminals. A synchronized predicate (or synchronized non-terminal) is protected by a mutex in order to allow for thread synchronization when proving a call to the predicate (or non-terminal). All predicates (and non-terminals) declared in the same synchronized directive share the same mutex. In order to use a separate mutex for each predicate (non-terminal) so that they are independently synchronized, a per-predicate synchronized directive must be used.

Declaring a predicate synchronized implicitly makes it deterministic. When using a single-threaded back-end Prolog compiler, calls to synchronized predicates behave as wrapped by the standard `once/1` meta-predicate.

Note that synchronized predicates cannot be declared dynamic (when necessary, declare the predicates updating the dynamic predicates as synchronized).

Template and modes

```
synchronized(+predicate_indicator_term)
synchronized(+non_terminal_indicator_term)
```

Examples

```
:- synchronized(db_update/1).

:- synchronized((write_stream/2, read_stream/2)).

:- synchronized([add_to_queue/2, remove_from_queue/2]).
```

See also:

predicate_property/2

---

uses/2

Description

```
uses(Object, Predicates)
uses(Object, PredicatesAndAliases)

uses(Object, NonTerminals)
uses(Object, NonTerminalsAndAliases)

uses(Object, Operators)
```

Declares that all calls made from predicates (or non-terminals) defined in the category or object containing the directive to the specified predicates (or non-terminals) are to be interpreted as messages to the specified object. Thus, this directive may be used to simplify writing of predicate definitions by allowing the programmer to omit the `Object::` prefix when using the predicates listed in the directive (as long as the calls do not occur as arguments for non-standard Prolog meta-predicates not declared on the adapter files). It is also possible to include operator declarations, `op(Precedence, Associativity, Operator)`, in the second argument.

This directive is also used when compiling calls to the database and reflection built-in methods by looking into these methods predicate arguments.

It is possible to specify a predicate alias using the notation `Name/Arity as Alias/Arity` or, in alternative, the notation `Name/Arity::Alias/Arity`. Aliases may be used either for avoiding conflicts between predicates specified in `use_module/2` and `uses/2` directives or for giving more meaningful names considering the using context of the predicates.

To enable the use of static binding, and thus optimal message sending performance, the objects should be loaded before compiling the entities that call their predicates.

Template and modes

```
uses(+object_identifier, +predicate_indicator_list)
uses(+object_identifier, +predicate_indicator_alias_list)

uses(+object_identifier, +non_terminal_indicator_list)
uses(+object_identifier, +non_terminal_indicator_alias_list)

uses(+object_identifier, +operator_list)
```

Examples

```
:- uses(list, [append/3, member/2]).
:- uses(store, [data/2]).

foo :-
    ...,
    findall(X, member(X, L), A),    % the same as findall(X, list::member(X, L), A)
    append(A, B, C),                % the same as list::append(A, B, C)
    assertz(data(X, C)),            % the same as store::assertz(data(X, C))
    ...
```

Another example, using the extended notation that allows us to define predicate aliases:

```
:- uses(btrees, [new/1 as new_btree/1]).
:- uses(queues, [new/1 as new_queue/1]).

btree_to_queue :-
    ...,
    new_btree(Tree),      % the same as btrees::new(Tree)
    new_queue(Queue),     % the same as queues::new(Queue)
    ...
```

See also:

use_module/2

use_module/2

Description

```
use_module(Module, Predicates)
use_module(Module, PredicatesAndAliases)

use_module(Module, NonTerminals)
use_module(Module, NonTerminalsAndAliases)

use_module(Module, Operators)
```

This directive declares that all calls (made from predicates defined in the category or object containing the directive) to the specified predicates (or non-terminals) are to be interpreted as calls to explicitly-qualified module predicates (or non-terminals). Thus, this directive may be used to simplify writing of predicate definitions by allowing the programmer to omit the `Module:` prefix when using the predicates listed in the directive (as long as the predicate calls do not occur as arguments for non-standard Prolog meta-predicates not declared on the adapter files). It is also possible to include operator declarations, `op(Precedence, Associativity, Operator)`, in the second argument.

This directive is also used when compiling calls to the database and reflection built-in methods by examining these methods predicate arguments.

It is possible to specify a predicate alias using the notation `Name/Arity as Alias/Arity` or, in alternative, the notation `Name/Arity:Alias/Arity`. Aliases may be used either for avoiding conflicts between predicates specified in `use_module/2` and *uses/2* directives or for giving more meaningful names considering the using context of the predicates.

Note that this directive differs from the directive with the same name found on some Prolog implementations by requiring the first argument to be a module name (an atom) instead of a file specification. In Logtalk, there's no mixing between *loading* a resource and (declaring the) *using* (of) a resource. As a consequence, this directive doesn't automatically load the module. Loading the module file is dependent of the used backend Prolog compiler and must be done separately (usually, using a source file `use_module/1` or `use_module/2` directive in the entity file or in the application loader file). Also, note that the name of the module may differ from the name of the module file.

The modules should be loaded prior to the compilation of entities that call the module predicates. This is required in general to allow the compiler to check if the called module predicate is a meta-predicate and retrieve its meta-predicate template to ensure proper call compilation.

Template and modes

```
use_module(+module_identifier, +predicate_indicator_list)
use_module(+module_identifier, +predicate_indicator_alias_list)

use_module(+module_identifier, +non_terminal_indicator_list)
use_module(+module_identifier, +non_terminal_indicator_alias_list)

use_module(+module_identifier, +operator_list)
```

Examples

```
:- use_module(lists, [append/3, member/2]).
:- use_module(store, [data/2]).
:- use_module(user,  [foo/1 as bar/1]).

foo :-
    ...,
    findall(X, member(X, L), A),    % same as findall(X, lists:member(X, L), A)
    append(A, B, C),                % same as lists:append(A, B, C)
    assertz(data(X, C)),            % same as assertz(store:data(X, C))
    retractall(bar(_)),             % same as retractall(user:foo(_))
    ...
```

See also:

uses/2

2.3 Built-in predicates

2.3.1 Enumerating objects, categories and protocols

current_category/1

Description

```
current_category(Category)
```

Enumerates, by backtracking, all currently defined categories. All categories are found, either static, dynamic, or built-in.

Template and modes

```
current_category(?category_identifier)
```

Errors

Category is neither a variable nor a valid category identifier: `type_error(category_identifier, Category)`

Examples

```
| ?- current_category(monitoring).
```

See also:

abolish_category/1, *category_property/2*, *create_category/4*, *complements_object/2*, *extends_category/2-3*, *imports_category/2-3*

current_object/1

Description

```
current_object(Object)
```

Enumerates, by backtracking, all currently defined objects. All objects are found, either static, dynamic or built-in.

Template and modes

```
current_object(?object_identifier)
```

Errors

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Examples

```
| ?- current_object(list).
```

See also:

abolish_object/1, *create_object/4*, *object_property/2*, *extends_object/2-3*, *instantiates_class/2-3*, *specializes_class/2-3*, *complements_object/2*

current_protocol/1

Description

```
current_protocol(Protocol)
```

Enumerates, by backtracking, all currently defined protocols. All protocols are found, either static, dynamic, or built-in.

Template and modes

```
current_protocol(?protocol_identifier)
```

Errors

Protocol is neither a variable nor a valid protocol identifier: `type_error(protocol_identifier, Protocol)`

Examples

```
| ?- current_protocol(listp).
```

See also:

abolish_protocol/1, *create_protocol/3*, *protocol_property/2*, *conforms_to_protocol/2-3*, *extends_protocol/2-3*, *implements_protocol/2-3*

2.3.2 Enumerating objects, categories and protocols properties

category_property/2

Description

```
category_property(Category, Property)
```

Enumerates, by backtracking, the properties associated with the defined categories. The valid category properties are listed in the language *Entity properties*.

Template and modes

```
category_property(?category_identifier, ?category_property)
```

Errors

Category is neither a variable nor a valid category identifier: `type_error(category_identifier, Category)`

Property is neither a variable nor a callable term: `type_error(callable, Property)`

Property is a callable term but not a valid category property: `domain_error(category_property, Property)`

Examples

```
| ?- category_property(Category, dynamic).
```

See also:

abolish_category/1, *create_category/4*, *current_category/1*, *complements_object/2*, *extends_category/2-3*, *imports_category/2-3*

---

object_property/2

Description

```
object_property(Object, Property)
```

Enumerates, by backtracking, the properties associated with the defined objects. The valid object properties are listed in the language *grammar*.

Template and modes

```
object_property(?object_identifier, ?object_property)
```

Errors

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Property is neither a variable nor a callable term: `type_error(callable, Property)`

Property is a callable term but not a valid object property: `domain_error(object_property, Property)`

Examples

```
| ?- object_property(list, Property).
```

See also:

abolish_object/1, *create_object/4*, *current_object/1*, *extends_object/2-3*, *instantiates_class/2-3*, *specializes_class/2-3*, *complements_object/2*

protocol_property/2

Description

```
protocol_property(Protocol, Property)
```

Enumerates, by backtracking, the properties associated with the currently defined protocols. The valid protocol properties are listed in the language *Entity properties*.

Template and modes

```
protocol_property(?protocol_identifier, ?protocol_property)
```

Errors

Protocol is neither a variable nor a valid protocol identifier: `type_error(protocol_identifier, Protocol)`

Property is neither a variable nor a callable term: `type_error(callable, Property)`

Property is a callable term but not a valid protocol property: `domain_error(protocol_property, Property)`

Examples

```
| ?- protocol_property(listp, Property).
```

See also:

abolish_protocol/1, *create_protocol/3*, *current_protocol/1*, *conforms_to_protocol/2-3*, *extends_protocol/2-3*, *implements_protocol/2-3*

2.3.3 Creating new objects, categories and protocols

create_category/4

Description

```
create_category(Identifier, Relations, Directives, Clauses)
```

Creates a new, dynamic category. This predicate is often used as a primitive to implement high-level category creation methods.

Note that, when opting for runtime generated category identifiers, it's possible to run out of identifiers when using a back-end Prolog compiler with bounded integer support. The portable solution, when creating a large number of dynamic category in long-running applications, is to recycle, whenever possible, the identifiers.

When using Logtalk multi-threading features, predicates calling this built-in predicate may need to be declared synchronized in order to avoid race conditions.

Template and modes

```
create_category(?category_identifier, @list(category_relation), @list(category_
↪directive), @list(clause))
```

Errors

Relations, Directives, or Clauses is a variable: `instantiation_error`

Identifier is neither a variable nor a valid category identifier: `type_error(category_identifier, Identifier)`

Identifier is already in use: `permission_error(modify, category, Identifier)` `permission_error(modify, object, Identifier)` `permission_error(modify, protocol, Identifier)`

Relations is neither a variable nor a proper list: `type_error(list, Relations)`

Repeated entity relation clause: `permission_error(repeat, entity_relation,`
    `implements/1)`   `permission_error(repeat, entity_relation, extends/1)`
    `permission_error(repeat, entity_relation, complements/1)`

Directives is neither a variable nor a proper list: `type_error(list, Directives)`

Clauses is neither a variable nor a proper list: `type_error(list, Clauses)`

Examples

```
| ?- create_category(
      tolerances,
      [implements(comparing)],
      [],
      [epsilon(1e-15), (equal(X, Y) :- epsilon(E), abs(X-Y) =< E)]
    ).
```

See also:

abolish_category/1,   *category_property/2*,   *current_category/1*,   *complements_object/2*,   *extends_category/2-3*, *imports_category/2-3*

create_object/4

Description

```
create_object(Identifier, Relations, Directives, Clauses)
```

Creates a new, dynamic object. The word *object* is used here as a generic term. This predicate can be used to create new prototypes, instances, and classes. This predicate is often used as a primitive to implement high-level object creation methods.

Note that, when opting for runtime generated object identifiers, it's possible to run out of identifiers when using a back-end Prolog compiler with bounded integer support. The portable solution, when creating a large number of dynamic objects in long-running applications, is to recycle, whenever possible, the identifiers.

When using Logtalk multi-threading features, predicates calling this built-in predicate may need to be declared synchronized in order to avoid race conditions.

Template and modes

```
create_object(?object_identifier, @list(object_relation), @list(object_directive), 
→@list(clause))
```

Errors

Relations, Directives, or Clauses is a variable: `instantiation_error`

Identifier is neither a variable nor a valid object identifier: `type_error(object_identifier, Identifier)`

Identifier is already in use: `permission_error(modify, category, Identifier)`
`permission_error(modify, object, Identifier)` `permission_error(modify, protocol, Identifier)`

Relations is neither a variable nor a proper list: `type_error(list, Relations)`

Repeated entity relation clause: `permission_error(repeat, entity_relation, implements/1)` `permission_error(repeat, entity_relation, imports/1)` `permission_error(repeat, entity_relation, extends/1)` `permission_error(repeat, entity_relation, instantiates/1)` `permission_error(repeat, entity_relation, specializes/1)`

Directives is neither a variable nor a proper list: `type_error(list, Directives)`

Clauses is neither a variable nor a proper list: `type_error(list, Clauses)`

Examples

Creating a simple, stand-alone object (a prototype): `| ?- create_object(translator, [], [public(int/2)], [int(0, zero)]).`

Creating a new prototype derived from a parent prototype: `| ?- create_object(mickey, [extends(mouse)], [public(alias/1)], [alias(mortimer)]).`

Creating a new class instance: `| ?- create_object(p1, [instantiates(person)], [], [name('Paulo Moura'), age(42)]).`

Creating a new class as a specialization of another class: `| ?- create_object(hovercraft, [specializes(vehicle)], [public([propeller/2, fan/2])], []).`

Creating a new object and defining its initialization goal: `| ?- create_object(runner, [instantiates(runners)], [initialization(start)], [length(22), time(60)]).`

Creating a new empty object with dynamic predicate declarations support: `| ?- create_object(database, [], [set_logtalk_flag(dynamic_declarations, allow)], []).`

See also:

abolish_object/1, *current_object/1*, *object_property/2*, *extends_object/2-3*, *instantiates_class/2-3*, *specializes_class/2-3*, *complements_object/2*

create_protocol/3

Description

```
create_protocol(Identifier, Relations, Directives)
```

Creates a new, dynamic, protocol. This predicate is often used as a primitive to implement high-level protocol creation methods.

Note that, when opting for runtime generated protocol identifiers, it's possible to run out of identifiers when using a back-end Prolog compiler with bounded integer support. The portable solution, when creating a large number of dynamic protocols in long-running applications, is to recycle, whenever possible, the identifiers.

When using Logtalk multi-threading features, predicates calling this built-in predicate may need to be declared synchronized in order to avoid race conditions.

Template and modes

```
create_protocol(?protocol_identifier, @list(protocol_relation), @list(protocol_
→directive))
```

Errors

Either Relations or Directives is a variable: `instantiation_error`

Identifier is neither a variable nor a valid protocol identifier: `type_error(protocol_identifier, Identifier)`

Identifier is already in use: `permission_error(modify, category, Identifier)` `permission_error(modify, object, Identifier)` `permission_error(modify, protocol, Identifier)`

Relations is neither a variable nor a proper list: `type_error(list, Relations)`

Repeated entity relation clause: `permission_error(repeat, entity_relation, extends/1)`

Directives is neither a variable nor a proper list: `type_error(list, Directives)`

Examples

```
| ?- create_protocol(
        logging,
        [extends(monitoring)],
        [public([log_file/1, log_on/0, log_off/0])]
    ).
```

See also:

abolish_protocol/1, *current_protocol/1*, *protocol_property/2*, *conforms_to_protocol/2-3*, *extends_protocol/2-3*, *implements_protocol/2-3*

2.3.4 Abolishing objects, categories and protocols

abolish_category/1

Description

```
abolish_category(Category)
```

Abolishes a dynamic category.

Template and modes

```
abolish_category(+category_identifier)
```

Errors

Category is a variable: `instantiation_error`

Category is neither a variable nor a valid category identifier: `type_error(category_identifier,`
`    Category)`

Category is an identifier of a static category: `permission_error(modify, static_category,`
`    Category)`

Category does not exist: `existence_error(category, Category)`

Examples

```
| ?- abolish_category(monitoring).
```

See also:

category_property/2,   *create_category/4*,   *current_category/1*   *complements_object/2*,   *extends_category/2-3*,
imports_category/2-3

abolish_object/1

Description

```
abolish_object(Object)
```

Abolishes a dynamic object.

Template and modes

```
abolish_object(+object_identifier)
```

Errors

Object is a variable: `instantiation_error`

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Object is an identifier of a static object: `permission_error(modify, static_object, Object)`

Object does not exist: `existence_error(object, Object)`

Examples

```
| ?- abolish_object(list).
```

See also:

create_object/4, *current_object/1*, *object_property/2*, *extends_object/2-3*, *instantiates_class/2-3*, *specializes_class/2-
3*, *complements_object/2*

---

abolish_protocol/1

Description

```
abolish_protocol(Protocol)
```

Abolishes a dynamic protocol.

Template and modes

```
abolish_protocol(@protocol_identifier)
```

Errors

Protocol is a variable: `instantiation_error`

Protocol is neither a variable nor a valid protocol identifier: `type_error(protocol_identifier, Protocol)`

Protocol is an identifier of a static protocol: `permission_error(modify, static_protocol, Protocol)`

Protocol does not exist: `existence_error(protocol, Protocol)`

Examples

```
| ?- abolish_protocol(listp).
```

See also:

create_protocol/3, *current_protocol/1*, *protocol_property/2*, *conforms_to_protocol/2-3*, *extends_protocol/2-3*, *implements_protocol/2-3*

2.3.5 Objects, categories, and protocols relations

extends_object/2-3

Description

```
extends_object(Prototype, Parent)
extends_object(Prototype, Parent, Scope)
```

Enumerates, by backtracking, all pairs of objects such that the first one extends the second. The relation scope is represented by the atoms `public`, `protected`, and `private`.

Template and modes

```
extends_object(?object_identifier, ?object_identifier)
extends_object(?object_identifier, ?object_identifier, ?scope)
```

Errors

Prototype is neither a variable nor a valid object identifier: `type_error(object_identifier, Prototype)`

Parent is neither a variable nor a valid object identifier: `type_error(object_identifier, Parent)`

Scope is neither a variable nor an atom: `type_error(atom, Scope)`

Scope is an atom but an invalid entity scope: `domain_error(scope, Scope)`

Examples

```
| ?- extends_object(Object, state_space).

| ?- extends_object(Object, list, public).
```

See also:

current_object/1, *instantiates_class/2-3*, *specializes_class/2-3*

extends_protocol/2-3

Description

```
extends_protocol(Protocol1, Protocol2)
extends_protocol(Protocol1, Protocol2, Scope)
```

Enumerates, by backtracking, all pairs of protocols such that the first one extends the second. The relation scope is represented by the atoms `public`, `protected`, and `private`.

Template and modes

```
extends_protocol(?protocol_identifier, ?protocol_identifier)
extends_protocol(?protocol_identifier, ?protocol_identifier, ?scope)
```

Errors

Protocol1 is neither a variable nor a valid protocol identifier: `type_error(protocol_identifier, Protocol1)`

Protocol2 is neither a variable nor a valid protocol identifier: `type_error(protocol_identifier, Protocol2)`

Scope is neither a variable nor an atom: `type_error(atom, Scope)`

Scope is an atom but an invalid entity scope: `domain_error(scope, Scope)`

Examples

```
| ?- extends_protocol(listp, Protocol).

| ?- extends_protocol(Protocol, termp, private).
```

See also:

current_protocol/1, *implements_protocol/2-3*, *conforms_to_protocol/2-3*

extends_category/2-3

Description

```
extends_category(Category1, Category2)
extends_category(Category1, Category2, Scope)
```

Enumerates, by backtracking, all pairs of categories such that the first one extends the second. The relation scope is represented by the atoms `public`, `protected`, and `private`.

Template and modes

```
extends_category(?category_identifier, ?category_identifier)
extends_category(?category_identifier, ?category_identifier, ?scope)
```

Errors

Category1 is neither a variable nor a valid protocol identifier: `type_error(category_identifier, Category1)`

Category2 is neither a variable nor a valid protocol identifier: `type_error(category_identifier, Category2)`

Scope is neither a variable nor an atom: `type_error(atom, Scope)`

Scope is an atom but an invalid entity scope: `domain_error(scope, Scope)`

Examples

```
| ?- extends_category(basic, Category).

| ?- extends_category(Category, extended, private).
```

See also:

current_category/1, *complements_object/2*, *imports_category/2-3*

implements_protocol/2-3

Description

```
implements_protocol(Object, Protocol)
implements_protocol(Category, Protocol)

implements_protocol(Object, Protocol, Scope)
implements_protocol(Category, Protocol, Scope)
```

Enumerates, by backtracking, all pairs of entities such that an object or a category implements a protocol. The relation scope is represented by the atoms `public`, `protected`, and `private`. This predicate only returns direct implementation relations; it does not implement a transitive closure.

Template and modes

```
implements_protocol(?object_identifier, ?protocol_identifier)
implements_protocol(?category_identifier, ?protocol_identifier)

implements_protocol(?object_identifier, ?protocol_identifier, ?scope)
implements_protocol(?category_identifier, ?protocol_identifier, ?scope)
```

Errors

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Category is neither a variable nor a valid category identifier: `type_error(category_identifier, Category)`

Protocol is neither a variable nor a valid protocol identifier: `type_error(protocol_identifier, Protocol)`

Scope is neither a variable nor an atom: `type_error(atom, Scope)`

Scope is an atom but an invalid entity scope: `domain_error(scope, Scope)`

Examples

```
| ?- implements_protocol(list, listp).

| ?- implements_protocol(list, listp, public).
```

See also:

current_protocol/1, *conforms_to_protocol/2-3*

conforms_to_protocol/2-3

Description

```
conforms_to_protocol(Object, Protocol)
conforms_to_protocol(Category, Protocol)

conforms_to_protocol(Object, Protocol, Scope)
conforms_to_protocol(Category, Protocol, Scope)
```

Enumerates, by backtracking, all pairs of entities such that an object or a category conforms to a protocol. The relation scope is represented by the atoms `public`, `protected`, and `private`. This predicate implements a transitive closure for the protocol implementation relation.

Template and modes

```
conforms_to_protocol(?object_identifier, ?protocol_identifier)
conforms_to_protocol(?category_identifier, ?protocol_identifier)

conforms_to_protocol(?object_identifier, ?protocol_identifier, ?scope)
conforms_to_protocol(?category_identifier, ?protocol_identifier, ?scope)
```

Errors

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Category is neither a variable nor a valid category identifier: `type_error(category_identifier, Category)`

Protocol is neither a variable nor a valid protocol identifier: `type_error(protocol_identifier, Protocol)`

Scope is neither a variable nor an atom: `type_error(atom, Scope)`

Scope is an atom but an invalid entity scope: `domain_error(scope, Scope)`

Examples

```
| ?- conforms_to_protocol(list, listp).

| ?- conforms_to_protocol(list, listp, public).
```

See also:

current_protocol/1, *implements_protocol/2-3*

complements_object/2

Description

```
complements_object(Category, Object)
```

Enumerates, by backtracking, all category–object pairs such that the category explicitly complements the object.

Template and modes

```
complements_object(?category_identifier, ?object_identifier)
```

Errors

Category is neither a variable nor a valid category identifier: `type_error(category_identifier,
    Prototype)`

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Parent)`

Examples

```
| ?- complements_object(logging, employee).
```

See also:

current_category/1, *imports_category/2-3*

imports_category/2-3

Description

```
imports_category(Object, Category)

imports_category(Object, Category, Scope)
```

Enumerates, by backtracking, importation relations between objects and categories. The relation scope is represented
by the atoms `public`, `protected`, and `private`.

Template and modes

```
imports_category(?object_identifier, ?category_identifier)

imports_category(?object_identifier, ?category_identifier, ?scope)
```

Errors

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Category is neither a variable nor a valid category identifier: `type_error(category_identifier,
    Category)`

Scope is neither a variable nor an atom: `type_error(atom, Scope)`

Scope is an atom but an invalid entity scope: `domain_error(scope, Scope)`

Examples

```
| ?- imports_category(debugger, monitoring).

| ?- imports_category(Object, monitoring, protected).
```

See also:

current_category/1, *complements_object/2*

instantiates_class/2-3

Description

```
instantiates_class(Instance, Class)
instantiates_class(Instance, Class, Scope)
```

Enumerates, by backtracking, all pairs of objects such that the first one instantiates the second. The relation scope is represented by the atoms `public`, `protected`, and `private`.

Template and modes

```
instantiates_class(?object_identifier, ?object_identifier)
instantiates_class(?object_identifier, ?object_identifier, ?scope)
```

Errors

Instance is neither a variable nor a valid object identifier: `type_error(object_identifier, Instance)`

Class is neither a variable nor a valid object identifier: `type_error(object_identifier, Class)`

Scope is neither a variable nor an atom: `type_error(atom, Scope)`

Scope is an atom but an invalid entity scope: `domain_error(scope, Scope)`

Examples

```
| ?- instantiates_class(water_jug, state_space).

| ?- instantiates_class(Space, state_space, public).
```

See also:

current_object/1, *extends_object/2-3*, *specializes_class/2-3*

specializes_class/2-3

Description

```
specializes_class(Class, Superclass)
specializes_class(Class, Superclass, Scope)
```

Enumerates, by backtracking, all pairs of objects such that the first one specializes the second. The relation scope is represented by the atoms `public`, `protected`, and `private`.

Template and modes

```
specializes_class(?object_identifier, ?object_identifier)
specializes_class(?object_identifier, ?object_identifier, ?scope)
```

Errors

Class is neither a variable nor a valid object identifier: `type_error(object_identifier, Class)`

Superclass is neither a variable nor a valid object identifier: `type_error(object_identifier, Superclass)`

Scope is neither a variable nor an atom: `type_error(atom, Scope)`

Scope is an atom but an invalid entity scope: `domain_error(scope, Scope)`

Examples

```
| ?- specializes_class(Subclass, state_space).

| ?- specializes_class(Subclass, state_space, public).
```

See also:

current_object/1, *extends_object/2-3*, *instantiates_class/2-3*

2.3.6 Event handling

abolish_events/5

Description

```
abolish_events(Event, Object, Message, Sender, Monitor)
```

Abolishes all matching events. The two types of events are represented by the atoms `before` and `after`. When the predicate is called with the first argument unbound, both types of events are abolished.

Template and modes

```
abolish_events(@term, @term, @term, @term, @term)
```

Errors

Event is neither a variable nor a valid event identifier: `type_error(event, Event)`

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Message is neither a variable nor a callable term: `type_error(callable, Message)`

Sender is neither a variable nor a valid object identifier: `type_error(object_identifier, Sender)`

Monitor is neither a variable nor a valid object identifier: `type_error(object_identifier, Monitor)`

Examples

```
| ?- abolish_events(_, list, _, _, debugger).
```

See also:

current_event/5, *define_events/5*, *before/3*, *after/3*

current_event/5

Description

```
current_event(Event, Object, Message, Sender, Monitor)
```

Enumerates, by backtracking, all defined events. The two types of events are represented by the atoms `before` and `after`.

Template and modes

```
current_event(?event, ?term, ?term, ?term, ?object_identifier)
```

Errors

Event is neither a variable nor a valid event identifier: `type_error(event, Event)`

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Message is neither a variable nor a callable term: `type_error(callable, Message)`

Sender is neither a variable nor a valid object identifier: `type_error(object_identifier, Sender)`

Monitor is neither a variable nor a valid object identifier: `type_error(object_identifier, Monitor)`

Examples

```
| ?- current_event(Event, Object, Message, Sender, debugger).
```

See also:

abolish_events/5, *define_events/5*, *before/3*, *after/3*

define_events/5

Description

```
define_events(Event, Object, Message, Sender, Monitor)
```

Defines a new set of events. The two types of events are represented by the atoms `before` and `after`. When the predicate is called with the first argument unbound, both types of events are defined. The object `Monitor` must define the event handler methods required by the `Event` argument.

Template and modes

```
define_events(@term, @term, @term, @term, +object_identifier)
```

Errors

Event is neither a variable nor a valid event identifier: `type_error(event, Event)`

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Message is neither a variable nor a callable term: `type_error(callable, Message)`

Sender is neither a variable nor a valid object identifier: `type_error(object_identifier, Sender)`

Monitor is a variable: `instantiation_error`

Monitor is neither a variable nor a valid object identifier: `existence_error(object_identifier, Monitor)`

Monitor does not define the required before/3 method: `existence_error(procedure, before/3)`

Monitor does not define the required after/3 method: `existence_error(procedure, after/3)`

Examples

```
| ?- define_events(_, list, member(_, _), _ , debugger).
```

See also:

abolish_events/5, *current_event/5*, *before/3*, *after/3*

2.3.7 Multi-threading

threaded/1

Description

```
threaded(Goals)

threaded(Conjunction)
threaded(Disjunction)
```

Proves each goal in a conjunction (disjunction) of goals in its own thread. This predicate is deterministic and opaque to cuts. The predicate argument is **not** flattened.

When the argument is a conjunction of goals, a call to this predicate blocks until either all goals succeed, one of the goals fail, or one of the goals generate an exception; the failure of one of the goals or an exception on the execution of one of the goals results in the termination of the remaining threads. The predicate call is true *iff* all goals are true.

When the argument is a disjunction of goals, a call to this predicate blocks until either one of the goals succeeds, all the goals fail, or one of the goals generate an exception; the success of one of the goals or an exception on the execution of one of the goals results in the termination of the remaining threads. The predicate call is true *iff* one of the goals is true.

When the predicate argument is neither a conjunction not a disjunction of goals, no threads are used. In this case, the predicate call is equivalent to a `once/1` predicate call.

Template and modes

```
threaded(+callable)
```

Errors

Goals is a variable: `instantiation_error`

A goal in Goals is a variable: `instantiation_error`

Goals is neither a variable nor a callable term: `type_error(callable, Goals)`

A goal Goal in Goals is neither a variable nor a callable term: `type_error(callable, Goal)`

Examples

Prove a conjunction of goals, each one in its own thread: `threaded((Goal, Goals))`

Prove a disjunction of goals, each one in its own thread: `threaded((Goal; Goals))`

See also:

threaded_call/1-2, *threaded_once/1-2*, *threaded_ignore/1*, *synchronized/1*

threaded_call/1-2

Description

```
threaded_call(Goal)
threaded_call(Goal, Tag)
```

Proves `Goal` asynchronously using a new thread. The argument can be a message sending goal. Calls to this predicate always succeeds and return immediately. The results (success, failure, or exception) are sent back to the message queue of the object containing the call (*this*); they can be retrieved by calling the *threaded_exit/1-2* predicate.

The variant `threaded_call/2` returns a threaded call identifier tag that can be used with the *threaded_exit/1-2* predicate. Tags shall be regarded as opaque terms; users shall not rely on its type.

Template and modes

```
threaded_call(@callable)
threaded_call(@callable, -nonvar)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Tag is not a variable: `type_error(variable, Goal)`

Examples

Prove `Goal` asynchronously in a new thread: `threaded_call(Goal)`

Prove `::Message` asynchronously in a new thread: `threaded_call(::Message)`

Prove `Object::Message` asynchronously in a new thread: `threaded_call(Object::Message)`

See also:

threaded_exit/1-2, *threaded_ignore/1*, *threaded_once/1-2*, *threaded_peek/1-2*, *threaded/1*, *synchronized/1*

threaded_once/1-2

Description

```
threaded_once(Goal)
threaded_once(Goal, Tag)
```

Proves `Goal` asynchronously using a new thread. Only the first goal solution is found. The argument can be a message sending goal. This call always succeeds. The result (success, failure, or exception) is sent back to the message queue of the object containing the call (*this*).

The variant `threaded_once/2` returns a threaded call identifier tag that can be used with the *threaded_exit/1-2* predicate. Tags shall be regarded as opaque terms; users shall not rely on its type.

Template and modes

```
threaded_once(@callable)
threaded_once(@callable, -nonvar)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Tag is not a variable: `type_error(variable, Goal)`

Examples

Prove `Goal` asynchronously in a new thread: `threaded_once(Goal)`

Prove `::Message` asynchronously in a new thread: `threaded_once(::Message)`

Prove `Object::Message` asynchronously in a new thread: `threaded_once(Object::Message)`

See also:

threaded_call/1-2, *threaded_exit/1-2*, *threaded_ignore/1*, *threaded_peek/1-2*, *threaded/1*, *synchronized/1*

threaded_ignore/1

Description

```
threaded_ignore(Goal)
```

Proves `Goal` asynchronously using a new thread. Only the first goal solution is found. The argument can be a message sending goal. This call always succeeds, independently of the result (success, failure, or exception), which is simply discarded instead of being sent back to the message queue of the object containing the call (*this*).

Template and modes

```
threaded_ignore(@callable)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Examples

Prove `Goal` asynchronously in a new thread: `threaded_ignore(Goal)`

Prove `::Message` asynchronously in a new thread: `threaded_ignore(::Message)`

Prove `Object::Message` asynchronously in a new thread: `threaded_ignore(Object::Message)`

See also:

threaded_call/1-2, *threaded_exit/1-2*, *threaded_once/1-2*, *threaded_peek/1-2*, *threaded/1*, *synchronized/1*

threaded_exit/1-2

Description

```
threaded_exit(Goal)
threaded_exit(Goal, Tag)
```

Retrieves the result of proving `Goal` in a new thread. This predicate blocks execution until the reply is sent to the *this* message queue by the thread executing the goal. When there is no thread proving the goal, the predicate generates an exception. This predicate is non-deterministic, providing access to any alternative solutions of its argument.

The argument of this predicate should be a *variant* of the argument of the corresponding *threaded_call/1-2* call. When the predicate argument is subsumed by the `threaded_call/1` call argument, the `threaded_exit/1` call will succeed iff its argument is a solution of the (more general) goal.

The variant `threaded_exit/2` accepts a threaded call identifier tag generated by the calls to the *threaded_call/1-2* and *threaded_once/1-2* predicates. Tags shall be regarded as an opaque term; users shall not rely on its type.

Template and modes

```
threaded_exit(+callable)
threaded_exit(+callable, +nonvar)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

no thread is running for proving Goal: `existence_error(goal_thread, Goal)`

Tag is a variable: `instantiation_error`

Examples

To retrieve an asynchronous goal proof result: `threaded_exit(Goal)`

To retrieve an asynchronous message to *self* result: `threaded_exit(::Goal)`

To retrieve an asynchronous message result: `threaded_exit(Object::Goal)`

See also:

threaded_call/1-2, *threaded_ignore/1*, *threaded_once/1-2*, *threaded_peek/1-2*, *threaded/1*

threaded_peek/1-2

Description

```
threaded_peek(Goal)
threaded_peek(Goal, Tag)
```

Checks if the result of proving `Goal` in a new thread is already available. This call succeeds or fails without blocking execution waiting for a reply to be available.

The argument of this predicate should be a *variant* of the argument of the corresponding *threaded_call/1-2* call. When the predicate argument is subsumed by the `threaded_call/1` call argument, the `threaded_peek/1` call will succeed iff its argument unifies with an already available solution of the (more general) goal.

The variant `threaded_peek/2` accepts a threaded call identifier tag generated by the calls to the *threaded_call/1-2* and *threaded_once/1-2* predicates. Tags shall be regarded as an opaque term; users shall not rely on its type.

Template and modes

```
threaded_peek(+callable)
threaded_peek(+callable, +nonvar)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Tag is a variable: `instantiation_error`

Examples

To check for an asynchronous goal proof result: `threaded_peek(Goal)`

To check for an asynchronous message to *self* result: `threaded_peek(::Goal)`

To check for an asynchronous message result: `threaded_peek(Object::Goal)`

See also:

threaded_call/1-2, *threaded_exit/1-2*, *threaded_ignore/1*, *threaded_once/1-2*, *threaded/1*

threaded_wait/1

Description

```
threaded_wait(Term)
threaded_wait([Term| Terms])
```

Suspends the thread making the call until a notification is received that unifies with `Term`. The call must be made within the same object (*this*) containing the calls to the *threaded_notify/1* predicate that will eventually send the notification. The argument may also be a list of notifications, `[Term| Terms]`. In this case, the thread making the call will suspend until all notifications in the list are received.

Template and modes

```
threaded_wait(?term)
threaded_wait(+list(term))
```

Errors

(none)

Examples

Wait until the `data_available` notification is received: `threaded_wait(data_available)`

See also:

threaded_notify/1

threaded_notify/1

Description

```
threaded_notify(Term)
threaded_notify([Term| Terms])
```

Sends `Term` as a notification to any thread suspended waiting for it in order to proceed. The call must be made within the same object (*this*) containing the calls to the *threaded_wait/1* predicate waiting for the notification. The argument may also be a list of notifications, `[Term| Terms]`. In this case, all notifications in the list will be sent to any threads suspended waiting for them in order to proceed.

Template and modes

```
threaded_notify(@term)
threaded_notify(@list(term))
```

Errors

(none)

Examples

Send the notification `data_available`: `threaded_notify(data_available)`

See also:

threaded_wait/1

2.3.8 Multi-threading engines

threaded_engine_create/3

Description

```
threaded_engine_create(AnswerTemplate, Goal, Engine)
```

Creates a new engine for proving the given goal and defines an answer template for retrieving the goal solution bindings. A message queue for passing arbitrary terms to the engine is also created. If the name for the engine is not given, a unique name is generated and returned. Engine names shall be regarded as opaque terms; users shall not rely on its type.

Template and modes

```
threaded_engine_create(@term, @callable, ?nonvar)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Engine is the name of an existing engine: `permission_error(create, engine, Engine)`

Examples

Create a new engine for finding members of a list: `threaded_engine_create(X, member(X, [1,2, 3]), worker_1)`

See also:

threaded_engine_destroy/1, *threaded_engine_self/1*, *threaded_engine/1*, *threaded_engine_next/2*, *threaded_engine_next_reified/2*

threaded_engine_destroy/1

Description

```
threaded_engine_destroy(Engine)
```

Stops and destroys an engine.

Template and modes

```
threaded_engine_destroy(@nonvar)
```

Errors

Engine is a variable: `instantiation_error`

Engine is neither a variable nor the name of an existing engine: `existence_error(engine, Engine)`

Examples

Stop an engine named `worker_1`: `threaded_engine_destroy(worker_1)`

Stop all engines: `forall(threaded_engine(Engine), threaded_engine_destroy(Engine))`

See also:

threaded_engine_create/3, *threaded_engine_self/1*, *threaded_engine/1*

threaded_engine/1

Description

```
threaded_engine(Engine)
```

Enumerates, by backtracking, all existing engines. Engine names shall be regarded as opaque terms; users shall not rely on its type.

Template and modes

```
threaded_engine(?nonvar)
```

Errors

(none)

Examples

Check that an engine named `worker_1` exists: `threaded_engine(worker_1)`

Write the names of all existing engines: `forall(threaded_engine(Engine), (writeq(Engine), nl))`

See also:

threaded_engine_create/3, *threaded_engine_self/1*, *threaded_engine_destroy/1*

threaded_engine_self/1

Description

```
threaded_engine_self(Engine)
```

Queries the name of engine calling the predicate.

Template and modes

```
threaded_engine_self(?nonvar)
```

Errors

(none)

Examples

Find the name of the engine making the query: `threaded_engine_self(Engine)`

Check if the name of the engine making the query is worker_1: `threaded_engine_self(worker_1)`

See also:

threaded_engine_create/3, *threaded_engine_destroy/1*, *threaded_engine/1*

threaded_engine_next/2

Description

```
threaded_engine_next(Engine, Answer)
```

Retrieves the next answer from an engine. This predicate blocks until an answer becomes available. The predicate fails when there are no more solutions to the engine goal. If the engine goal throws an exception, calling this predicate will re-throw the exception and subsequent calls will fail.

Template and modes

```
threaded_engine_next(@nonvar, ?term)
```

Errors

Engine is a variable: `instantiation_error`

Engine is neither a variable nor the name of an existing engine: `existence_error(engine, Engine)`

Examples

Gets the next engine answer: `threaded_engine_next(worker_1, Answer)`

See also:

threaded_engine_create/3, *threaded_engine_next_reified/2*, *threaded_engine_yield/1*

threaded_engine_next_reified/2

Description

```
threaded_engine_next_reified(Engine, Answer)
```

Retrieves the next reified answer from an engine. This predicate always succeeds and blocks until an answer becomes available. Answers are returned using the terms `the(Answer)`, `no`, and `exception(Error)`.

Template and modes

```
threaded_engine_next_reified(@nonvar, ?nonvar)
```

Errors

Engine is a variable: `instantiation_error`

Engine is neither a variable nor the name of an existing engine: `existence_error(engine, Engine)`

Examples

Gets the next engine answer: `threaded_engine_next_reified(worker_1, Answer)`

See also:

threaded_engine_create/3, *threaded_engine_next/2*, *threaded_engine_yield/1*

threaded_engine_yield/1

Description

```
threaded_engine_yield(Answer)
```

Returns an answer independent of the solutions of the engine goal. Fails if not called from within an engine. This predicate is usually used when the engine goal is call to a recursive predicate processing terms from the engine term queue.

This predicate blocks until the returned answer is consumed.

Note that this predicate should not be called as the last element of a conjunction resulting in an engine goal solution as, in this case, an answer will always be returned. For example, instead of `(threaded_engine_yield(ready); member(X,[1,2,3]))` use `(X=ready; member(X,[1,2,3]))`.

Template and modes

```
threaded_engine_yield(@term)
```

Errors

(none)

Examples

Returns the atom `ready` as an engine answer: `threaded_engine_yield(ready)`

See also:

threaded_engine_create/3, *threaded_engine_next/2*, *threaded_engine_next_reified/2*

threaded_engine_post/2

Description

```
threaded_engine_post(Engine, Term)
```

Posts a term to the engine term queue.

Template and modes

```
threaded_engine_post(@nonvar, @term)
```

Errors

Engine is a variable: `instantiation_error`

Engine is neither a variable nor the name of an existing engine: `existence_error(engine, Engine)`

Examples

Post the atom `ready` to the `worker_1` engine queue: `threaded_engine_post(worker_1, ready)`

See also:

threaded_engine_fetch/1

threaded_engine_fetch/1

Description

```
threaded_engine_fetch(Term)
```

Fetches a term from the engine term queue. Blocks until a term is available. Fails if not called from within an engine.

Template and modes

```
threaded_engine_fetch(?term)
```

Errors

(none)

Examples

Fetch a term from the engine term queue: `threaded_engine_fetch(Term)`

See also:

threaded_engine_post/2

2.3.9 Compiling and loading source files

logtalk_compile/1

Description

```
logtalk_compile(File)
logtalk_compile(Files)
```

Compiles to disk a *source file* or a list of source files using the current default compiler flag values. The Logtalk source file name extension (by default, `.lgt`) can be omitted. Source file paths can be absolute, relative to the current directory, or use library notation. This predicate can also be used to compile Prolog source files as Logtalk source code. When no recognized Logtalk or Prolog extension is specified, the compiler tries first to append a Logtalk source file extension and then a Prolog source file extension. If that fails, the compiler tries to use the file name as-is.

When this predicate is called from the top-level, relative source file paths are resolved using the current working directory. When the calls are made from a source file, relative source file paths are resolved using the source file directory.

Note that only the errors related to problems in the predicate argument are listed below. This predicate fails when errors are found during compilation of a source file.

Template and modes

```
logtalk_compile(@source_file_name)
logtalk_compile(@list(source_file_name))
```

Errors

File is a variable: `instantiation_error`

Files is a variable or a list with an element which is a variable: `instantiation_error`

File, or an element File of the Files list, is neither a variable nor a source file: `type_error(source_file_name, File)`

File, or an element File of the Files list, uses library notation but the library does not exist: `existence_error(library, Library)`

File or an element File of the Files list does not exist: `existence_error(file, File)`

Examples

```
| ?- logtalk_compile(set).

| ?- logtalk_load(types(tree)).

| ?- logtalk_compile([listp, list]).
```

See also:

logtalk_compile/2, *logtalk_load/1*, *logtalk_load/2*, *logtalk_make/0*, *logtalk_make/1*, *logtalk_library_path/2*

logtalk_compile/2

Description

```
logtalk_compile(File, Flags)
logtalk_compile(Files, Flags)
```

Compiles to disk a *source file* or a list of source files using a list of compiler flags. The Logtalk source file name extension (by default, `.lgt`) can be omitted. Source file paths can be absolute, relative to the current directory, or use library notation. This predicate can also be used to compile Prolog source files as Logtalk source code. When no recognized Logtalk or Prolog extension is specified, the compiler tries first to append a Logtalk source file extension and then a Prolog source file extension. If that fails, the compiler tries to use the file name as-is. Compiler flags are represented as *flag(value)*. For a description of the available compiler flags, please see *Compiler flags* in the User Manual.

When this predicate is called from the top-level, relative source file paths are resolved using the current working directory. When the calls are made from a source file, relative source file paths are resolved by default using the source file directory (unless a `relative_to/1` flag is passed).

Note that only the errors related to problems in the predicate argument are listed below. This predicate fails when errors are found during compilation of a source file.

> **Warning:** The compiler flags specified in the second argument only apply to the files listed in the first argument. Notably, if you are compiling a *loader file*, the flags only apply to the loader file itself.

Template and modes

```
logtalk_compile(@source_file_name, @list(compiler_flag))
logtalk_compile(@list(source_file_name), @list(compiler_flag))
```

Errors

File is a variable: `instantiation_error`

Files is a variable or a list with an element which is a variable: `instantiation_error`

File, or an element File of the Files list, is neither a variable nor a source file name:
    `type_error(source_file_name, File)`

File, or an element File of the Files list, uses library notation but the library does not exist:
    `existence_error(library, Library)`

File or an element File of the Files list, does not exist: `existence_error(file, File)`

Flags is a variable or a list with an element which is a variable: `instantiation_error`

Flags is neither a variable nor a proper list: `type_error(list, Flags)`

An element Flag of the Flags list is not a valid compiler flag: `type_error(compiler_flag, Flag)`

An element Flag of the Flags list defines a value for a read-only compiler flag: `permission_error(modify,`
    `flag, Flag)`

An element Flag of the Flags list defines an invalid value for a flag: `domain_error(flag_value,`
    `Flag+Value)`

Examples

```
| ?- logtalk_compile(list, []).

| ?- logtalk_compile(types(tree)).

| ?- logtalk_compile([listp, list], [source_data(off), portability(silent)]).
```

See also:

logtalk_compile/1, *logtalk_load/1*, *logtalk_load/2*, *logtalk_make/0*, *logtalk_make/1*, *logtalk_library_path/2*

logtalk_load/1

Description

```
logtalk_load(File)
logtalk_load(Files)
```

Compiles to disk and then loads to memory a *source file* or a list of source files using the current default compiler flag values. The Logtalk source file name extension (by default, `.lgt`) can be omitted. Source file paths can be absolute, relative to the current directory, or use library notation. This predicate can also be used to compile Prolog source files as Logtalk source code. When no recognized Logtalk or Prolog extension is specified, the compiler tries first to append a Logtalk source file extension and then a Prolog source file extension. If that fails, the compiler tries to use the file name as-is.

When this predicate is called from the top-level, relative source file paths are resolved using the current working directory. When the calls are made from a source file, relative source file paths are resolved using the source file directory.

Note that only the errors related to problems in the predicate argument are listed below. This predicate fails when errors are found during compilation of a source file.

Depending on the back-end Prolog compiler, the shortcuts `{File}` or `{File1, File2, ...}` may be used in alternative. Check the adapter files for the availability of these shortcuts as they are not part of the language (and thus should only be used at the top-level interpreter).

Template and modes

```
logtalk_load(@source_file_name)
logtalk_load(@list(source_file_name))
```

Errors

File is a variable: `instantiation_error`

Files is a variable or a list with an element which is a variable: `instantiation_error`

File, or an element File of the Files list, is neither a variable nor a source file name:
`type_error(source_file_name, File)`

File, or an element File of the Files list, uses library notation but the library does not exist:
`existence_error(library, Library)`

File or an element File of the Files list, does not exist: `existence_error(file, File)`

Examples

```
| ?- logtalk_load(set).

| ?- logtalk_load(types(tree)).

| ?- logtalk_load([listp, list]).
```

See also:

logtalk_compile/1, *logtalk_compile/2*, *logtalk_load/2*, *logtalk_make/0*, *logtalk_make/1*, *logtalk_library_path/2*

logtalk_load/2

Description

```
logtalk_load(File, Flags)
logtalk_load(Files, Flags)
```

Compiles to disk and then loads to memory a *source file* or a list of source files using a list of compiler flags. The Logtalk source file name extension (by default, `.lgt`) can be omitted. Compiler flags are represented as *flag(value)*. This predicate can also be used to compile Prolog source files as Logtalk source code. When no recognized Logtalk or Prolog extension is specified, the compiler tries first to append a Logtalk source file extension and then a Prolog source file extension. If that fails, the compiler tries to use the file name as-is. For a description of the available compiler flags, please consult the *Compiler flags* section in the User Manual. Source file paths can be absolute, relative to the current directory, or use library notation.

When this predicate is called from the top-level, relative source file paths are resolved using the current working directory. When the calls are made from a source file, relative source file paths are resolved by default using the source file directory (unless a `relative_to/1` flag is passed).

Note that only the errors related to problems in the predicate argument are listed below. This predicate fails when errors are found during compilation of a source file.

> **Warning:** The compiler flags specified in the second argument only apply to the files listed in the first argument and not to any files that those files may load or compile. Notably, if you are loading a *loader file*, the flags only apply to the loader file itself and not to the files loaded by it.

Template and modes

```
logtalk_load(@source_file_name, @list(compiler_flag))
logtalk_load(@list(source_file_name), @list(compiler_flag))
```

Errors

File is a variable: `instantiation_error`

Files is a variable or a list with an element which is a variable: `instantiation_error`

File, or an element File of the Files list, is neither a variable nor a source file name:
`type_error(source_file_name, File)`

File, or an element File of the Files list, uses library notation but the library does not exist:
`existence_error(library, Library)`

File or an element File of the Files list, does not exist: `existence_error(file, File)`

Flags is a variable or a list with an element which is a variable: `instantiation_error`

Flags is neither a variable nor a proper list: `type_error(list, Flags)`

An element Flag of the Flags list is not a valid compiler flag: `type_error(compiler_flag, Flag)`

An element Flag of the Flags list defines a value for a read-only compiler flag: `permission_error(modify, flag, Flag)`

An element Flag of the Flags list defines an invalid value for a flag: `domain_error(flag_value, Flag+Value)`

Examples

```
| ?- logtalk_load(list, []).

| ?- logtalk_load(types(tree)).

| ?- logtalk_load([listp, list], [source_data(off), portability(silent)]).
```

See also:

logtalk_compile/1, *logtalk_compile/2*, *logtalk_load/1*, *logtalk_make/0*, *logtalk_make/1*, *logtalk_library_path/2*

logtalk_make/0

Description

```
logtalk_make
```

Reloads all Logtalk source files that have been modified since the time they are last loaded. Only source files loaded using the `logtalk_load/1-2` predicates are reloaded. Non-modified files will also be reloaded when there is a change to the compilation mode (i.e. when the files were loaded without explicit `debug/1` or `optimize/1` flags and the default values of these flags changed after loading; no check is made, however, for other implicit compiler flags that may have changed since loading). When an included file is modified, this predicate reloads its main file (i.e. the file that contains the `include/1` directive).

Depending on the back-end Prolog compiler, the shortcut `{*}` may be used in alternative. Check the adapter files for the availability of these shortcuts as they are not part of the language (and thus should only be used at the top-level interpreter).

This predicate can be extended by the user by defining clauses for the `logtalk_make_target_action/1` multifile and dynamic hook predicate using the argument `all`. The additional user defined actions are run after the default one.

Template and modes

```
logtalk_make
```

Errors

(none)

Examples

```
| ?- logtalk_make.
```

See also:

logtalk_compile/1, *logtalk_compile/2*, *logtalk_load/1*, *logtalk_load/2*, *logtalk_make/1*, *logtalk_make_target_action/1*

logtalk_make/1

Description

```
logtalk_make(Target)
```

Allows reloading all Logtalk source files that have been modified since last loaded when called with the target `all`, deleting all intermediate files generated by the compilation of Logtalk source files when called with the target `clean`, checking for code issues when called with the target `check`, listing of circular dependencies between pairs or trios of objects when called with the target `circular`, generating documentation when called with the target `documentation`, and deleting the dynamic binding caches with the target `caches`.

There are also three variants of the `all` target: `debug`, `normal`, and `optimal`. These targets change the compilation mode (by changing the default value of the `debug` and `optimize` flags) and reload all files affected.

When using the `all` target, only source files loaded using the `logtalk_load/1-2` predicates are reloaded. Non-modified files will also be reloaded when there is a change to the compilation mode (i.e. when the files were loaded without explicit `debug/1` or `optimize/1` flags and the default values of these flags changed after loading; no check is made, however, for other implicit compiler flags that may have changed since loading). When an included file is modified, this target reloads its main file (i.e. the file that contains the `include/1` directive).

When using the `check` or `circular` targets, be sure to compile your source files with the `source_date` flag turned on for complete and detailed reports.

When using the `check` target, predicates for messages sent to objects that implement the `forwarding` built-in protocol are not reported. While this usually avoids only false positives, it may also result in failure to report true missing predicates in some cases.

When using the `circular` target, be prepared for a lengthy computation time for applications with a large combined number of objects and message calls. Only mutual and triangular dependencies are checked due to the computational cost.

The `documentation` target requires the `doclet` tool and a single *doclet* object to be loaded. See the `doclet` tool for more details.

Depending on the back-end Prolog compiler, the shortcuts `{*}`, `{!}`, `{?}`, `{@}`, `{#}`, and `{$}` may be used in alternative to, respectively, `logtalk_make(all)`, `logtalk_make(clean)`, `logtalk_make(check)`, `logtalk_make(circular)`, `logtalk_make(documentation)`, and `logtalk_make(caches)`. The shortcuts `{+d}`, `{+n}`, and `{+o}` may be used in alternative to, respectively, `logtalk_make(debug)`, `logtalk_make(normal)`, and `logtalk_make(optimal)`. Check the adapter files for the availability of these shortcuts as they are not part of the language (and thus should only be used at the top-level interpreter).

The target actions can be extended by defining clauses for the multifile and dynamic hook predicate `logtalk_make_target_action(Target)` where `Target` is one of the targets listed above. The additional user defined actions are run after the default ones.

Template and modes

```
logtalk_make(+atom)
```

Errors

(none)

Examples

```
| ?- logtalk_make(clean).
```

See also:

logtalk_compile/1, *logtalk_compile/2*, *logtalk_load/1*, *logtalk_load/2*, *logtalk_make/0*, *logtalk_make_target_action/1*

logtalk_make_target_action/1

Description

```
logtalk_make_target_action(Target)
```

Multifile and dynamic hook predicate that allows defining user actions for the `logtalk_make/1` targets. The user defined actions are run after the default ones using a failure driven loop.

Template and modes

```
logtalk_make_target_action(+atom)
```

Errors

(none)

Examples

```
% integrate the dead_code_scanner tool with logtalk_make/1

:- multifile(logtalk_make_target_action/1).
:- dynamic(logtalk_make_target_action/1).

logtalk_make_target_action(check) :-
    dead_code_scanner::all.
```

See also:

logtalk_make/1, *logtalk_make/0*

logtalk_library_path/2

Description

```
logtalk_library_path(Library, Path)
```

Dynamic and multifile user-defined predicate, allowing the declaration of aliases to library paths. Library aliases may also be used on the second argument (using the notation *alias(path)*). Paths must always end with the path directory separator character ('/').

Relative paths (e.g. `'../'` or `'./'`) should only be used within the *alias(path)*) notation so that library paths can always be expanded to absolute paths independently of the (usually unpredictable) current directory at the time the `logtalk_library_path/2` predicate is called.

When working with a relocatable application, the actual application installation directory can be retrieved by calling the *logtalk_load_context/2* predicate with the `directory` key and using the returned value to define the `logtalk_library_path/2` predicate. On a settings file, simply use an *initialization/1* directive to wrap the call to the `logtalk_load_context/2` predicate and the assert of the `logtalk_library_path/2` fact.

This predicate may also be used to override the default scratch directory by defining the library alias `scratch_directory` in a backend Prolog initialization file (assumed to be loaded prior to Logtalk loading). This allows e.g. Logtalk to be installed in a read-only directory by setting this alias to the operating-system directory for

temporary files. It also allows several Logtalk instances to run concurrently without conflict by using a unique scratch directory per instance (e.g. using a process ID or a UUID generator).

Template and modes

```
logtalk_library_path(?atom, -atom)
logtalk_library_path(?atom, -compound)
```

Errors

(none)

Examples

```
| ?- logtalk_library_path(viewpoints, Path).

Path = examples('viewpoints/')
yes

| ?- logtalk_library_path(Library, Path).

Library = home,
Path = '$HOME/' ;

Library = logtalk_home,
Path = '$LOGTALKHOME/' ;

Library = logtalk_user
Path = '$LOGTALKUSER/' ;

Library = examples
Path = logtalk_user('examples/') ;

Library = library
Path = logtalk_user('library/') ;

Library = viewpoints
Path = examples('viewpoints/')
yes
```

See also:

logtalk_compile/1, *logtalk_compile/2*, *logtalk_load/1*, *logtalk_load/2*

logtalk_load_context/2

Description

```
logtalk_load_context(Key, Value)
```

Provides access to the Logtalk compilation/loading context. The following keys are currently supported: `entity_identifier`, `entity_prefix`, `entity_type` (returns the value `module` when compiling a module as an object), `source`, `file` (the actual file being compiled, which is different from `source` only when processing an `include/1` directive), `basename`, `directory`, `stream`, `target` (the full path of the intermediate Prolog file), `flags` (the list of the explicit flags used for the compilation of the source file), `term` (the term being expanded), `term_position` (StartLine-EndLine), and `variable_names` ([Name1=Variable1, ...]). The `term_position` key is only supported in back-end Prolog compilers that provide access to the start and end lines of a read term.

The `logtalk_load_context/2` predicate can also be called *initialization/1* directives in a source file. A common scenario is to use the `directory` key to define library aliases.

Currently, any variables in the values of the `term` and `variable_names` keys are not shared with, respectively, the term and goal arguments of the `term_expansion/2` and `goal_expansion/2` methods.

Using the `variable_names` key requires calling the standard built-in predicate `term_variables/2` on the term read and unifying the term variables with the variables in the names list. This, however, may rise portability issues with those Prolog compilers that don't return the variables in the same order for the `term_variables/2` predicate and the option `variable_names/1` of the `read_term/3` built-in predicate, which is used by the Logtalk compiler to read source files.

Template and modes

```
logtalk_load_context(?atom, -nonvar)
```

Errors

(none)

Examples

```
term_expansion(Term, ExpandedTerms) :-
    ...
    logtalk_load_context(entity_identifier, Entity),
    ....

:- initialization((
    logtalk_load_context(directory, Directory),
    assertz(logtalk_library_path(my_app, Directory))
)).
```

See also:

goal_expansion/2, *term_expansion/2*

2.3.10 Flags

current_logtalk_flag/2

Description

```
current_logtalk_flag(Flag, Value)
```

Enumerates, by backtracking, the current Logtalk flag values.

Template and modes

```
current_logtalk_flag(?atom, ?atom)
```

Errors

Flag is neither a variable nor an atom: `type_error(atom, Flag)`

Flag is an atom but an invalid flag: `domain_error(flag, Value)`

Examples

```
| ?- current_logtalk_flag(source_data, Value).
```

See also:

create_logtalk_flag/3, *set_logtalk_flag/2*

set_logtalk_flag/2

Description

```
set_logtalk_flag(Flag, Value)
```

Sets Logtalk default, global, flag values. For local flag scope, use the corresponding *set_logtalk_flag/2* directive. To set a global flag value when compiling and loading a source file, wrap the calls to this built-in predicate with an *initialization/1* directive.

Template and modes

```
set_logtalk_flag(+atom, +nonvar)
```

Errors

Flag is a variable: `instantiation_error`

Value is a variable: `instantiation_error`

Flag is neither a variable nor an atom: `type_error(atom, Flag)`

Flag is an atom but an invalid flag: `domain_error(flag, Flag)`

Value is not a valid value for flag Flag: `domain_error(flag_value, Flag + Value)`

Flag is a read-only flag: `permission_error(modify, flag, Flag)`

Examples

```
| ?- set_logtalk_flag(unknown_entities, silent).
```

See also:

create_logtalk_flag/3, *current_logtalk_flag/2*

create_logtalk_flag/3

Description

```
create_logtalk_flag(Flag, Value, Options)
```

Creates a new Logtalk flag and sets its default value. User-defined flags can be queried and set in the same way as pre-defined flags by using, respectively, the *current_logtalk_flag/2* and *set_logtalk_flag/2* built-in predicates.

This predicate is based on the specification of the SWI-Prolog `create_prolog_flag/3` built-in predicate and supports the same options: `access(Access)`, where `Access` can be either `read_write` (the default) or `read_only`; `keep(Keep)`, where `Keep` can be either `false` (the default) or `true`, for deciding if an existing definition of the flag should be kept or replaced by the new one; and `type(Type)` for specifying the type of the flag, which can be `boolean`, `atom`, `integer`, `float`, or `term` (which only restricts the flag value to ground terms). When the `type/1` option is not specified, the type of the flag is inferred from its initial value.

Template and modes

```
create_logtalk_flag(+atom, +ground, +list(ground))
```

Errors

Flag is a variable: `instantiation_error`

Value is not a ground term: `instantiation_error`

Options is not a ground term: `instantiation_error`

Flag is neither a variable nor an atom: `type_error(atom, Flag)`

Options is neither a variable nor a list: `type_error(atom, Flag)`

Value is not a valid value for flag Flag: `domain_error(flag_value, Flag + Value)`

Flag is a system-defined flag: `permission_error(modify, flag, Flag)`

An element Option of the list Options is not a valid option domain_error(flag_option,Option)

The list Options contains a type(Type) option and Value is not a Type term type_error(Type, Value)

Examples

```
| ?- create_logtalk_flag(pretty_print_blobs, false, []).
```

See also:

current_logtalk_flag/2, *set_logtalk_flag/2*

2.4 Built-in methods

2.4.1 Execution context

context/1

Description

```
context(Context)
```

Returns the execution context for a predicate clause using the term `logtalk(Head, ExecutionContext)` where `Head` is the head of the clause containing the call. This private predicate is mainly used for providing a default error context when type-checking predicate arguments. The `ExecutionContext` term should be regarded as an opaque term, which can be decoded using the `logtalk::execution_context/7` predicate. Calls to this predicate are inlined at compilation time.

Template and modes

```
context(--callable)
```

Errors

Context is not a variable: `type_error(var, Context)`

Examples

```
foo(A, N) :-
    % type-check arguments
    context(Context),
    type::check(atom, A, Context),
    type::check(integer, N, Context),
    % arguments are fine; go ahead
    ... .
```

See also:

parameter/2, *self/1*, *sender/1*, *this/1*

parameter/2

Description

```
parameter(Number, Term)
```

Used in *parametric objects* (and parametric categories), this private method provides runtime access to the parameter values of the entity that contains the predicate clause whose body is being executed by using the argument number in the entity identifier. This predicate is implemented as a unification between its second argument and the corresponding implicit execution-context argument in the predicate containing the call. This unification occurs at the clause head when the second argument is not instantiated (the most common case). When the second argument is instantiated, the unification must be delayed to runtime and thus occurs at the clause body.

Entity parameters can also be accessed using *parametric variables*, which use the syntax _VariableName_. The compiler recognizes occurrences of these variables in entity clauses. Parametric variables allows us to abstract parameter positions thus simplifying code maintenance.

Template and modes

```
parameter(+integer, ?term)
```

Errors

Number is a variable: `instantiation_error`

Number is neither a variable nor an integer value: `type_error(integer, Number)`

Number is smaller than one or greater than the parametric entity identifier arity:
`domain_error(out_of_range, Number)`

Entity identifier is not a compound term: `type_error(compound, Entity)`

Examples

```
:- object(box(_Color, _Weight)).

    ...

    color(Color) :-              % this clause is translated into a fact
        parameter(1, Color).     % upon compilation

    heavy :-
        parameter(2, Weight),    % after compilation, the >/2 call will be
        Weight > 10.             % the first condition on the clause body

    ...
```

The same example using *parametric variables*:

```
:- object(box(_Color_, _Weight_)).

    ...
```

```
    color(_Color_).

    heavy :-
        _Weight_ > 10.

    ...
```

See also:

context/1, *self/1*, *sender/1*, *this/1*

self/1

Description

```
self(Self)
```

Returns the object that has received the message under processing. This private method is translated to a unification between its argument and the corresponding implicit context argument in the predicate containing the call. This unification occurs at the clause head when the argument is not instantiated (the most common case).

Template and modes

```
self(?object_identifier)
```

Errors

(none)

Examples

```
test :-
    self(Self),                              % after compilation, the write/1
    write('executing a method in behalf of '),  % call will be the first goal on
    writeq(Self), nl.                        % the clause body
```

See also:

context/1, *parameter/2*, *sender/1*, *this/1*

sender/1

Description

```
sender(Sender)
```

Returns the object that has sent the message under processing. This private method is translated into a unification between its argument and the corresponding implicit context argument in the predicate containing the call. This unification occurs at the clause head when the argument is not instantiated (the most common case).

Template and modes

```
sender(?object_identifier)
```

Errors

(none)

Examples

```
% after compilation, the write/1 call will be the first goal on the clause body:

test :-
    sender(Sender),
    write('executing a method to answer a message sent by '),
    writeq(Sender), nl.
```

See also:

context/1, *parameter/2*, *self/1*, *this/1*

this/1

Description

```
this(This)
```

Unifies its argument with the identifier of the object for which the predicate clause whose body is being executed is defined (or the object importing the category that contains the predicate clause). This private method is implemented as a unification between its argument and the corresponding implicit execution-context argument in the predicate containing the call. This unification occurs at the clause head when the argument is not instantiated (the most common case). This method is useful for avoiding hard-coding references to an object identifier or for retrieving all object parameters with a single call when using parametric objects.

Template and modes

```
this(?object_identifier)
```

Errors

(none)

Examples

```
% after compilation, the write/1 call will be the first goal on the clause body:

test :-
    this(This),
    write('Using a predicate clause contained in '),
    writeq(This), nl.
```

See also:

context/1, *parameter/2*, *self/1*, *sender/1*

2.4.2 Reflection

current_op/3

Description

```
current_op(Priority, Specifier, Operator)
```

Enumerates, by backtracking, the visible operators declared for an object. Operators not declared using a scope directive are not enumerated.

Template and modes

```
current_op(?operator_priority, ?operator_specifier, ?atom)
```

Errors

Priority is neither a variable nor an integer: `type_error(integer, Priority)`

Priority is an integer but not a valid operator priority: `domain_error(operator_priority, Priority)`

Specifier is neither a variable nor an atom: `type_error(atom, Specifier)`

Specifier is an atom but not a valid operator specifier: `domain_error(operator_specifier, Specifier)`

Operator is neither a variable nor an atom: `type_error(atom, Operator)`

Examples

To enumerate, by backtracking, the local operators or the operators visible in *this*:
    `current_op(Priority, Specifier, Operator)`

To enumerate, by backtracking, the public and protected operators visible in *self*:
    `::current_op(Priority, Specifier, Operator)`

To enumerate, by backtracking, the public operators visible for an explicit object:
    `Object::current_op(Priority, Specifier, Operator)`

See also:

current_predicate/1, *predicate_property/2*, *op/3*

current_predicate/1

Description

```
current_predicate(Predicate)
```

Enumerates, by backtracking, visible user predicates. When the predicate is declared in a `uses/2` or `use_module/2` directive, predicates are enumerated for the referenced object or module. Otherwise, predicates are enumerated for an object. In the case of objects, predicates not declared using a scope directive are not enumerated.

Template and modes

```
current_predicate(?predicate_indicator)
```

Errors

Predicate is neither a variable nor a valid predicate indicator: `type_error(predicate_indicator, Predicate)`

Predicate is a Name/Arity term but Functor is neither a variable nor an atom: `type_error(atom, Name)`

Predicate is a Name/Arity term but Arity is neither a variable nor an integer: `type_error(integer, Arity)`

Predicate is a Name/Arity term but Arity is a negative integer: `domain_error(not_less_than_zero, Arity)`

Examples

To enumerate, by backtracking, the locally visible user predicates or the user predicates visible in *this*:
    `current_predicate(Predicate)`

To enumerate, by backtracking, the public and protected user predicates visible in *self*:
    `::current_predicate(Predicate)`

To enumerate, by backtracking, the public user predicates visible for an explicit object:
    `Object::current_predicate(Predicate)`

See also:

current_op/3, *predicate_property/2*, *uses/2*, *use_module/2*

predicate_property/2

Description

```
predicate_property(Predicate, Property)
```

Enumerates, by backtracking, the properties of a visible predicate. When the predicate indicator for `Predicate` is declared in a `uses/2` or `use_module/2` directive, properties are enumerated for the referenced object or module predicate. Otherwise, properties are enumerated for an object predicate. In the case of objects, properties for predicates not declared using a scope directive are not enumerated. The valid predicate properties are listed in the language *grammar*.

Template and modes

```
predicate_property(+callable, ?predicate_property)
```

Errors

Predicate is a variable: `instantiation_error`

Predicate is neither a variable nor a callable term: `type_error(callable, Predicate)`

Property is neither a variable nor a valid predicate property: `domain_error(predicate_property, Property)`

Examples

To enumerate, by backtracking, the properties of a locally visible user predicate or a user predicate visible in *this*:
`predicate_property(foo(_), Property)`

To enumerate, by backtracking, the properties of a public or protected predicate visible in *self*:
`::predicate_property(foo(_), Property)`

To enumerate, by backtracking, the properties of a public predicate visible in an explicit object:
`Object::predicate_property(foo(_), Property)`

See also:

current_op/3, *current_predicate/1*, *uses/2*, *use_module/2*

2.4.3 Database

abolish/1

Description

```
abolish(Predicate)
abolish(Name/Arity)
```

Abolishes a runtime declared dynamic predicate or a local dynamic predicate. When the predicate indicator for `Head` is declared in a *uses/2* or *use_module/2* directive, the predicate is abolished in the referenced object or module. Otherwise, the predicate is abolished in an object's database. In the case of objects, only predicates that are dynamically declared (using a call to the `asserta/1` or `assertz/1` built-in methods) can be abolished.

Template and modes

```
abolish(+predicate_indicator)
```

Errors

Predicate is a variable: `instantiation_error`

Functor is a variable: `instantiation_error`

Arity is a variable: `instantiation_error`

Predicate is neither a variable nor a valid predicate indicator: `type_error(predicate_indicator, Predicate)`

Functor is neither a variable nor an atom: `type_error(atom, Functor)`

Arity is neither a variable nor an integer: `type_error(integer, Arity)`

Predicate is statically declared: `permission_error(modify, predicate_declaration, Name/Arity)`

Predicate is a private predicate: `permission_error(modify, private_predicate, Name/Arity)`

Predicate is a protected predicate: `permission_error(modify, protected_predicate, Name/Arity)`

Predicate is a static predicate: `permission_error(modify, static_predicate, Name/Arity)`

Predicate is not declared for the object receiving the message: `existence_error(predicate_declaration, Name/Arity)`

Examples

To abolish a local dynamic predicate or a dynamic predicate in *this*: `abolish(Predicate)`

To abolish a public or protected dynamic predicate in *self*: `::abolish(Predicate)`

To abolish a public dynamic predicate in an explicit object: `Object::abolish(Predicate)`

See also:

asserta/1, *assertz/1*, *clause/2*, *retract/1*, *retractall/1 dynamic/0*, *dynamic/1*, *uses/2*, *use_module/2*

asserta/1

Description

```
asserta(Head)
asserta((Head:-Body))
```

Asserts a clause as the first one for a dynamic predicate. When the predicate indicator for `Head` is declared in a *uses/2* or *use_module/2* directive, the clause is asserted in the referenced object or module. Otherwise, the clause is asserted for an object's dynamic predicate. If the predicate is not previously declared (using a scope directive), then a dynamic predicate declaration is added to the object (assuming that we are asserting locally or that the compiler flag `dynamic_declarations` was set to `allow` when the object was created or compiled).

This method may be used to assert clauses for predicates that are not declared dynamic for dynamic objects provided that the predicates are declared in *this*. This allows easy initialization of dynamically created objects when writing constructors.

Template and modes

```
asserta(+clause)
```

Errors

Head is a variable: `instantiation_error`

Head is a neither a variable nor a callable term: `type_error(callable, Head)`

Body cannot be converted to a goal: `type_error(callable, Body)`

The predicate indicator of Head, Name/Arity, is that of a private predicate: `permission_error(modify, private_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a protected predicate: `permission_error(modify, protected_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a static predicate: `permission_error(modify, static_predicate, Name/Arity)`

Target object was created/compiled with support for dynamic declaration of predicates turned off: `permission_error(create, predicate_declaration, Name/Arity)`

Examples

To assert a clause as the first one for a local dynamic predicate or a dynamic predicate in *this*: `asserta(Clause)`

To assert a clause as the first one for any public or protected dynamic predicate in *self*: `::asserta(Clause)`

To assert a clause as the first one for any public dynamic predicate in an explicit object: `Object::asserta(Clause)`

See also:

abolish/1, *assertz/1*, *clause/2*, *retract/1*, *retractall/1 dynamic/0*, *dynamic/1*, *uses/2*, *use_module/2*

assertz/1

Description

```
assertz(Head)
assertz((Head:-Body))
```

Asserts a clause as the last one for a dynamic predicate. When the predicate indicator for `Head` is declared in a *uses/2* or *use_module/2* directive, the clause is asserted in the referenced object or module. Otherwise, the clause is asserted for an object's dynamic predicate. If the predicate is not previously declared (using a scope directive), then a

dynamic predicate declaration is added to the object (assuming that we are asserting locally or that the compiler flag `dynamic_declarations` was set to `allow` when the object was created or compiled).

This method may be used to assert clauses for predicates that are not declared dynamic for dynamic objects provided that the predicates are declared in *this*. This allows easy initialization of dynamically created objects when writing constructors.

Template and modes

```
assertz(+clause)
```

Errors

Head is a variable: `instantiation_error`

Head is a neither a variable nor a callable term: `type_error(callable, Head)`

Body cannot be converted to a goal: `type_error(callable, Body)`

The predicate indicator of Head, Name/Arity, is that of a private predicate: `permission_error(modify, private_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a protected predicate: `permission_error(modify, protected_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a static predicate: `permission_error(modify, static_predicate, Name/Arity)`

Target object was created/compiled with support for dynamic declaration of predicates turned off: `permission_error(create, predicate_declaration, Name/Arity)`

Examples

To assert a clause as the last one for a local dynamic predicate or a dynamic predicate in *this*: `assertz(Clause)`

To assert a clause as the last one for any public or protected dynamic predicate in *self*: `::assertz(Clause)`

To assert a clause as the last one for any public dynamic predicate in an explicit object: `Object::assertz(Clause)`

See also:

abolish/1, *asserta/1*, *clause/2*, *retract/1*, *retractall/1 dynamic/0*, *dynamic/1*, *uses/2*, *use_module/2*

clause/2

Description

```
clause(Head, Body)
```

Enumerates, by backtracking, the clauses of a dynamic predicate. When the predicate indicator for `Head` is declared in a `uses/2` or `use_module/2` directive, the predicate enumerates the clauses in the referenced object or module. Otherwise, it enumerates the clauses for an object's dynamic predicate.

This method may be used to enumerate clauses for predicates that are not declared dynamic for dynamic objects provided that the predicates are declared in *this*.

Template and modes

```
clause(+callable, ?body)
```

Errors

Head is a variable: `instantiation_error`

Head is a neither a variable nor a callable term: `type_error(callable, Head)`

Body is a neither a variable nor a callable term: `type_error(callable, Body)`

The predicate indicator of Head, Name/Arity, is that of a private predicate: `permission_error(access, private_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a protected predicate: `permission_error(access, protected_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a static predicate: `permission_error(access, static_predicate, Name/Arity)`

Head is not a declared predicate: `existence_error(predicate_declaration, Name/Arity)`

Examples

To retrieve a matching clause of a local dynamic predicate or a dynamic predicate in *this*: `clause(Head, Body)`

To retrieve a matching clause of a public or protected dynamic predicate in *self*: `::clause(Head, Body)`

To retrieve a matching clause of a public dynamic predicate in an explicit object: `Object::clause(Head, Body)`

See also:

abolish/1, *asserta/1*, *assertz/1*, *retract/1*, *retractall/1 dynamic/0*, *dynamic/1*, *uses/2*, *use_module/2*

retract/1

Description

```
retract(Head)
retract((Head:-Body))
```

Retracts a clause for a dynamic predicate. When the predicate indicator for `Head` is declared in a *uses/2* or *use_module/2* directive, the clause is retracted in the referenced object or module. Otherwise, the clause is retracted in an object's dynamic predicate. On backtracking, the predicate retracts the next matching clause.

This method may be used to retract clauses for predicates that are not declared dynamic for dynamic objects provided that the predicates are declared in *this*.

Template and modes

```
retract(+clause)
```

Errors

Head is a variable: `instantiation_error`

Head is neither a variable nor a callable term: `type_error(callable, Head)`

The predicate indicator of Head, Name/Arity, is that of a private predicate: `permission_error(modify, private_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a protected predicate: `permission_error(modify, protected_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a static predicate: `permission_error(modify, static_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is not declared: `existence_error(predicate_declaration, Name/Arity)`

Examples

To retract a matching clause of a local dynamic predicate or a dynamic predicate in *this*:
`retract(Clause)`

To retract a matching clause of a public or protected dynamic predicate in *self*: `::retract(Clause)`

To retract a matching clause of a public dynamic predicate in an explicit object:
`Object::retract(Clause)`

See also:

abolish/1, *asserta/1*, *assertz/1*, *clause/2*, *retractall/1*, *dynamic/0*, *dynamic/1*, *uses/2*, *use_module/2*

retractall/1

Description

```
retractall(Head)
```

Retracts all clauses with a matching head for a dynamic predicate. When the predicate indicator for `Head` is declared in a *uses/2* or *use_module/2* directive, the clauses are retracted in the referenced object or module. Otherwise, the clauses are retracted in an object's dynamic predicate.

This method may be used to retract clauses for predicates that are not declared dynamic for dynamic objects provided that the predicates are declared in *this*.

Template and modes

```
retractall(+callable)
```

Errors

Head is a variable: `instantiation_error`

Head is neither a variable nor a callable term: `type_error(callable, Head)`

The predicate indicator of Head, Name/Arity, is that of a private predicate: `permission_error(modify, private_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a protected predicate: `permission_error(modify, protected_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is that of a static predicate: `permission_error(modify, static_predicate, Name/Arity)`

The predicate indicator of Head, Name/Arity, is not declared: `existence_error(predicate_declaration, Name/Arity)`

Examples

To retract all local predicate clauses or all predicate clauses in *this* with a matching head: `retractall(Head)`

To retract all public or protected predicate clauses with a matching head in *self*: `::retractall(Head)`

To retract all public predicate clauses with a matching head in an explicit object: `Object::retractall(Head)`

See also:

abolish/1, *asserta/1*, *assertz/1*, *clause/2*, *retract/1*, *dynamic/0*, *dynamic/1*, *uses/2*, *use_module/2*

2.4.4 Meta-calls

call/1-N

Description

```
call(Goal)
call(Closure, Arg1, ...)
```

Calls a goal, which might be constructed by appending additional arguments to a closure. The upper limit for `N` depends on the upper limit for the arity of a compound term of the back-end Prolog compiler. This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object. The `Closure` argument can also be a lambda expression or a Logtalk control construct. When using a back-end Prolog compiler supporting a module system, calls in the format `call(Module:Closure, Arg1, ...)` may also be used.

This meta-predicate is opaque to cuts in its arguments.

Template and modes

```
call(+callable)
call(+callable, ?term)
call(+callable, ?term, ?term)
...
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Closure is a variable: `instantiation_error`

Closure is neither a variable nor a callable term: `type_error(callable, Closure)`

Examples

Call a goal, constructed by appending additional arguments to a closure, in the context of the object or category containing the
    `call(Closure, Arg1, Arg2, ...)`

To send a goal, constructed by appending additional arguments to a closure, as a message to *self*:
    `call(::Closure, Arg1, Arg2, ...)`

To send a goal, constructed by appending additional arguments to a closure, as a message to an explicit object:
    `call(Object::Closure, Arg1, Arg2, ...)`

See also:

ignore/1, *once/1*, *\+/1*

ignore/1

Description

```
ignore(Goal)
```

This predicate succeeds whether its argument succeeds or fails and it is not re-executable. This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object.

This meta-predicate is opaque to cuts in its argument.

Template and modes

```
ignore(+callable)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Examples

Call a goal and succeeding even if it fails: `ignore(Goal)`

To send a message succeeding even if it fails to *self***:** `ignore(::Goal)`

To send a message succeeding even if it fails to an explicit object: `ignore(Object::Goal)`

See also:

call/1-N, *once/1*, *\+/1*

once/1

Description

```
once(Goal)
```

This predicate behaves as `call(Goal)` but it is not re-executable. This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object.

This meta-predicate is opaque to cuts in its argument.

Template and modes

```
once(+callable)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Examples

Call a goal deterministically in the context of the object or category containing the call: `once(Goal)`

To send a goal as a non-backtracable message to *self***:** `once(::Goal)`

To send a goal as a non-backtracable message to an explicit object: `once(Object::Goal)`

See also:

call/1-N, *ignore/1*, *\+/1*

\+/1

Description

```
\+ Goal
```

Not-provable meta-predicate. True iff `call(Goal)` is false. This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
\+ +callable
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Examples

Not-provable goal in the context of the object or category containing the call: `\+ Goal`

Not-provable goal sent as a message to *self*: `\+ ::Goal`

Not-provable goal sent as a message to an explicit object: `\+ Object::Goal`

See also:

call/1-N, *ignore/1*, *once/1*

2.4.5 Exception-handling

catch/3

Description

```
catch(Goal, Catcher, Recovery)
```

Catches exceptions thrown by a goal. See the ISO Prolog standard definition. This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
catch(?callable, ?term, ?term)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Examples

(none)

See also:

throw/1, *context/1*, *instantiation_error/0*, *type_error/2*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *evaluation_error/1*, *representation_error/1 resource_error/1*, *syntax_error/1*, *system_error/0*

throw/1

Description

```
throw(Exception)
```

Throws an exception. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
throw(+nonvar)
```

Errors

Exception is a variable: `instantiation_error`

Exception does not unify with the second argument of any call of *catch/3*: `system_error`

Examples

(none)

See also:

catch/3, *context/1*, *instantiation_error/0*, *type_error/2*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *evaluation_error/1*, *representation_error/1 resource_error/1*, *syntax_error/1*, *system_error/0*

instantiation_error/0

Description

```
instantiation_error
```

Throws an `error(instantiation_error, logtalk(Head,Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
instantiation_error
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *type_error/2*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *representation_error/1*, *evaluation_error/1*, *resource_error/1*, *syntax_error/1*, *system_error/0*

type_error/2

Description

```
type_error(Type, Culprit)
```

Throws an `error(type_error(Type,Culprit), logtalk(Head,Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
type_error(+nonvar,+term)
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *instantiation_error/0*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *representation_error/1* *evaluation_error/1*, *resource_error/1*, *syntax_error/1*, *system_error/0*,

domain_error/2

Description

```
domain_error(Domain, Culprit)
```

Throws an `error(domain_error(Domain,Culprit), logtalk(Head,Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
domain_error(+atom, +nonvar)
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *instantiation_error/0*, *type_error/2*, *existence_error/2*, *permission_error/3*, *representation_error/1*, *evaluation_error/1*, *resource_error/1*, *syntax_error/1*, *system_error/0*

existence_error/2

Description

```
existence_error(Thing, Culprit)
```

Throws an `error(existence_error(Thing,Culprit), logtalk(Head,Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
existence_error(+nonvar)
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *instantiation_error/0*, *type_error/2*, *domain_error/2*, *evaluation_error/1*, *permission_error/3*, *representation_error/1*, *resource_error/1*, *syntax_error/1*, *system_error/0*

permission_error/3

Description

```
permission_error(Operation, Permission, Culprit)
```

Throws an `error(permission_error(Operation,Permission,Culprit), logtalk(Head, Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
permission_error(+nonvar)
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *instantiation_error/0*, *type_error/2*, *domain_error/2*, *existence_error/2*, *representation_error/1*, *evaluation_error/1*, *resource_error/1*, *syntax_error/1*, *system_error/0*

representation_error/1

Description

```
representation_error(Flag)
```

Throws an `error(representation_error(Flag), logtalk(Head,Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
representation_error(+atom)
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *instantiation_error/0*, *type_error/2*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *evaluation_error/1*, *resource_error/1*, *syntax_error/1*, *system_error/0*

evaluation_error/1

Description

```
evaluation_error(Exception)
```

Throws an `error(evaluation_error(Exception), logtalk(Head,Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
evaluation_error(+nonvar)
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *instantiation_error/0*, *type_error/2*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *representation_error/1*, *resource_error/1*, *syntax_error/1*, *system_error/0*

resource_error/1

Description

```
resource_error(Resource)
```

Throws an `error(resource_error(Resource), logtalk(Head,Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
resource_error(+nonvar)
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *instantiation_error/0*, *type_error/2*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *representation_error/1*, *instantiation_error/0*, *syntax_error/1*, *system_error/0*

syntax_error/1

Description

```
syntax_error(Description)
```

Throws an `error(syntax_error(Description), logtalk(Head,Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
syntax_error(+nonvar)
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *instantiation_error/0*, *type_error/2*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *representation_error/1*, *instantiation_error/0*, *system_error/0 resource_error/1*

system_error/0

Description

```
system_error(Description)
```

Throws an `error(system_error(Description), logtalk(Head,Context))` exception term where `Head` is the head of the clause from where this predicate is called and `Context` is the execution context of the call. This built-in predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
system_error(+nonvar)
```

Errors

(none)

Examples

(none)

See also:

catch/3, *throw/1*, *context/1*, *instantiation_error/0*, *type_error/2*, *domain_error/2*, *existence_error/2*, *permission_error/3*, *representation_error/1 evaluation_error/1*, *resource_error/1*, *syntax_error/1*,

2.4.6 All solutions

bagof/3

Description

```
bagof(Template, Goal, List)
```

Collects a bag of solutions for the goal for each set of instantiations of the free variables in the goal. The order of the elements in the bag follows the order of the goal solutions. The free variables in the goal are the variables that occur in the goal but not in the template. Free variables can be ignored, however, by using the `^/2` existential qualifier. For example, if `T` is term containing all the free variables that we want to ignore, we can write `T^Goal`. Note that the term `T` can be written as `V1^V2^....`

When there are free variables, this method is re-executable on backtracking. This method fails when there are no solutions, never returning an empty list.

This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
bagof(@term, +callable, -list)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Goal is a call to a non-existing predicate: `existence_error(procedure, Predicate)`

Examples

To find a bag of solutions in the context of the object or category containing the call: `bagof(Template, Goal, List)`

To find a bag of solutions of sending a message to *self*: `bagof(Template, ::Message, List)`

To find a bag of solutions of sending a message to an explicit object: `bagof(Template, Object::Message, List)`

See also:

findall/3, *findall/4*, *forall/2*, *setof/3*

findall/3

Description

```
findall(Template, Goal, List)
```

Collects a list of solutions for the goal. The order of the elements in the list follows the order of the goal solutions. It succeeds returning an empty list when the goal has no solutions.

This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
findall(?term, +callable, ?list)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Goal is a call to a non-existing predicate: `existence_error(procedure, Predicate)`

Examples

To find all solutions in the context of the object or category containing the call: `findall(Template, Goal, List)`

To find all solutions of sending a message to *self*: `findall(Template, ::Message, List)`

To find all solutions of sending a message to an explicit object: `findall(Template, Object::Message, List)`

See also:

bagof/3, *findall/4*, *forall/2*, *setof/3*

findall/4

Description

```
findall(Template, Goal, List, Tail)
```

Variant of the `findall/3` method that allows passing the tail of the results list. It succeeds returning the tail argument when the goal has no solutions.

This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
findall(?term, +callable, ?list, +list)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Goal is a call to a non-existing predicate: `existence_error(procedure, Predicate)`

Examples

To find all solutions in the context of the object or category containing the call: `findall(Template, Goal, List, Tail)`

To find all solutions of sending a message to *self*: `findall(Template, ::Message, List, Tail)`

To find all solutions of sending a message to an explicit object: `findall(Template, Object::Message, List, Tail)`

See also:

bagof/3, *findall/3*, *forall/2*, *setof/3*

forall/2

Description

```
forall(Generator, Test)
```

For all solutions of `Generator`, `Test` is true. This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
forall(+callable, +callable)
```

Errors

Either Generator or Test is a variable: `instantiation_error`

Generator is neither a variable nor a callable term: `type_error(callable, Generator)`

Test is neither a variable nor a callable term: `type_error(callable, Test)`

Examples

To call both goals in the context of the object or category containing the call: `forall(Generator, Test)`

To send both goals as messages to *self*: `forall(::Generator, ::Test)`

To send both goals as messages to explicit objects: `forall(Object1::Generator, Object2::Test)`

See also:

bagof/3, *findall/3*, *findall/4*, *setof/3*

setof/3

Description

```
setof(Template, Goal, List)
```

Collects a set of solutions for the goal for each set of instantiations of the free variables in the goal. The solutions are sorted using standard term order. The free variables in the goal are the variables that occur in the goal but not in the template. Free variables can be ignored, however, by using the `^/2` existential qualifier. For example, if `T` is term containing all the free variables that we want to ignore, we can write `T^Goal`. Note that the term `T` can be written as `V1^V2^....`.

When there are free variables, this method is re-executable on backtracking. This method fails when there are no solutions, never returning an empty list.

This built-in meta-predicate is declared as a private method and thus cannot be used as a message to an object.

Template and modes

```
setof(@term, +callable, -list)
```

Errors

Goal is a variable: `instantiation_error`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Goal is a call to a non-existing predicate: `existence_error(procedure, Predicate)`

Examples

To find a set of solutions in the context of the object or category containing the call: `setof(Template, Goal, List)`

To find a set of solutions of sending a message to *self*: `setof(Template, ::Message, List)`

To find a set of solutions of sending a message to an explicit object: `setof(Template, Object::Message, List)`

See also:

bagof/3, *findall/3*, *findall/4*, *forall/2*

2.4.7 Event handling

before/3

Description

```
before(Object, Message, Sender)
```

User-defined method for handling `before` events. This method is declared in the `monitoring` built-in protocol as a public predicate. Note that you can make its scope protected or private by using, respectively, protected or private implementation of the `monitoring` protocol.

Template and modes

```
before(?object_identifier, ?callable, ?object_identifier)
```

Errors

(none)

Examples

```
:- object(...,
    implements(monitoring),
    ...).

    before(Object, Message, Sender) :-
        writeq(Object), write('::'), writeq(Message),
        write(' from '), writeq(Sender), nl.
```

See also:

after/3, *abolish_events/5*, *current_event/5*, *define_events/5*

after/3

Description

```
after(Object, Message, Sender)
```

User-defined method for handling `after` events. This method is declared in the `monitoring` built-in protocol as a public predicate. Note that you can make its scope protected or private by using, respectively, protected or private implementation of the `monitoring` protocol.

Template and modes

```
after(?object_identifier, ?callable, ?object_identifier)
```

Errors

(none)

Examples

```
:- object(...,
    implements(monitoring),
    ...).

    after(Object, Message, Sender) :-
        writeq(Object), write('::'), writeq(Message),
        write(' from '), writeq(Sender), nl.
```

See also:

before/3, *abolish_events/5*, *current_event/5*, *define_events/5*

2.4.8 Message forwarding

forward/1

Description

```
forward(Message)
```

User-defined method for forwarding unknown messages sent to an object (using the *::/2* control construct), auto-matically called by the runtime when defined. This method is declared in the `forwarding` built-in protocol as a public predicate. Note that you can make its scope protected or private by using, respectively, protected or private implementation of the `forwarding` protocol.

Template and modes

```
forward(+callable)
```

Errors

(none)

Examples

```
:- object(proxy,
    implements(forwarding),
    ...).

    forward(Message) :-
        % delegate the unknown message to other object
        [real::Message].
```

See also:

[]/1

2.4.9 Definite clause grammar rules

call//1-N

Description

```
call(Closure)
call(Closure, Arg1, ...)
call(Object::Closure, Arg1, ...)
call(::Closure, Arg1, ...)
```

This non-terminal takes a closure and is processed by appending the input list of tokens and the list of remaining tokens to the arguments of the closure. This built-in non-terminal is interpreted as a private non-terminal and thus cannot be used as a message to an object. When using a back-end Prolog compiler supporting a module system, calls in the format `call(Module:Closure)` may also be used. By using as argument a *lambda expression*, this built-in non-terminal can provide controlled access to the input list of tokens and to the list of the remaining tokens processed by the grammar rule containing the call.

Template and modes

```
call(+callable)
call(+callable, ?term)
call(+callable, ?term, ?term)
...
```

Errors

Closure is a variable: `instantiation_error`

Closure is neither a variable nor a callable term: `type_error(callable, Closure)`

Examples

Calls a goal, constructed by appending the input list of tokens and the list of remaining tokens to the arguments of the closure, i
> `call(Closure)`

To send a goal, constructed by appending the input list of tokens and the list of remaining tokens to the arguments of the closure
> `call(::Closure)`

To send a goal, constructed by appending the input list of tokens and the list of remaining tokens to the arguments of the closure
> `call(Object::Closure)`

See also:

eos//0, *phrase//1*, *phrase/2*, *phrase/3*

eos//0

Description

```
eos
```

This non-terminal matches the end-of-input. It is implemented by checking that the implicit difference list unifies with `[]-[]`.

Template and modes

```
eos
```

Errors

(none)

Examples

```
abc --> a, b, c, eos.
```

See also:

call//1-N, *phrase//1*, *phrase/2*, *phrase/3*

phrase//1

Description

```
phrase(NonTerminal)
```

This non-terminal takes a non-terminal or a grammar rule body and parses it using the current implicit list of tokens. A common use is to wrap what otherwise would be a naked variable in a grammar rule body.

Template and modes

```
phrase(+callable)
```

Errors

NonTerminal is a variable: `instantiation_error`

NonTerminal is neither a variable nor a callable term: `type_error(callable, NonTerminal)`

Examples

(none)

See also:

call//1-N, *phrase/2*, *phrase/3*

phrase/2

Description

```
phrase(GrammarRuleBody, Input)
phrase(::GrammarRuleBody, Input)
phrase(Object::GrammarRuleBody, Input)
```

True when the `GrammarRuleBody` grammar rule body can be applied to the `Input` list of tokens. In the most common case, `GrammarRuleBody` is a non-terminal defined by a grammar rule. This built-in method is declared private and thus cannot be used as a message to an object. When using a back-end Prolog compiler supporting a module system, calls in the format `phrase(Module:GrammarRuleBody, Input)` may also be used.

This method is opaque to cuts in the first argument. When the first argument is sufficiently instantiated at compile time, the method call is compiled in order to eliminate the implicit overheads of converting the grammar rule body into a goal and meta-calling it. For performance reasons, the second argument is only type-checked at compile time.

Template and modes

```
phrase(+callable, ?list)
```

Errors

NonTerminal is a variable: `instantiation_error`

NonTerminal is neither a variable nor a callable term: `type_error(callable, NonTerminal)`

Examples

To parse a list of tokens using a local non-terminal: `phrase(NonTerminal, Input)`

To parse a list of tokens using a non-terminal within the scope of *self*: `phrase(::NonTerminal, Input)`

To parse a list of tokens using a public non-terminal of an explicit object: `phrase(Object::NonTerminal, Input)`

See also:

call//1-N, *phrase//1*, *phrase/3*

phrase/3

Description

```
phrase(GrammarRuleBody, Input, Rest)
phrase(::GrammarRuleBody, Input, Rest)
phrase(Object::GrammarRuleBody, Input, Rest)
```

True when the `GrammarRuleBody` grammar rule body can be applied to the `Input-Rest` difference list of tokens. In the most common case, `GrammarRuleBody` is a non-terminal defined by a grammar rule. This built-in method is declared private and thus cannot be used as a message to an object. When using a back-end Prolog compiler supporting a module system, calls in the format `phrase(Module:GrammarRuleBody, Input, Rest)` may also be used.

This method is opaque to cuts in the first argument. When the first argument is sufficiently instantiated at compile time, the method call is compiled in order to eliminate the implicit overheads of converting the grammar rule body into a goal and meta-calling it. For performance reasons, the second and third arguments are only type-checked at compile time.

Template and modes

```
phrase(+callable, ?list, ?list)
```

Errors

NonTerminal is a variable: `instantiation_error`

NonTerminal is neither a variable nor a callable term: `type_error(callable, NonTerminal)`

Examples

To parse a list of tokens using a local non-terminal: `phrase(NonTerminal, Input, Rest)`

To parse a list of tokens using a non-terminal within the scope of *self* **:** `phrase(::NonTerminal, Input, Rest)`

To parse a list of tokens using a public non-terminal of an explicit object: `phrase(Object::NonTerminal, Input, Rest)`

See also:

call//1-N, *phrase/2*, *phrase/3*

2.4.10 Term and goal expansion

expand_term/2

Description

```
expand_term(Term, Expansion)
```

Expands a term. The most common use is to expand a grammar rule into a clause. Users may override the default Logtalk grammar rule translator by defining clauses for the *term_expansion/2* hook predicate.

The expansion works as follows: if the first argument is a variable, then it is unified with the second argument; if the first argument is not a variable and there are local or inherited clauses for the `term_expansion/2` hook predicate within scope, then this predicate is called to provide an expansion that is then unified with the second argument; if the `term_expansion/2` predicate is not used and the first argument is a compound term with functor `-->/2` then the default Logtalk grammar rule translator is used, with the resulting clause being unified with the second argument; when the translator is not used, the two arguments are unified. The `expand_term/2` predicate may return a single term or a list of terms.

This built-in method may be used to expand a grammar rule into a clause for use with the built-in database methods.

Automatic term expansion is only performed at compile time (to expand terms read from a source file) when using a *hook object*. This predicate can be used by the user to manually perform term expansion at runtime (for example, to convert a grammar rule into a clause).

Template and modes

```
expand_term(?term, ?term)
```

Errors

(none)

Examples

(none)

See also:

expand_goal/2, *goal_expansion/2*, *term_expansion/2*

term_expansion/2

Description

```
term_expansion(Term, Expansion)
```

Defines an expansion for a term. This predicate, when defined and within scope, is automatically called by the *expand_term/2* method. When that is not the case, the expand_term/2 method only uses the default expansions. Use of this predicate by the expand_term/2 method may be restricted by changing its default public scope.

The term_expansion/2 predicate may return a list of terms. Returning an empty list effectively suppresses the term.

Term expansion may be also be applied when compiling source files by defining the object providing access to the term_expansion/2 clauses as a *hook object*. Clauses for the term_expansion/2 predicate defined within an object or a category are **never** used in the compilation of the object or the category itself. Moreover, in this context, terms wrapped using the *{}/1* compiler bypass control construct are not expanded and any expanded term wrapped in this control construct will not be further expanded.

Objects and categories implementing this predicate should declare that they implement the expanding protocol if no ancestor already declares it. This protocol implementation relation can be declared as either protected or private to restrict the scope of this predicate.

Template and modes

```
term_expansion(+nonvar, -nonvar)
term_expansion(+nonvar, -list(nonvar))
```

Errors

(none)

Examples

```
term_expansion((:- license(default)), (:- license(gplv3))).
```

```
term_expansion(data(Millimeters), data(Meters)) :- Meters is Millimeters /
1000.
```

See also:

expand_goal/2, *expand_term/2*, *goal_expansion/2*, *logtalk_load_context/2*

expand_goal/2

Description

```
expand_goal(Goal, ExpandedGoal)
```

Expands a goal. The expansion works as follows: if the first argument is a variable, then it is unified with the second argument; if the first argument is not a variable and there are local or inherited clauses for the `goal_expansion/2` hook predicate within scope, then this predicate is recursively called until a fixed-point is reached to provide an expansion that is then unified with the second argument; if the `goal_expansion/2` predicate is not within scope, the two arguments are unified.

Automatic goal expansion is only performed at compile time (to expand the body of clauses and meta-directives read from a source file) when using *hook objects*. This predicate can be used by the user to manually perform goal expansion at runtime (for example, before asserting a clause).

Template and modes

```
expand_goal(?term, ?term)
```

Errors

(none)

Examples

(none)

See also:

expand_term/2, *goal_expansion/2*, *term_expansion/2*

goal_expansion/2

Description

```
goal_expansion(Goal, ExpandedGoal)
```

Defines an expansion for a goal. The first argument is the goal to be expanded. The expanded goal is returned in the second argument. This predicate is called recursively on the expanded goal until a fixed point is reached. Thus, care must be taken to avoid compilation loops. This predicate, when defined and within scope, is automatically called by the *expand_goal/2* method. Use of this predicate by the `expand_goal/2` method may be restricted by changing its default public scope.

Goal expansion may be also be applied when compiling source files by defining the object providing access to the `goal_expansion/2` clauses as a *hook object*. Clauses for the `goal_expansion/2` predicate defined within an object or a category are **never** used in the compilation of the object or the category itself. Moreover, in this context, goals wrapped using the *{}/1* compiler bypass control construct are not expanded and any expanded goal wrapped in this control construct will not be further expanded.

Objects and categories implementing this predicate should declare that they implement the `expanding` protocol if no ancestor already declares it. This protocol implementation relation can be declared as either protected or private to restrict the scope of this predicate.

Template and modes

```
goal_expansion(+callable, -callable)
```

Errors

(none)

Examples

```
goal_expansion(write(Term), (write_term(Term, []), nl)).
```

```
goal_expansion(read(Term), (write('Input: '), {read(Term)})).
```

See also:

expand_goal/2, *expand_term/2*, *term_expansion/2*, *logtalk_load_context/2*

2.4.11 Coinduction hooks

coinductive_success_hook/1-2

Description

```
coinductive_success_hook(Head, Hypothesis)
coinductive_success_hook(Head)
```

User-defined hook predicates that are automatically called in case of coinductive success when proving a query for a coinductive predicates. The hook predicates are called with the head of the coinductive predicate on coinductive success and, optionally, with the hypothesis used that to reach coinductive success.

When both hook predicates are defined, the `coinductive_success_hook/1` clauses are only used if no `coinductive_success_hook/2` clause applies. The compiler ensures zero performance penalties when defining coinductive predicates without a corresponding definition for the coinductive success hook predicates.

The compiler assumes that these hook predicates are defined as static predicates in order to optimize their use.

Template and modes

```
coinductive_success_hook(+callable, +callable)
coinductive_success_hook(+callable)
```

Errors

(none)

Examples

(none)

See also:

coinductive/1

2.4.12 Message printing

print_message/3

Description

```
print_message(Kind, Component, Term)
```

Built-in method for printing a message represented by a term, which is converted to the message text using the *logtalk::message_tokens(Term, Component)* hook non-terminal. This method is declared in the `logtalk` built-in object as a public predicate. The line prefix and the output stream used for each `Kind-Component` pair can be found using the *logtalk::message_prefix_stream(Kind, Component, Prefix, Stream)* hook predicate.

This predicate starts by converting the message term to a list of tokens and by calling the *logtalk::message_hook(Message, Kind, Component, Tokens)* hook predicate. If this predicate succeeds, the `print_message/3` predicate assumes that the message have been successfully printed.

Template and modes

```
print_message(+nonvar, +nonvar, +nonvar)
```

Errors

(none)

Examples

```
..., logtalk::print_message(information, core, redefining_entity(object, foo)), ...
```

See also:

message_hook/4, *message_prefix_stream/4*, *message_tokens//2*, *print_message_tokens/3*, *print_message_token/4*, *ask_question/5*, *question_hook/6*, *question_prompt_stream/4*

message_tokens//2

Description

```
message_tokens(Message, Component)
```

User-defined non-terminal hook used to rewrite a message term into a list of tokens and declared in the `logtalk` built-in object as a public, multifile, and dynamic non-terminal. The list of tokens can be printed by calling the *print_message_tokens/3* method. This non-terminal hook is automatically called by the *print_message/3* method.

Template and modes

```
message_tokens(+nonvar, +nonvar)
```

Errors

(none)

Examples

```
:- multifile(logtalk::message_tokens//2).
:- dynamic(logtalk::message_tokens//2).

logtalk::message_tokens(redefining_entity(Type, Entity), core) -->
    ['Redefining ~w ~q'-[Type, Entity], nl].
```

See also:

message_hook/4, *message_prefix_stream/4*, *print_message/3*, *print_message_tokens/3*, *print_message_token/4*, *ask_question/5*, *question_hook/6*, *question_prompt_stream/4*

message_hook/4

Description

```
message_hook(Message, Kind, Component, Tokens)
```

User-defined hook method for intercepting printing of a message, declared in the `logtalk` built-in object as a public, multifile, and dynamic predicate. This hook method is automatically called by the *print_message/3* method. When the call succeeds, the `print_message/3` method assumes that the message have been successfully printed.

Template and modes

```
message_hook(@nonvar, @nonvar, @nonvar, @list(nonvar))
```

Errors

(none)

Examples

```
:- multifile(logtalk::message_hook/4).
:- dynamic(logtalk::message_hook/4).

% print silent messages instead of discarding them as default
logtalk::message_hook(_, silent, core, Tokens) :-
    logtalk::message_prefix_stream(silent, core, Prefix, Stream),
    logtalk::print_message_tokens(Stream, Prefix, Tokens).
```

See also:

message_prefix_stream/4, *message_tokens//2*, *print_message/3*, *print_message_tokens/3*, *print_message_token/4*, *ask_question/5*, *question_hook/6*, *question_prompt_stream/4*

message_prefix_stream/4

Description

```
message_prefix_stream(Kind, Component, Prefix, Stream)
```

User-defined hook method for specifying the default prefix and stream for printing a message for a given kind and component. This method is declared in the `logtalk` built-in object as a public, multifile, and dynamic predicate.

Template and modes

```
message_prefix_stream(?nonvar, ?nonvar, ?atom, ?stream_or_alias)
```

Errors

(none)

Examples

```
:- multifile(logtalk::message_prefix_stream/4).
:- dynamic(logtalk::message_prefix_stream/4).

logtalk::message_prefix_stream(information, core, '% ', user_output).
```

See also:

message_hook/4, *message_tokens//2*, *print_message/3*, *print_message_tokens/3*, *print_message_token/4*, *ask_question/5*, *question_hook/6*, *question_prompt_stream/4*

print_message_tokens/3

Description

```
print_message_tokens(Stream, Prefix, Tokens)
```

Built-in method for printing a list of message tokens, declared in the `logtalk` built-in object as a public predicate. This method is automatically called by the *print_message/3* method (assuming that the message was not intercepted by a *message_hook/4* definition) and calls the user-defined hook predicate *print_message_token/4* for each token. When a call to this hook predicate succeeds, the `print_message_tokens/3` predicate assumes that the token have been printed. When the call fails, the `print_message_tokens/3` predicate uses a default printing procedure for the token.

Template and modes

```
print_message_tokens(@stream_or_alias, +atom, @list(nonvar))
```

Errors

(none)

Examples

```
...,
logtalk::print_message_tokens(user_error, '% ', ['Redefining ~w ~q'-[object,foo],
→nl]),
...
```

See also:

message_hook/4, *message_prefix_stream/4*, *message_tokens//2*, *print_message/3*, *print_message_token/4*, *ask_question/5*, *question_hook/6*, *question_prompt_stream/4*

print_message_token/4

Description

```
print_message_token(Stream, Prefix, Token, Tokens)
```

User-defined hook method for printing a message token, declared in the `logtalk` built-in object as a public, multifile, and dynamic predicate. It allows the user to intercept the printing of a message token. This hook method is automatically called by the *print_message_tokens/3* built-in method for each token.

Template and modes

```
print_message_token(@stream_or_alias, @atom, @nonvar, @list(nonvar))
```

Errors

(none)

Examples

```
:- multifile(logtalk::print_message_token/4).
:- dynamic(logtalk::print_message_token/4).

% ignore all flush tokens
logtalk::print_message_token(_Stream, _Prefix, flush, _Tokens).
```

See also:

message_hook/4,   *message_prefix_stream/4*,   *message_tokens//2*,   *print_message/3*,   *print_message_tokens/3*,
ask_question/5, *question_hook/6*, *question_prompt_stream/4*

2.4.13 Question asking

ask_question/5

Description

```
ask_question(Question, Kind, Component, Check, Answer)
```

Built-in method for asking a question represented by a term, Question, which is converted to the question text using
the *logtalk::message_tokens(Question, Component)* hook predicate. This method is declared in the logtalk built-in
object as a public predicate. The default question prompt and the input stream used for each Kind-Component
pair can be found using the *logtalk::question_prompt_stream(Kind, Component, Prompt, Stream)* hook predicate. The
Check argument is a closure that is converted into a checking goal by extending it with the user supplied answer. This
predicate implements a read-loop that terminates when the checking predicate succeeds.

This predicate starts by calling the *logtalk::question_hook(Question, Kind, Component, Check, Answer)* hook predi-
cate. If this predicate succeeds, the ask_question/5 predicate assumes that the question have been successfully
asked and replied.

Template and modes

```
ask_question(+nonvar, +nonvar, +nonvar, +callable, -term)
```

Meta-predicate template

```
ask_question(*, *, *, 1, *)
```

Errors

(none)

Examples

```
...,
logtalk::ask_question(enter_age, question, my_app, integer, Age),
...
```

See also:

question_hook/6, *question_prompt_stream/4*, *message_hook/4*, *message_prefix_stream/4*, *message_tokens//2*, *print_message/3*, *print_message_tokens/3*, *print_message_token/4*

question_hook/6

Description

```
question_hook(Question, Kind, Component, Tokens, Check, Answer)
```

User-defined hook method for intercepting asking a question, declared in the `logtalk` built-in object as a public, multifile, and dynamic predicate. This hook method is automatically called by the *ask_question/5* method. When the call succeeds, the `ask_question/5` method assumes that the question have been successfully asked and replied.

Template and modes

```
question_hook(+nonvar, +nonvar, +nonvar, +list(nonvar), +callable, -term)
```

Meta-predicate template

```
question_hook(*, *, *, *, 1, *)
```

Errors

(none)

Examples

```
:- multifile(logtalk::question_hook/6).
:- dynamic(logtalk::question_hook/6).

% use a pre-defined answer instead of asking the user
logtalk::question_hook(upper_limit, question, my_app, _, float, 3.7).
```

See also:

ask_question/5, *question_prompt_stream/4* *message_hook/4*, *message_prefix_stream/4*, *message_tokens//2*, *print_message/3*, *print_message_tokens/3*, *print_message_token/4*,

---

question_prompt_stream/4

Description

```
question_prompt_stream(Kind, Component, Prompt, Stream)
```

User-defined hook method for specifying the default prompt and input stream for asking a question for a given kind and component. This method is declared in the `logtalk` built-in object as a public, multifile, and dynamic predicate.

Template and modes

```
question_prompt_stream(?nonvar, ?nonvar, ?atom, ?stream_or_alias)
```

Errors

(none)

Examples

```
:- multifile(logtalk::question_prompt_stream/4).
:- dynamic(logtalk::question_prompt_stream/4).

logtalk::question_prompt_stream(question, debugger, '    > ', user_input).
```

See also:

ask_question/5, *question_hook/6*, *message_hook/4*, *message_prefix_stream/4*, *message_tokens//2*, *print_message/3*, *print_message_tokens/3*, *print_message_token/4*

2.5 Control constructs

2.5.1 Message sending

::/2

Description

```
Object::Message
{Proxy}::Message
```

Sends a message to an object. The message argument must match a public predicate of the receiver object. When the message corresponds to a protected or private predicate, the call is only valid if the *sender* matches the *predicate scope container*. When the predicate is declared but not defined, the message simply fails (as per the closed-world assumption).

The `{Proxy}::Message` syntax allows simplified access to parametric object *proxies*. Its operational semantics is equivalent to the goal conjunction `(call(Proxy), Proxy::Message)`. I.e. `Proxy` is proved within the context of the pseudo-object *user* and, if successful, the goal term is used as a parametric object identifier. Exceptions

thrown when proving `Proxy` are handled by the `::/2` control construct. This syntax construct supports backtracking over the `{Proxy}` goal.

The lookups for the message declaration and the corresponding method are performed using a depth-first strategy. Depending on the value of the `optimize` flag, these lookups are performed at compile time whenever sufficient information is available. When the lookups are performed at runtime, a caching mechanism is used to improve performance in subsequent messages.

Template and modes

```
+object_identifier::+callable
{+object_identifier}::+callable
```

Errors

Either Object or Message is a variable: `instantiation_error`

Object is neither a variable nor a valid object identifier: `type_error(object_identifier, Object)`

Message is neither a variable nor a callable term: `type_error(callable, Message)`

Message, with predicate indicator Name/Arity, is declared private: `permission_error(access, private_predicate, Name/Arity)`

Message, with predicate indicator Name/Arity, is declared protected: `permission_error(access, protected_predicate, Name/Arity)`

Message, with predicate indicator Name/Arity, is not declared: `existence_error(predicate_declaration, Name/Arity)`

Object does not exist: `existence_error(object, Object)`

Proxy is a variable: `instantiation_error`

Proxy is neither a variable nor a callable term: `type_error(callable, Proxy)`

Proxy, with predicate indicator Name/Arity, does not exist in the *user* pseudo-object: `existence_error(procedure, Name/Arity)`

Examples

```
| ?- list::member(X, [1, 2, 3]).

X = 1 ;
X = 2 ;
X = 3
yes
```

See also:

::/1, *^^/1*, *[]/1*

::/1

Description

```
::Message
```

Sends a message to *self*. Can only used in the body of a predicate definition. The argument should match a public or protected predicate of *self*. It may also match a private predicate if the predicate is within the scope of the object where the method making the call is defined, if imported from a category, if used from within a category, or when using private inheritance. When the predicate is declared but not defined, the message simply fails (as per the closed-world assumption).

The lookups for the message declaration and the corresponding method are performed using a depth-first strategy. A message to *self* necessarily implies the use of dynamic binding but a caching mechanism is used to improve performance in subsequent messages.

Template and modes

```
::+callable
```

Errors

Message is a variable: `instantiation_error`

Message is neither a variable nor a callable term: `type_error(callable, Message)`

Message, with predicate indicator Name/Arity, is declared private: `permission_error(access, private_predicate, Name/Arity)`

Message, with predicate indicator Name/Arity, is not declared: `existence_error(predicate_declaration, Name/Arity)`

Examples

```
area(Area) :-
    ::width(Width),
    ::height(Height),
    Area is Width*Height.
```

See also:

::/2, *^^/1*, *[]/1*

2.5.2 Message delegation

[]/1

Description

```
[Object::Message]
[{Proxy}::Message]
```

This control construct allows the programmer to send a message to an object while preserving the original sender. It is mainly used in the definition of object handlers for unknown messages. This functionality is usually known as *delegation* but be aware that this is an overloaded word that can mean different things in different object-oriented programming languages.

To prevent using of this control construct to break object encapsulation, an attempt to delegate a message to the original sender results in an error. The remaining error conditions are the same as the *::/2* control construct.

Note that, despite the correct functor for this control construct being (traditionally) `'.'/2`, we refer to it as `[]/1` simply to emphasize that the syntax is a list with a single element.

Template and modes

```
[+object_identifier::+callable]
[{+object_identifier}::+callable]
```

Errors

Object and the original sender are the same object: `permission_error(access, object, Sender)`

Either Object or Message is a variable: `instantiation_error`

Object is neither a variable nor an object identifier: `type_error(object_identifier, Object)`

Message is neither a variable nor a callable term: `type_error(callable, Message)`

Message, with predicate indicator Name/Arity, is declared private: `permission_error(access, private_predicate, Name/Arity)`

Message, with predicate indicator Name/Arity, is declared protected: `permission_error(access, protected_predicate, Name/Arity)`

Message, with predicate indicator Name/Arity, is not declared: `existence_error(predicate_declaration, Name/Arity)`

Object does not exist: `existence_error(object, Object)`

Proxy is a variable: `instantiation_error`

Proxy is neither a variable nor an object identifier: `type_error(object_identifier, Proxy)`

Proxy, with predicate indicator Name/Arity, does not exist in the *user* pseudo-object: `existence_error(procedure, Name/Arity)`

Examples

```
forward(Message) :-
    [backup::Message].
```

See also:

::/2, *::/1*, *^^/1*, *forward/1*

2.5.3 Calling imported and inherited predicates

^^/1

Description

```
^^Predicate
```

Calls an imported or inherited predicate definition. The call fails if the predicate is declared but there is no imported or inherited predicate definition (as per the closed-world assumption). This control construct may be used within objects or categories in the body of a predicate definition.

This control construct preserves the implicit execution context *self* and *sender* arguments (plus the meta-call context and coinduction stack when applicable) when calling the inherited (or imported) predicate definition.

The lookups for the predicate declaration and the predicate definition are performed using a depth-first strategy. Depending on the value of the `optimize` flag, these lookups are performed at compile time when the predicate is static and sufficient information is available. When the lookups are performed at runtime, a caching mechanism is used to improve performance in subsequent calls.

When the call is made from within an object, the lookup for the predicate definition starts at the imported categories, if any. If an imported predicate definition is not found, the lookup proceeds to the ancestor objects. Calls from predicates defined in complementing categories lookup inherited definitions as if the calls were made from the complemented object, thus allowing more comprehensive object patching. For other categories, the predicate definition lookup is restricted to the extended categories.

The called predicate should be declared public or protected. It may also be declared private if within the scope of the entity where the method making the call is defined.

This control construct is a generalization of the Smalltalk *super* keyword to take into account Logtalk support for prototypes and categories besides classes.

Template and modes

```
^^+callable
```

Errors

Predicate is a variable: `instantiation_error`

Predicate is neither a variable nor a callable term: `type_error(callable, Predicate)`

Predicate, with predicate indicator Name/Arity, is declared private: `permission_error(access, private_predicate, Name/Arity)`

Predicate, with predicate indicator Name/Arity, is not declared: `existence_error(predicate_declaration, Name/Arity)`

Examples

```
init :-
    assertz(counter(0)),
    ^^init.
```

See also:

::/2, *::/1*, *[]/1*

2.5.4 Calling external predicates

{}/1

Description

```
{Term}
{Goal}
```

This control construct allows the programmer to bypass the Logtalk compiler. It can also be used to wrap a source file term (either a clause or a directive) to bypass the term-expansion mechanism. Similarly, it can also be used to wrap a goal to bypass the goal-expansion mechanism. When used to wrap a goal, it is opaque to cuts and the argument is called within the context of the pseudo-object *user*. It is also possible to use {Closure} as the first argument of *call/1-N* calls. In this case, `Closure` will be extended with the remaining arguments of the `call/2-N` call in order to construct a goal that will be called within the context of `user`. It can also be used as a message to any object. This is useful when the message is e.g. a conjunction of messages, some of which being calls to Prolog built-in predicates.

This control construct may also be used in place of an object identifier when sending a message. In this case, the result of proving its argument as a goal (within the context of the pseudo-object `user`) is used as an object identifier in the message sending call. This feature is mainly used with *parametric objects* when their identifiers correspond to predicates defined in `user`.

Template and modes

```
{+callable}
```

Errors

Term or Goal is a variable: `instantiation_error`

Term is neither a variable nor a callable term: `type_error(callable, Term)`

Goal is neither a variable nor a callable term: `type_error(callable, Goal)`

Examples

```
{:- load_foreign_resource(file)}.

N1/D1 < N2/D2 :-
    {N1*D2 < N2*D1}.

call_in_user(F, X, Y, Z) :-
    call({F}, X, Y, Z).

| ?- {circle(Id, Radius, Color)}::area(Area).
...
```

```
| ?- logtalk::{write('Hello world!'), nl}.
Hello world!
yes
```

2.5.5 Context switching calls

<</2

Description

```
Object<<Goal
{Proxy}<<Goal
```

Debugging control construct. Calls a goal within the context of the specified object. The goal is called with the execution context (*sender*, *this*, and *self*) set to the object. The goal may need to be written between parenthesis to avoid parsing errors due to operator conflicts. This control construct should only be used for debugging or for writing unit tests. This control construct can only be used for objects compiled with the compiler flag context_switching_calls set to allow. Set this compiler flag to deny to disable this control construct and thus preventing using it to break encapsulation.

The {Proxy}<<Goal syntax allows simplified access to *parametric object proxies*. Its operational semantics is equivalent to the goal conjunction (call(Proxy), Proxy<<Goal). I.e. Proxy is proved within the context of the pseudo-object *user* and, if successful, the goal term is used as a parametric object identifier. Exceptions thrown when proving Proxy are handled by the <</2 control construct. This syntax construct supports backtracking over the {Proxy} goal.

Caveat: although the goal argument is fully compiled before calling, some necessary information for the second compiler pass may not be available at runtime.

Template and modes

```
+object_identifier<<+callable
{+object_identifier}<<+callable
```

Errors

Either Object or Goal is a variable: instantiation_error

Object is neither a variable nor a valid object identifier: type_error(object_identifier, Object)

Goal is neither a variable nor a callable term: type_error(callable, Goal)

Object does not contain a local definition for the Goal predicate: existence_error(procedure, Goal)

Object does not exist: existence_error(object, Object)

Object was created/compiled with support for context switching calls turned off: permission_error(access, database, Goal)

Proxy is a variable: instantiation_error

Proxy is neither a variable nor an object identifier: `type_error(object_identifier, Proxy)`

The predicate Proxy does not exist in the *user* **pseudo-object:** `existence_error(procedure,`
     `ProxyFunctor/ProxyArity)`

Examples

```
test(member) :-
    list << member(1, [1]).
```

TUTORIAL

3.1 List predicates

In this example, we will illustrate the use of:

- objects
- protocols

by using common list utility predicates.

3.1.1 Defining a list object

We will start by defining an object, `list`, containing predicate definitions for some common list predicates like `append/3`, `length/2`, and `member/2`:

```logtalk
:- object(list).

    :- public(append/3).
    :- public(length/2).
    :- public(member/2).

    append([], List, List).
    append([Head| Tail], List, [Head| Tail2]) :-
        append(Tail, List, Tail2).

    length(List, Length) :-
        length(List, 0, Length).

    length([], Length, Length).
    length([_| Tail], Acc, Length) :-
        Acc2 is Acc + 1,
        length(Tail, Acc2, Length).

    member(Element, [Element| _]).
    member(Element, [_| List]) :-
        member(Element, List).

:- end_object.
```

What is different here from a regular Prolog program? The definitions of the list predicates are the usual ones. We have two new directives, *object/1-5* and *end_object/0*, that encapsulate the object's code. In Logtalk, by default, all object predicates are private; therefore, we have to explicitly declare all predicates that we want to be public, that is, that we want to call from outside the object. This is done using the *public/1* scope directive.

After we copy the object code to a text file and saved it under the name `list.lgt`, we need to change the Prolog working directory to the one used to save our file (consult your Prolog compiler reference manual). Then, after starting Logtalk (see the Installing and running Logtalk section on the User Manual), we can compile and load the object using the *logtalk_load/1* Logtalk built-in predicate:

```
| ?- logtalk_load(list).

object list loaded
yes
```

We can now try goals like:

```
| ?- list::member(X, [1, 2, 3]).

X = 1;
X = 2;
X = 3;
no
```

or:

```
| ?- list::length([1, 2, 3], L).

L = 3
yes
```

The infix operator *::/2* is used in Logtalk to send a message to an object. The message must match a public object predicate. If we try to call a non-public predicate such as the `length/3` auxiliary predicate an exception will be generated:

```
| ?- list::length([1, 2, 3], 0, L).

uncaught exception:
    error(
        existence_error(predicate_declaration, length([1, 2, 3], 0, L)),
        list::length([1, 2, 3], 0, L),
        user)
```

The error term describes the type of error, the message that caused the exception, and the sender of the message (in this case, the pseudo-object `user` because we are sending the message from the top-level interpreter).

3.1.2 Defining a list protocol

As we saw in the above example, a Logtalk object may contain predicate directives and predicate definitions (clauses). The set of predicate directives defines what we call the object's *protocol* or interface. An interface may have several implementations. For instance, we may want to define a new object that implements the list predicates using difference lists. However, we do not want to repeat the predicate directives in the new object. Therefore, what we need is to split the object's protocol from the object's implementation by defining a new Logtalk entity known as a protocol. Logtalk protocols are compilations units, at the same level as objects and categories. That said, let us define a `listp` protocol:

```
:- protocol(listp).

    :- public(append/3).
    :- public(length/2).
    :- public(member/2).
```

```
:- end_protocol.
```

Similar to what we have done for objects, we use the *protocol/1-2* and *end_protocol/0* directives to encapsulate the predicate directives. We can improve this protocol by documenting the call/return modes and the number of proofs of each predicate using the *mode/2* directive:

```
:- protocol(listp).

    :- public(append/3).
    :- mode(append(?list, ?list, ?list), zero_or_more).

    :- public(length/2).
    :- mode(length(?list, ?integer), zero_or_more).

    :- public(member/2).
    :- mode(member(?term, ?list), zero_or_more).

:- end_protocol.
```

We now need to change our definition of the list object by removing the predicate directives and by declaring that the object implements the listp protocol:

```
:- object(list,
    implements(listp)).

    append([], List, List).
    append([Head| Tail], List, [Head| Tail2]) :-
        append(Tail, List, Tail2).

    ...

:- end_object.
```

The protocol declared in listp may now be alternatively implemented using difference lists by defining a new object, difflist:

```
:- object(difflist,
    implements(listp).

    append(L1-X, X-L2, L1-L2).
    ...

:- end_object.
```

3.1.3 Summary

- It is easy to define a simple object: just put your Prolog code inside starting and ending object directives and add the necessary scope directives. The object will be self-defining and ready to use.

- Define a protocol when you may want to provide or enable several alternative definitions to a given set of predicates. This way we avoid needless repetition of predicate directives.

3.2 Dynamic object attributes

In this example, we will illustrate the use of:

- categories
- category predicates
- dynamic predicates

by defining a category that implements a set of predicates for handling dynamic object attributes.

3.2.1 Defining a category

We want to define a set of predicates to handle dynamic object attributes. We need public predicates to set, get, and delete attributes, and a private dynamic predicate to store the attributes values. Let us name these predicates `set_attribute/2` and `get_attribute/2`, for getting and setting an attribute value, `del_attribute/2` and `del_attributes/2`, for deleting attributes, and `attribute_/2`, for storing the attributes values.

But we do not want to encapsulate these predicates in an object. Why? Because they are a set of useful, closely related, predicates that may be used by several, unrelated, objects. If defined at an object level, we would be constrained to use inheritance in order to have the predicates available to other objects. Furthermore, this could force us to use multi-inheritance or to have some kind of generic root object containing all kinds of possible useful predicates.

For this kind of situation, Logtalk enables the programmer to encapsulate the predicates in a *category*, so that they can be used in any object. A category is a Logtalk entity, at the same level as objects and protocols. It can contain predicates directives and/or definitions. Category predicates can be imported by any object, without code duplication and without resorting to inheritance.

When defining category predicates, we need to remember that a category can be imported by more than one object. Thus, the calls to the built-in methods that handle the private dynamic predicate (such as *assertz/1* or *retract/1*) must be made either in the context of *self*, using the *message to self* control structure, *::/1*, or in the context of *this* (i.e. in the context of the object importing the category). This way, we ensure that when we call one of the attribute predicates on an object, the intended object own definition of `attribute_/2` will be used. The predicates definitions are straightforward. For example, if opting for storing the attributes in *self*:

```
:- category(attributes).

    :- public(set_attribute/2).
    :- mode(set_attribute(+nonvar, +nonvar), one).

    :- public(get_attribute/2).
    :- mode(get_attribute(?nonvar, ?nonvar), zero_or_more).

    :- public(del_attribute/2).
    :- mode(del_attribute(?nonvar, ?nonvar), zero_or_more).

    :- public(del_attributes/2).
    :- mode(del_attributes(@term, @term), one).

    :- private(attribute_/2).
    :- mode(attribute_(?nonvar, ?nonvar), zero_or_more).
    :- dynamic(attribute_/2).

    set_attribute(Attribute, Value):-
        ::retractall(attribute_(Attribute, _)),
        ::assertz(attribute_(Attribute, Value)).
```

(continues on next page)

```
    get_attribute(Attribute, Value):-
        ::attribute_(Attribute, Value).

    del_attribute(Attribute, Value):-
        ::retract(attribute_(Attribute, Value)).

    del_attributes(Attribute, Value):-
        ::retractall(attribute_(Attribute, Value)).

:- end_category.
```

The alternative, opting for storing the attributes on *this*, is similar: just delete the uses of the `::/1` control structure from the code above.

We have two new directives, *category/1-3* and *end_category/0*, that encapsulate the category code. If needed, we can put the predicates directives inside a protocol that will be implemented by the category:

```
:- category(attributes,
    implements(attributes_protocol)).

    ...

:- end_category.
```

Any protocol can be implemented by either an object, a category, or both.

3.2.2 Importing the category

We reuse a category's predicates by importing them into an object:

```
:- object(person,
    imports(attributes)).

    ...

:- end_object.
```

After compiling and loading this object and our category, we can now try queries like:

```
| ?- person::set_attribute(name, paulo).

yes

| ?- person::set_attribute(gender, male).

yes

| ?- person::get_attribute(Attribute, Value).

Attribute = name, Value = paulo ;
Attribute = gender, Value = male ;
no
```

3.2.3 Summary

- Categories are similar to objects: we just write our predicate directives and definitions bracketed by opening and ending category directives.

- An object reuses a category by importing it. The imported predicates behave as if they have been defined in the object itself.

- When do we use a category instead of an object? Whenever we have a set of closely related predicates that we want to reuse in several, unrelated, objects without being constrained by inheritance relations. Thus, categories can be interpreted as object building components.

3.3 A reflective class-based system

When compiling an object, Logtalk distinguishes prototypes from instance or classes by examining the object relations. If an object instantiates and/or specializes another object, then it is compiled as an instance or class, otherwise it is compiled as a prototype. A consequence of this is that, in order to work with instance or classes, we always have to define root objects for the instantiation and specialization hierarchies (however, we are not restricted to a single hierarchy). The best solution is often to define a reflective class-based system [Maes87], where every class is also an object and, as such, an instance of some class.

In this example, we are going to define the basis for a reflective class-based system, based on an extension of the ideas presented in [Cointe87]. This extension provides, along with root objects for the instantiation and specialization hierarchies, explicit support for abstract classes [Moura94].

3.3.1 Defining the base classes

We will start by defining three classes: `object`, `abstract_class`, and `class`. The class `object` will contain all predicates common to all objects. It will be the root of the inheritance graph:

```
:- object(object,
    instantiates(class)).

    % predicates common to all objects

:- end_object.
```

The class `abstract_class` specializes `object` by adding predicates common to all classes. It will be the default meta-class for abstract classes:

```
:- object(abstract_class,
    instantiates(class),
    specializes(object)).

    % predicates common to all classes

:- end_object.
```

The class `class` specializes `abstract_class` by adding predicates common to all instantiable classes. It will be the root of the instantiation graph and the default meta-class for instantiable classes:

```
:- object(class,
    instantiates(class),
    specializes(abstract_class)).
```

```
    % predicates common to all instantiable classes

:- end_object.
```

Note that all three objects are instances of class `class`. The instantiation and specialization relationships are chosen so that each object may use the predicates defined in itself and in the other two objects, with no danger of method lookup endless loops.

3.3.2 Summary

- An object that does not instantiate or specialize other objects is always compiled as a prototype.

- An instance must instantiate at least one object (its class). Similarly, a class must at least specialize or instantiate other object.

- The distinction between abstract classes and instantiable classes is an operational one, depending on the class inherited methods. A class is instantiable if inherits methods for creating instances. Conversely, a class is abstract if does not inherit any instance creation method.

3.4 Profiling programs

In this example, we will illustrate the use of:

- events

- monitors

by defining a simple profiler that prints the starting and ending time for processing a message sent to an object.

3.4.1 Messages as events

In a pure object-oriented system, all computations start by sending messages to objects. We can thus define an *event* as the sending of a message to an object. An event can then be specified by the tuple (`Object, Message, Sender`). This definition can be refined by interpreting the sending of a message and the return of the control to the object that has sent the message as two distinct events. We call these events respectively `before` and `after`. Therefore, we end up by representing an event by the tuple (`Event, Object, Message, Sender`). For instance, if we send the message:

```
| ?- foo::bar(X).

X = 1
yes
```

the two corresponding events will be:

```
(before, foo, bar(X), user)
(after, foo, bar(1), user)
```

Note that the second event is only generated if the message succeeds. If the message as a goal have multiple solutions, then one `after` event will be generated for each solution.

Events are automatically generated by the message sending mechanisms for each public message sent using the *::/2* operator.

3.4.2 Profilers as monitors

A monitor is an object that reacts whenever a spied event occurs. The monitor actions are defined by two event handlers: *before/3* for `before` events and *after/3* for `after` events. These predicates are automatically called by the message sending mechanisms when an event registered for the monitor occurs. These event handlers are declared as public predicates in the `monitoring` built-in protocol.

In our example, we need a way to get the current time before and after we process a message. We will assume that we have a `time` object implementing a `cpu_time/1` predicate that returns the current CPU time for the Prolog session:

```
:- object(time).

    :- public(cpu_time/1).
    :- mode(cpu_time(-number), one).

    ...

:- end_object.
```

Our profiler will be named `stop_watch`. It must define event handlers for the `before` and `after` events that will print the event description (object, message, and sender) and the current time:

```
:- object(stop_watch,
    implements(monitoring)).    % event handler predicate declarations

    :- uses(time).

    before(Object, Message, Sender) :-
        write(Object), write(' <-- '), writeq(Message),
        write(' from '), write(Sender), nl, write('STARTING at '),
        time::cpu_time(Seconds), write(Seconds), write(' seconds'), nl.

    after(Object, Message, Sender) :-
        write(Object), write(' <-- '), writeq(Message),
        write(' from '), write(Sender), nl, write('ENDING at '),
        time::cpu_time(Seconds), write(Seconds), write(' seconds'), nl.

:- end_object.
```

After compiling and loading the `stop_watch` object (and the objects that we want to profile), we can use the *define_events/5* built-in predicate to set up our profiler. For example, to profile all messages that are sent to the object `foo`, we need to call the goal:

```
| ?- define_events(_, foo, _, _, stop_watch).

yes
```

This call will register `stop_watch` as a monitor to all messages sent to object `foo`, for both `before` and `after` events. Note that we say "as a monitor", not "the monitor": we can have any number of monitors over the same events.

From now on, every time we sent a message to `foo`, the `stop_watch` monitor will print the starting and ending times for the message execution. For instance:

```
| ?- foo::bar(X).

foo <-- bar(X) from user
STARTING at 12.87415 seconds
foo <-- bar(1) from user
```

```
ENDING at 12.87419 seconds

X = 1
yes
```

To stop profiling the messages sent to `foo` we use the *abolish_events/5* built-in predicate:

```
| ?- abolish_events(_, foo, _, _, stop_watch).

yes
```

This call will abolish all events defined over the object `foo` assigned to the `stop_watch` monitor.

3.4.3 Summary

- An event is defined as the sending of a (public) message to an object.

- There are two kinds of events: `before` events, generated before a message is processed, and `after` events, generated after the message processing completed successfully.

- Any object can be declared as a monitor to any event. A monitor shall reference the `monitoring` built-in protocol in the object opening directive.

- A monitor defines event handlers, the predicates *before/3* and *after/3*, that are automatically called by the runtime engine when a spied event occurs.

- Three built-in predicates, *define_events/5*, *current_event/5*, and *abolish_events/5*, enables us define, query, and abolish both events and monitors.

FAQ

4.1 General

- *Why are all versions of Logtalk numbered 2.x or 3.x?*
- *Why do I need a Prolog compiler to use Logtalk?*
- *Is the Logtalk implementation based on Prolog modules?*
- *Does the Logtalk implementation use term-expansion?*

4.1.1 Why are all versions of Logtalk numbered 2.x or 3.x?

The numbers "2" and "3" in the Logtalk version string refers to, respectively, the second and the third generations of the Logtalk language. Development of Logtalk 2 started on January 1998, with the first alpha release for registered users on July and the first public beta on October. The first stable version of Logtalk 2 was released on February 9, 1999. Development of Logtalk 3 started on April 2012, with the first public alpha released on August 21, 2012. The first stable version of Logtalk 3 was released on January 7, 2015.

4.1.2 Why do I need a Prolog compiler to use Logtalk?

Currently, the Logtalk language is implemented as a Prolog extension instead of as a standalone compiler. Compilation of Logtalk source files is performed in two steps. First, the Logtalk compiler converts a source file to a Prolog file. Second, the chosen Prolog compiler is called by Logtalk to compile the intermediate Prolog file generated on the first step. The implementation of Logtalk as a Prolog extension allows users to use Logtalk together with features only available on specific Prolog compilers.

4.1.3 Is the Logtalk implementation based on Prolog modules?

No. Logtalk is (currently) implemented is plain Prolog code. Only a few Prolog compilers include a module system, with several compatibility problems between them. Moreover, the current ISO Prolog standard for modules is next to worthless and is ignored by most of the Prolog community. Nevertheless, the Logtalk compiler is able to compile simple modules (using a common subset of module directives) as objects for backward-compatibility with existing code (see the *Prolog Integration and Migration Guide* for details).

4.1.4 Does the Logtalk implementation use term-expansion?

No. Term-expansion mechanisms are not standard and are not available in all supported Prolog compilers.

4.2 Compatibility

- *What are the backend Prolog compiler requirements to run Logtalk?*
- *Can I use constraint-based packages with Logtalk?*
- *Can I use Logtalk objects and Prolog modules at the same time?*

4.2.1 What are the backend Prolog compiler requirements to run Logtalk?

See the wiki page on backend Prolog compiler requirements.

4.2.2 Can I use constraint-based packages with Logtalk?

Usually, yes. Some constraint-based packages may define operators which clash with the ones defined by Logtalk. In these cases, compatibility with Logtalk depends on the constraint-based packages providing an alternative for accessing the functionality provided by those operators. When the constraint solver is encapsulated using a Prolog module, a possible workaround is to use either explicit module qualification or encapsulate the call using the *{}/1* control construct (thus bypassing the Logtalk compiler).

4.2.3 Can I use Logtalk objects and Prolog modules at the same time?

Yes. In order to call a module predicate from within an object (or category) you may use an *use_module/2* directive or use explicit module qualification (possibly wrapping the call using the Logtalk control construct *{}/1* that allows bypassing of the Logtalk compiler when compiling a predicate call). Logtalk also allows modules to be compiled as objects (see the *Prolog Integration and Migration Guide* for details).

4.3 Installation

- *The integration scripts/shortcuts are not working!*
- *I get errors when starting up Logtalk after upgrading to the latest version!*

4.3.1 The integration scripts/shortcuts are not working!

Check that the `LOGTALKHOME` and `LOGTALKUSER` environment variables are defined, that the Logtalk user folder is available on the location pointed by `LOGTALKUSER` (you can create this folder by running the `logtalk_user_setup` shell script), and that the Prolog compilers that you want to use are supported and available from the system path. If the problem persists, run the shell script that creates the integration script or shortcut manually and check for any error message or additional instructions. For some Prolog compilers such as XSB and Ciao, the first call of the integration script or shortcut must be made by an administrator user. If you are using Windows, make sure that any anti-virus or other security software that you might have installed is not silently blocking some of the installer tasks.

4.3.2 I get errors when starting up Logtalk after upgrading to the latest version!

Changes in the Logtalk compiler between releases may render Prolog adapter files from older versions incompatible with new ones. You may need to update your local Logtalk user files by running e.g. the `logtalk_user_setup`

shell script. Check the UPGRADING.md file on the root of the Logtalk installation directory and the release notes for any incompatible changes to the adapter files.

4.4 Portability

- *Are my Logtalk applications portable across Prolog compilers?*
- *Are my Logtalk applications portable across operating systems?*

4.4.1 Are my Logtalk applications portable across Prolog compilers?

Yes, as long you don't use built-in predicates or special features only available on some Prolog compilers. There is a compiler flag (portability) that you can set to instruct Logtalk to print a warning for each occurrence of non-ISO Prolog standard features such as proprietary built-in predicates. In addition, it is advisable that you constrain, if possible, the use of platform or compiler dependent code to a small number of objects with clearly defined protocols. You may also use Logtalk support for conditional compilation to compile different entity or predicate definitions depending on the back-end Prolog compiler being used.

4.4.2 Are my Logtalk applications portable across operating systems?

Yes, as long you don't use built-in predicates or special features that your chosen backend Prolog compiler only supports in some operating-systems. You may need to change the end-of-lines characters of your source files to match the ones on the target operating system and the expectations of your Prolog compiler. Some Prolog compilers silently fail to compile source files with the wrong end-of-lines characters.

4.5 Programming

- *Should I use prototypes or classes in my application?*
- *Can I use both classes and prototypes in the same application?*
- *Can I mix classes and prototypes in the same hierarchy?*
- *Can I use a protocol or a category with both prototypes and classes?*
- *What support is provided in Logtalk for defining and using components?*
- *What support is provided in Logtalk for reflective programming?*

4.5.1 Should I use prototypes or classes in my application?

Prototypes and classes provide different patterns of code reuse. A prototype encapsulates code that can be used by itself and by its descendent prototypes. A class encapsulates code to be used by its descendent instances. Prototypes provide the best replacement to the use of modules as encapsulation units, avoiding the need to instantiate a class in order to access its code.

4.5.2 Can I use both classes and prototypes in the same application?

Yes. In addition, you may freely exchange messages between prototypes, classes, and instances.

4.5.3 Can I mix classes and prototypes in the same hierarchy?

No. However, you may define as many prototype hierarchies and class hierarchies and classes as needed by your application.

4.5.4 Can I use a protocol or a category with both prototypes and classes?

Yes. A protocol may be implemented by both prototypes and classes in the same application. Likewise, a category may be imported by both prototypes and classes in the same application.

4.5.5 What support is provided in Logtalk for defining and using components?

Logtalk supports component-based programming (since its inception on January 1998), by using *categories* (which are first-class entities like objects and protocols). Logtalk categories can be used with both classes and prototypes and are inspired on the Smalltalk-80 (documentation-only) concept of method categories and on Objective-C categories, hence the name. For more information, please consult the *Categories* section and the examples provided with the distribution.

4.5.6 What support is provided in Logtalk for reflective programming?

Logtalk supports meta-classes, behavioral reflection through the use of event-driven programming, and structural reflection through the use of a set of built-in predicates and built-in methods which allow us to query the system about existing entities, entity relations, and entity predicates.

4.6 Troubleshooting

- *Using compiler options on calls to the Logtalk compiling and loading predicates do not work!*
- *Gecko-based browsers (e.g. Firefox) show non-rendered HTML entities when browsing XML documenting files!*
- *Compiling a source file results in errors or warnings but the Logtalk compiler reports a successful compilation with zero errors and zero warnings!*

4.6.1 Using compiler options on calls to the Logtalk compiling and loading predicates do not work!

Using compiler options on calls to the Logtalk *logtalk_compile/2* and *logtalk_load/2* built-in predicates only apply the file being compiled. If the first argument is a *loader file*, the compiler options will only be used in the compilation of the loader file itself, not in the compilation of the files loaded by the loader file. The solution is to edit the loader file and add the compiler options to the calls that compile/load the individual files.

4.6.2 Gecko-based browsers (e.g. Firefox) show non-rendered HTML entities when browsing XML documenting files!

Using Gecko-based browsers (e.g. Firefox) show non-rendered HTML entities (e.g. `–`) when browsing XML documenting files after running the `lgt2xml` shell script in the directory containing the XML documenting files. This is a consequence of the lack of support for the `disable-output-escaping` attribute in the browser XSLT processor. The workaround is to use other browser (e.g. Safari or Opera) or to use instead the `lgt2html` shell script in the directory containing the XML documenting files to convert them to (X)HTML files for browsing.

4.6.3 Compiling a source file results in errors or warnings but the Logtalk compiler reports a successful compilation with zero errors and zero warnings!

This may happen when your Prolog compiler implementation of the ISO Prolog standard `write_canonical/2` built-in predicate is buggy and writes terms that cannot be read back when consulting the intermediate Prolog files generated by the Logtalk compiler. Often, syntax errors found when consulting result in error messages but not in exceptions as the Prolog compiler tries to continue the compilation despite the problems found. As the Logtalk compiler relies on the exception mechanisms to catch compilation problems, it may report zero errors and zero warnings despite the error messages. Send a bug report to the Prolog compiler developers asking them to fix the `write_canonical/2` buggy implementation.

4.7 Usability

- *Is there a shortcut for compiling and loading source files?*

- *Is there an equivalent directive to the ensure_loaded/1 Prolog directive?*

- *Are there shortcuts for the make functionality?*

4.7.1 Is there a shortcut for compiling and loading source files?

Yes. With most back-end Prolog compilers, you can use `{File}` as a shortcut for `logtalk_load(File)`. For compiling and loading multiple files simply use `{File1, File2, ...}`. See the documentation of the *logtalk_load/1* predicate for details.

4.7.2 Is there an equivalent directive to the `ensure_loaded/1` Prolog directive?

You can use the goal `logtalk_load(File, [reload(skip)])` to ensure that `File` is only loaded once. See the documentation of the *logtalk_load/2* predicate for details.

4.7.3 Are there shortcuts for the make functionality?

Yes. With most back-end Prolog compilers, you can use `{*}` as a shortcut for `logtalk_make(all)` to reload all files modified since last compiled and loaded, `{!}` as a shortcut for `logtalk_make(clean)` to delete all intermediate Prolog files generated by the compilation of Logtalk source files, `{?}` as a shortcut for `logtalk_make(missing)` to list missing entities and predicates, and `{@}` as a shortcut for `logtalk_make(circular)` to list circular references. See the documentation of the *logtalk_make/1* predicate for details.

4.8 Deployment

- *Can I create standalone applications with Logtalk?*

4.8.1 Can I create standalone applications with Logtalk?

It depends on the Prolog compiler that you use to run Logtalk. Assuming that your Prolog compiler supports the creation of standalone executables, your application must include the adapter file for your compiler and the Logtalk

compiler and runtime. The distribution includes embedding scripts for selected backend Prolog compilers and embedding examples.

For instructions on how to embed Logtalk and Logtalk applications see the wiki section on embedding.

4.9 Performance

- *Is Logtalk implemented as a meta-interpreter?*
- *What kind of code Logtalk generates when compiling objects? Dynamic code? Static code?*
- *How about message-sending performance? Does Logtalk use static binding or dynamic binding?*
- *How does Logtalk performance compare with plain Prolog and with Prolog modules?*

4.9.1 Is Logtalk implemented as a meta-interpreter?

No. Objects and their encapsulated predicates are compiled, not meta-interpreted. In particular, inheritance relations are pre-compiled for improved performance. Moreover, no meta-interpreter is used even for objects compiled in debug mode.

4.9.2 What kind of code Logtalk generates when compiling objects? Dynamic code? Static code?

Static objects are compiled to static code. Static objects containing dynamic predicates are also compiled to static code, except, of course, for the dynamic predicates themselves. Dynamic objects are necessarily compiled to dynamic code. As in Prolog programming, for best performance, dynamic object predicates and dynamic objects should only be used when truly needed.

4.9.3 How about message-sending performance? Does Logtalk use static binding or dynamic binding?

Logtalk supports both static binding and dynamic binding. When static binding is not possible, Logtalk uses dynamic binding coupled with a caching mechanism that avoids repeated lookups of predicate declarations and predicate definitions. This is a solution common to other programming languages supporting dynamic binding. Message lookups are automatically cached the first time a message is sent. Cache entries are automatically removed when loading entities or using Logtalk dynamic features that invalidate the cached lookups. Whenever static binding is used, message sending performance is essentially the same as a predicate call in plain Prolog. Performance of dynamic binding when lookups are cached is close to the performance that would be achieved with static binding. See the wiki section on performance for more details.

4.9.4 Which Prolog-dependent factors are most crucial for good Logtalk performance?

Logtalk compiles objects assuming first-argument indexing for static code. First-argument indexing of dynamic code, when available, helps to improve performance due to the automatic caching of method lookups and the necessary use of book-keeping tables by the runtime engine (this is specially important when using event-driven programming). Dynamic objects and static objects containing dynamic predicates also benefit from first-argument indexing of dynamic predicates. The availability of multi-argument indexing, notably for dynamic predicates, also benefits dynamic binding performance.

4.9.5 How does Logtalk performance compare with plain Prolog and with Prolog modules?

Plain Prolog, Prolog modules, and Logtalk objects provide different trade-offs between performance and features. In general, for a given predicate definition, the best performance will be attained using plain Prolog, second will be Prolog modules (assuming no explicitly qualified calls are used), and finally Logtalk objects. Whenever static binding is used, the performance of Logtalk is equal or close to that of plain Prolog (depending on the Prolog virtual machine implementation and compiler optimizations). See the simple benchmark test results using some popular Prolog compilers.

4.10 Licensing

- *What's the Logtalk distribution license?*
- *Can Logtalk be used in commercial applications?*
- *What's the final license for a combination of Logtalk with a Prolog compiler?*

4.10.1 What's the Logtalk distribution license?

Logtalk follows the Apache License 2.0.

4.10.2 Can Logtalk be used in commercial applications?

Yes, the Apache License 2.0 allows commercial use. See e.g. the Apache License and Distribution FAQ.

4.10.3 What's the final license for a combination of Logtalk with a Prolog compiler?

See the wiki page on licensing for details and relevant resources.

4.11 Support

- *Are there professional consulting, training and supporting services?*

4.11.1 Are there professional consulting, training and supporting services?

Yes. Please visit logtalk.pt for professional consulting, developing, training, and other supporting services.

FIVE

GLOSSARY

abstract class A class that cannot be instantiated. Usually used to contain common predicates that are inherited by other classes.

abstract method A method implementing an algorithm whose step corresponds to calls to methods defined in the descendants of the object (or category) containing it.

adapter file A Prolog source file defining a minimal abstraction layer between the Logtalk compiler/runtime and a specific backend Prolog compiler.

ancestor A class or a parent prototype that contributes (via inheritance) to the definition of an object. For class-based hierarchies, the ancestors of an instance are its class(es) and all the superclasses of its class(es). For prototype-based hierarchies, the ancestors of a prototype are its parent(s) and the ancestors of its parent(s).

built-in method A predefined method that can be called from within any object or category. Built-in methods cannot be redefined.

built-in predicate A predefined predicate that can be called from anywhere. Built-in predicates can be redefined within objects and categories.

category A set of predicates directives and clauses that can be (virtually) imported by any object. Categories support composing objects using fine-grained units of code reuse and also hot patching of existing objects. A category should be functionally-cohesive, defining a single functionality.

class An object that specializes another object, interpreted as its superclass. Classes define the common predicates of a set of objects that instantiates it. An object can also be interpreted as a class when it instantiates itself.

closed-world assumption The assumption that what cannot be proved true is false. Therefore, sending a message corresponding to a *declared* but not *defined* predicate, or calling a declared predicate with no clauses, fails. But messages or calls to undeclared predicates generate an error.

coinductive predicate A predicate whose calls are proved using greatest fixed point semantics. Coinductive predicates allows reasoning about infinite rational entities such as cyclic terms and ω-automata.

complementing category A category used for hot patching an existing object (or a set of objects).

directive A source file term that affects the interpretation of source code. Directives use the `(:-)/1` prefix operator as functor.

doclet file A source file whose main purpose is to generate documentation for e.g. a library or an application.

doclet object An object specifying the steps necessary to (re)generate the API documentation for a project. See the `doclet` and `lgtdoc` tool for details.

dynamic binding Runtime lookup of a predicate declaration and definition to verify the validity of a message (or a super call) and find the predicate definition that will be used to answer the message (or the super call). Also known as *late binding*. See also *static binding*.

encapsulation The hiding of an object implementation. This promotes software reuse by isolating the object clients from its implementation details. Encapsulation is enforced in Logtalk by using predicate scope directives.

entity Generic name for Logtalk compilation units: objects, categories, and protocols. Entities share a single namespace (i.e. entity identifiers must be unique).

entity directive A directive that affects how Logtalk entities (objects, protocols, or categories) are used or compiled.

event The sending of a message to an object. An event can be expressed as an ordered tuple: (`Event, Object, Message, Sender`). Logtalk distinguish between the sending of a message — `before` event — and the return of control to the sender — `after` event.

grammar rule An alternative notation for predicates used to parse or generate sentences on some language. This notation hides the arguments used to pass the sequences of tokens being processed, thus simplifying the representation of grammars. Grammar rules are represented using as functor the infix operator (`-->`)/2 instead of the (`:-`)/2 operator used with predicate clauses.

grammar rule non-terminal A syntactic category of words or phrases. A non-terminal is identified by its *non-terminal indicator*, i.e. by its name and number of arguments using the notation `Name//Arity`.

grammar rule terminal A word or basic symbol of a language.

hook object An object, implementing the `expanding` built-in protocol, defining term- and goal-expansion clauses, used in the compilation of Logtalk or Prolog source files. A hook object can be specified using the compiler flag `hook/1`. It can also be specified using a `set_logtalk_flag/2` directive in the source files to be expanded.

identity Property of an entity that distinguishes it from every other entity. Object and category identifiers can be atoms or compound terms. Protocol identities must be atoms. All Logtalk entities (objects, protocols, and categories) share the same namespace.

inheritance An entity inherits predicate directives and clauses from related entities. In the particular case of objects, when an object extends other object, we have prototype-based inheritance. When an object specializes or instantiates another object, we have class-based inheritance.

instance An object that instantiates another object, interpreted as its class.

instantiation The process of creating a new class instance. In Logtalk, this does not necessarily imply dynamic creation of an object at runtime; an instance may also be defined as a static object in a source file.

lambda expression A compound term that can be used in place of a goal or closure meta-argument and that abstracts a predicate definition by listing its variables and a callable term that implements the definition. Lambda expressions help avoiding the need of naming and defining auxiliary predicates.

lambda free variable A variable that is global to a *lambda expression*. All used global variables must be explicitly listed in a lambda expression.

lambda parameter A term (usually a variable or a non-ground compound term) that is local to a *lambda expression*. All lambda parameters must be explicitly enumerated in a lambda expression.

library A directory containing source files. The library name can be used as an alias to the directory path when compiling and loading source files using the notation `library_name('source_file_relative_path')`.

loader file A source file whose main purpose is to load a set of source files.

local predicate A predicate that is defined in an object (or in a category) but that is not listed in a scope directive. These predicates behave like private predicates but are invisible to the reflection built-in methods. Local predicates are usually auxiliary predicates and only relevant to the entity where they are defined.

message A query sent to an object. In logical terms, a message can be seen as a request for proof construction using an object's database and the databases of related entities.

message to self A message sent to the object that received the original message under processing. Messages to self require *dynamic binding* as the value of self is only know at runtime.

meta-argument A predicate argument that is called as a goal, used as a closure to construct a goal that will be called, or that is handled in a way that requires awareness of the predicate calling context.

meta-interpreter A program capable of running other programs written in the same language.

meta-predicate A predicate with one or more *meta-arguments*. For example, `call/1` and `findall/3` are built-in meta-predicates.

metaclass The class of a class, when interpreted as an instance. Metaclass instances are themselves classes.

method The predicate definition used to answer a message sent to an object. Logtalk supports both *static binding* and dynamic binding to find which method to run to answer a message.

module A Prolog entity characterized by an identity and a set of predicate directives and clauses. Prolog modules are usually static although some Prolog systems allow the creation of dynamic modules at runtime. Prolog modules can be seen as prototypes.

monitor Any object, implementing the `monitoring` built-in protocol, that is notified by the runtime when a spied event occurs. The spied *events* can be set by the monitor itself or by any other object.

multifile predicate A predicate whose clauses can be defined in multiple entities. The object or category holding the directive without an entity prefix qualifying the predicate holds the multifile predicate *primary declaration*, which consists of both a scope directive and a `multifile/1` directive for the predicate.

object An entity characterized by an identity and a set of predicate directives and clauses. Logtalk objects can be either static or dynamic. Logtalk objects can play the role of classes, instances, or prototypes. The role or roles an object plays depends on its relations with other objects.

object database The set of predicates locally defined inside an object.

parameter An argument of a parametric object or a parametric category identifier. Parameters are logical variables implicitly shared by all the entity predicate clauses.

parameter variable A variable used as parameter in a parametric object or a parametric category using the syntax `_VariableName_`. Occurrences of parameter variables in entity clauses are implicitly unified with the corresponding entity parameters.

parametric category See parametric entity.

parametric entity An object or category whose identifier is a compound term containing free variables that can be used to parameterize the entity predicates. Parameters are *logical variables* implicitly shared by all the entity clauses.

parametric object See parametric entity.

parametric object proxy A compound term (usually represented as a plain Prolog fact) with the same functor and with the same number of arguments as the identifier of a parametric object.

parent A prototype that is extended by another prototype.

predicate Predicates describe what is true about the application domain. A predicate is identified by its *predicate indicator*, i.e. by its name and number of arguments using the notation `Name/Arity`.

predicate alias An alternative functor (`Name/Arity`) for a predicate. Predicate aliases can be defined for any inherited predicate using the `alias/2` directive and for predicates listed in `uses/2` and `use_module` directives. Predicate aliases can be used to solve inheritance conflicts and to improve code clarity by using alternative names that are more meaningful in the calling context.

predicate directive A directive that affects how predicates are called or compiled.

predicate scope container The object that inherits a predicate declaration from an imported category or an implemented protocol.

primary predicate declaration See *multifile predicate*.

private inheritance All public and protected predicates are inherited as private predicates.

private predicate A predicate that can only be called from the object that contains the scope directive.

profiler A program that collects data about other program performance.

protected inheritance All public predicates are inherited as protected. No change for protected or private predicates.

protected predicate A predicate that can only be called from the object containing the scope directive or from an object that inherits the predicate.

protocol An entity that contains predicate declarations. A predicate is declared using a scope directive. It may be further specified by additional predicate directives. Protocols support the separation between interface and implementation, can be implemented by both objects and categories, and can be extended by other protocols. A protocol should be functionally-cohesive, specifying a single functionality.

prototype A self-describing object that may extend or be extended by other objects. An object with no instantiation or specialization relations with other objects is always interpreted as a prototype.

public inheritance All inherited predicates maintain their declared scope.

public predicate A predicate that can be called from any object.

self The object that received the message under processing.

sender An object that sends a message to other object. When a message is sent from within a category, the *sender* is the object importing the category.

settings file A source file, compiled and loaded at Logtalk startup, mainly defining default values for compiler flags that override the defaults found on the backend Prolog compiler adapter files.

singleton method A method defined in an instance itself. Singleton methods are supported in Logtalk and can also be found in other object-oriented programming languages.

source file A text file defining Logtalk and/or Prolog code. Multiple Logtalk entities may be defined in a single source file. Prolog code may be intermixed with entity definitions.

source file directive A directive that affects how a source file is compiled.

specialization A class is specialized by defining a new class that inherit its predicates and possibly add new ones.

static binding Compile time lookup of a predicate declaration and definition when compiling a message sending call (or a super call). Dynamic binding is used whenever static binding is not possible (e.g. due to the predicate being dynamic or due to lack of enough information). Also known as *early binding*. See also *dynamic binding*.

subclass A class that is a specialization, direct or indirectly, of another class.

super call Call of an inherited (or imported) predicate definition. Mainly used when redefining an inherited (or imported) predicate to call the overridden definition while making additional calls. Super calls preserve self and may require *dynamic binding* if the predicate is dynamic.

superclass A class from which another class is a specialization (directly or indirectly via another class).

synchronized predicate A synchronized predicate is protected by a mutex ensuring that, in a multi-threaded application, it can only be called by a single thread at a time.

tester file A source file whose main purpose is to load and a run a set of unit tests.

this The object that contains the predicate clause under execution. When the predicate clause is contained in a category, *this* is a reference to the object importing the category for which the predicate clause is being executed.

visible predicate A predicate that is declared for an object, a built-in method, a Logtalk built-in predicate, or a Prolog built-in predicate.

- genindex

BIBLIOGRAPHY

[Alexiev93] Mutable Object State for Object-Oriented Logic Programming: A Survey Alexiev, V. Technical Report TR 93-15, Department of Computing Science, University of Alberta, Canada

[Belli_et_al_92] Object-oriented programming in Prolog: rationale and a case study Belli, F., Jack, O., Naish, L. Technical Report 92/2, Department of Electrical and Electronics Engineering, University of Paderborn, Germany URL: http://www.cs.mu.oz.au/~lee/papers/oolp/

[Block89] An Extended Frame Language Block, F. P., Chan, N. C. Proceedings OOPLSLA 89(10):151-157, ACM

[Bobrow_et_al_88] Common Lisp Object System Specification Bobrow, D. G., Michiel, L. G., Gabriel, R. P., Keene, S. E., Kiczales, G., Moon, D. A. ACM SIGPLAN Notices(23)

[Bratko90] Prolog Programming for Artificial Intelligence Bratko, I. Addison Wesley, 2º edition, 1990

[Champaux92] A comparative Study of Object-Oriented Analysis Methods Champaux, D., Faure, P. Journal of Object-Oriented Programming, Vol. 5, N.1, 1992

[Clocksin87] Programming in Prolog Clocksin, W.F., Mellish, C.S. Springer-Verlag, New York, 1987

[Cointe87] Metaclasses are First Class: the ObjVlisp Model Cointe, P. Proceedings OOPLSLA 87(10):156-167, ACM

[Cordes91] The Literate Programming Paradigm Cordes, D., Brown, M. IEEE Computer, June 1991:52-61

[Covington94] ISO Prolog: A Summary of the Draft Proposed Standard Covington, M. A. URL: ftp://ai.uga.edu/pub/prolog.standard/

[Cox86] Object-Oriented Programming: An Evolutionary Approach Cox, Brad J. Addison-Wesley Publishing Company, Don Mills, Ontario

[Davison89] Polka: A Parlog Object oriented language Davison, A. Ph.D. Thesis, Imperial College, London, 1989

[Davison92] A survey of logic programming-based object oriented languages Davison, A. Tech Report 92/3, Dept. of Computer Science, University of Melbourne, Australia URL: http://www.cs.mu.oz.au/tr_db/mu_92_03.ps.gz

[Davison93] The deductive and object oriented features of BeBOP Davison, A. Tech Report 93/6, Dept. of Computer Science, University of Melbourne, Australia URL:http://www.cs.mu.oz.au/tr_db/mu_93_06.ps.gz

[Delzanno97] Logic and Object-Oriented Programming in Linear Logic Delzanno, G. Ph.D. Thesis, University of Pisa, Italy URL:http://www.mpi-sb.mpg.de/~delzanno/

[Dony90] Exception Handling and Object-Oriented Programming: Towards a Synthesis Dony, C. Proceedings OOPLSLA 90:322-330, ACM

[Fornarino_et_al_89] An Original Object-Oriented Approach for Relation Management Fornarino, M., Pinna, A.-M.,Trousse, B. Proceedings of the 4th Portuguese Conference on Artificial Intelligence Lecture Notes in Artificial Intelligence, Springer-Verlag (390):13-26

[Fromherz93] OL(P): Object Layer for Prolog Fromherz, M. URL: ftp://parcftp.xerox.com/ftp/pub/ol/

[Fukunaga86] An Experience with a Prolog-based Object-Oriented Language Fukunaga, K., Hirose, S. Proceedings OOPLSLA 86, 21(11):224-231, ACM

[Goldberg83] Smalltalk-80 The language and its implementation Goldberg, A., Robson, D. Addison-Wesley Series in Computer Science

[Joy_et_al_00] The Java Language Specification, Second Edition Joy, B., Steele, G., Gosling, J., Bracha, G. Addison-Wesley, 2000

[ISO95] ISO/IEC DIS 13211-1 - Programming Language Prolog Part 1: General Core Joint Technical Committee ISO/IEC JTC 1 URL: https://www.iso.org/standard/21413.html

[Knuth84] Literate Programming Knuth, D. E. Computer Journal, May 84, 27(2):97-111

[Lieberman86] Using Prototypical Objects to Implement Shared Behaviour in Object Oriented Systems Lieberman, H. Proceedings OOPLSLA 86:189-214, ACM

[Maes87] Concepts and Experiments in Computational Reflection Maes, P. Proceedings OOPLSLA 87, ACM

[McCabe92] Logic and Objects McCabe, F. G. Prentice Hall Series in Computer Science

[Moon86] Object-Oriented Programming in Flavors Moon, D. Proceedings OOPLSLA 86:1-8, ACM

[Moss94] Prolog++ The Power of Object-Oriented and Logic Programming Moss, C. Addison-Wesley International Series in Logic Programming, 1994

[Moura94] Logtalk: Programação Orientada para Objectos em Prolog Moura, P., Costa, E. 2ª Conferência e Exposição Portuguesa de Tecnologia Orientada por Objectos 3i Consultores, Lisboa

[Moura99] Porting Prolog: Notes on porting a Prolog program to 22 Prolog compilers or the relevance of the ISO Prolog standard Moura, P. ALP Newsletter, Vol. 12/2, May 1999

[Moura00] Logtalk 2.6 Documentation Moura, P. Technical Report DMI 2000/1 University of Beira Interior, Portugal

[Razek92] Combining Objects and Relations Razek, G. Comunications of the ACM, 27(12):66-70

[Rumbaugh87] Relations as Semantic Constructs in an Object-Oriented Language Rumbaugh, J. Proceedings OOPLSLA 87:466-481, ACM

[Rumbaugh88] Controlling Propagation of Operations using Attributes on Relations Rumbaugh, J. Proceedings OOPLSLA 88:285-296, ACM

[Schachte95] Efficient Object-Oriented Programming in Prolog Schachte, P., Saab, G. Logic Programming: Formal Methods and Pratical Applications Studies in Computer Science and Artificial Intelligence, 11 Elsevier Science B.V. North-Holland, Amsterdam, 1995

[SICStus95] SICStus Prolog Manual SICStus URL: http://www.sics.se/ps/sicstus.html

[Shan_et_al_93] Is Multiple Inheritance Essential to OOP? (Panel) Shan, Y., Cargill, T., Cox, B., Cook, W., Loomis, M., Snyder, A. Proceedings OOPLSLA 93:360-363

[Stefik_et_al_86] Integrating Acess-Oriented Programming into a Multiparadigm Environment Stefik, M. J., Bobrow, D. G. , Kahn, K. M. IEEE Software, January 1986:10-18

[Stroustrup86] The C++ Programming Language Stroustrup, B. Addison-Wesley Series in Computer Science

[Taenzer89] Problems in Object-Oriented Software Reuse Taenzer, D., Ganti, M., Podar, S. Proceedings of ECOOP 89 British Computer Society Workshop Series, Cambridge University Press

[Tanzer95] Remarks on Object-Oriented Modeling of Associations Tanzer, C. Journal of Object-Oriented Programming, February 1995, SIGS Publications

[Tanenbaum87] Operating Systems - Design and Implementation Tanenbaum, A. Prentice-Hall Software Series, 1987

[Welsch89] Reasoning Objects with Dynamic Knowledge Bases Welsch, C., Barth, G. Proceedings of the 4th Portuguese Conference on Artificial Intelligence(390):257-268 Lecture Notes in Artificial Intelligence, Springer-Verlag, 1989